

What is rustc?

Welcome to "The rustc book"! `rustc` is the compiler for the Rust programming language, provided by the project itself. Compilers take your source code and produce binary code, either as a library or executable.

Most Rust programmers don't invoke `rustc` directly, but instead do it through Cargo. It's all in service of `rustc` though! If you want to see how Cargo calls `rustc`, you can

```
$ cargo build --verbose
```

And it will print out each `rustc` invocation. This book can help you understand what each of these options does. Additionally, while most Rustaceans use Cargo, not all do: sometimes they integrate `rustc` into other build systems. This book should provide a guide to all of the options you'd need to do so.

Basic usage

Let's say you've got a little hello world program in a file `hello.rs`:

```
fn main() {
    println!("Hello, world!");
}
```

To turn this source code into an executable, you can use `rustc`:

```
$ rustc hello.rs
$ ./hello # on a *NIX
$ .\hello.exe # on Windows
```

Note that we only ever pass `rustc` the *crate root*, not every file we wish to compile. For example, if we had a `main.rs` that looked like this:

```
mod foo;

fn main() {
    foo::hello();
}
```

And a `foo.rs` that had this:

```
pub fn hello() {
    println!("Hello, world!");
}
```

To compile this, we'd run this command:

```
$ rustc main.rs
```

No need to tell `rustc` about `foo.rs`; the `mod` statements give it everything that it needs. This is different than how you would use a C compiler, where you invoke the compiler on each file, and then link everything together. In other words, the `crate` is a translation unit, not a particular module.

Command-line Arguments

Here's a list of command-line arguments to `rustc` and what they do.

-h/**--help**: get help

This flag will print out help information for `rustc`.

--cfg: configure the compilation environment

This flag can turn on or off various `#[cfg]` settings for [conditional compilation](#).

The value can either be a single identifier or two identifiers separated by `=`.

For example, `--cfg 'verbose'` or `--cfg 'feature="serde"`. These correspond to `#[cfg(verbose)]` and `#[cfg(feature = "serde")]` respectively.

-L: add a directory to the library search path

The `-L` flag adds a path to search for external crates and libraries.

The kind of search path can optionally be specified with the form `-L KIND=PATH` where `KIND` may be one of:

- `dependency` – Only search for transitive dependencies in this directory.
- `crate` – Only search for this crate's direct dependencies in this directory.
- `native` – Only search for native libraries in this directory.
- `framework` – Only search for macOS frameworks in this directory.
- `all` – Search for all library kinds in this directory. This is the default if `KIND` is not specified.

-l: link the generated crate to a native library

Syntax: `-l [KIND[:MODIFIERS]=]NAME[:RENAME]`.

This flag allows you to specify linking to a specific native library when building a crate.

The kind of library can optionally be specified with the form `-l KIND=lib` where `KIND` may be one of:

- `dylib` – A native dynamic library.
- `static` – A native static library (such as a `.a` archive).
- `framework` – A macOS framework.

If the kind is specified, then linking modifiers can be attached to it. Modifiers are specified as a comma-delimited string with each modifier prefixed with either a `+` or `-` to indicate that the modifier is enabled or disabled, respectively.

Specifying multiple `modifiers` arguments in a single `link` attribute, or multiple identical modifiers in the same `modifiers` argument is not currently supported.

Example: `-l static:+whole-archive=mylib`.

The kind of library and the modifiers can also be specified in a `##[link] attribute`. If the kind is not specified in the `link` attribute or on the command-line, it will link a dynamic library if available, otherwise it will use a static library. If the kind is specified on the command-line, it will override the kind specified in a `link` attribute.

The name used in a `link` attribute may be overridden using the form `-l ATTR_NAME:LINK_NAME` where `ATTR_NAME` is the name in the `link` attribute, and `LINK_NAME` is the name of the actual library that will be linked.

Linking modifiers: `whole-archive`

This modifier is only compatible with the `static` linking kind. Using any other kind will result in a compiler error.

`+whole-archive` means that the static library is linked as a whole archive without throwing any object files away.

This modifier translates to `--whole-archive` for `ld`-like linkers, to `/WHOLEARCHIVE` for `link.exe`, and to `-force_load` for `ld64`. The modifier does nothing for linkers that don't support it.

The default for this modifier is `-whole-archive`.

NOTE: The default may currently be different in some cases for backward compatibility, but it is not guaranteed. If you need whole archive semantics use `+whole-archive` explicitly.

Linking modifiers: `bundle`

This modifier is only compatible with the `static` linking kind. Using any other kind will result in a compiler error.

When building a `rlib` or `staticlib` `+bundle` means that the native static library will be packed into the `rlib` or `staticlib` archive, and then retrieved from there during linking of the final binary.

When building a rlib `-bundle` means that the native static library is registered as a dependency of that rlib "by name", and object files from it are included only during linking of the final binary, the file search by that name is also performed during final linking.

When building a staticlib `-bundle` means that the native static library is simply not included into the archive and some higher level build system will need to add it later during linking of the final binary.

This modifier has no effect when building other targets like executables or dynamic libraries.

The default for this modifier is `+bundle`.

Linking modifiers: `verbatim`

This modifier is compatible with all linking kinds.

`+verbatim` means that rustc itself won't add any target-specified library prefixes or suffixes (like `lib` or `.a`) to the library name, and will try its best to ask for the same thing from the linker.

For `ld`-like linkers supporting GNU extensions rustc will use the `-l:filename` syntax (note the colon) when passing the library, so the linker won't add any prefixes or suffixes to it. See `-l namespec` in `ld` documentation for more details.

For linkers not supporting any verbatim modifiers (e.g. `link.exe` or `ld64`) the library name will be passed as is. So the most reliable cross-platform use scenarios for this option are when no linker is involved, for example bundling native libraries into rlibs.

`-verbatim` means that rustc will either add a target-specific prefix and suffix to the library name before passing it to linker, or won't prevent linker from implicitly adding it.

In case of `raw-dylib` kind in particular `.dll` will be added to the library name on Windows.

The default for this modifier is `-verbatim`.

NOTE: Even with `+verbatim` and `-l:filename` syntax `ld`-like linkers do not typically support passing absolute paths to libraries. Usually such paths need to be passed as input files without using any options like `-l`, e.g. `ld /my/absolute/path`.

`-Clink-arg=/my/absolute/path` can be used for doing this from stable `rustc`.

--crate-type: a list of types of crates for the compiler to emit

This instructs `rustc` on which crate type to build. This flag accepts a comma-separated list of values, and may be specified multiple times. The valid crate types are:

- `lib` – Generates a library kind preferred by the compiler, currently defaults to `rlib`.
- `rlib` – A Rust static library.
- `staticlib` – A native static library.
- `dylib` – A Rust dynamic library.
- `cdylib` – A native dynamic library.
- `bin` – A runnable executable program.
- `proc-macro` – Generates a format suitable for a procedural macro library that may be loaded by the compiler.

The crate type may be specified with the `crate_type` attribute. The `--crate-type` command-line value will override the `crate_type` attribute.

More details may be found in the [linkage chapter](#) of the reference.

--crate-name: specify the name of the crate being built

This informs `rustc` of the name of your crate.

--edition: specify the edition to use

This flag takes a value of `2015`, `2018` or `2021`. The default is `2015`. More information about editions may be found in the [edition guide](#).

--emit: specifies the types of output files to generate

This flag controls the types of output files generated by the compiler. It accepts a comma-separated list of values, and may be specified multiple times. The valid emit kinds are:

- `asm` – Generates a file with the crate's assembly code. The default output filename is `CRATE_NAME.s`.
- `dep-info` – Generates a file with Makefile syntax that indicates all the source files that were loaded to generate the crate. The default output filename is `CRATE_NAME.d`.

- `link` – Generates the crates specified by `--crate-type`. The default output filenames depend on the crate type and platform. This is the default if `--emit` is not specified.
- `llvm-bc` – Generates a binary file containing the LLVM bitcode. The default output filename is `CRATE_NAME.bc`.
- `llvm-ir` – Generates a file containing LLVM IR. The default output filename is `CRATE_NAME.ll`.
- `metadata` – Generates a file containing metadata about the crate. The default output filename is `libCRATE_NAME.rmeta`.
- `mir` – Generates a file containing rustc's mid-level intermediate representation. The default output filename is `CRATE_NAME.mir`.
- `obj` – Generates a native object file. The default output filename is `CRATE_NAME.o`.

The output filename can be set with the `-o` flag. A suffix may be added to the filename with the `-c extra-filename` flag. The files are written to the current directory unless the `--out-dir` flag is used. Each emission type may also specify the output filename with the form `KIND=PATH`, which takes precedence over the `-o` flag.

`--print`: print compiler information

This flag prints out various information about the compiler. This flag may be specified multiple times, and the information is printed in the order the flags are specified. Specifying a `--print` flag will usually disable the `--emit` step and will only print the requested information. The valid types of print values are:

- `crate-name` – The name of the crate.
- `file-names` – The names of the files created by the `link` emit kind.
- `sysroot` – Path to the sysroot.
- `target-libdir` – Path to the target libdir.
- `cfg` – List of cfg values. See [conditional compilation](#) for more information about cfg values.
- `target-list` – List of known targets. The target may be selected with the `--target` flag.
- `target-cpus` – List of available CPU values for the current target. The target CPU may be selected with the `-C target-cpu=val` flag.
- `target-features` – List of available target features for the current target. Target features may be enabled with the `-C target-feature=val` flag. This flag is unsafe. See [known issues](#) for more details.
- `relocation-models` – List of relocation models. Relocation models may be selected with the `-C relocation-model=val` flag.
- `code-models` – List of code models. Code models may be selected with the `-C code-model=val` flag.
- `tls-models` – List of Thread Local Storage models supported. The model may be selected with the `-Z tls-model=val` flag.

- `native-static-libs` – This may be used when creating a `staticlib` crate type. If this is the only flag, it will perform a full compilation and include a diagnostic note that indicates the linker flags to use when linking the resulting static library. The note starts with the text `native-static-libs:` to make it easier to fetch the output.
- `link-args` – This flag does not disable the `--emit` step. When linking, this flag causes `rustc` to print the full linker invocation in a human-readable form. This can be useful when debugging linker options. The exact format of this debugging output is not a stable guarantee, other than that it will include the linker executable and the text of each command-line argument passed to the linker.

-g: include debug information

A synonym for `-C debuginfo=2`.

-O: optimize your code

A synonym for `-C opt-level=2`.

-o: filename of the output

This flag controls the output filename.

--out-dir: directory to write the output in

The outputted crate will be written to this directory. This flag is ignored if the `-p` flag is used.

--explain: provide a detailed explanation of an error message

Each error of `rustc`'s comes with an error code; this will print out a longer explanation of a given error.

--test: build a test harness

When compiling this crate, `rustc` will ignore your `main` function and instead produce a test harness. See the [Tests chapter](#) for more information about tests.

--target: select a target triple to build

This controls which `target` to produce.

-W: set lint warnings

This flag will set which lints should be set to the [warn level](#).

Note: The order of these lint level arguments is taken into account, see [lint level via compiler flag](#) for more information.

--force-warn: force a lint to warn

This flag sets the given lint to the [forced warn level](#) and the level cannot be overridden, even ignoring the [lint caps](#).

-A: set lint allowed

This flag will set which lints should be set to the [allow level](#).

Note: The order of these lint level arguments is taken into account, see [lint level via compiler flag](#) for more information.

-D: set lint denied

This flag will set which lints should be set to the [deny level](#).

Note: The order of these lint level arguments is taken into account, see [lint level via compiler flag](#) for more information.

-F: set lint forbidden

This flag will set which lints should be set to the [forbid level](#).

Note: The order of these lint level arguments is taken into account, see [lint level via compiler flag](#) for more information.

-Z: set unstable options

This flag will allow you to set unstable options of rustc. In order to set multiple options, the -Z flag can be used multiple times. For example: `rustc -Z verbose -Z time-passes`. Specifying options with -Z is only available on nightly. To view all available options run: `rustc -Z help`, or see [The Unstable Book](#).

--cap-lints: set the most restrictive lint level

This flag lets you 'cap' lints, for more, see [here](#).

-C/--codegen: code generation options

This flag will allow you to set [codegen options](#).

-V/--version: print a version

This flag will print out `rustc`'s version.

-v/--verbose: use verbose output

This flag, when combined with other flags, makes them produce extra output.

--extern: specify where an external library is located

This flag allows you to pass the name and location for an external crate of a direct dependency. Indirect dependencies (dependencies of dependencies) are located using the `-L` flag. The given crate name is added to the [extern prelude](#), similar to specifying `extern crate` within the root module. The given crate name does not need to match the name the library was built with.

Specifying `--extern` has one behavior difference from `extern crate`: `--extern` merely makes the crate a *candidate* for being linked; it does not actually link it unless it's actively used. In rare occasions you may wish to ensure a crate is linked even

if you don't actively use it from your code: for example, if it changes the global allocator or if it contains `#[no_mangle]` symbols for use by other programming languages. In such cases you'll need to use `extern crate`.

This flag may be specified multiple times. This flag takes an argument with either of the following formats:

- `CRATENAME=PATH` – Indicates the given crate is found at the given path.
- `CRATENAME` – Indicates the given crate may be found in the search path, such as within the sysroot or via the `-L` flag.

The same crate name may be specified multiple times for different crate types. If both an `rlib` and `dylib` are found, an internal algorithm is used to decide which to use for linking. The `-C prefer-dynamic` flag may be used to influence which is used.

If the same crate name is specified with and without a path, the one with the path is used and the pathless flag has no effect.

--sysroot: Override the system root

The "sysroot" is where `rustc` looks for the crates that come with the Rust distribution; this flag allows that to be overridden.

--error-format: control how errors are produced

This flag lets you control the format of messages. Messages are printed to stderr. The valid options are:

- `human` – Human-readable output. This is the default.
- `json` – Structured JSON output. See the [JSON chapter](#) for more detail.
- `short` – Short, one-line messages.

--color: configure coloring of output

This flag lets you control color settings of the output. The valid options are:

- `auto` – Use colors if output goes to a tty. This is the default.
- `always` – Always use colors.
- `never` – Never colorize output.

--diagnostic-width: specify the terminal width for diagnostics

This flag takes a number that specifies the width of the terminal in characters. Formatting of diagnostics will take the width into consideration to make them better fit on the screen.

--remap-path-prefix: remap source names in output

Remap source path prefixes in all output, including compiler diagnostics, debug information, macro expansions, etc. It takes a value of the form `FROM=TO` where a path prefix equal to `FROM` is rewritten to the value `TO`. The `FROM` may itself contain an `=` symbol, but the `TO` value may not. This flag may be specified multiple times.

This is useful for normalizing build products, for example by removing the current directory out of pathnames emitted into the object files. The replacement is purely textual, with no consideration of the current system's pathname syntax. For example `--remap-path-prefix foo=bar` will match `foo/lib.rs` but not `./foo/lib.rs`.

When multiple remappings are given and several of them match, the `last` matching one is applied.

--json: configure json messages printed by the compiler

When the `--error-format=json` option is passed to `rustc` then all of the compiler's diagnostic output will be emitted in the form of JSON blobs. The `--json` argument can be used in conjunction with `--error-format=json` to configure what the JSON blobs contain as well as which ones are emitted.

With `--error-format=json` the compiler will always emit any compiler errors as a JSON blob, but the following options are also available to the `--json` flag to customize the output:

- `diagnostic-short` - json blobs for diagnostic messages should use the "short" rendering instead of the normal "human" default. This means that the output of `--error-format=short` will be embedded into the JSON diagnostics instead of the default `--error-format=human`.
- `diagnostic-rendered-ansi` - by default JSON blobs in their `rendered` field will contain a plain text rendering of the diagnostic. This option instead indicates that the diagnostic should have embedded ANSI color codes intended to be used to colorize the message in the manner `rustc` typically already does for terminal outputs. Note that this is usefully combined with crates like `fwdansi` to translate these ANSI codes on Windows to console commands or `strip-ansi-escapes` if you'd like to optionally remove the ansi colors afterwards.

- **artifacts** - this instructs rustc to emit a JSON blob for each artifact that is emitted. An artifact corresponds to a request from the **--emit CLI argument**, and as soon as the artifact is available on the filesystem a notification will be emitted.
- **future-incompat** - includes a JSON message that contains a report if the crate contains any code that may fail to compile in the future.

Note that it is invalid to combine the **--json** argument with the **--color** argument, and it is required to combine **--json** with **--error-format=json**.

See [the JSON chapter](#) for more detail.

@path: load command-line flags from a path

If you specify **@path** on the command-line, then it will open **path** and read command line options from it. These options are one per line; a blank line indicates an empty option. The file can use Unix or Windows style line endings, and must be encoded as UTF-8.

Codegen Options

All of these options are passed to `rustc` via the `-C` flag, short for "codegen." You can see a version of this list for your exact compiler by running `rustc -C help`.

ar

This option is deprecated and does nothing.

code-model

This option lets you choose which code model to use.

Code models put constraints on address ranges that the program and its symbols may use.

With smaller address ranges machine instructions may be able to use more compact addressing modes.

The specific ranges depend on target architectures and addressing modes available to them.

For x86 more detailed description of its code models can be found in [System V Application Binary Interface](#) specification.

Supported values for this option are:

- `tiny` - Tiny code model.
- `small` - Small code model. This is the default model for majority of supported targets.
- `kernel` - Kernel code model.
- `medium` - Medium code model.
- `large` - Large code model.

Supported values can also be discovered by running `rustc --print code-models`.

codegen-units

This flag controls how many code generation units the crate is split into. It takes an integer greater than 0.

When a crate is split into multiple codegen units, LLVM is able to process them in parallel. Increasing parallelism may speed up compile times, but may also produce slower code. Setting this to 1 may improve the performance of generated code, but may be slower to compile.

The default value, if not specified, is 16 for non-incremental builds. For incremental builds the default is 256 which allows caching to be more granular.

control-flow-guard

This flag controls whether LLVM enables the Windows [Control Flow Guard](#) platform security feature. This flag is currently ignored for non-Windows targets. It takes one of the following values:

- `y`, `yes`, `on`, `true`, `checks`, or no value: enable Control Flow Guard.
- `nochecks`: emit Control Flow Guard metadata without runtime enforcement checks (this should only be used for testing purposes as it does not provide security enforcement).
- `n`, `no`, `off`, `false`: do not enable Control Flow Guard (the default).

debug-assertions

This flag lets you turn `cfg(debug_assertions)` conditional compilation on or off. It takes one of the following values:

- `y`, `yes`, `on`, `true`, or no value: enable debug-assertions.
- `n`, `no`, `off` or `false`: disable debug-assertions.

If not specified, debug assertions are automatically enabled only if the [opt-level](#) is 0.

debuginfo

This flag controls the generation of debug information. It takes one of the following values:

- `0` or `none`: no debug info at all (the default).
- `line-directives-only`: line info directives only. For the nvptx* targets this enables [profiling](#). For other use cases, `line-tables-only` is the better, more compatible choice.
- `line-tables-only`: line tables only. Generates the minimal amount of debug info for backtraces with filename/line number info, but not anything else, i.e. no variable or function parameter info.
- `1` or `limited`: debug info without type or variable-level information.
- `2` or `full`: full debug info.

Note: The `-g` flag is an alias for `-C debuginfo=2`.

default-linker-libraries

This flag controls whether or not the linker includes its default libraries. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: include default libraries (the default).
- `n`, `no`, `off` or `false`: exclude default libraries.

For example, for gcc flavor linkers, this issues the `-nodefaultlibs` flag to the linker.

embed-bitcode

This flag controls whether or not the compiler embeds LLVM bitcode into object files. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: put bitcode in rlibs (the default).
- `n`, `no`, `off` or `false`: omit bitcode from rlibs.

LLVM bitcode is required when rustc is performing link-time optimization (LTO). It is also required on some targets like iOS ones where vendors look for LLVM bitcode. Embedded bitcode will appear in rustc-generated object files inside of a section whose name is defined by the target platform. Most of the time this is `.llvmbc`.

The use of `-C embed-bitcode=no` can significantly improve compile times and reduce generated file sizes if your compilation does not actually need bitcode (e.g. if you're not compiling for iOS or you're not performing LTO). For these reasons, Cargo uses `-C embed-bitcode=no` whenever possible. Likewise, if you are building directly with `rustc` we recommend using `-C embed-bitcode=no` whenever you are not using LTO.

If combined with `-C lto`, `-C embed-bitcode=no` will cause `rustc` to abort at start-up, because the combination is invalid.

Note: if you're building Rust code with LTO then you probably don't even need the `embed-bitcode` option turned on. You'll likely want to use `-clinker-plugin-lto` instead which skips generating object files entirely and simply replaces object files with LLVM bitcode. The only purpose for `-C embed-bitcode` is when you're generating an rlib that is both being used with and without LTO. For example Rust's standard library ships with embedded bitcode since users link to it both with and without LTO.

This also may make you wonder why the default is `yes` for this option. The reason for that is that it's how it was for rustc 1.44 and prior. In 1.45 this option was added to turn off what had always been the default.

extra-filename

This option allows you to put extra data in each output filename. It takes a string to add as a suffix to the filename. See the `--emit` flag for more information.

force-frame-pointers

This flag forces the use of frame pointers. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: force-enable frame pointers.
- `n`, `no`, `off` or `false`: do not force-enable frame pointers. This does not necessarily mean frame pointers will be removed.

The default behaviour, if frame pointers are not force-enabled, depends on the target.

force-unwind-tables

This flag forces the generation of unwind tables. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: Unwind tables are forced to be generated.
- `n`, `no`, `off` or `false`: Unwind tables are not forced to be generated. If unwind tables are required by the target an error will be emitted.

The default if not specified depends on the target.

incremental

This flag allows you to enable incremental compilation, which allows `rustc` to save information after compiling a crate to be reused when recompiling the crate, improving re-compile times. This takes a path to a directory where incremental files will be stored.

inline-threshold

This option lets you set the default threshold for inlining a function. It takes an unsigned integer as a value. Inlining is based on a cost model, where a higher threshold will allow more inlining.

The default depends on the opt-level:

opt-level	Threshold
0	N/A, only inlines always-inline functions
1	N/A, only inlines always-inline functions and LLVM lifetime intrinsics
2	225
3	275
s	75
z	25

instrument-coverage

This option enables instrumentation-based code coverage support. See the chapter on [instrumentation-based code coverage](#) for more information.

Note that while the `-C instrument-coverage` option is stable, the profile data format produced by the resulting instrumentation may change, and may not work with coverage tools other than those built and shipped with the compiler.

link-arg

This flag lets you append a single extra argument to the linker invocation.

"Append" is significant; you can pass this flag multiple times to add multiple arguments.

link-args

This flag lets you append multiple extra arguments to the linker invocation. The options should be separated by spaces.

link-dead-code

This flag controls whether the linker will keep dead code. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: keep dead code.
- `n`, `no`, `off` or `false`: remove dead code (the default).

An example of when this flag might be useful is when trying to construct code coverage metrics.

link-self-contained

On `windows-gnu`, `linux-musl`, and `wasi` targets, this flag controls whether the linker will use libraries and objects shipped with Rust instead or those in the system. It takes one of the following values:

- no value: `rustc` will use heuristic to disable self-contained mode if system has necessary tools.
- `y`, `yes`, `on`, `true`: use only libraries/objects shipped with Rust.
- `n`, `no`, `off` or `false`: rely on the user or the linker to provide non-Rust libraries/objects.

This allows overriding cases when detection fails or user wants to use shipped libraries.

linker

This flag controls which linker `rustc` invokes to link your code. It takes a path to the linker executable. If this flag is not specified, the linker will be inferred based on the target. See also the `linker-flavor` flag for another way to specify the linker.

linker-flavor

This flag controls the linker flavor used by `rustc`. If a linker is given with the `-C linker` flag, then the linker flavor is inferred from the value provided. If no linker is given then the linker flavor is used to determine the linker to use. Every `rustc` target defaults to some linker flavor. Valid options are:

- `em`: use `Emscripten emcc`.
- `gcc`: use the `cc` executable, which is typically `gcc` or `clang` on many systems.
- `ld`: use the `ld` executable.
- `msvc`: use the `link.exe` executable from Microsoft Visual Studio MSVC.
- `ptx-linker`: use `rust-ptx-linker` for Nvidia NVPTX GPGPU support.
- `bpf-linker`: use `bpf-linker` for eBPF support.
- `wasm-ld`: use the `wasm-ld` executable, a port of LLVM `lld` for WebAssembly.
- `ld64.lld`: use the LLVM `lld` executable with the `-flavor darwin` flag for Apple's `ld`.
- `ld.lld`: use the LLVM `lld` executable with the `-flavor gnu` flag for GNU binutils' `ld`.

- `lld-link`: use the LLVM `lld` executable with the `-flavor link` flag for Microsoft's `link.exe`.

linker-plugin-lto

This flag defers LTO optimizations to the linker. See [linker-plugin-LTO](#) for more details. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: enable linker plugin LTO.
- `n`, `no`, `off` or `false`: disable linker plugin LTO (the default).
- A path to the linker plugin.

More specifically this flag will cause the compiler to replace its typical object file output with LLVM bitcode files. For example an rlib produced with `-clinker-plugin-lto` will still have `*.o` files in it, but they'll all be LLVM bitcode instead of actual machine code. It is expected that the native platform linker is capable of loading these LLVM bitcode files and generating code at link time (typically after performing optimizations).

Note that rustc can also read its own object files produced with `-clinker-plugin-lto`. If an rlib is only ever going to get used later with a `-clto` compilation then you can pass `-clinker-plugin-lto` to speed up compilation and avoid generating object files that aren't used.

llvm-args

This flag can be used to pass a list of arguments directly to LLVM.

The list must be separated by spaces.

Pass `--help` to see a list of options.

lto

This flag controls whether LLVM uses [link time optimizations](#) to produce better optimized code, using whole-program analysis, at the cost of longer linking time. It takes one of the following values:

- `y`, `yes`, `on`, `true`, `fat`, or no value: perform "fat" LTO which attempts to perform optimizations across all crates within the dependency graph.
- `n`, `no`, `off`, `false`: disables LTO.
- `thin`: perform ["thin" LTO](#). This is similar to "fat", but takes substantially less time to run while still achieving performance gains similar to "fat".

If `-c lto` is not specified, then the compiler will attempt to perform "thin local LTO" which performs "thin" LTO on the local crate only across its [codegen units](#). When `-c lto` is not specified, LTO is disabled if codegen units is 1 or optimizations are disabled (`-C opt-level=0`). That is:

- When `-c lto` is not specified:
 - `codegen-units=1`: disable LTO.
 - `opt-level=0`: disable LTO.
- When `-c lto` is specified:
 - `lto`: 16 codegen units, perform fat LTO across crates.
 - `codegen-units=1 + lto`: 1 codegen unit, fat LTO across crates.

See also [linker-plugin-lto](#) for cross-language LTO.

metadata

This option allows you to control the metadata used for symbol mangling. This takes a space-separated list of strings. Mangled symbols will incorporate a hash of the metadata. This may be used, for example, to differentiate symbols between two different versions of the same crate being linked.

no-prepopulate-passes

This flag tells the pass manager to use an empty list of passes, instead of the usual pre-populated list of passes.

no-redzone

This flag allows you to disable [the red zone](#). It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: disable the red zone.
- `n`, `no`, `off` or `false`: enable the red zone.

The default behaviour, if the flag is not specified, depends on the target.

no-stack-check

This option is deprecated and does nothing.

no-vectorize-loops

This flag disables [loop vectorization](#).

no-vectorize-slp

This flag disables vectorization using [superword-level parallelism](#).

opt-level

This flag controls the optimization level.

- `0`: no optimizations, also turns on `cfg(debug_assertions)` (the default).
- `1`: basic optimizations.
- `2`: some optimizations.
- `3`: all optimizations.
- `s`: optimize for binary size.
- `z`: optimize for binary size, but also turn off loop vectorization.

Note: The `-O` flag is an alias for `-C opt-level=2`.

The default is `0`.

overflow-checks

This flag allows you to control the behavior of [runtime integer overflow](#). When overflow-checks are enabled, a panic will occur on overflow. This flag takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: enable overflow checks.
- `n`, `no`, `off` or `false`: disable overflow checks.

If not specified, overflow checks are enabled if `debug-assertions` are enabled, disabled otherwise.

panic

This option lets you control what happens when the code panics.

- `abort`: terminate the process upon panic
- `unwind`: unwind the stack upon panic

If not specified, the default depends on the target.

passes

This flag can be used to add extra LLVM passes to the compilation.

The list must be separated by spaces.

See also the `no-prepopulate-passes` flag.

prefer-dynamic

By default, `rustc` prefers to statically link dependencies. This option will indicate that dynamic linking should be used if possible if both a static and dynamic versions of a library are available. There is an internal algorithm for determining whether or not it is possible to statically or dynamically link with a dependency. For example, `cdylib` crate types may only use static linkage. This flag takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: use dynamic linking.
- `n`, `no`, `off` or `false`: use static linking (the default).

profile-generate

This flag allows for creating instrumented binaries that will collect profiling data for use with profile-guided optimization (PGO). The flag takes an optional argument which is the path to a directory into which the instrumented binary will emit the collected data. See the chapter on [profile-guided optimization](#) for more information.

profile-use

This flag specifies the profiling data file to be used for profile-guided optimization (PGO). The flag takes a mandatory argument which is the path to a valid `.profdata` file. See the chapter on [profile-guided optimization](#) for more information.

relocation-model

This option controls generation of [position-independent code \(PIC\)](#).

Supported values for this option are:

Primary relocation models

- `static` - non-relocatable code, machine instructions may use absolute addressing modes.
- `pic` - fully relocatable position independent code, machine instructions need to use relative addressing modes.
Equivalent to the "uppercase" `-fPIC` or `-fPIE` options in other compilers, depending on the produced crate types.
This is the default model for majority of supported targets.
- `pie` - position independent executable, relocatable code but without support for symbol interpositioning (replacing symbols by name using `LD_PRELOAD` and similar). Equivalent to the "uppercase" `-fPIE` option in other compilers. `pie` code cannot be linked into shared libraries (you'll get a linking error on attempt to do this).

Special relocation models

- `dynamic-no-pic` - relocatable external references, non-relocatable code.
Only makes sense on Darwin and is rarely used.
If StackOverflow tells you to use this as an opt-out of PIC or PIE, don't believe it, use `-C relocation-model=static` instead.
- `ropi`, `rwpri` and `ropi-rwpri` - relocatable code and read-only data, relocatable read-write data, and combination of both, respectively.
Only makes sense for certain embedded ARM targets.
- `default` - relocation model default to the current target.
Only makes sense as an override for some other explicitly specified relocation model previously set on the command line.

Supported values can also be discovered by running `rustc --print relocation-models`.

Linking effects

In addition to codegen effects, `relocation-model` has effects during linking.

If the relocation model is `pic` and the current target supports position-independent executables (PIE), the linker will be instructed (`-pie`) to produce one.

If the target doesn't support both position-independent and statically linked executables, then `-C target-feature=+crt-static` "wins" over `-C relocation-model=pic`, and the linker is instructed (`-static`) to produce a statically linked but not position-independent executable.

remark

This flag lets you print remarks for optimization passes.

The list of passes should be separated by spaces.

`all` will remark on every pass.

rpath

This flag controls whether `rpath` is enabled. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: enable rpath.
- `n`, `no`, `off` or `false`: disable rpath (the default).

save-temp

This flag controls whether temporary files generated during compilation are deleted once compilation finishes. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: save temporary files.
- `n`, `no`, `off` or `false`: delete temporary files (the default).

soft-float

This option controls whether `rustc` generates code that emulates floating point instructions in software. It takes one of the following values:

- `y`, `yes`, `on`, `true` or no value: use soft floats.
- `n`, `no`, `off` or `false`: use hardware floats (the default).

split-debuginfo

This option controls the emission of "split debuginfo" for debug information that `rustc` generates. The default behavior of this option is platform-specific, and not all possible values for this option work on all platforms. Possible values are:

- `off` - This is the default for platforms with ELF binaries and windows-gnu (not Windows MSVC and not macOS). This typically means that DWARF debug information can be found in the final artifact in sections of the executable. This option is not supported on Windows MSVC. On macOS this option prevents the final execution of `dsymutil` to generate debuginfo.
- `packed` - This is the default for Windows MSVC and macOS. The term "packed" here means that all the debug information is packed into a separate file from the

main executable. On Windows MSVC this is a `*.pdb` file, on macOS this is a `*.dSYM` folder, and on other platforms this is a `*.dwp` file.

- **unpacked** - This means that debug information will be found in separate files for each compilation unit (object file). This is not supported on Windows MSVC. On macOS this means the original object files will contain debug information. On other Unix platforms this means that `*.dwo` files will contain debug information.

Note that all three options are supported on Linux and Apple platforms, `packed` is supported on Windows-MSVC, and all other platforms support `off`. Attempting to use an unsupported option requires using the nightly channel with the `-Z unstable-options` flag.

strip

The option `-C strip=val` controls stripping of debuginfo and similar auxiliary data from binaries during linking.

Supported values for this option are:

- `none` - debuginfo and symbols (if they exist) are copied to the produced binary or separate files depending on the target (e.g. `.pdb` files in case of MSVC).
- `debuginfo` - debuginfo sections and debuginfo symbols from the symbol table section are stripped at link time and are not copied to the produced binary or separate files.
- `symbols` - same as `debuginfo`, but the rest of the symbol table section is stripped as well if the linker supports it.

symbol-mangling-version

This option controls the `name mangling` format for encoding Rust item names for the purpose of generating object code and linking.

Supported values for this option are:

- `v0` - The "v0" mangling scheme. The specific format is not specified at this time.

The default, if not specified, will use a compiler-chosen default which may change in the future.

target-cpu

This instructs `rustc` to generate code specifically for a particular processor.

You can run `rustc --print target-cpus` to see the valid options to pass here. Each target has a default base CPU. Special values include:

- `native` can be passed to use the processor of the host machine.
- `generic` refers to an LLVM target with minimal features but modern tuning.

target-feature

Individual targets will support different features; this flag lets you control enabling or disabling a feature. Each feature should be prefixed with a `+` to enable it or `-` to disable it.

Features from multiple `-C target-feature` options are combined.

Multiple features can be specified in a single option by separating them with commas `-C target-feature=+x,-y`.

If some feature is specified more than once with both `+` and `-`, then values passed later override values passed earlier.

For example, `-C target-feature=+x,-y,+z -Ctarget-feature=-x,+y` is equivalent to `-C target-feature=-x,+y,+z`.

To see the valid options and an example of use, run `rustc --print target-features`.

Using this flag is unsafe and might result in [undefined runtime behavior](#).

See also the `target_feature` attribute for controlling features per-function.

This also supports the feature `+crt-static` and `-crt-static` to control [static C runtime linkage](#).

Each target and `target-cpu` has a default set of enabled features.

tune-cpu

This instructs `rustc` to schedule code specifically for a particular processor. This does not affect the compatibility (instruction sets or ABI), but should make your code slightly more efficient on the selected CPU.

The valid options are the same as those for `target-cpu`. The default is `None`, which LLVM translates as the `target-cpu`.

This is an unstable option. Use `-z tune-cpu=machine` to specify a value.

Due to limitations in LLVM (12.0.0-git9218f92), this option is currently effective only for x86 targets.

Lints

In software, a "lint" is a tool used to help improve your source code. The Rust compiler contains a number of lints, and when it compiles your code, it will also run the lints. These lints may produce a warning, an error, or nothing at all, depending on how you've configured things.

Here's a small example:

```
$ cat main.rs
fn main() {
    let x = 5;
}
$ rustc main.rs
warning: unused variable: `x`
--> main.rs:2:9
2 |     let x = 5;
   |
   |
= note: `#[warn(unused_variables)]` on by default
= note: to avoid this warning, consider using `_x` instead
```

This is the `unused_variables` lint, and it tells you that you've introduced a variable that you don't use in your code. That's not *wrong*, so it's not an error, but it might be a bug, so you get a warning.

Future-incompatible lints

Sometimes the compiler needs to be changed to fix an issue that can cause existing code to stop compiling. "Future-incompatible" lints are issued in these cases to give users of Rust a smooth transition to the new behavior. Initially, the compiler will continue to accept the problematic code and issue a warning. The warning has a description of the problem, a notice that this will become an error in the future, and a link to a tracking issue that provides detailed information and an opportunity for feedback. This gives users some time to fix the code to accommodate the change. After some time, the warning may become an error.

The following is an example of what a future-incompatible looks like:

```
warning: borrow of packed field is unsafe and requires unsafe function or block (error E0133)
--> lint_example.rs:11:13
11 |     let y = &x.data.0;
   |           ^^^^^^^^^^
   |
   = note: `#[warn(safe_packed_borrows)]` on by default
   = warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
   = note: for more information, see issue #46043 <https://github.com/rust-lang/rust-lang/issues/46043>
   = note: fields of packed structs might be misaligned: dereferencing a misaligned pointer or
even just creating a misaligned reference is undefined behavior
```

For more information about the process and policy of future-incompatible changes,
see [RFC 1589](#).

Lint Levels

In `rustc`, lints are divided into five levels:

1. allow
2. warn
3. force-warn
4. deny
5. forbid

Each lint has a default level (explained in the lint listing later in this chapter), and the compiler has a default warning level. First, let's explain what these levels mean, and then we'll talk about configuration.

allow

These lints exist, but by default, do nothing. For example, consider this source:

```
pub fn foo() {}
```

Compiling this file produces no warnings:

```
$ rustc lib.rs --crate-type=lib
$
```

But this code violates the `missing_docs` lint.

These lints exist mostly to be manually turned on via configuration, as we'll talk about later in this section.

warn

The 'warn' lint level will produce a warning if you violate the lint. For example, this code runs afoul of the `unused_variables` lint:

```
pub fn foo() {
    let x = 5;
}
```

This will produce this warning:

```
$ rustc lib.rs --crate-type=lib
warning: unused variable: `x`
--> lib.rs:2:9
2 |     let x = 5;
  |
  = note: #[warn(unused_variables)]` on by default
  = note: to avoid this warning, consider using `_x` instead
```

force-warn

'force-warn' is a special lint level. It's the same as 'warn' in that a lint at this level will produce a warning, but unlike the 'warn' level, the 'force-warn' level cannot be overridden. If a lint is set to 'force-warn', it is guaranteed to warn: no more, no less. This is true even if the overall lint level is capped via cap-lints.

deny

A 'deny' lint produces an error if you violate it. For example, this code runs into the `exceeding_bitshifts` lint.

```
fn main() {
    100u8 << 10;
}
```

```
$ rustc main.rs
error: bitshift exceeds the type's number of bits
--> main.rs:2:13
2 |     100u8 << 10;
  |     ^^^^^^^^^^
  |
  = note: #[deny(exceeding_bitshifts)]` on by default
```

What's the difference between an error from a lint and a regular old error? Lints are configurable via levels, so in a similar way to 'allow' lints, warnings that are 'deny' by default let you allow them. Similarly, you may wish to set up a lint that is `warn` by default to produce an error instead. This lint level gives you that.

forbid

'forbid' is a special lint level that fills the same role for 'deny' that 'force-warn' does for 'warn'. It's the same as 'deny' in that a lint at this level will produce an error, but unlike the 'deny' level, the 'forbid' level can not be

overridden to be anything lower than an error. However, lint levels may still be capped with `--cap-lints` (see below) so `rustc --cap-lints warn` will make lints set to 'forbid' just warn.

Configuring warning levels

Remember our `missing_docs` example from the 'allow' lint level?

```
$ cat lib.rs
pub fn foo() {}
$ rustc lib.rs --crate-type=lib
$
```

We can configure this lint to operate at a higher level, both with compiler flags, as well as with an attribute in the source code.

You can also "cap" lints so that the compiler can choose to ignore certain lint levels. We'll talk about that last.

Via compiler flag

The `-A`, `-W`, `--force-warn`, `-D`, and `-F` flags let you turn one or more lints into allowed, warning, force-warn, deny, or forbid levels, like this:

```
$ rustc lib.rs --crate-type=lib -W missing-docs
warning: missing documentation for crate
--> lib.rs:1:1
  |
1 | pub fn foo() {}
  | ^^^^^^^^^^^^
  |
  = note: requested on the command line with `'-W missing-docs`'

warning: missing documentation for a function
--> lib.rs:1:1
  |
1 | pub fn foo() {}
  | ^^^^^^^^^^^^
```

```
$ rustc lib.rs --crate-type=lib -D missing-docs
error: missing documentation for crate
--> lib.rs:1:1
1 | pub fn foo() {}
| ^^^^^^^^^^^^^^
|
= note: requested on the command line with `'-D missing-docs`

error: missing documentation for a function
--> lib.rs:1:1
1 | pub fn foo() {}
| ^^^^^^^^^^^^^^

error: aborting due to 2 previous errors
```

You can also pass each flag more than once for changing multiple lints:

```
$ rustc lib.rs --crate-type=lib -D missing-docs -D unused-variables
```

And of course, you can mix these five flags together:

```
$ rustc lib.rs --crate-type=lib -D missing-docs -A unused-variables
```

The order of these command line arguments is taken into account. The following allows the `unused-variables` lint, because it is the last argument for that lint:

```
$ rustc lib.rs --crate-type=lib -D unused-variables -A unused-variables
```

You can make use of this behavior by overriding the level of one specific lint out of a group of lints. The following example denies all the lints in the `unused` group, but explicitly allows the `unused-variables` lint in that group (forbid still trumps everything regardless of ordering):

```
$ rustc lib.rs --crate-type=lib -D unused -A unused-variables
```

Since `force-warn` and `forbid` cannot be overridden, setting one of them will prevent any later level for the same lint from taking effect.

Via an attribute

You can also modify the lint level with a crate-wide attribute:

```
$ cat lib.rs
#![warn(missing_docs)]  
  
pub fn foo() {}  
$ rustc lib.rs --crate-type=lib  
warning: missing documentation for crate  
--> lib.rs:1:1  
| / #![warn(missing_docs)]  
2 |  
3 | | pub fn foo() {}  
| |-----^  
|  
note: lint level defined here  
--> lib.rs:1:9  
|  
1 | #![warn(missing_docs)]  
| |-----^  
  
warning: missing documentation for a function  
--> lib.rs:3:1  
|  
3 | pub fn foo() {}  
| |-----^
```

`warn`, `allow`, `deny`, and `forbid` all work this way. There is no way to set a lint to `force-warn` using an attribute.

You can also pass in multiple lints per attribute:

```
#![warn(missing_docs, unused_variables)]  
  
pub fn foo() {}
```

And use multiple attributes together:

```
#![warn(missing_docs)]  
#![deny(unused_variables)]  
  
pub fn foo() {}
```

Capping lints

`rustc` supports a flag, `--cap-lints LEVEL` that sets the "lint cap level." This is the maximum level for all lints. So for example, if we take our code sample from the "deny" lint level above:

```
fn main() {  
    100u8 << 10;  
}
```

And we compile it, capping lints to warn:

```
$ rustc lib.rs --cap-lints warn
warning: bitshift exceeds the type's number of bits
  → lib.rs:2:5
  |
2 |     100u8 << 10;
  |     ^^^^^^^^^^
  |
= note: `#[warn(exceeding_bitshifts)]` on by default

warning: this expression will panic at run-time
  → lib.rs:2:5
  |
2 |     100u8 << 10;
  |     ^^^^^^^^^^ attempt to shift left with overflow
```

It now only warns, rather than errors. We can go further and allow all lints:

```
$ rustc lib.rs --cap-lints allow
$
```

This feature is used heavily by Cargo; it will pass `--cap-lints allow` when compiling your dependencies, so that if they have any warnings, they do not pollute the output of your build.

Lint Groups

`rustc` has the concept of a "lint group", where you can toggle several warnings through one name.

For example, the `nonstandard-style` lint sets `non-camel-case-types`, `non-snake-case`, and `non-upper-case-globals` all at once. So these are equivalent:

```
$ rustc -D nonstandard-style
$ rustc -D non-camel-case-types -D non-snake-case -D non-upper-case-globals
```

Here's a list of each lint group, and the lints that they are made up of:

Group	Description	Lints
warnings	All lints that are set to issue warnings	See warn-by-default for the default set of warnings
future-incompatible	Lints that detect code that has future-compatibility problems	<code>ambiguous-associated-items</code> , <code>byte-slice-in-packed-struct-with-derive</code> , <code>cenum-impl-drop-cast</code> , <code>coherence-leak-check</code> , <code>conflicting-repr-hints</code> , <code>const-evaluatable-unchecked</code> , <code>deprecated-cfg-attr-crate-type-name</code> , <code>deref-into-dyn-supertrait</code> , <code>forbidden-lint-groups</code> , <code>ill-formed-attribute-input</code> , <code>illegal-floating-point-literal-pattern</code> , <code>implied-bounds-entailment</code> , <code>indirect-structural-match</code> , <code>invalid-doc-attributes</code> , <code>invalid-type-param-default</code> , <code>late-bound-lifetime-arguments</code> , <code>legacy-derive-helpers</code> , <code>macro-expanded-macro-exports-accessed-by-absolute-paths</code> , <code>missing-fragment-specifier</code> , <code>nontrivial-structural-match</code> , <code>order-dependent-trait-objects</code> , <code>patterns-in-fns-without-body</code> , <code>pointer-structural-match</code> , <code>private-in-public</code> , <code>proc-macro-back-compat</code> , <code>proc-macro-derive-resolution-fallback</code> , <code>pub-use-of-private-extern-crate</code> , <code>repr-transparent-external-private-fields</code> , <code>semicolon-in-expressions-from-macros</code> , <code>soft-unstable</code> , <code>suspicious-auto-trait-impls</code> , <code>uninhabited-static</code> , <code>unstable-name-collisions</code> , <code>unstable-syntax-pre-expansion</code> , <code>unsupported-calling-conventions</code> , <code>where-clauses-object-safety</code>
let-underscore	Lints that detect wildcard let bindings that	<code>let-underscore-drop</code> , <code>let-underscore-lock</code>

Group	Description	Lints
	are likely to be invalid	
nonstandard-style	Violation of standard naming conventions	non-camel-case-types, non-snake-case, non-upper-case-globals
rust-2018-compatibility	Lints used to transition code from the 2015 edition to 2018	absolute-paths-not-starting-with-crate, anonymous-parameters, keyword-idents, tyvar-behind-raw-pointer
rust-2018-idioms	Lints to nudge you toward idiomatic features of Rust 2018	bare-trait-objects, elided-lifetimes-in-paths, ellipsis-inclusive-range-patterns, explicit-outlives-requirements, unused-extern-crates
rust-2021-compatibility	Lints used to transition code from the 2018 edition to 2021	array-into-iter, bare-trait-objects, ellipsis-inclusive-range-patterns, non-fmt-panics, rust-2021-incompatible-closure-captures, rust-2021-incompatible-or-patterns, rust-2021-prefixes-incompatible-syntax, rust-2021-prelude-collisions
unused	Lints that detect things being declared but not used, or excess syntax	dead-code, map-unit-fn, path-statements, redundant-semicolons, unreachable-code, unreachable-patterns, unused-allocation, unused-assignments, unused-attributes, unused-braces, unused-doc-comments, unused-extern-crates, unused-features, unused-imports, unused-labels, unused-macro-rules, unused-macros, unused-must-use, unused-mut, unused-parens, unused-unsafe, unused-variables

Additionally, there's a `bad-style` lint group that's a deprecated alias for `nonstandard-style`.

Finally, you can also see the table above by invoking `rustc -W help`. This will give you the exact values for the specific compiler you have installed.

Lint Listing

This section lists out all of the lints, grouped by their default lint levels.

You can also see this list by running `rustc -W help`.

Allowed-by-default Lints

These lints are all set to the 'allow' level by default. As such, they won't show up unless you set them to a higher lint level with a flag or attribute.

- `absolute_paths_not_starting_with_crate`
- `box_pointers`
- `elided_lifetimes_in_paths`
- `explicit_outlives_requirements`
- `ffi_unwind_calls`
- `fuzzy_provenance_casts`
- `keyword_idents`
- `let_underscore_drop`
- `lossy_provenance_casts`
- `macro_use_extern_crate`
- `meta_variable_misuse`
- `missing_abi`
- `missing_copy_implementations`
- `missing_debug_implementations`
- `missing_docs`
- `multiple_supertrait_upcastable`
- `must_not_suspend`
- `non_ascii_idents`
- `non_exhaustive_omitted_patterns`
- `noop_method_call`
- `pointer_structural_match`
- `rust_2021_incompatible_closure_captures`
- `rust_2021_incompatible_or_patterns`
- `rust_2021_prefixes_incompatible_syntax`
- `rust_2021_prelude_collisions`
- `single_use_lifetimes`
- `trivial_casts`
- `trivial_numeric_casts`
- `unreachable_pub`
- `unsafe_code`
- `unsafe_op_in_unsafe_fn`
- `unstable_features`
- `unused_crate_dependencies`
- `unused_extern_crates`
- `unused_import_braces`
- `unused_lifetimes`
- `unused_macro_rules`
- `unused_qualifications`
- `unused_results`

- [unused_tuple_struct_fields](#)
- [variant_size_differences](#)

absolute-paths-not-starting-with-crate

The `absolute_paths_not_starting_with_crate` lint detects fully qualified paths that start with a module name instead of `crate`, `self`, or an extern crate name

Example

```
#![deny(absolute_paths_not_starting_with_crate)]

mod foo {
    pub fn bar() {}
}

fn main() {
    ::foo::bar();
}
```

This will produce:

```
error: absolute paths must start with `self`, `super`, `crate`, or an external crate name in
the 2018 edition
--> lint_example.rs:8:5
8 |     ::foo::bar();
|     ^^^^^^^^^^^ help: use `crate`:: `crate::foo::bar`
|
|= warning: this is accepted in the current edition (Rust 2015) but is a hard error in Rust
2018!
= note: for more information, see issue #53130 <https://github.com/rust-lang/rust-lang/issues/53130>
note: the lint level is defined here
--> lint_example.rs:1:9
|
1 | #![deny(absolute_paths_not_starting_with_crate)]
|     ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

Rust [editions](#) allow the language to evolve without breaking backwards compatibility. This lint catches code that uses absolute paths in the style of the 2015 edition. In the 2015 edition, absolute paths (those starting with `::`) refer to either the crate root or an external crate. In the 2018 edition it was changed so that they only refer to external crates. The path prefix `crate::` should be used instead to reference items from the crate root.

If you switch the compiler from the 2015 to 2018 edition without updating the code, then it will fail to compile if the old style paths are used. You can manually change the paths to use the `crate::` prefix to transition to the 2018 edition.

This lint solves the problem automatically. It is "allow" by default because the code is perfectly valid in the 2015 edition. The `cargo fix` tool with the `--edition` flag will switch this lint to "warn" and automatically apply the suggested fix from the compiler. This provides a completely automated way to update old code to the 2018 edition.

box-pointers

The `box_pointers` lints use of the `Box` type.

Example

```
#![deny(box_pointers)]
struct Foo {
    x: Box<isize>,
}
```

This will produce:

```
error: type uses owned (Box type) pointers: Box<isize>
--> lint_example.rs:4:5
|
4 |     x: Box<isize>,
|     ^^^^^^^^^^^^^^
|
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #![deny(box_pointers)]
|     ^^^^^^^^^^^^^^
```

Explanation

This lint is mostly historical, and not particularly useful. `Box<T>` used to be built into the language, and the only way to do heap allocation. Today's Rust can call into other allocators, etc.

elided-lifetimes-in-paths

The `elided_lifetimes_in_paths` lint detects the use of hidden lifetime parameters.

Example

```
#![deny(elided_lifetimes_in_paths)]
struct Foo<'a> {
    x: &'a u32
}

fn foo(x: &Foo) {
```

This will produce:

```
error: hidden lifetime parameters in types are deprecated
--> lint_example.rs:7:12
|
7 | fn foo(x: &Foo) {
|         ^^^ expected lifetime parameter
|
note: the lint level is defined here
--> lint_example.rs:1:9
|
1 | #![deny(elided_lifetimes_in_paths)]
|          ^^^^^^^^^^^^^^^^^^^^^^^^^^
help: indicate the anonymous lifetime
|
7 | fn foo(x: &Foo<'_>) {
|         +++
```

Explanation

Elided lifetime parameters can make it difficult to see at a glance that borrowing is occurring. This lint ensures that lifetime parameters are always explicitly stated, even if it is the `'_ placeholder lifetime`.

This lint is "allow" by default because it has some known issues, and may require a significant transition for old code.

explicit-outlives-requirements

The `explicit_outlives_requirements` lint detects unnecessary lifetime bounds that can be inferred.

Example

```
#![deny(explicit_outlives_requirements)]  
  
struct SharedRef<'a, T>  
where  
    T: 'a,  
{  
    data: &'a T,  
}
```

This will produce:

```
error: outlives requirements can be inferred  
--> lint_example.rs:5:24  
5 |     struct SharedRef<'a, T>  
5 |     ^-----  
6 |     where  
6 |         T: 'a,  
6 |         ^ help: remove this bound  
|  
note: the lint level is defined here  
--> lint_example.rs:2:9  
2 | #![deny(explicit_outlives_requirements)]  
   | ^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

If a `struct` contains a reference, such as `&'a T`, the compiler requires that `T` outlives the lifetime `'a`. This historically required writing an explicit lifetime bound to indicate this requirement. However, this can be overly explicit, causing clutter and unnecessary complexity. The language was changed to automatically infer the bound if it is not specified. Specifically, if the struct contains a reference, directly or indirectly, to `T` with lifetime `'x`, then it will infer that `T: 'x` is a requirement.

This lint is "allow" by default because it can be noisy for existing code that already had these requirements. This is a stylistic choice, as it is still valid to explicitly state the bound. It also has some false positives that can cause confusion.

See [RFC 2093](#) for more details.

ffi-unwind-calls

The `ffi_unwind_calls` lint detects calls to foreign functions or function pointers with `c-unwind` or other FFI-unwind ABIs.

Example

```
#![feature(c_unwind)]
#![warn(ffi_unwind_calls)]

extern "C-unwind" {
    fn foo();
}

fn bar() {
    unsafe { foo(); }
    let ptr: unsafe extern "C-unwind" fn() = foo;
    unsafe { ptr(); }
}
```

This will produce:

```
warning: call to foreign function with FFI-unwind ABI
--> lint_example.rs:10:14
  |
10 |     unsafe { foo(); }
  |          ^^^^^^ call to foreign function with FFI-unwind ABI
  |
note: the lint level is defined here
--> lint_example.rs:2:9
  |
2 | #![warn(ffi_unwind_calls)]
  |          ^^^^^^^^^^^^^^

warning: call to function pointer with FFI-unwind ABI
--> lint_example.rs:12:14
  |
12 |     unsafe { ptr(); }
  |          ^^^^^^ call to function pointer with FFI-unwind ABI
```

Explanation

For crates containing such calls, if they are compiled with `-C panic=unwind` then the produced library cannot be linked with crates compiled with `-C panic=abort`. For crates that desire this ability it is therefore necessary to avoid such calls.

fuzzy-provenance-casts

The `fuzzy_provenance_casts` lint detects an `as` cast between an integer and a pointer.

Example

```
#![feature(strict_provenance)]
#![warn(fuzzy_provenance_casts)]

fn main() {
    let _dangling = 16_usize as *const u8;
}
```

This will produce:

```
warning: strict provenance disallows casting integer `usize` to pointer `*const u8`
--> lint_example.rs:5:21
5 |     let _dangling = 16_usize as *const u8;
   |     ^^^^^^^^^^^^^^^^^^^^^^^^^^

= help: if you can't comply with strict provenance and don't have a pointer with the correct
provenance you can use `std::ptr::from_exposed_addr()` instead
note: the lint level is defined here
--> lint_example.rs:2:9
2 | #![warn(fuzzy_provenance_casts)]
   |     ^^^^^^^^^^^^^^^^^^^^^^

help: use `with_addr()` to adjust a valid pointer in the same allocation, to this address
5 |     let _dangling = (...).with_addr(16_usize);
   |     ++++++~
```

Explanation

This lint is part of the strict provenance effort, see [issue #95228](#). Casting an integer to a pointer is considered bad style, as a pointer contains, besides the *address* also a *provenance*, indicating what memory the pointer is allowed to read/write. Casting an integer, which doesn't have provenance, to a pointer requires the compiler to assign (guess) provenance. The compiler assigns "all exposed valid" (see the docs of `ptr::from_exposed_addr` for more information about this "exposing"). This penalizes the optimiser and is not well suited for dynamic analysis/dynamic program verification (e.g. Miri or CHERI platforms).

It is much better to use `ptr::with_addr` instead to specify the provenance you want. If using this function is not possible because the code relies on exposed provenance then there is as an escape hatch `ptr::from_exposed_addr`.

keyword-idents

The `keyword_idents` lint detects edition keywords being used as an identifier.

Example

```
#![deny(keyword_idents)]
// edition 2015
fn dyn() {}
```

This will produce:

```
error: `dyn` is a keyword in the 2018 edition
--> lint_example.rs:4:4
   |
4 | fn dyn() {}
   |     ^^^ help: you can use a raw identifier to stay compatible: `r#dyn`
   |
 = warning: this is accepted in the current edition (Rust 2015) but is a hard error in Rust
2018!
 = note: for more information, see issue #49716 <https://github.com/rust-lang/rust/issues/49716>
note: the lint level is defined here
--> lint_example.rs:1:9
   |
1 | #![deny(keyword_idents)]
   | ^^^^^^^^^^^^^^^^^^
```

Explanation

Rust [editions](#) allow the language to evolve without breaking backwards compatibility. This lint catches code that uses new keywords that are added to the language that are used as identifiers (such as a variable name, function name, etc.). If you switch the compiler to a new edition without updating the code, then it will fail to compile if you are using a new keyword as an identifier.

You can manually change the identifiers to a non-keyword, or use a [raw identifier](#), for example `r#dyn`, to transition to a new edition.

This lint solves the problem automatically. It is "allow" by default because the code is perfectly valid in older editions. The [cargo fix](#) tool with the `--edition` flag will switch this lint to "warn" and automatically apply the suggested fix from the compiler (which is to use a raw identifier). This provides a completely automated way to update old code for a new edition.

let-underscore-drop

The `let_underscore_drop` lint checks for statements which don't bind an expression which has a non-trivial Drop implementation to anything, causing the expression to be dropped immediately instead of at end of scope.

Example

```
struct SomeStruct;
impl Drop for SomeStruct {
    fn drop(&mut self) {
        println!("Dropping SomeStruct");
    }
}

fn main() {
    #[warn(let_underscore_drop)]
    // SomeStruct is dropped immediately instead of at end of scope,
    // so "Dropping SomeStruct" is printed before "end of main".
    // The order of prints would be reversed if SomeStruct was bound to
    // a name (such as "_foo").
    let _ = SomeStruct;
    println!("end of main");
}
```

This will produce:

```
warning: non-binding let on a type that implements `Drop`
--> lint_example.rs:14:5
  |
14 |     let _ = SomeStruct;
  |     ^^^^^^^^^^^^^^^^^^

note: the lint level is defined here
--> lint_example.rs:9:11
  |
9  |     #[warn(let_underscore_drop)]
  |     ^^^^^^^^^^^^^^^^^^

help: consider binding to an unused variable to avoid immediately dropping the value
14 |     let _unused = SomeStruct;
  |     ~~~~~

help: consider immediately dropping the value
14 |     drop(SomeStruct);
  |     ~~~~      +
```

Explanation

Statements which assign an expression to an underscore causes the expression to immediately drop instead of extending the expression's lifetime to the end of the scope. This is usually unintended, especially for types like `MutexGuard`, which are typically used to lock a mutex for the duration of an entire scope.

If you want to extend the expression's lifetime to the end of the scope, assign an underscore-prefixed name (such as `_foo`) to the expression. If you do actually want to drop the expression immediately, then calling `std::mem::drop` on the expression is clearer and helps convey intent.

lossy-provenance-casts

The `lossy_provenance_casts` lint detects an `as` cast between a pointer and an integer.

Example

```
#![feature(strict_provenance)]
#![warn(lossy_provenance_casts)]

fn main() {
    let x: u8 = 37;
    let _addr: usize = &x as *const u8 as usize;
}
```

This will produce:

```
warning: under strict provenance it is considered bad style to cast pointer `*const u8` to
integer `usize`
--> lint_example.rs:6:24
|
6 |     let _addr: usize = &x as *const u8 as usize;
|           ^^^^^^^^^^^^^^^^^^^^^^^^^^
|
= help: if you can't comply with strict provenance and need to expose the pointer provenance
you can use `*.expose_addr()` instead
note: the lint level is defined here
--> lint_example.rs:2:9
|
2 | #[warn(lossy_provenance_casts)]
|           ^^^^^^^^^^^^^^^^^^
help: use `*.addr()` to obtain the address of a pointer
|
6 |     let _addr: usize = (&x as *const u8).addr();
|           +             ~~~~~~
```

Explanation

This lint is part of the strict provenance effort, see [issue #95228](#). Casting a pointer to an integer is a lossy operation, because beyond just an *address* a pointer may be associated with a particular *provenance*. This information is used by the optimiser and for dynamic analysis/dynamic program verification (e.g. Miri or CHERI platforms).

Since this cast is *lossy*, it is considered good style to use the `ptr::addr` method instead, which has a similar effect, but doesn't "expose" the pointer provenance. This improves optimisation potential. See the docs of `ptr::addr` and `ptr::expose_addr` for more information about exposing pointer provenance.

If your code can't comply with strict provenance and needs to expose the provenance, then there is `ptr::expose_addr` as an escape hatch, which preserves the behaviour of `as usize` casts while being explicit about the semantics.

macro-use-extern-crate

The `macro_use_extern_crate` lint detects the use of the `macro_use` attribute.

Example

```
#![deny(macro_use_extern_crate)]  
  
#[macro_use]  
extern crate serde_json;  
  
fn main() {  
    let _ = json!{{}};  
}
```

This will produce:

```
error: deprecated `#[macro_use]` attribute used to import macros should be replaced at use  
sites with a `use` item to import the macro instead  
--> src/main.rs:3:1  
|  
3 | #[macro_use]  
| ^^^^^^^^^^  
|  
note: the lint level is defined here  
--> src/main.rs:1:9  
|  
1 | #![deny(macro_use_extern_crate)]  
| ^^^^^^^^^^
```

Explanation

The `macro_use` attribute on an `extern crate` item causes macros in that external crate to be brought into the prelude of the crate, making the macros in scope everywhere. As part of the efforts to simplify handling of dependencies in the [2018 edition](#), the use of `extern crate` is being phased out. To bring macros from extern crates into scope, it is recommended to use a `use import`.

This lint is "allow" by default because this is a stylistic choice that has not been settled, see [issue #52043](#) for more information.

meta-variable-misuse

The `meta_variable_misuse` lint detects possible meta-variable misuse in macro definitions.

Example

```
#![deny(meta_variable_misuse)]

macro_rules! foo {
    () => {};
    ($($i:ident = $($j:ident),+ );*) => { $($($i = $k; )+ )* };
}

fn main() {
    foo!();
}
```

This will produce:

```
error: unknown macro variable `k`
--> lint_example.rs:5:55
5 |     ($($i:ident = $($j:ident),+ );*) => { $($($i = $k; )+ )* }; ^
   |
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #![deny(meta_variable_misuse)]
   | ^^^^^^^^^^^^^^^^^^
```

Explanation

There are quite a few different ways a `macro_rules!` macro can be improperly defined. Many of these errors were previously only detected when the macro was expanded or not at all. This lint is an attempt to catch some of these problems when the macro is *defined*.

This lint is "allow" by default because it may have false positives and other issues. See [issue #61053](#) for more details.

missing-abi

The `missing_abi` lint detects cases where the ABI is omitted from extern declarations.

Example

```
#![deny(missing_abi)]

extern fn foo() {}
```

This will produce:

```
error: extern declarations without an explicit ABI are deprecated
--> lint_example.rs:4:1
|
4 | extern fn foo() {}
| ^^^^^^^^^^^^^^ ABI should be specified here
|
= help: the default ABI is C
note: the lint level is defined here
--> lint_example.rs:1:9
|
1 | #![deny(missing_abi)]
| ^^^^^^^^^^^^^^
```

Explanation

Historically, Rust implicitly selected C as the ABI for extern declarations. We expect to add new ABIs, like `C-unwind`, in the future, though this has not yet happened, and especially with their addition seeing the ABI easily will make code review easier.

missing-copy-implementations

The `missing_copy_implementations` lint detects potentially-forgotten implementations of `Copy`.

Example

```
#![deny(missing_copy_implementations)]
pub struct Foo {
    pub field: i32
}
```

This will produce:

```
error: type could implement `Copy`; consider adding `impl Copy`
--> lint_example.rs:2:1
|
2 | / pub struct Foo {
3 | |     pub field: i32
4 | | }
| |_ ^
|
note: the lint level is defined here
--> lint_example.rs:1:9
|
1 | #![deny(missing_copy_implementations)]
| ^^^^^^^^^^^^^^
```

Explanation

Historically (before 1.0), types were automatically marked as `Copy` if possible. This was changed so that it required an explicit opt-in by implementing the `Copy` trait. As part of this change, a lint was added to alert if a copyable type was not marked `Copy`.

This lint is "allow" by default because this code isn't bad; it is common to write newtypes like this specifically so that a `Copy` type is no longer `Copy`. `Copy` types can result in unintended copies of large data which can impact performance.

missing-debug-implementations

The `missing_debug_implementations` lint detects missing implementations of `fmt::Debug`.

Example

```
#![deny(missing_debug_implementations)]
pub struct Foo;
```

This will produce:

```
error: type does not implement `Debug`; consider adding `#[derive(Debug)]` or a manual
implementation
--> lint_example.rs:2:1
 |
2 | pub struct Foo;
| ^^^^^^^^^^^^^^^^
|
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #![deny(missing_debug_implementations)]
| ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

Having a `Debug` implementation on all types can assist with debugging, as it provides a convenient way to format and display a value. Using the `#[derive(Debug)]` attribute will automatically generate a typical implementation, or a custom implementation can be added by manually implementing the `Debug` trait.

This lint is "allow" by default because adding `Debug` to all types can have a negative impact on compile time and code size. It also requires boilerplate to be added to every type, which can be an impediment.

missing-docs

The `missing_docs` lint detects missing documentation for public items.

Example

```
#![deny(missing_docs)]
pub fn foo() {}
```

This will produce:

```
error: missing documentation for the crate
--> lint_example.rs:1:1
|
1 | / #![deny(missing_docs)]
2 | | fn main() {
3 | | pub fn foo() {}
4 | |
|_| ^
|
note: the lint level is defined here
--> lint_example.rs:1:9
|
1 | #![deny(missing_docs)]
| ^^^^^^^^^^
```

Explanation

This lint is intended to ensure that a library is well-documented. Items without documentation can be difficult for users to understand how to use properly.

This lint is "allow" by default because it can be noisy, and not all projects may want to enforce everything to be documented.

multiple-supertrait-upcastable

The `multiple_supertrait_upcastable` lint detects when an object-safe trait has multiple supertraits.

Example

```
trait A {}
trait B {}

#[warn(multiple_supertrait_upcastable)]
trait C: A + B {}
```

This will produce:

```
warning: unknown lint: `multiple_supertrait_upcastable`
--> lint_example.rs:5:1
  |
5 | #[warn(multiple_supertrait_upcastable)]
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  |
  = note: the `multiple_supertrait_upcastable` lint is unstable
  = help: add `#![feature(multiple_supertrait_upcastable)]` to the crate attributes to enable
  = note: `#[warn(unknown_lints)]` on by default
```

Explanation

To support upcasting with multiple supertraits, we need to store multiple vtables and this can result in extra space overhead, even if no code actually uses upcasting. This lint allows users to identify when such scenarios occur and to decide whether the additional overhead is justified.

must-not-suspend

The `must_not_suspend` lint guards against values that shouldn't be held across suspend points (`.await`)

Example

```
#![feature(must_not_suspend)]
#![warn(must_not_suspend)]

#[must_not_suspend]
struct SyncThing {}

async fn yield_now() {}

pub async fn uhoh() {
    let guard = SyncThing {};
    yield_now().await;
}
```

This will produce:

```
warning: `SyncThing` held across a suspend point, but should not be
--> lint_example.rs:11:9
11 |     let guard = SyncThing {};
|     ^^^^^^
12 |     yield_now().await;
|             ----- the value is held across this suspend point
|
help: consider using a block (`{ ... }`) to shrink the value's scope, ending before the
suspend point
--> lint_example.rs:11:9
11 |     let guard = SyncThing {};
|     ^^^^^^
note: the lint level is defined here
--> lint_example.rs:2:9
2 | #[warn(must_not_suspend)]
| ^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

The `must_not_suspend` lint detects values that are marked with the `#[must_not_suspend]` attribute being held across suspend points. A "suspend" point is usually a `.await` in an `async` function.

This attribute can be used to mark values that are semantically incorrect across suspends (like certain types of timers), values that have `async` alternatives, and values that regularly cause problems with the `send`-ness of `async` fn's returned futures (like `MutexGuard`'s)

non-ascii-idents

The `non_ascii_idents` lint detects non-ASCII identifiers.

Example

```
#![deny(non_ascii_idents)]
fn main() {
    let föö = 1;
}
```

This will produce:

```
error: identifier contains non-ASCII characters
--> lint_example.rs:4:9
4 |     let föö = 1;
   |     ^
|
note: the lint level is defined here
--> lint_example.rs:2:9
2 | #[deny(non_ascii_idents)]
   | ^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

This lint allows projects that wish to retain the limit of only using ASCII characters to switch this lint to "forbid" (for example to ease collaboration or for security reasons). See [RFC 2457](#) for more details.

non-exhaustive-omitted-patterns

The `non_exhaustive_omitted_patterns` lint detects when a wildcard (`_` or `..`) in a pattern for a `#[non_exhaustive]` struct or enum is reachable.

Example

```
// crate A
#[non_exhaustive]
pub enum Bar {
    A,
    B, // added variant in non breaking change
}

// in crate B
#![feature(non_exhaustive_omitted_patterns_lint)]

match Bar::A {
    Bar::A => {},
    #[warn(non_exhaustive_omitted_patterns)]
    _ => {},
}
```

This will produce:

```
warning: reachable patterns not covered of non exhaustive enum
--> $DIR/reachable-patterns.rs:70:9
LL |         - => {}
|           ^ pattern `B` not covered
|
note: the lint level is defined here
--> $DIR/reachable-patterns.rs:69:16
LL |     #[warn(non_exhaustive_omitted_patterns)]
|           ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
= help: ensure that all possible cases are being handled by adding the suggested match arms
= note: the matched value is of type `Bar` and the `non_exhaustive_omitted_patterns` attribute was found
```

Explanation

Structs and enums tagged with `#[non_exhaustive]` force the user to add a (potentially redundant) wildcard when pattern-matching, to allow for future addition of fields or variants. The `non_exhaustive_omitted_patterns` lint detects when such a wildcard happens to actually catch some fields/variants. In other words, when the match without the wildcard would not be exhaustive. This lets the user be informed if new fields/variants were added.

noop-method-call

The `noop_method_call` lint detects specific calls to noop methods such as a calling `<&T as Clone>::clone` where `T: !Clone`.

Example

```
#![warn(noop_method_call)]
struct Foo;
let foo = &Foo;
let clone: &Foo = foo.clone();
```

This will produce:

```
warning: call to `clone()` on a reference in this situation does nothing
--> lint_example.rs:6:22
6 | let clone: &Foo = foo.clone();
   |          ^^^^^^ unnecessary method call
   |
   = note: the type `&Foo` which `clone` is being called on is the same as the type returned
from `clone`, so the method call does not do anything and can be removed
note: the lint level is defined here
--> lint_example.rs:2:9
2 | #[warn(noop_method_call)]
   |          ^^^^^^^^^^^^^^
```

Explanation

Some method calls are noops meaning that they do nothing. Usually such methods are the result of blanket implementations that happen to create some method invocations that end up not doing anything. For instance, `clone` is implemented on all `&T`, but calling `clone` on a `&T` where `T` does not implement `clone`, actually doesn't do anything as references are copy. This lint detects these calls and warns the user about them.

pointer-structural-match

The `pointer_structural_match` lint detects pointers used in patterns whose behaviour cannot be relied upon across compiler versions and optimization levels.

Example

```
#![deny(pointer_structural_match)]
fn foo(a: usize, b: usize) -> usize { a + b }
const FOO: fn(usize, usize) -> usize = foo;
fn main() {
    match FOO {
        FOO => {},
        _ => {},
    }
}
```

This will produce:

```
error: function pointers and unsized pointers in patterns behave unpredictably and should not
be relied upon. See https://github.com/rust-lang/rust/issues/70861 for details.
--> lint_example.rs:6:9
|           FOO => {},
|           ^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #62411 <https://github.com/rust-
lang/rust/issu
```

note: the lint level is defined here
--> lint_example.rs:1:9

```
| #![deny(pointer_structural_match)]
| ^^^^^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

Previous versions of Rust allowed function pointers and wide raw pointers in patterns. While these work in many cases as expected by users, it is possible that due to optimizations pointers are "not equal to themselves" or pointers to different functions compare as equal during runtime. This is because LLVM optimizations can deduplicate functions if their bodies are the same, thus also making pointers to these functions point to the same location. Additionally functions may get duplicated if they are instantiated in different crates and not deduplicated again via LTO.

rust-2021-incompatible-closure-captures

The `rust_2021_incompatible_closure_captures` lint detects variables that aren't completely captured in Rust 2021, such that the `Drop` order of their fields may differ between Rust 2018 and 2021.

It can also detect when a variable implements a trait like `Send`, but one of its fields does not, and the field is captured by a closure and used with the assumption that said field implements the same trait as the root variable.

Example of drop reorder

```
#![deny(rust_2021_incompatible_closure_captures)]

struct FancyInteger(i32);

impl Drop for FancyInteger {
    fn drop(&mut self) {
        println!("Just dropped {}", self.0);
    }
}

struct Point { x: FancyInteger, y: FancyInteger }

fn main() {
    let p = Point { x: FancyInteger(10), y: FancyInteger(20) };

    let c = || {
        let x = p.x;
    };

    c();
    // ... More code ...
}
```

This will produce:

```
error: changes to closure capture in Rust 2021 will affect drop order
--> lint_example.rs:17:11
   |
17 |     let c = || {
   |     ^
18 |         let x = p.x;
   |         --- in Rust 2018, this closure captures all of `p`, but in Rust 2021, it
will only capture `p.x`
...
24 |     | - in Rust 2018, `p` is dropped here, but in Rust 2021, only `p.x` will be dropped here as
part of the closure
   |
   = note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/disjoint-capture-in-closures.html>
note: the lint level is defined here
--> lint_example.rs:1:9
   |
1 | #![deny(rust_2021_incompatible_closure_captures)]
   | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
help: add a dummy let to cause `p` to be fully captured
   |
17 ~     let c = || {
18 +         let _ = &p;
   |
```

Explanation

In the above example, `p.y` will be dropped at the end of `f` instead of with `c` in Rust 2021.

Example of auto-trait

```
#![deny(rust_2021_incompatible_closure_captures)]
use std::thread;

struct Pointer(*mut i32);
unsafe impl Send for Pointer {}

fn main() {
    let mut f = 10;
    let fptr = Pointer(&mut f as *mut i32);
    thread::spawn(move || unsafe {
        *fptr.0 = 20;
    });
}
```

This will produce:

```
error: changes to closure capture in Rust 2021 will affect which traits the closure implements
--> lint_example.rs:10:19
10 |     thread::spawn(move || unsafe {
|     |         ^^^^^^ in Rust 2018, this closure implements `Send` as `fptr`
| implements `Send`, but in Rust 2021, this closure will no longer implement `Send` because
| `fptr` is not fully captured and `fptr.0` does not implement `Send`
11 |         *fptr.0 = 20;
|         ----- in Rust 2018, this closure captures all of `fptr`, but in Rust 2021, it
will only capture `fptr.0`
|
= note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/disjoint-capture-in-closures.html>
note: the lint level is defined here
--> lint_example.rs:1:9
|
1 | #![deny(rust_2021_incompatible_closure_captures)]
| ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
help: add a dummy let to cause `fptr` to be fully captured
|
10 ~    thread::spawn(move || { let _ = &fptr; unsafe {
11 |         *fptr.0 = 20;
12 ~     } });
|
```

Explanation

In the above example, only `fptr.0` is captured in Rust 2021. The field is of type `*mut i32`, which doesn't implement `Send`, making the code invalid as the field cannot be sent between threads safely.

rust-2021-incompatible-or-patterns

The `rust_2021_incompatible_or_patterns` lint detects usage of old versions of or-patterns.

Example

```
#!/[deny(rust_2021_incompatible_or_patterns)]  
  
macro_rules! match_any {  
    ( $expr:expr , $($($pat:pat)|+ => $expr_arm:expr ),+ ) => {  
        match $expr {  
            $($($pat => $expr_arm, )+  
            )+  
        }  
    };  
}  
  
fn main() {  
    let result: Result<i64, i32> = Err(42);  
    let int: i64 = match_any!(result, Ok(i) | Err(i) => i.into());  
    assert_eq!(int, 42);  
}
```

This will produce:

```
error: the meaning of the `pat` fragment specifier is changing in Rust 2021, which may affect  
this macro  
--> lint_example.rs:4:26  
|     ( $expr:expr , $($($pat:pat)|+ => $expr_arm:expr ),+ ) => {  
|         ^^^^^^^^^ help: use pat_param to preserve semantics:  
`$pat:pat_param`  
|  
= warning: this is accepted in the current edition (Rust 2018) but is a hard error in Rust  
2021!  
= note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/or-patterns-macro-rules.html>  
note: the lint level is defined here  
--> lint_example.rs:1:9  
|  
1 | #![deny(rust_2021_incompatible_or_patterns)]  
|         ^^^^^^^^^^^^^^^^^^
```

Explanation

In Rust 2021, the `pat` matcher will match additional patterns, which include the `|` character.

rust-2021-prefixes-incompatible-syntax

The `rust_2021_prefixes_incompatible_syntax` lint detects identifiers that will be parsed as a prefix instead in Rust 2021.

Example

```
#![deny(rust_2021_prefixes_incompatible_syntax)]  
  
macro_rules! m {  
    (z $x:expr) => ();  
}  
  
m!(z"hey");
```

This will produce:

```
error: prefix `z` is unknown  
--> lint_example.rs:8:4  
|  
8 |     m!(z"hey");  
|     ^ unknown prefix  
|  
= warning: this is accepted in the current edition (Rust 2018) but is a hard error in Rust  
2021!  
= note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/reserving-syntax.html>  
note: the lint level is defined here  
--> lint_example.rs:1:9  
|  
1 | #![deny(rust_2021_prefixes_incompatible_syntax)]  
|           ^^^^^^^^^^^^^^^^^^^^^^^^^^  
help: insert whitespace here to avoid this being parsed as a prefix in Rust 2021  
|  
8 |     m!(z "hey");  
|     +
```

Explanation

In Rust 2015 and 2018, `z"hey"` is two tokens: the identifier `z` followed by the string literal `"hey"`. In Rust 2021, the `z` is considered a prefix for `"hey"`.

This lint suggests to add whitespace between the `z` and `"hey"` tokens to keep them separated in Rust 2021.

rust-2021-prelude-collisions

The `rust_2021_prelude_collisions` lint detects the usage of trait methods which are ambiguous with traits added to the prelude in future editions.

Example

```
#![deny(rust_2021_prelude_collisions)]  
  
trait Foo {  
    fn try_into(self) -> Result<String, !>;  
}  
  
impl Foo for &str {  
    fn try_into(self) -> Result<String, !> {  
        Ok(String::from(self))  
    }  
}  
  
fn main() {  
    let x: String = "3".try_into().unwrap();  
    // ^^^^^^  
    // This call to try_into matches both Foo::try_into and TryInto::try_into as  
    // `TryInto` has been added to the Rust prelude in 2021 edition.  
    println!("{}");  
}
```

This will produce:

```
error: trait method `try_into` will become ambiguous in Rust 2021  
--> lint_example.rs:14:21  
|  
14 |     let x: String = "3".try_into().unwrap();  
|     ^^^^^^ help: disambiguate the associated function:  
`Foo::try_into(&"3")'  
|  
= warning: this is accepted in the current edition (Rust 2018) but is a hard error in Rust  
2021!  
= note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/prelude.html>  
note: the lint level is defined here  
--> lint_example.rs:1:9  
|  
1 | #![deny(rust_2021_prelude_collisions)]  
| ^^^^^^
```

Explanation

In Rust 2021, one of the important introductions is the [prelude changes](#), which add `TryFrom`, `TryInto`, and `FromIterator` into the standard library's prelude. Since this results in an ambiguity as to which method/function to call when an existing `try_into` method is called via dot-call syntax or a `try_from`/`from_iter` associated function is called directly on a type.

single-use-lifetimes

The `single_use_lifetimes` lint detects lifetimes that are only used once.

Example

```
#![deny(single_use_lifetimes)]  
  
fn foo<'a>(x: &'a u32) {}
```

This will produce:

```
error: lifetime parameter ``'a` only used once  
--> lint_example.rs:4:8  
|  
4 | fn foo<'a>(x: &'a u32) {}  
|     ^^      -- ...is used only here  
|  
|     this lifetime...  
  
note: the lint level is defined here  
--> lint_example.rs:1:9  
|  
1 | #![deny(single_use_lifetimes)]  
|     ^^^^^^^^^^^^^^^^^^  
help: elide the single-use lifetime  
|  
4 - fn foo<'a>(x: &'a u32) {}  
4 + fn foo(x: &u32) {}  
|
```

Explanation

Specifying an explicit lifetime like `'a` in a function or `impl` should only be used to link together two things. Otherwise, you should just use `'_` to indicate that the lifetime is not linked to anything, or elide the lifetime altogether if possible.

This lint is "allow" by default because it was introduced at a time when `'_` and elided lifetimes were first being introduced, and this lint would be too noisy. Also, there are some known false positives that it produces. See [RFC 2115](#) for historical context, and [issue #44752](#) for more details.

trivial-casts

The `trivial_casts` lint detects trivial casts which could be replaced with coercion, which may require a temporary variable.

Example

```
#![deny(trivial_casts)]  
let x: &u32 = &42;  
let y = x as *const u32;
```

This will produce:

```
error: trivial cast: `&u32` as `*const u32`
--> lint_example.rs:4:9
4 | let y = x as *const u32;
   | ^^^^^^^^^^^^^^^^^^
|
= help: cast can be replaced by coercion; this might require a temporary variable
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #[deny(trivial_casts)]
   | ^^^^^^^^^^^^^^
```

Explanation

A trivial cast is a cast `e as T` where `e` has type `U` and `U` is a subtype of `T`. This type of cast is usually unnecessary, as it can be usually be inferred.

This lint is "allow" by default because there are situations, such as with FFI interfaces or complex type aliases, where it triggers incorrectly, or in situations where it will be more difficult to clearly express the intent. It may be possible that this will become a warning in the future, possibly with an explicit syntax for coercions providing a convenient way to work around the current issues. See [RFC 401 \(coercions\)](#), [RFC 803 \(type ascription\)](#) and [RFC 3307 \(remove type ascription\)](#) for historical context.

trivial-numeric-casts

The `trivial_numeric_casts` lint detects trivial numeric casts of types which could be removed.

Example

```
#![deny(trivial_numeric_casts)]
let x = 42_i32 as i32;
```

This will produce:

```

error: trivial numeric cast: `i32` as `i32`
--> lint_example.rs:3:9
|   let x = 42_i32 as i32;
|   ^^^^^^^^^^^^^^
|
= help: cast can be replaced by coercion; this might require a temporary variable
note: the lint level is defined here
--> lint_example.rs:1:9
|
1 | #[deny(trivial_numeric_casts)]
| ^^^^^^^^^^^^^^^^^^^^^^^^^^

```

Explanation

A trivial numeric cast is a cast of a numeric type to the same numeric type. This type of cast is usually unnecessary.

This lint is "allow" by default because there are situations, such as with FFI interfaces or complex type aliases, where it triggers incorrectly, or in situations where it will be more difficult to clearly express the intent. It may be possible that this will become a warning in the future, possibly with an explicit syntax for coercions providing a convenient way to work around the current issues. See [RFC 401 \(coercions\)](#), [RFC 803 \(type ascription\)](#) and [RFC 3307 \(remove type ascription\)](#) for historical context.

unreachable-pub

The `unreachable_pub` lint triggers for `pub` items not reachable from the crate root.

Example

```

#[deny(unreachable_pub)]
mod foo {
    pub mod bar {
        }
    }
}

```

This will produce:

```
error: unreachable `pub` item
--> lint_example.rs:4:5
4 |     pub mod bar {
|     ^^^^^^^^^^
|     |
|     help: consider restricting its visibility: `pub(crate)`  
= help: or consider exporting it for use by other crates
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #[deny(unreachable_pub)]
| ^^^^^^^^^^^^^^^^^^
```

Explanation

A bare `pub` visibility may be misleading if the item is not actually publicly exported from the crate. The `pub(crate)` visibility is recommended to be used instead, which more clearly expresses the intent that the item is only visible within its own crate.

This lint is "allow" by default because it will trigger for a large amount existing Rust code, and has some false-positives. Eventually it is desired for this to become warn-by-default.

unsafe-code

The `unsafe_code` lint catches usage of `unsafe` code.

Example

```
#![deny(unsafe_code)]
fn main() {
    unsafe {
        }
}
```

This will produce:

```
error: usage of an `unsafe` block
--> lint_example.rs:3:5
3 | /     unsafe {
4 | |
5 | |     }
|-----^
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #[deny(unsafe_code)]
|           ^^^^^^^^^^
```

Explanation

This lint is intended to restrict the usage of `unsafe`, which can be difficult to use correctly.

unsafe-op-in-unsafe-fn

The `unsafe_op_in_unsafe_fn` lint detects unsafe operations in unsafe functions without an explicit unsafe block.

Example

```
#![deny(unsafe_op_in_unsafe_fn)]  
  
unsafe fn foo() {}  
  
unsafe fn bar() {  
    foo();  
}  
  
fn main() {}
```

This will produce:

```
error: call to unsafe function is unsafe and requires unsafe block (error E0133)
--> lint_example.rs:6:5
6 |     foo();
|     ^^^^^^ call to unsafe function
|
= note: consult the function's documentation for information on how to avoid undefined
behavior
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #[deny(unsafe_op_in_unsafe_fn)]
|           ^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

Currently, an `unsafe fn` allows any `unsafe` operation within its body. However, this can increase the surface area of code that needs to be scrutinized for proper behavior. The `unsafe block` provides a convenient way to make it clear exactly which parts of the code are performing unsafe operations. In the future, it is desired to change it so that unsafe operations cannot be performed in an `unsafe fn` without an `unsafe` block.

The fix to this is to wrap the unsafe code in an `unsafe` block.

This lint is "allow" by default since this will affect a large amount of existing code, and the exact plan for increasing the severity is still being considered. See [RFC #2585](#) and [issue #71668](#) for more details.

unstable-features

The `unstable_features` is deprecated and should no longer be used.

unused-crate-dependencies

The `unused_crate_dependencies` lint detects crate dependencies that are never used.

Example

```
#![deny(unused_crate_dependencies)]
```

This will produce:

```
error: external crate `regex` unused in `lint_example`: remove the dependency or add `use
regex as _;`
|
note: the lint level is defined here
--> src/lib.rs:1:9
|
1 | #[deny(unused_crate_dependencies)]
  ^^^^^^^^^^^^^^^^^^
```

Explanation

After removing the code that uses a dependency, this usually also requires removing the dependency from the build configuration. However, sometimes that step can be missed, which leads to time wasted building dependencies that are no longer used. This lint can be enabled to detect dependencies that are never used (more specifically, any dependency passed with the `--extern` command-line flag that is never referenced via `use`, `extern crate`, or in any `path`).

This lint is "allow" by default because it can provide false positives depending on how the build system is configured. For example, when using Cargo, a "package" consists of multiple crates (such as a library and a binary), but the dependencies are defined for the package as a whole. If there is a dependency that is only used in the binary, but not the library, then the lint will be incorrectly issued in the library.

unused-extern-crates

The `unused_extern_crates` lint guards against `extern crate` items that are never used.

Example

```
#![deny(unused_extern_crates)]
extern crate proc_macro;
```

This will produce:

```
error: unused extern crate
--> lint_example.rs:3:1
  |
3 | extern crate proc_macro;
  | ^^^^^^^^^^^^^^^^^^^^^^ help: remove it
  |
note: the lint level is defined here
--> lint_example.rs:1:9
  |
1 | #[deny(unused_extern_crates)]
  | ^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

`extern crate` items that are unused have no effect and should be removed. Note that there are some cases where specifying an `extern crate` is desired for the side effect of ensuring the given crate is linked, even though it is not otherwise directly referenced. The lint can be silenced by aliasing the crate to an underscore, such as `extern crate foo as _`. Also note that it is no longer idiomatic to use `extern crate` in the [2018 edition](#), as extern crates are now automatically added in scope.

This lint is "allow" by default because it can be noisy, and produce false-positives. If a dependency is being removed from a project, it is recommended to remove it from the build configuration (such as `Cargo.toml`) to ensure stale build entries aren't left behind.

unused-import-braces

The `unused_import_braces` lint catches unnecessary braces around an imported item.

Example

```
#![deny(unused_import_braces)]
use test:::{A};

pub mod test {
    pub struct A;
}
```

This will produce:

```

error: braces around A is unnecessary
--> lint_example.rs:2:1
2 | use test::{A};
  ^^^^^^^^^^
  |
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #[deny(unused_import_braces)]
  ^^^^^^^^^^^^^^^^^^^^^^
  
```

Explanation

If there is only a single item, then remove the braces (`use test::A;` for example).

This lint is "allow" by default because it is only enforcing a stylistic choice.

unused-lifetimes

The `unused_lifetimes` lint detects lifetime parameters that are never used.

Example

```

#[deny(unused_lifetimes)]

pub fn foo<'a>() {}
  
```

This will produce:

```

error: lifetime parameter ``a` never used
--> lint_example.rs:4:12
4 | pub fn foo<'a>() {}
  |          -^- help: elide the unused lifetime
  |
note: the lint level is defined here
--> lint_example.rs:2:8
2 | #[deny(unused_lifetimes)]
  |          ^^^^^^^^^^
  
```

Explanation

Unused lifetime parameters may signal a mistake or unfinished code. Consider removing the parameter.

unused-macro-rules

The `unused_macro_rules` lint detects macro rules that were not used.

Note that the lint is distinct from the `unused_macros` lint, which fires if the entire macro is never called, while this lint fires for single unused rules of the macro that is otherwise used. `unused_macro_rules` fires only if `unused_macros` wouldn't fire.

Example

```
#[warn(unused_macro_rules)]
macro_rules! unused_empty {
    (hello) => { println!("Hello, world!"); }; // This rule is unused
    () => { println!("empty"); }; // This rule is used
}

fn main() {
    unused_empty!(hello);
}
```

This will produce:

```
warning: 2nd rule of macro `unused_empty` is never used
--> lint_example.rs:4:5
4 |     () => { println!("empty"); }; // This rule is used
   |     ^
|
note: the lint level is defined here
--> lint_example.rs:1:8
1 | #[warn(unused_macro_rules)]
   | ^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

Unused macro rules may signal a mistake or unfinished code. Furthermore, they slow down compilation. Right now, silencing the warning is not supported on a single rule level, so you have to add an `allow` to the entire macro definition.

If you intended to export the macro to make it available outside of the crate, use the `macro_export` attribute.

unused-qualifications

The `unused_qualifications` lint detects unnecessarily qualified names.

Example

```
#![deny(unused_qualifications)]
mod foo {
    pub fn bar() {}
}

fn main() {
    use foo::bar;
    foo::bar();
}
```

This will produce:

```
error: unnecessary qualification
--> lint_example.rs:8:5
  |
8 |     foo::bar();
  |     ^^^^^^^^^^
  |
note: the lint level is defined here
--> lint_example.rs:1:9
  |
1 | #![deny(unused_qualifications)]
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

If an item from another module is already brought into scope, then there is no need to qualify it in this case. You can call `bar()` directly, without the `foo::`.

This lint is "allow" by default because it is somewhat pedantic, and doesn't indicate an actual problem, but rather a stylistic choice, and can be noisy when refactoring or moving around code.

unused-results

The `unused_results` lint checks for the unused result of an expression in a statement.

Example

```
#![deny(unused_results)]
fn foo<T>() -> T { panic!() }

fn main() {
    foo::<usize>();
}
```

This will produce:

```
error: unused result of type `usize`
--> lint_example.rs:5:5
5 |     foo::<usize>();
  | ^^^^^^^^^^^^^^

note: the lint level is defined here
--> lint_example.rs:1:9
| 
1 | #[deny(unused_results)]
  | ^^^^^^^^^^
```

Explanation

Ignoring the return value of a function may indicate a mistake. In cases where it is almost certain that the result should be used, it is recommended to annotate the function with the `must_use` attribute. Failure to use such a return value will trigger the `unused_must_use` lint which is warn-by-default. The `unused_results` lint is essentially the same, but triggers for all return values.

This lint is "allow" by default because it can be noisy, and may not be an actual problem. For example, calling the `remove` method of a `Vec` or `HashMap` returns the previous value, which you may not care about. Using this lint would require explicitly ignoring or discarding such values.

unused-tuple-struct-fields

The `unused_tuple_struct_fields` lint detects fields of tuple structs that are never read.

Example

```
#[warn(unused_tuple_struct_fields)]
struct S(i32, i32, i32);
let s = S(1, 2, 3);
let _ = (s.0, s.2);
```

This will produce:

```
warning: field `1` is never read
--> lint_example.rs:3:15
3 | struct S(i32, i32, i32);
  |         ^^^
  |
  |     field in this struct

note: the lint level is defined here
--> lint_example.rs:2:8
2 | #[warn(unused_tuple_struct_fields)]
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

help: consider changing the field to be of unit type to suppress this warning while preserving
the field numbering, or remove the field
3 | struct S(i32, (), i32);
  |      ~~
```

Explanation

Tuple struct fields that are never read anywhere may indicate a mistake or unfinished code. To silence this warning, consider removing the unused field(s) or, to preserve the numbering of the remaining fields, change the unused field(s) to have unit type.

variant-size-differences

The `variant_size_differences` lint detects enums with widely varying variant sizes.

Example

```
#![deny(variant_size_differences)]
enum En {
    V0(u8),
    VBig([u8; 1024]),
}
```

This will produce:

```
error: enum variant is more than three times larger (1024 bytes) than the next largest
--> lint_example.rs:5:5
5 |     VBig([u8; 1024]),
  | ^^^^^^^^^^^^^^^^^^
|
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #[deny(variant_size_differences)]
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

It can be a mistake to add a variant to an enum that is much larger than the other variants, bloating the overall size required for all variants. This can impact performance and memory usage. This is triggered if one variant is more than 3 times larger than the second-largest variant.

Consider placing the large variant's contents on the heap (for example via [Box](#)) to keep the overall size of the enum itself down.

This lint is "allow" by default because it can be noisy, and may not be an actual problem. Decisions about this should be guided with profiling and benchmarking.

Warn-by-default Lints

These lints are all set to the 'warn' level by default.

- `ambiguous_glob_reexports`
- `anonymous_parameters`
- `array_into_iter`
- `asm_sub_register`
- `bad_asm_style`
- `bare_trait_objects`
- `break_with_label_and_loop`
- `byte_slice_in_packed_struct_with_derive`
- `clashing_extern_declarations`
- `coherence_leak_check`
- `confusable_idents`
- `const_evaluatable_unchecked`
- `const_item_mutation`
- `dead_code`
- `deprecated`
- `deprecated_where_clause_location`
- `deref_into_dyn_supertrait`
- `deref_nullptr`
- `drop_bounds`
- `duplicate_macro_attributes`
- `dyn_drop`
- `ellipsis_inclusive_range_patterns`
- `exported_private_dependencies`
- `for_loops_over_fallibles`
- `forbidden_lint_groups`
- `function_item_references`
- `illegal_floating_point_literal_pattern`
- `improper_ctypes`
- `improper_ctypes_definitions`
- `incomplete_features`
- `indirect_structural_match`
- `inline_no_sanitize`
- `invalid_doc_attributes`
- `invalid_macro_export_arguments`
- `invalid_value`
- `irrefutable_let_patterns`
- `large_assignments`
- `late_bound_lifetime_arguments`
- `legacy_derive_helpers`
- `map_unit_fn`

- `mixed_script_confusables`
- `named_arguments_used_positionally`
- `no_mangle_generic_items`
- `non_camel_case_types`
- `non_fmt_panics`
- `non_shorthand_field_patterns`
- `non_snake_case`
- `non_upper_case_globals`
- `nontrivial_structural_match`
- `opaque_hidden_inferred_bound`
- `overlapping_range_endpoints`
- `path_statements`
- `private_in_public`
- `redundant_semicolons`
- `renamed_and_removed_lints`
- `repr_transparent_external_private_fields`
- `semicolon_in_expressions_from_macros`
- `special_module_name`
- `stable_features`
- `suspicious_auto_trait_impls`
- `temporary_cstring_as_ptr`
- `trivial_bounds`
- `type_alias_bounds`
- `tyvar_behind_raw_pointer`
- `uncommon_codepoints`
- `unconditional_recursion`
- `undefined_naked_function_abi`
- `unexpected_cfgs`
- `unfulfilled_lint_expectations`
- `ungated_async_fn_track_caller`
- `uninhabited_static`
- `unknown_lints`
- `unnameable_test_items`
- `unreachable_code`
- `unreachable_patterns`
- `unstable_name_collisions`
- `unstable_syntax_pre_expansion`
- `unsupported_calling_conventions`
- `unused_allocation`
- `unused_assignments`
- `unused_attributes`
- `unused_braces`
- `unused_comparisons`
- `unused_doc_comments`
- `unused_features`
- `unused_imports`

- [unused_labels](#)
- [unused_macros](#)
- [unused_must_use](#)
- [unused_mut](#)
- [unused_parens](#)
- [unused_unsafe](#)
- [unused_variables](#)
- [warnings](#)
- [where_clauses_object_safety](#)
- [while_true](#)

ambiguous-glob-reexports

The `ambiguous_glob_reexports` lint detects cases where names re-exported via globs collide. Downstream users trying to use the same name re-exported from multiple globs will receive a warning pointing out redefinition of the same name.

Example

```
#![deny(ambiguous_glob_reexports)]
pub mod foo {
    pub type X = u8;
}

pub mod bar {
    pub type Y = u8;
    pub type X = u8;
}

pub use foo::*;

pub use bar::*;

pub fn main() {}
```

This will produce:

```
error: ambiguous glob re-exports
--> lint_example.rs:11:9
11 | pub use foo::::*;
   |          ^^^^^^ the name `X` in the type namespace is first re-exported here
12 | pub use bar::::*;
   |          ----- but the name `X` in the type namespace is also re-exported here
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #![deny(ambiguous_glob_reexports)]
   |          ^^^^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

This was previously accepted but it could silently break a crate's downstream users code. For example, if `foo::*` and `bar::*` were re-exported before `bar::X` was added to the re-exports, downstream users could use `this_crate::X` without problems. However, adding `bar::X` would cause compilation errors in downstream crates because `X` is defined multiple times in the same namespace of `this_crate`.

anonymous-parameters

The `anonymous_parameters` lint detects anonymous parameters in trait definitions.

Example

```
#![deny(anonymous_parameters)]
// edition 2015
pub trait Foo {
    fn foo(usize);
}
fn main() {}
```

This will produce:

```

error: anonymous parameters are deprecated and will be removed in the next edition
--> lint_example.rs:4:12
4 |     fn foo(usize);
|         ^^^^^ help: try naming the parameter or explicitly ignoring it: `_: usize`
|
= warning: this is accepted in the current edition (Rust 2015) but is a hard error in Rust
2018!
= note: for more information, see issue #41686 <https://github.com/rust-lang/rust/issues/41686>
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #[deny(anonymous_parameters)]
|          ^^^^^^^^^^^^^^^^^^^^^^

```

Explanation

This syntax is mostly a historical accident, and can be worked around quite easily by adding an `_` pattern or a descriptive identifier:

```

trait Foo {
    fn foo(_: usize);
}

```

This syntax is now a hard error in the 2018 edition. In the 2015 edition, this lint is "warn" by default. This lint enables the `cargo fix` tool with the `--edition` flag to automatically transition old code from the 2015 edition to 2018. The tool will run this lint and automatically apply the suggested fix from the compiler (which is to add `_` to each parameter). This provides a completely automated way to update old code for a new edition. See [issue #41686](#) for more details.

array-into-iter

The `array_into_iter` lint detects calling `into_iter` on arrays.

Example

```
[1, 2, 3].into_iter().for_each(|n| { *n; });
```

This will produce:

```
warning: this method call resolves to `<&[T; N] as IntoIterator>::into_iter` (due to backwards compatibility), but will resolve to <[T; N] as IntoIterator>::into_iter in Rust 2021
--> lint_example.rs:3:11
3 | [1, 2, 3].into_iter().for_each(|n| { *n; });
   |
   = warning: this changes meaning in Rust 2021
   = note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/IntoIterator-for-arrays.html>
   = note: `#[warn(array_into_iter)]` on by default
help: use `.'.iter()` instead of `.'.into_iter()` to avoid ambiguity
3 | [1, 2, 3].iter().for_each(|n| { *n; });
   |
   help: or use `IntoIterator::into_iter(..)` instead of `.'.into_iter()` to explicitly iterate by value
3 | IntoIterator::into_iter([1, 2, 3]).for_each(|n| { *n; });
   ++++++++~
```

Explanation

Since Rust 1.53, arrays implement `IntoIterator`. However, to avoid breakage, `array.into_iter()` in Rust 2015 and 2018 code will still behave as `(&array).into_iter()`, returning an iterator over references, just like in Rust 1.52 and earlier. This only applies to the method call syntax `array.into_iter()`, not to any other syntax such as `for _ in array` or `IntoIterator::into_iter(array)`.

asm-sub-register

The `asm_sub_register` lint detects using only a subset of a register for inline asm inputs.

Example

```
#[cfg(target_arch="x86_64")]
use std::arch::asm;

fn main() {
    #[cfg(target_arch="x86_64")]
    unsafe {
        asm!("mov {0}, {0}", in(reg) 0i16);
    }
}
```

This will produce:

```
warning: formatting may not be suitable for sub-register argument
--> src/main.rs:7:19
7 |     asm!("mov {0}, {0}", in(reg) 0i16);
   |           ^^^ ^^^                   ---- for this argument
= note: `#[warn(asm_sub_register)]` on by default
= help: use the `x` modifier to have the register formatted as `ax`
= help: or use the `r` modifier to keep the default formatting of `rax`
```

Explanation

Registers on some architectures can use different names to refer to a subset of the register. By default, the compiler will use the name for the full register size. To explicitly use a subset of the register, you can override the default by using a modifier on the template string operand to specify when subregister to use. This lint is issued if you pass in a value with a smaller data type than the default register size, to alert you of possibly using the incorrect width. To fix this, add the suggested modifier to the template, or cast the value to the correct size.

See [register template modifiers](#) in the reference for more details.

bad-asm-style

The `bad_asm_style` lint detects the use of the `.intel_syntax` and `.att_syntax` directives.

Example

```
#[cfg(target_arch="x86_64")]
use std::arch::asm;

fn main() {
    #[cfg(target_arch="x86_64")]
    unsafe {
        asm!(
            ".att_syntax",
            "movq %{0}, %{0}", in(reg) 0usize
        );
    }
}
```

This will produce:

```
warning: avoid using `.`att_syntax` , prefer using `options(att_syntax)` instead
--> src/main.rs:8:14
8 |         ".att_syntax",
|         ^^^^^^^^^^^^
= note: `#[warn(bad_asm_style)]` on by default
```

Explanation

On x86, `asm!` uses the intel assembly syntax by default. While this can be switched using assembler directives like `.att_syntax`, using the `att_syntax` option is recommended instead because it will also properly prefix register placeholders with `%` as required by AT&T syntax.

bare-trait-objects

The `bare_trait_objects` lint suggests using `dyn Trait` for trait objects.

Example

```
trait Trait { }

fn takes_trait_object(_: Box<Trait>) {
```

This will produce:

```
warning: trait objects without an explicit `dyn` are deprecated
--> lint_example.rs:4:30
4 | fn takes_trait_object(_: Box<Trait>) {
|         ^^^^^^
|
= warning: this is accepted in the current edition (Rust 2018) but is a hard error in Rust
2021!
= note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/warnings-promoted-to-error.html>
= note: `#[warn(bare_trait_objects)]` on by default
help: use `dyn`
4 | fn takes_trait_object(_: Box<dyn Trait>) {
|         +++
```

Explanation

Without the `dyn` indicator, it can be ambiguous or confusing when reading code as to whether or not you are looking at a trait object. The `dyn` keyword makes it explicit, and adds a symmetry to contrast with `impl Trait`.

break-with-label-and-loop

The `break_with_label_and_loop` lint detects labeled `break` expressions with an unlabeled loop as their value expression.

Example

```
'label: loop {  
    break 'label loop { break 42; };  
};
```

This will produce:

```
warning: this labeled break expression is easy to confuse with an unlabeled break with a  
labeled value expression  
--> lint_example.rs:3:5  
3 |     break 'label loop { break 42; };  
|     ^^^^^^  
| = note: `#[warn(break_with_label_and_loop)]` on by default  
help: wrap this expression in parentheses  
3 |     break 'label (loop { break 42; });  
|           +  
|           +
```

Explanation

In Rust, loops can have a label, and `break` expressions can refer to that label to break out of specific loops (and not necessarily the innermost one). `break` expressions can also carry a value expression, which can be another loop. A labeled `break` with an unlabeled loop as its value expression is easy to confuse with an unlabeled `break` with a labeled loop and is thus discouraged (but allowed for compatibility); use parentheses around the loop expression to silence this warning. Unlabeled `break` expressions with labeled loops yield a hard error, which can also be silenced by wrapping the expression in parentheses.

byte-slice-in-packed-struct-with-derive

The `byte_slice_in_packed_struct_with_derive` lint detects cases where a byte slice field (`[u8]`) or string slice field (`str`) is used in a `packed` struct that derives one or more built-in traits.

Example

```
#[repr(packed)]
#[derive(Hash)]
struct FlexZeroSlice {
    width: u8,
    data: [u8],
}
```

This will produce:

```
warning: byte slice in a packed struct that derives a built-in trait
--> lint_example.rs:6:5
  |
3 |     #[derive(Hash)]
  |         ----- in this derive macro expansion
...
6 |     data: [u8],
  | ^^^^^^^^^^^^
  |
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #107457 <https://github.com/rust-lang/rust/issues/107457>
= help: consider implementing the trait by hand, or remove the `packed` attribute
= note: `#[warn(byte_slice_in_packed_struct_with_derive)]` on by default
= note: this warning originates in the derive macro `Hash` (in Nightly builds, run with -Z
macro-backtrace for more info)
```

Explanation

This was previously accepted but is being phased out, because fields in packed structs are now required to implement `Copy` for `derive` to work. Byte slices and string slices are a temporary exception because certain crates depended on them.

clashing-extern-declarations

The `clashing_extern_declarations` lint detects when an `extern fn` has been declared with the same name but different types.

Example

```
mod m {
    extern "C" {
        fn foo();
    }
}

extern "C" {
    fn foo(_: u32);
}
```

This will produce:

```
warning: `foo` redeclared with a different signature
--> lint_example.rs:9:5
|
4 |     fn foo();
|     ----- `foo` previously declared here
...
9 |     fn foo(_: u32);
|     ^^^^^^^^^^^^^^^ this signature doesn't match the previous declaration
|
= note: expected `unsafe extern "C" fn()`
      found `unsafe extern "C" fn(u32)`
= note: `#[warn(clashing_extern_declarations)]` on by default
```

Explanation

Because two symbols of the same name cannot be resolved to two different functions at link time, and one function cannot possibly have two types, a clashing extern declaration is almost certainly a mistake. Check to make sure that the `extern` definitions are correct and equivalent, and possibly consider unifying them in one location.

This lint does not run between crates because a project may have dependencies which both rely on the same extern function, but declare it in a different (but valid) way. For example, they may both declare an opaque type for one or more of the arguments (which would end up distinct types), or use types that are valid conversions in the language the `extern fn` is defined in. In these cases, the compiler can't say that the clashing declaration is incorrect.

coherence-leak-check

The `coherence_leak_check` lint detects conflicting implementations of a trait that are only distinguished by the old leak-check code.

Example

```
trait SomeTrait { }
impl SomeTrait for for<'a> fn(&'a u8) { }
impl<'a> SomeTrait for fn(&'a u8) { }
```

This will produce:

```
warning: conflicting implementations of trait `SomeTrait` for type `for<'a> fn(&'a u8)`
--> lint_example.rs:4:1
  |
3 | impl SomeTrait for for<'a> fn(&'a u8) { }
  |----- first implementation here
4 | impl<'a> SomeTrait for fn(&'a u8) { }
  |^^^^^^^^^^^^^^^^^^^^^^^^^^^^^ conflicting implementation for `for<'a> fn(&'a u8)`
  |
  = warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
  = note: for more information, see issue #56105 <https://github.com/rust-lang/rust/issues/56105>
  = note: this behavior recently changed as a result of a bug fix; see rust-lang/rust#56105
for details
  = note: `#[warn(coherence_leak_check)]` on by default
```

Explanation

In the past, the compiler would accept trait implementations for identical functions that differed only in where the lifetime binder appeared. Due to a change in the borrow checker implementation to fix several bugs, this is no longer allowed. However, since this affects existing code, this is a [future-incompatible](#) lint to transition this to a hard error in the future.

Code relying on this pattern should introduce "newtypes", like `struct Foo(for<'a> fn(&'a u8))`.

See [issue #56105](#) for more details.

confusable-idents

The `confusable_idents` lint detects visually confusable pairs between identifiers.

Example

```
// Latin Capital Letter E With Caron
pub const Ě: i32 = 1;
// Latin Capital Letter E With Breve
pub const Ĕ: i32 = 2;
```

This will produce:

```
warning: identifier pair considered confusable between `Ě` and `Ę`
--> lint_example.rs:5:11
3 | pub const Ě: i32 = 1;
   |          ^ this is where the previous identifier occurred
4 | // Latin Capital Letter E With Breve
5 | pub const Ę: i32 = 2;
   |
= note: `#[warn(confusable_ids)]` on by default
```

Explanation

This lint warns when different identifiers may appear visually similar, which can cause confusion.

The confusable detection algorithm is based on [Unicode® Technical Standard #39 Unicode Security Mechanisms Section 4 Confusable Detection](#). For every distinct identifier X execute the function `skeleton(X)`. If there exist two distinct identifiers X and Y in the same crate where `skeleton(X) = skeleton(Y)` report it. The compiler uses the same mechanism to check if an identifier is too similar to a keyword.

Note that the set of confusable characters may change over time. Beware that if you "forbid" this lint that existing code may fail in the future.

const-evaluatable-unchecked

The `const_evaluatable_unchecked` lint detects a generic constant used in a type.

Example

```
const fn foo<T>() -> usize {
    if std::mem::size_of::<*mut T>() < 8 { // size of *mut T does not depend on T
        4
    } else {
        8
    }
}

fn test<T>() {
    let _ = [0; foo::<T>()];
}
```

This will produce:

```
warning: cannot use constants which depend on generic parameters in types
--> lint_example.rs:11:17
11 |     let _ = [0; foo::<T>()];
      ^^^^^^^^^^
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #76200 <https://github.com/rust-lang/rust/issues/76200>
= note: `#[warn(const_evaluatable_unchecked)]` on by default
```

Explanation

In the 1.43 release, some uses of generic parameters in array repeat expressions were accidentally allowed. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #76200](#) for a more detailed description and possible fixes.

const-item-mutation

The `const_item_mutation` lint detects attempts to mutate a `const` item.

Example

```
const FOO: [i32; 1] = [0];

fn main() {
    FOO[0] = 1;
    // This will print "[0]".
    println!("{:?}", FOO);
}
```

This will produce:

```
warning: attempting to modify a `const` item
--> lint_example.rs:4:5
4 |     FOO[0] = 1;
  |     ^^^^^^^^^^
  |
  = note: each usage of a `const` item creates a new temporary; the original `const` item will
not be modified
note: `const` item defined here
--> lint_example.rs:1:1
1 | const FOO: [i32; 1] = [0];
  | ^^^^^^^^^^^^^^^^^^^^^^
  = note: `#[warn(const_item_mutation)]` on by default
```

Explanation

Trying to directly mutate a `const` item is almost always a mistake. What is happening in the example above is that a temporary copy of the `const` is mutated, but the original `const` is not. Each time you refer to the `const` by name (such as `FOO` in the example above), a separate copy of the value is inlined at that location.

This lint checks for writing directly to a field (`FOO.field = some_value`) or array entry (`FOO[0] = val`), or taking a mutable reference to the `const` item (`&mut FOO`), including through an autoderef (`FOO.some_mut_self_method()`).

There are various alternatives depending on what you are trying to accomplish:

- First, always reconsider using mutable globals, as they can be difficult to use correctly, and can make the code more difficult to use or understand.
- If you are trying to perform a one-time initialization of a global:
 - If the value can be computed at compile-time, consider using const-compatible values (see [Constant Evaluation](#)).
 - For more complex single-initialization cases, consider using a third-party crate, such as `lazy_static` or `once_cell`.
 - If you are using the `nightly` channel, consider the new `lazy` module in the standard library.
- If you truly need a mutable global, consider using a `static`, which has a variety of options:
 - Simple data types can be directly defined and mutated with an `atomic` type.
 - More complex types can be placed in a synchronization primitive like a `Mutex`, which can be initialized with one of the options listed above.
 - A `mutable static` is a low-level primitive, requiring unsafe. Typically This should be avoided in preference of something higher-level like one of the above.

dead-code

The `dead_code` lint detects unused, unexported items.

Example

```
fn foo() {}
```

This will produce:

```
warning: function `foo` is never used
--> lint_example.rs:2:4
  |
2 | fn foo() {}
  |     ^
  |
= note: `#[warn(dead_code)]` on by default
```

Explanation

Dead code may signal a mistake or unfinished code. To silence the warning for individual items, prefix the name with an underscore such as `_foo`. If it was intended to expose the item outside of the crate, consider adding a visibility modifier like `pub`. Otherwise consider removing the unused code.

deprecated

The `deprecated` lint detects use of deprecated items.

Example

```
#[deprecated]
fn foo() {}

fn bar() {
    foo();
}
```

This will produce:

```
warning: use of deprecated function `main::foo`
--> lint_example.rs:6:5
6 |     foo();
  |     ^
  |
= note: `#[warn(deprecated)]` on by default
```

Explanation

Items may be marked "deprecated" with the `deprecated` attribute to indicate that they should no longer be used. Usually the attribute should include a note on what to use instead, or check the documentation.

deprecated_where_clause_location

The `deprecated_where_clause_location` lint detects when a where clause is in front of the equals in an associated type.

Example

```
trait Trait {
    type Assoc<'a> where Self: 'a;
}

impl Trait for () {
    type Assoc<'a> where Self: 'a = ();
}
```

This will produce:

```
warning: where clause not allowed here
--> lint_example.rs:7:18
7 |     type Assoc<'a> where Self: 'a = ();
  |     ^^^^^^
  |
= note: see issue #89122 <https://github.com/rust-lang/rust/issues/89122> for more
information
= note: `#[warn(deprecated_where_clause_location)]` on by default
help: move it to the end of the type declaration
|
7 -     type Assoc<'a> where Self: 'a = ();
7 +     type Assoc<'a> = () where Self: 'a;
|
```

Explanation

The preferred location for where clauses on associated types in impls is after the type. However, for most of generic associated types development, it was only accepted before the equals. To provide a transition period and further evaluate this change, both are currently accepted. At some point in the future, this may be disallowed at an edition boundary; but, that is undecided currently.

deref-into-dyn-supertrait

The `deref_into_dyn_supertrait` lint is output whenever there is a use of the `Deref` implementation with a `dyn SuperTrait` type as `Output`.

These implementations will become shadowed when the `trait_upcasting` feature is stabilized. The `deref` functions will no longer be called implicitly, so there might be behavior change.

Example

```
#![deny(deref_into_dyn_supertrait)]
#![allow(dead_code)]

use core::ops::Deref;

trait A {}
trait B: A {}

impl<'a> Deref for dyn 'a + B {
    type Target = dyn A;
    fn deref(&self) -> &Self::Target {
        todo!()
    }
}

fn take_a(_: &dyn A) {}

fn take_b(b: &dyn B) {
    take_a(b);
}
```

This will produce:

```
error: `(dyn B + 'a)` implements `Deref` with supertrait `A` as target
--> lint_example.rs:9:1
   |
9 |     impl<'a> Deref for dyn 'a + B {
   |     ^^^^^^^^^^
10|         type Target = dyn A;
           ----- target type is set here
   |
   = warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
   = note: for more information, see issue #89460 <https://github.com/rust-lang/rust-lang/issues/89460>
note: the lint level is defined here
--> lint_example.rs:1:9
   |
1 | #[deny(deref_into_dyn_supertrait)]
   | ^^^^^^^^^^
```

Explanation

The dynamic upcasting coercion feature adds new coercion rules, taking priority over certain other coercion rules, which will cause some behavior change.

deref-nullptr

The `deref_nullptr` lint detects when an null pointer is dereferenced, which causes undefined behavior.

Example

```
use std::ptr;
unsafe {
    let x = &*ptr::null::<i32>();
    let x = ptr::addr_of!(*ptr::null::<i32>());
    let x = *(0 as *const i32);
}
```

This will produce:

```

warning: dereferencing a null pointer
--> lint_example.rs:5:14
5 |     let x = &*ptr::null::<i32>();
|           ^^^^^^^^^^^^^^^^^^ this code causes undefined behavior when executed
|
= note: `#[warn(deref_nullptr)]` on by default

warning: dereferencing a null pointer
--> lint_example.rs:6:27
6 |     let x = ptr::addr_of!(*ptr::null::<i32>());
|           ^^^^^^^^^^^^^^^^^^ this code causes undefined behavior when
executed

warning: dereferencing a null pointer
--> lint_example.rs:7:13
7 |     let x = *(0 as *const i32);
|           ^^^^^^^^^^^^^^^^^^ this code causes undefined behavior when executed

```

Explanation

Dereferencing a null pointer causes **undefined behavior** even as a place expression, like `&*(0 as *const i32)` or `addr_of!(*(0 as *const i32))`.

drop-bounds

The `drop_bounds` lint checks for generics with `std::ops::Drop` as bounds.

Example

```
fn foo<T: Drop>() {}
```

This will produce:

```

warning: bounds on `T: Drop` are most likely incorrect, consider instead using
`std::mem::needs_drop` to detect whether a type can be trivially dropped
--> lint_example.rs:2:11
2 | fn foo<T: Drop>() {}
|           ^
|
= note: `#[warn(drop_bounds)]` on by default

```

Explanation

A generic trait bound of the form `T: Drop` is most likely misleading and not what the programmer intended (they probably should have used `std::mem::needs_drop` instead).

`Drop` bounds do not actually indicate whether a type can be trivially dropped or not, because a composite type containing `Drop` types does not necessarily implement `Drop` itself. Naïvely, one might be tempted to write an implementation that assumes that a type can be trivially dropped while also supplying a specialization for `T: Drop` that actually calls the destructor. However, this breaks down e.g. when `T` is `String`, which does not implement `Drop` itself but contains a `Vec`, which does implement `Drop`, so assuming `T` can be trivially dropped would lead to a memory leak here.

Furthermore, the `Drop` trait only contains one method, `Drop::drop`, which may not be called explicitly in user code (E0040), so there is really no use case for using `Drop` in trait bounds, save perhaps for some obscure corner cases, which can use `#[allow(drop_bounds)]`.

duplicate-macro-attributes

The `duplicate_macro_attributes` lint detects when a `#[test]`-like built-in macro attribute is duplicated on an item. This lint may trigger on `bench`, `cfg_eval`, `test` and `test_case`.

Example

```
#[test]
#[test]
fn foo() {}
```

This will produce:

```
warning: duplicated attribute
--> src/lib.rs:2:1
2 | #[test]
  | ^^^^^^
  |
= note: `#[warn(duplicate_macro_attributes)]` on by default
```

Explanation

A duplicated attribute may erroneously originate from a copy-paste and the effect of it being duplicated may not be obvious or desirable.

For instance, doubling the `#[test]` attributes registers the test to be run twice with no change to its environment.

dyn-drop

The `dyn_drop` lint checks for trait objects with `std::ops::Drop`.

Example

```
fn foo(_x: Box<dyn Drop>) {}
```

This will produce:

```
warning: types that do not implement `Drop` can still have drop glue, consider instead using
`std::mem::needs_drop` to detect whether a type is trivially dropped
--> lint_example.rs:2:20
  |
2 | fn foo(_x: Box<dyn Drop>) {}
  |          ^^^^
  |
= note: `#[warn(dyn_drop)]` on by default
```

Explanation

A trait object bound of the form `dyn Drop` is most likely misleading and not what the programmer intended.

`Drop` bounds do not actually indicate whether a type can be trivially dropped or not, because a composite type containing `Drop` types does not necessarily implement `Drop` itself. Naively, one might be tempted to write a deferred drop system, to pull cleaning up memory out of a latency-sensitive code path, using `dyn Drop` trait objects. However, this breaks down e.g. when `T` is `String`, which does not implement `Drop`, but should probably be accepted.

To write a trait object bound that accepts anything, use a placeholder trait with a blanket implementation.

```
trait Placeholder {}
impl<T> Placeholder for T {}
fn foo(_x: Box<dyn Placeholder>) {}
```

ellipsis-inclusive-range-patterns

The `ellipsis_inclusive_range_patterns` lint detects the `... range pattern`, which is deprecated.

Example

```
let x = 123;
match x {
    0...100 => {}
    _ => {}
}
```

This will produce:

```
warning: `...` range patterns are deprecated
--> lint_example.rs:4:6
4 |     0...100 => {}
  |     ^^^ help: use `..=` for an inclusive range
= warning: this is accepted in the current edition (Rust 2018) but is a hard error in Rust
2021!
= note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/warnings-promoted-to-error.html>
= note: `#[warn(ellipsis_inclusive_range_patterns)]` on by default
```

Explanation

The `... range pattern` syntax was changed to `..=` to avoid potential confusion with the `... range expression`. Use the new form instead.

exported-private-dependencies

The `exported_private_dependencies` lint detects private dependencies that are exposed in a public interface.

Example

```
pub fn foo() -> Option<some_private_dependency::Thing> {
    None
}
```

This will produce:

```
warning: type `bar::Thing` from private dependency 'bar' in public interface
--> src/lib.rs:3:1
3 | pub fn foo() -> Option<bar::Thing> {
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
= note: `#[warn(exported_private_dependencies)]` on by default
```

Explanation

Dependencies can be marked as "private" to indicate that they are not exposed in the public interface of a crate. This can be used by Cargo to independently resolve those dependencies because it can assume it does not need to unify them with other packages using that same dependency. This lint is an indication of a violation of that contract.

To fix this, avoid exposing the dependency in your public interface. Or, switch the dependency to a public dependency.

Note that support for this is only available on the nightly channel. See [RFC 1977](#) for more details, as well as the [Cargo documentation](#).

for-loops-over-fallibles

The `for_loops_over_fallibles` lint checks for `for` loops over `Option` or `Result` values.

Example

```
let opt = Some(1);
for x in opt { /* ... */}
```

This will produce:

```
warning: for loop over an `Option`. This is more readably written as an `if let` statement
--> lint_example.rs:3:10
3 | for x in opt { /* ... */}
  | ^^^
= note: `#[warn(for_loops_over_fallibles)]` on by default
help: to check pattern in a loop use `while let`
|
3 | while let Some(x) = opt { /* ... */}
  | ~~~~~ ~~~
help: consider using `if let` to clear intent
|
3 | if let Some(x) = opt { /* ... */}
  | ~~~~~ ~~~
```

Explanation

Both `Option` and `Result` implement `IntoIterator` trait, which allows using them in a `for` loop. `for` loop over `Option` or `Result` will iterate either 0 (if the value is `None`/`Err(_)`) or 1 time (if the value is `Some(_)`/`Ok(_)`). This is not very useful and is more clearly expressed via `if let`.

`for` loop can also be accidentally written with the intention to call a function multiple times, while the function returns `Some(_)`, in these cases `while let` loop should be used instead.

The "intended" use of `IntoIterator` implementations for `Option` and `Result` is passing them to generic code that expects something implementing `IntoIterator`. For example using `.chain(option)` to optionally add a value to an iterator.

forbidden-lint-groups

The `forbidden_lint_groups` lint detects violations of `forbid` applied to a lint group. Due to a bug in the compiler, these used to be overlooked entirely. They now generate a warning.

Example

```
#![forbid(warnings)]
#![deny(bad_style)]

fn main() {}
```

This will produce:

```
warning: deny(bad_style) incompatible with previous forbid
--> lint_example.rs:2:9
  |
1 | #![forbid(warnings)]
  |          ----- `forbid` level set here
2 | #![deny(bad_style)]
  |          ^^^^^^^^^^ overruled by previous forbid
  |
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #81670 <https://github.com/rust-lang/rust/issues/81670>
= note: `#[warn(forbidden_lint_groups)]` on by default
```

Recommended fix

If your crate is using `#![forbid(warnings)]`, we recommend that you change to `#![deny(warnings)]`.

Explanation

Due to a compiler bug, applying `forbid` to lint groups previously had no effect. The bug is now fixed but instead of enforcing `forbid` we issue this future-compatibility warning to avoid breaking existing crates.

function-item-references

The `function_item_references` lint detects function references that are formatted with `fmt::Pointer` or transmuted.

Example

```
fn foo() { }

fn main() {
    println!("{:p}", &foo);
}
```

This will produce:

```
warning: taking a reference to a function item does not give a function pointer
--> lint_example.rs:4:22
  |
4 |     println!("{:p}", &foo);
      ^^^^ help: cast `foo` to obtain a function pointer: `foo as fn()`
  |
= note: `#[warn(function_item_references)]` on by default
```

Explanation

Taking a reference to a function may be mistaken as a way to obtain a pointer to that function. This can give unexpected results when formatting the reference as a pointer or transmuting it. This lint is issued when function references are formatted as pointers, passed as arguments bound by `fmt::Pointer` or transmuted.

illegal-floating-point-literal-pattern

The `illegal_floating_point_literal_pattern` lint detects floating-point literals used in patterns.

Example

```
let x = 42.0;

match x {
    5.0 => {}
    _   => {}
}
```

This will produce:

```
warning: floating-point types cannot be used in patterns
--> lint_example.rs:5:5
  |
5 |     5.0 => {}
  |     ^^^
  |
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #41620 <https://github.com/rust-lang/rust/issues/41620>
= note: `#[warn(illegal_floating_point_literal_pattern)]` on by default
```

Explanation

Previous versions of the compiler accepted floating-point literals in patterns, but it was later determined this was a mistake. The semantics of comparing floating-point values may not be clear in a pattern when contrasted with "structural equality". Typically you can work around this by using a `match guard`, such as:

```
match x {
    y if y == 5.0 => {}
    _   => {}
}
```

This is a `future-incompatible` lint to transition this to a hard error in the future. See [issue #41620](#) for more details.

improper-ctypes

The `improper_ctypes` lint detects incorrect use of types in foreign modules.

Example

```
extern "C" {
    static STATIC: String;
}
```

This will produce:

```
warning: `extern` block uses type `String`, which is not FFI-safe
--> lint_example.rs:3:20
  |
3 |     static STATIC: String;
  |             ^^^^^^ not FFI-safe
  |
= help: consider adding a `#[repr(C)]` or `#[repr(transparent)]` attribute to this struct
= note: this struct has unspecified layout
= note: `#[warn(improper_ctypes)]` on by default
```

Explanation

The compiler has several checks to verify that types used in `extern` blocks are safe and follow certain rules to ensure proper compatibility with the foreign interfaces. This lint is issued when it detects a probable mistake in a definition. The lint usually should provide a description of the issue, along with possibly a hint on how to resolve it.

improper-ctypes-definitions

The `improper_ctypes_definitions` lint detects incorrect use of `extern function` definitions.

Example

```
pub extern "C" fn str_type(p: &str) { }
```

This will produce:

```
warning: `extern` fn uses type `str`, which is not FFI-safe
--> lint_example.rs:3:31
  |
3 | pub extern "C" fn str_type(p: &str) { }
      ^^^^^ not FFI-safe
  |
  = help: consider using `*const u8` and a length instead
  = note: string slices have no C equivalent
  = note: `#[warn(improper_ctypes_definitions)]` on by default
```

Explanation

There are many parameter and return types that may be specified in an `extern` function that are not compatible with the given ABI. This lint is an alert that these types should not be used. The lint usually should provide a description of the issue, along with possibly a hint on how to resolve it.

incomplete-features

The `incomplete_features` lint detects unstable features enabled with the `feature attribute` that may function improperly in some or all cases.

Example

```
#![feature(generic_const_exprs)]
```

This will produce:

```
warning: the feature `generic_const_exprs` is incomplete and may not be safe to use and/or
cause compiler crashes
--> lint_example.rs:1:12
  |
1 | #![feature(generic_const_exprs)]
      ^^^^^^^^^^^^^^^^^^^^^^
  |
  = note: see issue #76560 <https://github.com/rust-lang/rust/issues/76560> for more
information
  = note: `#[warn(incomplete_features)]` on by default
```

Explanation

Although it is encouraged for people to experiment with unstable features, some of them are known to be incomplete or faulty. This lint is a signal that the feature has not yet been finished, and you may experience problems with it.

indirect-structural-match

The `indirect_structural_match` lint detects a `const` in a pattern that manually implements `PartialEq` and `Eq`.

Example

```
#![deny(indirect_structural_match)]

struct NoDerive(i32);
impl PartialEq for NoDerive { fn eq(&self, _: &Self) -> bool { false } }
impl Eq for NoDerive { }
#[derive(PartialEq, Eq)]
struct WrapParam<T>(T);
const WRAP_INDIRECT_PARAM: & &WrapParam<NoDerive> = & &WrapParam(NoDerive(0));
fn main() {
    match WRAP_INDIRECT_PARAM {
        WRAP_INDIRECT_PARAM => { }
        _ => { }
    }
}
```

This will produce:

```
error: to use a constant of type `NoDerive` in a pattern, `NoDerive` must be annotated with `#`  
[derive(PartialEq, Eq)]`  
--> lint_example.rs:11:9  
|  
11 |     WRAP_INDIRECT_PARAM => { }  
|     ^^^^^^^^^^  
  
= warning: this was previously accepted by the compiler but is being phased out; it will  
become a hard error in a future release!  
= note: for more information, see issue #62411 <https://github.com/rust-lang/rust-lang/issues/62411>  
= note: the traits must be derived, manual `impl`s are not sufficient  
= note: see https://doc.rust-lang.org/stable/std/marker/trait.StructuralEq.html for details  
note: the lint level is defined here  
--> lint_example.rs:1:9  
|  
1 | #[deny(indirect_structural_match)]  
|     ^^^^^^^^^^
```

Explanation

The compiler unintentionally accepted this form in the past. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #62411](#) for a complete description of the problem, and some possible solutions.

inline-no-sanitize

The `inline_no_sanitze` lint detects incompatible use of `#[inline(always)]` and `#[no_sanitze(...)]`.

Example

```
#![feature(no_sanitze)]  
  
#[inline(always)]  
#[no_sanitze(address)]  
fn x() {}  
  
fn main() {  
    x()  
}
```

This will produce:

```
warning: `no_sanitze` will have no effect after inlining  
--> lint_example.rs:4:1  
4 | #[no_sanitze(address)]  
| ^^^^^^^^^^^^^^^^^^  
|  
note: inlining requested here  
--> lint_example.rs:3:1  
3 | #[inline(always)]  
| ^^^^^^^^^^  
= note: `#[warn(inline_no_sanitze)]` on by default
```

Explanation

The use of the `#[inline(always)]` attribute prevents the the `#[no_sanitze(...)]` attribute from working. Consider temporarily removing `inline` attribute.

invalid-doc-attributes

The `invalid_doc_attributes` lint detects when the `#[doc(...)]` is misused.

Example

```
#![deny(warnings)]  
  
pub mod submodule {  
    #![doc(test(no_crate_inject))]  
}
```

This will produce:

```
error: this attribute can only be applied at the crate level  
--> lint_example.rs:5:12  
|  
5 |     #![doc(test(no_crate_inject))]  
|           ^^^^^^^^^^^^^^^^^^  
|  
= warning: this was previously accepted by the compiler but is being phased out; it will  
become a hard error in a future release!  
= note: for more information, see issue #82730 <https://github.com/rust-lang/rust/issues/82730>  
= note: read <https://doc.rust-lang.org/nightly/rustdoc/the-doc-attribute.html#at-the-crate-level> for more information  
note: the lint level is defined here  
--> lint_example.rs:1:9  
|  
1 | #![deny(warnings)]  
|       ^^^^^^  
= note: `#[deny(invalid_doc_attributes)]` implied by `#[deny(warnings)]`
```

Explanation

Previously, incorrect usage of the `#[doc(..)]` attribute was not being validated. Usually these should be rejected as a hard error, but this lint was introduced to avoid breaking any existing crates which included them.

This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #82730](#) for more details.

invalid-macro-export-arguments

The `invalid_macro_export_arguments` lint detects cases where `#[macro_export]` is being used with invalid arguments.

Example

```
#![deny(invalid_macro_export_arguments)]  
  
#[macro_export(invalid_parameter)]  
macro_rules! myMacro {  
    () => {  
        // [...]  
    }  
}  
  
#[macro_export(too, many, items)]
```

This will produce:

```
warning: `invalid_parameter` isn't a valid `#[macro_export]` argument  
--> lint_example.rs:4:16  
4 | #[macro_export(invalid_parameter)]  
| ^^^^^^^^^^^^^^^^^^  
= note: `#[warn(invalid_macro_export_arguments)]` on by default
```

Explanation

The only valid argument is `#[macro_export(local_inner_macros)]` or no argument (`#[macro_export]`). You can't have multiple arguments in a `#[macro_export(..)]`, or mention arguments other than `local_inner_macros`.

invalid-value

The `invalid_value` lint detects creating a value that is not valid, such as a null reference.

Example

```
unsafe {  
    let x: &'static i32 = std::mem::zeroed();  
}
```

This will produce:

```
warning: the type `&i32` does not permit zero-initialization
--> lint_example.rs:4:27
4 |     let x: &'static i32 = std::mem::zeroed();
   |     ^^^^^^^^^^^^^^^^^^
   |
   |     this code causes undefined behavior when executed
   |     help: use `MaybeUninit<T>` instead, and only call `assume_init`
after initialization is done
   |
   |= note: references must be non-null
   |= note: `#[warn(invalid_value)]` on by default
```

Explanation

In some situations the compiler can detect that the code is creating an invalid value, which should be avoided.

In particular, this lint will check for improper use of `mem::zeroed`, `mem::uninitialized`, `mem::transmute`, and `MaybeUninit::assume_init` that can cause [undefined behavior](#). The lint should provide extra information to indicate what the problem is and a possible solution.

irrefutable-let-patterns

The `irrefutable_let_patterns` lint detects [irrefutable patterns](#) in `if let`s, `while let`s, and `if let` guards.

Example

```
if let _ = 123 {
    println!("always runs!");
}
```

This will produce:

```
warning: irrefutable `if let` pattern
--> lint_example.rs:2:4
2 | if let _ = 123 {
   |     ^^^^^^^^^^
   |
   |= note: this pattern will always match, so the `if let` is useless
   |= help: consider replacing the `if let` with a `let`
   |= note: `#[warn(irrefutable_let_patterns)]` on by default
```

Explanation

There usually isn't a reason to have an irrefutable pattern in an `if let` or `while let` statement, because the pattern will always match successfully. A `let` or `loop` statement will suffice. However, when generating code with a macro, forbidding irrefutable patterns would require awkward workarounds in situations where the macro doesn't know if the pattern is refutable or not. This lint allows macros to accept this form, while alerting for a possibly incorrect use in normal code.

See [RFC 2086](#) for more details.

large-assignments

The `large_assignments` lint detects when objects of large types are being moved around.

Example

```
let x = [0; 50000];
let y = x;
```

produces:

```
warning: moving a large value
--> $DIR/move-large.rs:1:3
let y = x;
    - Copied large value here
```

Explanation

When using a large type in a plain assignment or in a function argument, idiomatic code can be inefficient. Ideally appropriate optimizations would resolve this, but such optimizations are only done in a best-effort manner. This lint will trigger on all sites of large moves and thus allow the user to resolve them in code.

late-bound-lifetime-arguments

The `late_bound_lifetime_arguments` lint detects generic lifetime arguments in path segments with late bound lifetime parameters.

Example

```
struct S;

impl S {
    fn late(self, _: &u8, _: &u8) {}
}

fn main() {
    S.late::<'static>(&0, &0);
}
```

This will produce:

```
warning: cannot specify lifetime arguments explicitly if late bound lifetime parameters are
present
--> lint_example.rs:8:14
4 |     fn late(self, _: &u8, _: &u8) {}
|             ^ the late bound lifetime parameter is introduced here
...
8 |     S.late::<'static>(&0, &0);
|             ^^^^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #42868 <https://github.com/rust-lang/rust/issues/42868>
= note: `#[warn(late_bound_lifetime_arguments)]` on by default
```

Explanation

It is not clear how to provide arguments for early-bound lifetime parameters if they are intermixed with late-bound parameters in the same list. For now, providing any explicit arguments will trigger this lint if late-bound parameters are present, so in the future a solution can be adopted without hitting backward compatibility issues. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #42868](#) for more details, along with a description of the difference between early and late-bound parameters.

legacy-derive-helpers

The `legacy_derive_helpers` lint detects derive helper attributes that are used before they are introduced.

Example

```
#[serde(rename_all = "camelCase")]
#[derive(Deserialize)]
struct S { /* fields */ }
```

produces:

```
warning: derive helper attribute is used before it is introduced
--> $DIR/legacy-derive-helpers.rs:1:3
  |
1 | #[serde(rename_all = "camelCase")]
  | ^^^^^^
...
2 | #[derive(Deserialize)]
  | ----- the attribute is introduced here
```

Explanation

Attributes like this work for historical reasons, but attribute expansion works in left-to-right order in general, so, to resolve `#[serde]`, compiler has to try to "look into the future" at not yet expanded part of the item , but such attempts are not always reliable.

To fix the warning place the helper attribute after its corresponding derive.

```
#[derive(Deserialize)]
#[serde(rename_all = "camelCase")]
struct S { /* fields */ }
```

map-unit-fn

The `map_unit_fn` lint checks for `Iterator::map` receive a callable that returns `()`.

Example

```
fn foo(items: &mut Vec<u8>) {
    items.sort();
}

fn main() {
    let mut x: Vec<Vec<u8>> = vec![
        vec![0, 2, 1],
        vec![5, 4, 3],
    ];
    x.iter_mut().map(foo);
}
```

This will produce:

```
warning: `Iterator::map` call that discard the iterator's values
--> lint_example.rs:10:18
  |
1 | fn foo(items: &mut Vec<u8>) {
  | ----- this function returns `()`, which is likely not what you
wanted
...
10|     x.iter_mut().map(foo);
   |           ^^^^^^--^
   |           |
   |           called `Iterator::map` with callable that returns `()`
   |           after this call to map, the resulting iterator is `impl Iterator<Item = ()>`, which means the only information carried by the iterator is the number of items
   |
   = note: `Iterator::map`, like many of the methods on `Iterator`, gets executed lazily,
meaning that its effects won't be visible until it is iterated
   = note: `#[warn(map_unit_fn)]` on by default
help: you might have meant to use `Iterator::for_each`
10|     x.iter_mut().for_each(foo);
   |           ~~~~~~
```

Explanation

Mapping to `()` is almost always a mistake.

mixed-script-confusables

The `mixed_script_confusables` lint detects visually confusable characters in identifiers between different [scripts](#).

Example

```
// The Japanese katakana character イ can be confused with the Han character 工.
const イ: &'static str = "アイウ";
```

This will produce:

```
warning: the usage of Script Group `Japanese, Katakana` in this crate consists solely of mixed
script confusables
--> lint_example.rs:3:7
  |
3 | const エ: &'static str = "アイウ";
  |     ^^
  |
= note: the usage includes 'エ' (U+30A8)
= note: please recheck to make sure their usages are indeed what you want
= note: `#[warn(mixed_script_confusables)]` on by default
```

Explanation

This lint warns when characters between different scripts may appear visually similar, which can cause confusion.

If the crate contains other identifiers in the same script that have non-confusable characters, then this lint will *not* be issued. For example, if the example given above has another identifier with katakana characters (such as `let カタカナ = 123;`), then this indicates that you are intentionally using katakana, and it will not warn about it.

Note that the set of confusable characters may change over time. Beware that if you "forbid" this lint that existing code may fail in the future.

named-arguments-used-positionally

The `named_arguments_used_positionally` lint detects cases where named arguments are only used positionally in format strings. This usage is valid but potentially very confusing.

Example

```
#![deny(named_arguments_used_positionally)]
fn main() {
    let _x = 5;
    println!("{}", _x); // Prints 1, will trigger lint

    println!("{}", _x); // Prints 5, no lint emitted
    println!("{{_x}}", _x = _x); // Prints 5, no lint emitted
}
```

This will produce:

```
error: named argument `_x` is not used by name
--> lint_example.rs:4:20
4 |     println!("{}", _x = 1); // Prints 1, will trigger lint
   |           ^^^ this named argument is referred to by position in formatting string
   |           |
   |           this formatting argument uses named argument `_x` by position

note: the lint level is defined here
--> lint_example.rs:1:9
1 | #[deny(named_arguments_used_positionally)]
   | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

help: use the named argument by name to avoid ambiguity
4 |     println!("{_x}", _x = 1); // Prints 1, will trigger lint
   |           ++
```

Explanation

Rust formatting strings can refer to named arguments by their position, but this usage is potentially confusing. In particular, readers can incorrectly assume that the declaration of named arguments is an assignment (which would produce the unit type). For backwards compatibility, this is not a hard error.

no-mangle-generic-items

The `no_mangle_generic_items` lint detects generic items that must be mangled.

Example

```
#[no_mangle]
fn foo<T>(t: T) {  
}
```

This will produce:

```
warning: functions generic over types or consts must be mangled
--> lint_example.rs:3:1
2 |     #[no_mangle]
|----- help: remove this attribute
3 | / fn foo<T>(t: T) {
4 | |
5 | |
|_-^
= note: `#[warn(no_mangle_generic_items)]` on by default
```

Explanation

A function with generics must have its symbol mangled to accommodate the generic parameter. The `no_mangle` attribute has no effect in this situation, and should be removed.

non-camel-case-types

The `non_camel_case_types` lint detects types, variants, traits and type parameters that don't have camel case names.

Example

```
struct my_struct;
```

This will produce:

```
warning: type `my_struct` should have an upper camel case name
--> lint_example.rs:2:8
2 | struct my_struct;
| ^^^^^^^^^^^ help: convert the identifier to upper camel case: `MyStruct`
|-
= note: `#[warn(non_camel_case_types)]` on by default
```

Explanation

The preferred style for these identifiers is to use "camel case", such as `MyStruct`, where the first letter should not be lowercase, and should not use underscores between letters. Underscores are allowed at the beginning and end of the identifier, as well as between non-letters (such as `x86_64`).

non-fmt-panics

The `non_fmt_panics` lint detects `panic!(..)` invocations where the first argument is not a formatting string.

Example

```
panic!("{}");
panic!(123);
```

This will produce:

```
warning: panic message contains an unused formatting placeholder
--> lint_example.rs:2:9
  |
2 | panic!("{}");
  ^^^
  |
= note: this message is not used as a format string when given without arguments, but will
be in Rust 2021
= note: `#[warn(non_fmt_panics)]` on by default
help: add the missing argument
  |
2 | panic!("{}", ...);
  |      ++++++
help: or add a "{}" format string to use the message literally
  |
2 | panic!("{}", "{}");
  |      +++++

warning: panic message is not a string literal
--> lint_example.rs:3:8
  |
3 | panic!(123);
  ^^^
  |
= note: this usage of `panic!()` is deprecated; it will be a hard error in Rust 2021
= note: for more information, see <https://doc.rust-lang.org/nightly/edition-guide/rust-2021/panic-macro-consistency.html>
help: add a "{}" format string to `Display` the message
  |
3 | panic!("{}", 123);
  |      +++++
help: or use std::panic::panic_any instead
  |
3 | std::panic::panic_any(123);
  | ~~~~~
```

Explanation

In Rust 2018 and earlier, `panic!(x)` directly uses `x` as the message. That means that `panic!("{}")` panics with the message `"{}"` instead of using it as a formatting string, and `panic!(123)` will panic with an `i32` as message.

Rust 2021 always interprets the first argument as format string.

non-shorthand-field-patterns

The `non_shorthand_field_patterns` lint detects using `Struct { x: x }` instead of `Struct { x }` in a pattern.

Example

```
struct Point {
    x: i32,
    y: i32,
}

fn main() {
    let p = Point {
        x: 5,
        y: 5,
    };

    match p {
        Point { x: x, y: y } => (),
    }
}
```

This will produce:

```
warning: the `x:` in this pattern is redundant
--> lint_example.rs:14:17
14 |         Point { x: x, y: y } => (),
|             ^^^^ help: use shorthand field pattern: `x`
| =
= note: `#[warn(non_shorthand_field_patterns)]` on by default

warning: the `y:` in this pattern is redundant
--> lint_example.rs:14:23
14 |         Point { x: x, y: y } => (),
|             ^^^^ help: use shorthand field pattern: `y`
```

Explanation

The preferred style is to avoid the repetition of specifying both the field name and the binding name if both identifiers are the same.

non-snake-case

The `non_snake_case` lint detects variables, methods, functions, lifetime parameters and modules that don't have snake case names.

Example

```
let MY_VALUE = 5;
```

This will produce:

```
warning: variable `MY_VALUE` should have a snake case name
--> lint_example.rs:2:5
2 | let MY_VALUE = 5;
| ^^^^^^^^^ help: convert the identifier to snake case: `my_value`
| = note: `#[warn(non_snake_case)]` on by default
```

Explanation

The preferred style for these identifiers is to use "snake case", where all the characters are in lowercase, with words separated with a single underscore, such as `my_value`.

non-upper-case-globals

The `non_upper_case_globals` lint detects static items that don't have uppercase identifiers.

Example

```
static max_points: i32 = 5;
```

This will produce:

```
warning: static variable `max_points` should have an upper case name
--> lint_example.rs:2:8
2 | static max_points: i32 = 5;
| ^^^^^^^^^^^ help: convert the identifier to upper case: `MAX_POINTS`
| = note: `#[warn(non_upper_case_globals)]` on by default
```

Explanation

The preferred style is for static item names to use all uppercase letters such as `MAX_POINTS`.

nontrivial-structural-match

The `nontrivial_structural_match` lint detects constants that are used in patterns, whose type is not structural-match and whose initializer body actually uses values that are not structural-match. So `Option<NotStructuralMatch>` is ok if the constant is just `None`.

Example

```
#![deny(nontrivial_structural_match)]

#[derive(Copy, Clone, Debug)]
struct NoDerive(u32);
impl PartialEq for NoDerive { fn eq(&self, _: &Self) -> bool { false } }
impl Eq for NoDerive { }
fn main() {
    const INDEX: Option<NoDerive> = [None, Some(NoDerive(10))][0];
    match None { Some(_) => panic!("whoops"), INDEX => dbg!(INDEX), };
}
```

This will produce:

```
error: to use a constant of type `NoDerive` in a pattern, the constant's initializer must be
trivial or `NoDerive` must be annotated with `#[derive(PartialEq, Eq)]`
--> lint_example.rs:9:47
 |
9 |     match None { Some(_) => panic!("whoops"), INDEX => dbg!(INDEX), };
   |           ^^^^^^
   |
   = warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
  = note: for more information, see issue #73448 <https://github.com/rust-lang/rust-lang/issues/73448>
  = note: the traits must be derived, manual `impl`s are not sufficient
  = note: see https://doc.rust-lang.org/stable/std/marker/trait.StructuralEq.html for details
note: the lint level is defined here
--> lint_example.rs:1:9
 |
1 | #![deny(nontrivial_structural_match)]
   | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

Previous versions of Rust accepted constants in patterns, even if those constants' types did not have `PartialEq` derived. Thus the compiler falls back to runtime execution of `PartialEq`, which can report that two constants are not equal even if they are bit-equivalent.

opaque-hidden-inferred-bound

The `opaque_hidden_inferred_bound` lint detects cases in which nested `impl Trait` in associated type bounds are not written generally enough to satisfy the bounds of the associated type.

Explanation

This functionality was removed in #97346, but then rolled back in #99860 because it caused regressions.

We plan on reintroducing this as a hard error, but in the mean time, this lint serves to warn and suggest fixes for any use-cases which rely on this behavior.

Example

```
#![feature(type_alias_impl_trait)]  
  
trait Duh {}  
  
impl Duh for i32 {}  
  
trait Trait {  
    type Assoc: Duh;  
}  
  
struct Struct;  
  
impl<F: Duh> Trait for F {  
    type Assoc = F;  
}  
  
type Tait = impl Sized;  
  
fn test() -> impl Trait<Assoc = Tait> {  
    42  
}
```

This will produce:

```
warning: opaque type `impl Trait<Assoc = Tait>` does not satisfy its associated type bounds
--> lint_example.rs:20:25
9 |     type Assoc: Duh;
|             --- this associated type bound is unsatisfied for `Tait`
...
20 | fn test() -> impl Trait<Assoc = Tait> {
|             ^^^^^^^^^^
|
|= note: `#[warn(opaque_hidden_inferred_bound)]` on by default
help: add this bound
18 | type Tait = impl Sized + Duh;
|         +++++
```

In this example, `test` declares that the associated type `Assoc` for `impl Trait` is `impl Sized`, which does not satisfy the `Send` bound on the associated type.

Although the hidden type, `i32` does satisfy this bound, we do not consider the return type to be well-formed with this lint. It can be fixed by changing `Tait = impl Sized` into `Tait = impl Sized + Send`.

overlapping-range-endpoints

The `overlapping_range_endpoints` lint detects `match` arms that have `range patterns` that overlap on their endpoints.

Example

```
let x = 123u8;
match x {
    0..=100 => { println!("small"); }
    100..=255 => { println!("large"); }
}
```

This will produce:

```
warning: multiple patterns overlap on their endpoints
--> lint_example.rs:5:5
4 |     0..=100 => { println!("small"); }
|             ----- this range overlaps on `100..=255`...
5 |     100..=255 => { println!("large"); }
|             ^^^^^^ ... with this range
|
|= note: you likely meant to write mutually exclusive ranges
|= note: `#[warn(overlapping_range_endpoints)]` on by default
```

Explanation

It is likely a mistake to have range patterns in a match expression that overlap in this way. Check that the beginning and end values are what you expect, and keep in mind that with `...=` the left and right bounds are inclusive.

path-statements

The `path_statements` lint detects path statements with no effect.

Example

```
let x = 42;  
x;
```

This will produce:

```
warning: path statement with no effect  
--> lint_example.rs:4:1  
4 | x;  
  | ^  
  = note: `#[warn(path_statements)]` on by default
```

Explanation

It is usually a mistake to have a statement that has no effect.

private-in-public

The `private_in_public` lint detects private items in public interfaces not caught by the old implementation.

Example

```
struct SemiPriv;

mod m1 {
    struct Priv;
    impl super::SemiPriv {
        pub fn f(_: Priv) {}
    }
}
```

This will produce:

```
warning: private type `Priv` in public interface (error E0446)
--> lint_example.rs:7:9
7 |     pub fn f(_: Priv) {}
|     ^^^^^^^^^^^^^^^^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #34537 <https://github.com/rust-lang/rust/issues/34537>
= note: `#[warn(private_in_public)]` on by default
```

Explanation

The visibility rules are intended to prevent exposing private items in public interfaces. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #34537](#) for more details.

redundant-semicolons

The `redundant_semicolons` lint detects unnecessary trailing semicolons.

Example

```
let _ = 123;;
```

This will produce:

```
warning: unnecessary trailing semicolon
--> lint_example.rs:2:13
2 | let _ = 123;;
   |          ^ help: remove this semicolon
= note: `#[warn(redundant_semicolons)]` on by default
```

Explanation

Extra semicolons are not needed, and may be removed to avoid confusion and visual clutter.

renamed-and-removed-lints

The `renamed_and_removed_lints` lint detects lints that have been renamed or removed.

Example

```
#![deny(raw_pointer_derive)]
```

This will produce:

```
warning: lint `raw_pointer_derive` has been removed: using derive with raw pointers is ok
--> lint_example.rs:1:9
1 | #![deny(raw_pointer_derive)]
   |          ^^^^^^^^^^^^^^^^^^
= note: `#[warn(renamed_and_removed_lints)]` on by default
```

Explanation

To fix this, either remove the lint or use the new name. This can help avoid confusion about lints that are no longer valid, and help maintain consistency for renamed lints.

repr-transparent-external-private-fields

The `repr_transparent_external_private_fields` lint detects types marked `#[repr(transparent)]` that (transitively) contain an external ZST type marked `#[non_exhaustive]` or containing private fields

Example

```
#![deny(repr_transparent_external_private_fields)]
use foo::NonExhaustiveZst;

#[repr(transparent)]
struct Bar(u32, ([u32; 0], NonExhaustiveZst));
```

This will produce:

```
error: zero-sized fields in repr(transparent) cannot contain external non-exhaustive types
--> src/main.rs:5:28
   |
5 | struct Bar(u32, ([u32; 0], NonExhaustiveZst));
   |          ^^^^^^^^^^^^^^^^^^
   |
note: the lint level is defined here
--> src/main.rs:1:9
   |
1 | #![deny(repr_transparent_external_private_fields)]
   |          ^^^^^^^^^^^^^^^^^^
   = warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
   = note: for more information, see issue #78586 <https://github.com/rust-lang/rust/issues/78586>
   = note: this struct contains `NonExhaustiveZst`, which is marked with `#[non_exhaustive]`, and makes it not a breaking change to become non-zero-sized in the future.
```

Explanation

Previously, Rust accepted fields that contain external private zero-sized types, even though it should not be a breaking change to add a non-zero-sized field to that private type.

This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #78586](#) for more details.

semicolon-in-expressions-from-macros

The `semicolon_in_expressions_from_macros` lint detects trailing semicolons in macro bodies when the macro is invoked in expression position. This was previously accepted, but is being phased out.

Example

```
#![deny(semicolon_in_expressions_from_macros)]
macro_rules! foo {
    () => { true; }
}

fn main() {
    let val = match true {
        true => false,
        _ => foo!()
    };
}
```

This will produce:

```
error: trailing semicolon in macro used in expression position
--> lint_example.rs:3:17
3 |     () => { true; }
   |          ^
...
9 |         _ => foo!()
   |             ----- in this macro invocation
   |
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #79813 <https://github.com/rust-lang/rust/issues/79813>
note: the lint level is defined here
--> lint_example.rs:1:9
  |
1 | #![deny(semicolon_in_expressions_from_macros)]
  | ^^^^^^^^^^
= note: this error originates in the macro `foo` (in Nightly builds, run with -Z macro-
backtrace for more info)
```

Explanation

Previously, Rust ignored trailing semicolon in a macro body when a macro was invoked in expression position. However, this makes the treatment of semicolons in the language inconsistent, and could lead to unexpected runtime behavior in some circumstances (e.g. if the macro author expects a value to be dropped).

This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #79813](#) for more details.

special-module-name

The `special_module_name` lint detects module declarations for files that have a special meaning.

Example

```
mod lib;

fn main() {
    lib::run();
}
```

This will produce:

```
warning: found module declaration for lib.rs
--> lint_example.rs:1:1
|
1 | mod lib;
| ^^^^^^^^
|
= note: lib.rs is the root of this crate's library target
= help: to refer to it from other targets, use the library's name as the path
= note: `#[warn(special_module_name)]` on by default
```

Explanation

Cargo recognizes `lib.rs` and `main.rs` as the root of a library or binary crate, so declaring them as modules will lead to miscompilation of the crate unless configured explicitly.

To access a library from a binary target within the same crate, use `your_crate_name::` as the path instead of `lib::`:

```
// bar/src/lib.rs
fn run() {
    // ...
}

// bar/src/main.rs
fn main() {
    bar::run();
}
```

Binary targets cannot be used as libraries and so declaring one as a module is not allowed.

stable-features

The `stable_features` lint detects a `feature` attribute that has since been made stable.

Example

```
#![feature(test_accepted_feature)]
fn main() {}
```

This will produce:

```
warning: the feature `test_accepted_feature` has been stable since 1.0.0 and no longer
requires an attribute to enable
--> lint_example.rs:1:12
  |
1 | #![feature(test_accepted_feature)]
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^
  |
= note: `#[warn(stable_features)]` on by default
```

Explanation

When a feature is stabilized, it is no longer necessary to include a `#![feature]` attribute for it. To fix, simply remove the `#![feature]` attribute.

suspicious-auto-trait-impls

The `suspicious_auto_trait_impls` lint checks for potentially incorrect implementations of auto traits.

Example

```
struct Foo<T>(T);

unsafe impl<T> Send for Foo<*const T> {}
```

This will produce:

```
warning: cross-crate traits with a default impl, like `Send`, should not be specialized
--> lint_example.rs:4:1
| unsafe impl<T> Send for Foo<*const T> {}
| ^^^^^^^^^^
|
= warning: this will change its meaning in a future release!
= note: for more information, see issue #93367 <https://github.com/rust-lang/rust/issues/93367>
= note: `*const T` is not a generic parameter
note: try using the same sequence of generic parameters as the struct definition
--> lint_example.rs:2:1
|
2 | struct Foo<T>(T);
| ^^^^^^
= note: `#[warn(suspicious_auto_trait_impls)]` on by default
```

Explanation

A type can implement auto traits, e.g. `Send`, `Sync` and `Unpin`, in two different ways: either by writing an explicit `impl` or if all fields of the type implement that auto trait.

The compiler disables the automatic implementation if an explicit one exists for given type constructor. The exact rules governing this are currently unsound, quite subtle, and will be modified in the future. This change will cause the automatic implementation to be disabled in more cases, potentially breaking some code.

temporary-cstring-as-ptr

The `temporary_cstring_as_ptr` lint detects getting the inner pointer of a temporary `CString`.

Example

```
let c_str = CString::new("foo").unwrap().as_ptr();
```

This will produce:

```
warning: getting the inner pointer of a temporary `CString`
--> lint_example.rs:4:42
4 | let c_str = CString::new("foo").unwrap().as_ptr();
   |                                     ----- ^^^^^^ this pointer will be invalid
   |                                     |
   |                                     this `CString` is deallocated at the end of the statement, bind it to a
variable to extend its lifetime
   |
   = note: pointers do not have a lifetime; when calling `as_ptr` the `CString` will be
deallocated at the end of the statement because nothing is referencing it as far as the type
system is concerned
   = help: for more information, see https://doc.rust-lang.org/reference/destructors.html
   = note: `#[warn(temporary_cstring_as_ptr)]` on by default
```

Explanation

The inner pointer of a `CString` lives only as long as the `cstring` it points to. Getting the inner pointer of a *temporary* `CString` allows the `CString` to be dropped at the end of the statement, as it is not being referenced as far as the typesystem is concerned. This means outside of the statement the pointer will point to freed memory, which causes undefined behavior if the pointer is later dereferenced.

trivial-bounds

The `trivial_bounds` lint detects trait bounds that don't depend on any type parameters.

Example

```
#![feature(trivial_bounds)]
pub struct A where i32: Copy;
```

This will produce:

```
warning: trait bound i32: Copy does not depend on any type or lifetime parameters
--> lint_example.rs:3:25
3 | pub struct A where i32: Copy;
   |           ^
   |
   = note: `#[warn(trivial_bounds)]` on by default
```

Explanation

Usually you would not write a trait bound that you know is always true, or never true. However, when using macros, the macro may not know whether or not the constraint would hold or not at the time when generating the code. Currently, the compiler does not alert you if the constraint is always true, and generates an error if it is never true. The `trivial_bounds` feature changes this to be a warning in both cases, giving macros more freedom and flexibility to generate code, while still providing a signal when writing non-macro code that something is amiss.

See [RFC 2056](#) for more details. This feature is currently only available on the nightly channel, see [tracking issue #48214](#).

type-alias-bounds

The `type_alias_bounds` lint detects bounds in type aliases.

Example

```
type SendVec<T: Send> = Vec<T>;
```

This will produce:

```
warning: bounds on generic parameters are not enforced in type aliases
--> lint_example.rs:2:17
  |
2 | type SendVec<T: Send> = Vec<T>;          ^^^^
  |
  = note: `#[warn(type_alias_bounds)]` on by default
help: the bound will not be checked when the type alias is used, and should be removed
  |
2 - type SendVec<T: Send> = Vec<T>;
2 + type SendVec<T> = Vec<T>;
  |
```

Explanation

The trait bounds in a type alias are currently ignored, and should not be included to avoid confusion. This was previously allowed unintentionally; this may become a hard error in the future.

tyvar-behind-raw-pointer

The `tyvar_behind_raw_pointer` lint detects raw pointer to an inference variable.

Example

```
// edition 2015
let data = std::ptr::null();
let _ = &data as *const *const ();

if data.is_null() {}
```

This will produce:

```
warning: type annotations needed
--> lint_example.rs:6:9
  |
6 | if data.is_null() {}
  | ^^^^^^^^
  |
= warning: this is accepted in the current edition (Rust 2015) but is a hard error in Rust
2018!
= note: for more information, see issue #46906 <https://github.com/rust-lang/rust/issu...>
= note: `#[warn(tyvar_behind_raw_pointer)]` on by default
```

Explanation

This kind of inference was previously allowed, but with the future arrival of [arbitrary self types](#), this can introduce ambiguity. To resolve this, use an explicit type instead of relying on type inference.

This is a [future-incompatible](#) lint to transition this to a hard error in the 2018 edition. See [issue #46906](#) for more details. This is currently a hard-error on the 2018 edition, and is "warn" by default in the 2015 edition.

uncommon-codepoints

The `uncommon_codepoints` lint detects uncommon Unicode codepoints in identifiers.

Example

```
const μ: f64 = 0.000001;
```

This will produce:

```
warning: identifier contains uncommon Unicode codepoints
--> lint_example.rs:3:7
3 | const μ: f64 = 0.000001;
  |
  = note: `#[warn(uncommon_codepoints)]` on by default
```

Explanation

This lint warns about using characters which are not commonly used, and may cause visual confusion.

This lint is triggered by identifiers that contain a codepoint that is not part of the set of "Allowed" codepoints as described by [Unicode® Technical Standard #39 Unicode Security Mechanisms Section 3.1 General Security Profile for Identifiers](#).

Note that the set of uncommon codepoints may change over time. Beware that if you "forbid" this lint that existing code may fail in the future.

unconditional-recursion

The `unconditional_recursion` lint detects functions that cannot return without calling themselves.

Example

```
fn foo() {
    foo();
}
```

This will produce:

```
warning: function cannot return without recursing
--> lint_example.rs:2:1
2 | fn foo() {
  | ^^^^^^^^^ cannot return without recursing
3 |     foo();
  |     ----- recursive call site
  |
  = help: a `loop` may express intention better if this is on purpose
  = note: `#[warn(unconditional_recursion)]` on by default
```

Explanation

It is usually a mistake to have a recursive call that does not have some condition to cause it to terminate. If you really intend to have an infinite loop, using a `loop` expression is recommended.

undefined-naked-function-abi

The `undefined_naked_function_abi` lint detects naked function definitions that either do not specify an ABI or specify the Rust ABI.

Example

```
#![feature(asm_experimental_arch, naked_functions)]  
  
use std::arch::asm;  
  
#[naked]  
pub fn default_abi() -> u32 {  
    unsafe { asm!("", options(noreturn)); }  
}  
  
#[naked]  
pub extern "Rust" fn rust_abi() -> u32 {  
    unsafe { asm!("", options(noreturn)); }  
}
```

This will produce:

```
warning: Rust ABI is unsupported in naked functions  
--> lint_example.rs:7:1  
|  
7 | pub fn default_abi() -> u32 {  
| ^^^^^^^^^^^^^^^^^^  
|= note: `#[warn(undefined_naked_function_abi)]` on by default  
  
warning: Rust ABI is unsupported in naked functions  
--> lint_example.rs:12:1  
|  
12 | pub extern "Rust" fn rust_abi() -> u32 {  
| ^^^^^^^^^^
```

Explanation

The Rust ABI is currently undefined. Therefore, naked functions should specify a non-Rust ABI.

unexpected-cfgs

The `unexpected_cfgs` lint detects unexpected conditional compilation conditions.

Example

```
rustc --check-cfg 'names()'
```

```
#[cfg(widnows)]
fn foo() {}
```

This will produce:

```
warning: unknown condition name used
--> lint_example.rs:1:7
  |
1 | #[cfg(widnows)]
  | ^^^^^^
  |
= note: `#[warn(unexpected_cfgs)]` on by default
```

Explanation

This lint is only active when a `--check-cfg='names(...)' option` has been passed to the compiler and triggers whenever an unknown condition name or value is used. The known condition include names or values passed in `--check-cfg`, `--cfg`, and some well-known names and values built into the compiler.

unfulfilled-lint-expectations

The `unfulfilled_lint_expectations` lint detects lint trigger expectations that have not been fulfilled.

Example

```
#![feature(lint_reasons)]
#[expect(unused_variables)]
let x = 10;
println!("{}", x);
```

This will produce:

```
warning: this lint expectation is unfulfilled
--> lint_example.rs:4:10
   |
4 | #[expect(unused_variables)]
   | ^^^^^^^^^^^^^^^^^^
   |
= note: `#[warn(unfulfilled_lint_expectations)]` on by default
```

Explanation

It was expected that the marked code would emit a lint. This expectation has not been fulfilled.

The `expect` attribute can be removed if this is intended behavior otherwise it should be investigated why the expected lint is no longer issued.

In rare cases, the expectation might be emitted at a different location than shown in the shown code snippet. In most cases, the `#[expect]` attribute works when added to the outer scope. A few lints can only be expected on a crate level.

Part of RFC 2383. The progress is being tracked in [#54503](#)

ungated-async-fn-track-caller

The `ungated_async_fn_track_caller` lint warns when the `#[track_caller]` attribute is used on an async function, method, or closure, without enabling the corresponding unstable feature flag.

Example

```
#[track_caller]
async fn foo() {}
```

This will produce:

```
warning: `#[track_caller]` on async functions is a no-op
--> lint_example.rs:2:1
2 | #[track_caller]
  ^^^^^^^^^^^^^^
3 | async fn foo() {}  
----- this function will not propagate the caller location
|  
= note: see issue #87417 <https://github.com/rust-lang/rust/issues/87417> for more
information
= help: add `#![feature(closure_track_caller)]` to the crate attributes to enable
= note: `#[warn(ungated_async_fn_track_caller)]` on by default
```

Explanation

The attribute must be used in conjunction with the [closure_track_caller](#) feature flag. Otherwise, the `#[track_caller]` annotation will function as a no-op.

uninhabited-static

The `uninhabited_static` lint detects uninhabited statics.

Example

```
enum Void {}
extern {
    static EXTERN: Void;
}
```

This will produce:

```
warning: static of uninhabited type
--> lint_example.rs:4:5
4 |     static EXTERN: Void;
  ^^^^^^^^^^^^^^^^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #74840 <https://github.com/rust-lang/rust/issues/74840>
= note: uninhabited statics cannot be initialized, and any access would be an immediate
error
= note: `#[warn(uninhabited_static)]` on by default
```

Explanation

Statics with an uninhabited type can never be initialized, so they are impossible to define. However, this can be side-stepped with an `extern static`, leading to problems later in the compiler which assumes that there are no initialized uninhabited places (such as locals or statics). This was accidentally allowed, but is being phased out.

unknown-lints

The `unknown_lints` lint detects unrecognized lint attributes.

Example

```
#![allow(not_a_real_lint)]
```

This will produce:

```
warning: unknown lint: `not_a_real_lint`
--> lint_example.rs:1:10
  |
1 | #![allow(not_a_real_lint)]
  | ^^^^^^^^^^^^^^^^^^
  |
= note: `#[warn(unknown_lints)]` on by default
```

Explanation

It is usually a mistake to specify a lint that does not exist. Check the spelling, and check the lint listing for the correct name. Also consider if you are using an old version of the compiler, and the lint is only available in a newer version.

unameable-test-items

The `unameable_test_items` lint detects `#[test]` functions that are not able to be run by the test harness because they are in a position where they are not nameable.

Example

```
fn main() {
    #[test]
    fn foo() {
        // This test will not fail because it does not run.
        assert_eq!(1, 2);
    }
}
```

This will produce:

```
warning: cannot test inner items
--> lint_example.rs:2:5
2 |     #[test]
|     ^^^^^^
|
= note: `#[warn(unnameable_test_items)]` on by default
= note: this warning originates in the attribute macro `test` (in Nightly builds, run with -Z macro-backtrace for more info)
```

Explanation

In order for the test harness to run a test, the test function must be located in a position where it can be accessed from the crate root. This generally means it must be defined in a module, and not anywhere else such as inside another function. The compiler previously allowed this without an error, so a lint was added as an alert that a test is not being used. Whether or not this should be allowed has not yet been decided, see [RFC 2471](#) and [issue #36629](#).

unreachable-code

The `unreachable_code` lint detects unreachable code paths.

Example

```
panic!("we never go past here!");

let x = 5;
```

This will produce:

```
warning: unreachable statement
--> lint_example.rs:4:1
2 |     panic!("we never go past here!");
|----- any code following this expression is unreachable
3 |
4 | let x = 5;
| ^^^^^^^^^^^ unreachable statement
|
= note: `#[warn(unreachable_code)]` on by default
```

Explanation

Unreachable code may signal a mistake or unfinished code. If the code is no longer in use, consider removing it.

unreachable-patterns

The `unreachable_patterns` lint detects unreachable patterns.

Example

```
let x = 5;
match x {
    y => (),
    5 => (),
}
```

This will produce:

```
warning: unreachable pattern
--> lint_example.rs:5:5
4 |     y => (),
|     - matches any value
5 |     5 => (),
|     ^ unreachable pattern
|
= note: `#[warn(unreachable_patterns)]` on by default
```

Explanation

This usually indicates a mistake in how the patterns are specified or ordered. In this example, the `y` pattern will always match, so the five is impossible to reach. Remember, match arms match in order, you probably wanted to put the `5` case above the `y` case.

unstable-name-collisions

The `unstable_name_collisions` lint detects that you have used a name that the standard library plans to add in the future.

Example

```
trait MyIterator : Iterator {
    // is_sorted is an unstable method that already exists on the Iterator trait
    fn is_sorted(self) -> bool where Self: Sized {true}
}

impl<T: ?Sized> MyIterator for T where T: Iterator { }

let x = vec![1, 2, 3];
let _ = x.iter().is_sorted();
```

This will produce:

```
warning: a method with this name may be added to the standard library in the future
--> lint_example.rs:10:18
10 |     let _ = x.iter().is_sorted();
   |     ^^^^^^^^^^
   |
   = warning: once this associated item is added to the standard library, the ambiguity may
cause an error or change in behavior!
  = note: for more information, see issue #48919 <https://github.com/rust-lang/rust-lang/issues/48919>
  = help: call with fully qualified syntax `MyIterator::is_sorted(...)` to keep using the
current method
  = help: add `#![feature(is_sorted)]` to the crate attributes to enable `is_sorted`
  = note: `#[warn(unstable_name_collisions)]` on by default
```

Explanation

When new methods are added to traits in the standard library, they are usually added in an "unstable" form which is only available on the [nightly channel](#) with a [feature attribute](#). If there is any pre-existing code which extends a trait to have a method with the same name, then the names will collide. In the future, when the method is stabilized, this will cause an error due to the ambiguity. This lint is an early-warning to let you know that there may be a collision in the future. This can be avoided by adding type annotations to disambiguate which trait method you intend to call, such as `MyIterator::is_sorted(my_iter)` or renaming or removing the method.

unstable-syntax-pre-expansion

The `unstable_syntax_pre_expansion` lint detects the use of unstable syntax that is discarded during attribute expansion.

Example

```
#[cfg(FALSE)]
macro foo() {}
```

This will produce:

```
warning: `macro` is experimental
--> lint_example.rs:3:1
  |
3 | macro foo() {}
  | ^^^^^^^^^^^^^^
  |
= note: see issue #39412 <https://github.com/rust-lang/rust/issues/39412> for more
information
= help: add `#![feature(decl_macro)]` to the crate attributes to enable
= warning: unstable syntax can change at any point in the future, causing a hard error!
= note: for more information, see issue #65860 <https://github.com/rust-
lang/rust/issu
```

Explanation

The input to active attributes such as `#[cfg]` or procedural macro attributes is required to be valid syntax. Previously, the compiler only gated the use of unstable syntax features after resolving `#[cfg]` gates and expanding procedural macros.

To avoid relying on unstable syntax, move the use of unstable syntax into a position where the compiler does not parse the syntax, such as a functionlike macro.

```
macro_rules! identity {
    ( $($tokens:tt)* ) => { $($tokens)* }
}

#[cfg(FALSE)]
identity! {
    macro foo() {}
}
```

This is a `future-incompatible` lint to transition this to a hard error in the future. See [issue #65860](#) for more details.

unsupported-calling-conventions

The `unsupported_calling_conventions` lint is output whenever there is a use of the `stdcall`, `fastcall`, `thiscall`, `vectorcall` calling conventions (or their unwind variants) on targets that cannot meaningfully be supported for the requested target.

For example `stdcall` does not make much sense for a `x86_64` or, more apparently, `powerpc` code, because this calling convention was never specified for those targets.

Historically MSVC toolchains have fallen back to the regular C calling convention for targets other than `x86`, but Rust doesn't really see a similar need to introduce a similar hack across many more targets.

Example

```
extern "stdcall" fn stdcall() {}
```

This will produce:

```
warning: use of calling convention not supported on this target
--> $DIR/unsupported.rs:39:1
   |
LL | extern "stdcall" fn stdcall() {}
   | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
   |
= note: `#[warn(unsupported_calling_conventions)]` on by default
= warning: this was previously accepted by the compiler but is being phased out;
          it will become a hard error in a future release!
= note: for more information, see issue ...
```

Explanation

On most of the targets the behaviour of `stdcall` and similar calling conventions is not defined at all, but was previously accepted due to a bug in the implementation of the compiler.

unused-allocation

The `unused_allocation` lint detects unnecessary allocations that can be eliminated.

Example

```
fn main() {
    let a = Box::new([1, 2, 3]).len();
}
```

This will produce:

```
warning: unnecessary allocation, use `&` instead
--> lint_example.rs:2:13
2 |     let a = Box::new([1, 2, 3]).len();
   |           ^^^^^^^^^^
= note: `#[warn(unused_allocation)]` on by default
```

Explanation

When a `Box` expression is immediately coerced to a reference, then the allocation is unnecessary, and a reference (using `&` or `&mut`) should be used instead to avoid the allocation.

unused-assignments

The `unused_assignments` lint detects assignments that will never be read.

Example

```
let mut x = 5;
x = 6;
```

This will produce:

```
warning: value assigned to `x` is never read
--> lint_example.rs:3:1
3 | x = 6;
  |
= help: maybe it is overwritten before being read?
= note: `#[warn(unused_assignments)]` on by default
```

Explanation

Unused assignments may signal a mistake or unfinished code. If the variable is never used after being assigned, then the assignment can be removed. Variables with an underscore prefix such as `_x` will not trigger this lint.

unused-attributes

The `unused_attributes` lint detects attributes that were not used by the compiler.

Example

```
#![ignore]
```

This will produce:

```
warning: `#[ignore]` only has an effect on functions
--> lint_example.rs:1:1
  |
1 | #[ignore]
  ^^^^^^^^^^
= note: `#[warn(unused_attributes)]` on by default
```

Explanation

Unused [attributes](#) may indicate the attribute is placed in the wrong position. Consider removing it, or placing it in the correct position. Also consider if you intended to use an *inner attribute* (with a `!` such as `#![allow(unused)]`) which applies to the item the attribute is within, or an *outer attribute* (without a `!` such as `#[allow(unused)]`) which applies to the item *following* the attribute.

unused-braces

The `unused_braces` lint detects unnecessary braces around an expression.

Example

```
if { true } {
    // ...
}
```

This will produce:

```
warning: unnecessary braces around `if` condition
--> lint_example.rs:2:4
2 | if { true } {
|   ^ ^
|
|= note: `#[warn(unused_braces)]` on by default
help: remove these braces
|
2 - if { true } {
2 + if true {
|
```

Explanation

The braces are not needed, and should be removed. This is the preferred style for writing these expressions.

unused-comparisons

The `unused_comparisons` lint detects comparisons made useless by limits of the types involved.

Example

```
fn foo(x: u8) {
    x >= 0;
}
```

This will produce:

```
warning: comparison is useless due to type limits
--> lint_example.rs:3:5
3 |     x >= 0;
|     ^^^^^^
|
|= note: `#[warn(unused_comparisons)]` on by default
```

Explanation

A useless comparison may indicate a mistake, and should be fixed or removed.

unused-doc-comments

The `unused_doc_comments` lint detects doc comments that aren't used by `rustdoc`.

Example

```
/// docs for x
let x = 12;
```

This will produce:

```
warning: unused doc comment
--> lint_example.rs:2:1
   |
2 | /// docs for x
   | ^^^^^^^^^^^^^^
3 | let x = 12;
   | ----- rustdoc does not generate documentation for statements
   |
= help: use `//` for a plain comment
= note: `#[warn(unused_doc_comments)]` on by default
```

Explanation

`rustdoc` does not use doc comments in all positions, and so the doc comment will be ignored. Try changing it to a normal comment with `//` to avoid the warning.

unused-features

The `unused_features` lint detects unused or unknown features found in crate-level [feature attributes](#).

Note: This lint is currently not functional, see [issue #44232](#) for more details.

unused-imports

The `unused_imports` lint detects imports that are never used.

Example

```
use std::collections::HashMap;
```

This will produce:

```
warning: unused import: `std::collections::HashMap`  
--> lint_example.rs:2:5  
2 | use std::collections::HashMap;  
  | ^^^^^^^^^^^^^^^^^^^^^^  
= note: `#[warn(unused_imports)]` on by default
```

Explanation

Unused imports may signal a mistake or unfinished code, and clutter the code, and should be removed. If you intended to re-export the item to make it available outside of the module, add a visibility modifier like `pub`.

unused-labels

The `unused_labels` lint detects `labels` that are never used.

Example

```
'unused_label: loop {}
```

This will produce:

```
warning: unused label  
--> lint_example.rs:2:1  
2 | 'unused_label: loop {}  
  | ^^^^^^^^^^  
= note: `#[warn(unused_labels)]` on by default
```

Explanation

Unused labels may signal a mistake or unfinished code. To silence the warning for the individual label, prefix it with an underscore such as `'_my_label:'`.

unused-macros

The `unused_macros` lint detects macros that were not used.

Note that this lint is distinct from the `unused_macro_rules` lint, which checks for single rules that never match of an otherwise used macro, and thus never expand.

Example

```
macro_rules! unused {
    () => {};
}

fn main() {
}
```

This will produce:

```
warning: unused macro definition: `unused`
--> lint_example.rs:1:14
  |
1 | macro_rules! unused {
  |         ^^^^^^
  |
= note: `#[warn(unused_macros)]` on by default
```

Explanation

Unused macros may signal a mistake or unfinished code. To silence the warning for the individual macro, prefix the name with an underscore such as `_my_macro`. If you intended to export the macro to make it available outside of the crate, use the `macro_export` attribute.

unused-must-use

The `unused_must_use` lint detects unused result of a type flagged as `#[must_use]`.

Example

```
fn returns_result() -> Result<(), ()> {
    Ok(())
}

fn main() {
    returns_result();
}
```

This will produce:

```
warning: unused `Result` that must be used
--> lint_example.rs:6:5
6 |     returns_result();
  | ^^^^^^^^^^^^^^^^^^
|
= note: this `Result` may be an `Err` variant, which should be handled
= note: `#[warn(unused_must_use)]` on by default
help: use `let _ = ...` to ignore the resulting value
6 |     let _ = returns_result();
  | ++++++
```

Explanation

The `#[must_use]` attribute is an indicator that it is a mistake to ignore the value. See [the reference](#) for more details.

unused-mut

The `unused_mut` lint detects `mut` variables which don't need to be mutable.

Example

```
let mut x = 5;
```

This will produce:

```
warning: variable does not need to be mutable
--> lint_example.rs:2:5
2 | let mut x = 5;
  | ^----^
  |
  | help: remove this `mut`
  |
= note: `#[warn(unused_mut)]` on by default
```

Explanation

The preferred style is to only mark variables as `mut` if it is required.

unused-parens

The `unused_parens` lint detects `if`, `match`, `while` and `return` with parentheses; they do not need them.

Examples

```
if(true) {}
```

This will produce:

```
warning: unnecessary parentheses around `if` condition
--> lint_example.rs:2:3
2 | if(true) {}
  | ^   ^
  |
= note: `#[warn(unused_parens)]` on by default
help: remove these parentheses
|
2 - if(true) {}
2 + if true {}
```

Explanation

The parentheses are not needed, and should be removed. This is the preferred style for writing these expressions.

unused-unsafe

The `unused_unsafe` lint detects unnecessary use of an `unsafe` block.

Example

```
unsafe {}
```

This will produce:

```
warning: unnecessary `unsafe` block
--> lint_example.rs:2:1
2 | unsafe {}
| ^^^^^^ unnecessary `unsafe` block
|
= note: `#[warn(unused_unsafe)]` on by default
```

Explanation

If nothing within the block requires `unsafe`, then remove the `unsafe` marker because it is not required and may cause confusion.

unused-variables

The `unused_variables` lint detects variables which are not used in any way.

Example

```
let x = 5;
```

This will produce:

```
warning: unused variable: `x`
--> lint_example.rs:2:5
2 | let x = 5;
|   ^ help: if this is intentional, prefix it with an underscore: `_x`
|
= note: `#[warn(unused_variables)]` on by default
```

Explanation

Unused variables may signal a mistake or unfinished code. To silence the warning for the individual variable, prefix it with an underscore such as `_x`.

warnings

The `warnings` lint allows you to change the level of other lints which produce warnings.

Example

```
#![deny(warnings)]
fn foo() {}
```

This will produce:

```
error: function `foo` is never used
--> lint_example.rs:3:4
  |
3 | fn foo() {}
  ^^^

note: the lint level is defined here
--> lint_example.rs:1:9
  |
1 | #![deny(warnings)]
  ^^^^^^^^^^
= note: `#[deny(dead_code)]` implied by `#[deny(warnings)]`
```

Explanation

The `warnings` lint is a bit special; by changing its level, you change every other warning that would produce a warning to whatever value you'd like. As such, you won't ever trigger this lint in your code directly.

where-clauses-object-safety

The `where_clauses_object_safety` lint detects for `object safety` of `where clauses`.

Example

```
trait Trait {}

trait X { fn foo(&self) where Self: Trait; }

impl X for () { fn foo(&self) {} }

impl Trait for dyn X {}

// Segfault at opt-level 0, SIGILL otherwise.
pub fn main() { <dyn X as X>::foo(&()); }
```

This will produce:

```
warning: the trait `X` cannot be made into an object
--> lint_example.rs:3:14
| trait X { fn foo(&self) where Self: Trait; }
|     ^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #51443 <https://github.com/rust-lang/rust/issues/51443>
note: for a trait to be "object safe" it needs to allow building a vtable to allow the call to
be resolvable dynamically; for more information visit <https://doc.rust-lang.org/reference/items/traits.html#object-safety>
--> lint_example.rs:3:14
|
3 | trait X { fn foo(&self) where Self: Trait; }
|     -          ^^^ ...because method `foo` references the `Self` type in its `where` clause
|     |
|     this trait cannot be made into an object...
= help: consider moving `foo` to another trait
= note: `#[warn(where_clauses_object_safety)]` on by default
```

Explanation

The compiler previously allowed these object-unsafe bounds, which was incorrect. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #51443](#) for more details.

while-true

The `while_true` lint detects `while true { }`.

Example

```
while true {  
}
```

This will produce:

```
warning: denote infinite loops with `loop { ... }`
--> lint_example.rs:2:1
| while true {  
| ^^^^^^^^^^ help: use `loop`  
|
= note: `#[warn(while_true)]` on by default
```

Explanation

`while true` should be replaced with `loop`. A `loop` expression is the preferred way to write an infinite loop because it more directly expresses the intent of the loop.

Deny-by-default Lints

These lints are all set to the 'deny' level by default.

- `ambiguous_associated_items`
- `arithmetic_overflow`
- `bindings_with_variant_name`
- `cenum_impl_drop_cast`
- `conflicting_repr_hints`
- `deprecated_cfg_attr_crate_type_name`
- `enum_intrinsics_non_enums`
- `ill-formed_attribute_input`
- `implied_bounds_entailment`
- `incomplete_include`
- `ineffective_unstable_trait_impl`
- `invalid_alignment`
- `invalid_atomic_ordering`
- `invalid_type_param_default`
- `let_underscore_lock`
- `macro_expanded_macro_exports_accessed_by_absolute_paths`
- `missing_fragment_specifier`
- `mutable_transmutes`
- `named_asm_labels`
- `no_mangle_const_items`
- `order_dependent_trait_objects`
- `overflowing_literals`
- `patterns_in_fns_without_body`
- `proc_macro_back_compat`
- `proc_macro_derive_resolution_fallback`
- `pub_use_of_private_extern_crate`
- `soft_unstable`
- `test_unstable_lint`
- `text_direction_codepoint_in_comment`
- `text_direction_codepoint_in_literal`
- `unconditional_panic`
- `unknown_crate_types`
- `useless_DEPRECATED`

ambiguous-associated-items

The `ambiguous_associated_items` lint detects ambiguity between `associated items` and `enum variants`.

Example

```
enum E {
    V
}

trait Tr {
    type V;
    fn foo() -> Self::V;
}

impl Tr for E {
    type V = u8;
    // `Self::V` is ambiguous because it may refer to the associated type or
    // the enum variant.
    fn foo() -> Self::V { 0 }
}
```

This will produce:

```
error: ambiguous associated item
--> lint_example.rs:15:17
15 |     fn foo() -> Self::V { 0 }
|           ^^^^^^^ help: use fully-qualified syntax: `<E as Tr>::V`
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #57644 <https://github.com/rust-lang/rust-lang/issues/57644>
note: `V` could refer to the variant defined here
--> lint_example.rs:3:5
3 |     V
|   ^
note: `V` could also refer to the associated type defined here
--> lint_example.rs:7:5
7 |     type V;
|   ^^^^^^
= note: `#[deny(ambiguous_associated_items)]` on by default
```

Explanation

Previous versions of Rust did not allow accessing enum variants through [type aliases](#). When this ability was added (see [RFC 2338](#)), this introduced some situations where it can be ambiguous what a type was referring to.

To fix this ambiguity, you should use a [qualified path](#) to explicitly state which type to use. For example, in the above example the function can be written as `fn f() -> <Self as Tr>::V { 0 }` to specifically refer to the associated type.

This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #57644](#) for more details.

arithmetic-overflow

The `arithmetic_overflow` lint detects that an arithmetic operation will **overflow**.

Example

```
1_i32 << 32;
```

This will produce:

```
error: this arithmetic operation will overflow
--> lint_example.rs:2:1
  |
2 | 1_i32 << 32;
  | ^^^^^^^^^^^ attempt to shift left by `32_i32`, which would overflow
  |
= note: `#[deny(arithmetic_overflow)]` on by default
```

Explanation

It is very likely a mistake to perform an arithmetic operation that overflows its value. If the compiler is able to detect these kinds of overflows at compile-time, it will trigger this lint. Consider adjusting the expression to avoid overflow, or use a data type that will not overflow.

bindings-with-variant-name

The `bindings_with_variant_name` lint detects pattern bindings with the same name as one of the matched variants.

Example

```
pub enum Enum {
    Foo,
    Bar,
}

pub fn foo(x: Enum) {
    match x {
        Foo => {}
        Bar => {}
    }
}
```

This will produce:

```
error[E0170]: pattern binding `Foo` is named the same as one of the variants of the type
`main::Enum`
--> lint_example.rs:9:9
9 |     Foo => {}
|     ^^^ help: to match on the variant, qualify the path: `main::Enum::Foo`
|
= note: `#[deny(bindings_with_variant_name)]` on by default
```

Explanation

It is usually a mistake to specify an enum variant name as an [identifier pattern](#). In the example above, the `match` arms are specifying a variable name to bind the value of `x` to. The second arm is ignored because the first one matches all values. The likely intent is that the arm was intended to match on the enum variant.

Two possible solutions are:

- Specify the enum variant using a [path pattern](#), such as `Enum::Foo`.
- Bring the enum variants into local scope, such as adding `use Enum::*;` to the beginning of the `foo` function in the example above.

cenum-impl-drop-cast

The `cenum_impl_drop_cast` lint detects an `as` cast of a field-less `enum` that implements `Drop`.

Example

```
enum E {
    A,
}

impl Drop for E {
    fn drop(&mut self) {
        println!("Drop");
    }
}

fn main() {
    let e = E::A;
    let i = e as u32;
}
```

This will produce:

```

error: cannot cast enum `E` into integer `u32` because it implements `Drop`
--> lint_example.rs:14:13
14 |     let i = e as u32;
      ^^^^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #73333 <https://github.com/rust-lang/rust/issues/73333>
= note: `#[deny(cenum_impl_drop_cast)]` on by default

```

Explanation

Casting a field-less `enum` that does not implement `Copy` to an integer moves the value without calling `drop`. This can result in surprising behavior if it was expected that `drop` should be called. Calling `drop` automatically would be inconsistent with other move operations. Since neither behavior is clear or consistent, it was decided that a cast of this nature will no longer be allowed.

This is a `future-incompatible` lint to transition this to a hard error in the future. See [issue #73333](#) for more details.

conflicting-repr-hints

The `conflicting_repr_hints` lint detects `repr` attributes with conflicting hints.

Example

```

#[repr(u32, u64)]
enum Foo {
    Variant1,
}

```

This will produce:

```

error[E0566]: conflicting representation hints
--> lint_example.rs:2:8
2 | #[repr(u32, u64)]
  | ^^^ ^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #68585 <https://github.com/rust-lang/rust/issues/68585>
= note: `#[deny(conflicting_repr_hints)]` on by default

```

Explanation

The compiler incorrectly accepted these conflicting representations in the past. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #68585](#) for more details.

To correct the issue, remove one of the conflicting hints.

deprecated-cfg-attr-crate-type-name

The `deprecated_cfg_attr_crate_type_name` lint detects uses of the `#[cfg_attr(..., crate_type = "...")]` and `#[cfg_attr(..., crate_name = "...")]` attributes to conditionally specify the crate type and name in the source code.

Example

```
#[cfg_attr(debug_assertions, crate_type = "lib")]
```

This will produce:

```
error: `crate_type` within an `#[cfg_attr]` attribute is deprecated
--> lint_example.rs:1:31
  |
1 | #[cfg_attr(debug_assertions, crate_type = "lib")]
  | ^^^^^^^^^^^^^^^^^^
  |
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #91632 <https://github.com/rust-lang/rust/issues/91632>
= note: `#[deny(deprecated_cfg_attr_crate_type_name)]` on by default
```

Explanation

The `#[crate_type]` and `#[crate_name]` attributes require a hack in the compiler to be able to change the used crate type and crate name after macros have been expanded. Neither attribute works in combination with Cargo as it explicitly passes `--crate-type` and `--crate-name` on the commandline. These values must match the value used in the source code to prevent an error.

To fix the warning use `--crate-type` on the commandline when running `rustc` instead of `#[cfg_attr(..., crate_type = "...")]` and `--crate-name` instead of `#[cfg_attr(..., crate_name = "...")]`.

enum-intrinsics-nonEnums

The `enum_intrinsics_nonEnums` lint detects calls to intrinsic functions that require an enum (`core::mem::discriminant`, `core::mem::variant_count`), but are called with a non-enum type.

Example

```
#![deny(enum_intrinsics_nonEnums)]
core::mem::discriminant::<i32>(&123);
```

This will produce:

```
error: the return value of `mem::discriminant` is unspecified when called with a non-enum type
--> lint_example.rs:3:1
  |
3 | core::mem::discriminant::<i32>(&123);
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  |
note: the argument to `discriminant` should be a reference to an enum, but it was passed a
reference to a `i32`, which is not an enum.
--> lint_example.rs:3:32
  |
3 | core::mem::discriminant::<i32>(&123);
  |      ^^^^
  note: the lint level is defined here
--> lint_example.rs:1:9
  |
1 | #![deny(enum_intrinsics_nonEnums)]
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^
```

Explanation

In order to accept any enum, the `mem::discriminant` and `mem::variant_count` functions are generic over a type `T`. This makes it technically possible for `T` to be a non-enum, in which case the return value is unspecified.

This lint prevents such incorrect usage of these functions.

ill-formed-attribute-input

The `ill_formed_attribute_input` lint detects ill-formed attribute inputs that were previously accepted and used in practice.

Example

```
#[inline = "this is not valid"]
fn foo() {}
```

This will produce:

```
error: attribute must be of the form `#[inline]` or `#[inline(always|never)]`
--> lint_example.rs:2:1
2 | #[inline = "this is not valid"]
   ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
   |
   = warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
   = note: for more information, see issue #57571 <https://github.com/rust-lang/rust/issues/57571>
   = note: `#[deny(ill-formed_attribute_input)]` on by default
```

Explanation

Previously, inputs for many built-in attributes weren't validated and nonsensical attribute inputs were accepted. After validation was added, it was determined that some existing projects made use of these invalid forms. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #57571](#) for more details.

Check the [attribute reference](#) for details on the valid inputs for attributes.

implied-bounds-entailment

The `implied_bounds_entailment` lint detects cases where the arguments of an `impl` method have stronger implied bounds than those from the trait method it's implementing.

Example

```
#![deny(implied_bounds_entailment)]

trait Trait {
    fn get<'s>(s: &'s str, _: &'static &'static ()) -> &'static str;
}

impl Trait for () {
    fn get<'s>(s: &'s str, _: &'static &'s ()) -> &'static str {
        s
    }
}

let val = <() as Trait>::get(&String::from("blah blah blah"), &&());
println!("{}", val);
```

This will produce:

```
error: impl method assumes more implied bounds than the corresponding trait method
--> lint_example.rs:9:31
9 |     fn get<'s>(s: &'s str, _: &'static &'s ()) -> &'static str {
|           ^^^^^^^^^^^^^^^^^^ help: replace this type to make the impl
signature compatible: `&'static &'static ()`
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #105572 <https://github.com/rust-lang/rust-lang/issues/105572>
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #![deny(implied_bounds_entailment)]
| ^^^^^^^^^^^^^^^^^^
```

Explanation

Neither the trait method, which provides no implied bounds about `'s`, nor the impl, requires the main function to prove that `'s: 'static`, but the impl method is allowed to assume that `'s: 'static` within its own body.

This can be used to implement an unsound API if used incorrectly.

incomplete-include

The `incomplete_include` lint detects the use of the `include!` macro with a file that contains more than one expression.

Example

```
fn main() {
    include!("foo.txt");
}
```

where the file `foo.txt` contains:

```
println!("hi!");
```

produces:

```
error: include macro expected single expression in source
--> foo.txt:1:14
1 | println!("1");
   |
   = note: `#[deny(incomplete_include)]` on by default
```

Explanation

The `include!` macro is currently only intended to be used to include a single [expression](#) or multiple [items](#). Historically it would ignore any contents after the first expression, but that can be confusing. In the example above, the `println!` expression ends just before the semicolon, making the semicolon "extra" information that is ignored. Perhaps even more surprising, if the included file had multiple print statements, the subsequent ones would be ignored!

One workaround is to place the contents in braces to create a [block expression](#). Also consider alternatives, like using functions to encapsulate the expressions, or use [proc-macros](#).

This is a lint instead of a hard error because existing projects were found to hit this error. To be cautious, it is a lint for now. The future semantics of the `include!` macro are also uncertain, see [issue #35560](#).

ineffective-unstable-trait-impl

The `ineffective_unstable_trait_impl` lint detects `#[unstable]` attributes which are not used.

Example

```
#![feature(staged_api)]

#[derive(Clone)]
#[stable(feature = "x", since = "1")]
struct S {}

#[unstable(feature = "y", issue = "none")]
impl Copy for S {}
```

This will produce:

```
error: an `#[unstable]` annotation here has no effect
--> lint_example.rs:8:1
   |
8 | #[unstable(feature = "y", issue = "none")]
   | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
   |
   = note: see issue #55436 <https://github.com/rust-lang/rust/issues/55436> for more
information
   = note: `#[deny(ineffective_unstable_trait_impl)]` on by default
```

Explanation

`staged_api` does not currently support using a stability attribute on `impl` blocks. `impls` are always stable if both the type and trait are stable, and always unstable otherwise.

invalid-alignment

The `invalid_alignment` lint detects dereferences of misaligned pointers during constant evaluation.

Example

```
#![feature(const_mut_refs)]
const FOO: () = unsafe {
    let x = &[0_u8; 4];
    let y = x.as_ptr().cast::<u32>();
    let mut z = 123;
    y.copy_to_nonoverlapping(&mut z, 1); // the address of a `u8` array is unknown
                                         // and thus we don't know if it is aligned enough for copying a `u32`.
};
```

This will produce:

```

error: accessing memory with alignment 1, but alignment 4 is required
--> /checkout/library/core/src/intrinsics.rs:2637:9
2637 |     copy_nonoverlapping(src, dst, count)
|     ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #68585 <https://github.com/rust-lang/rust/issues/104616>
note: inside `copy_nonoverlapping::<u32>`
--> /checkout/library/core/src/intrinsics.rs:2637:9
2637 |     copy_nonoverlapping(src, dst, count)
|     ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
note: inside `ptr::const_ptr::<impl *const u32>::copy_to_nonoverlapping`
--> /checkout/library/core/src/ptr/const_ptr.rs:1279:18
1279 |     unsafe { copy_nonoverlapping(self, dest, count) }
|     ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
note: inside `F00`
--> lint_example.rs:7:5
7 |     y.copy_to_nonoverlapping(&mut z, 1); // the address of a `u8` array is unknown
|     ^^^^^^^^^^^^^^^^^^
= note: `#[deny(invalid_alignment)]` on by default

```

Explanation

The compiler allowed dereferencing raw pointers irrespective of alignment during const eval due to the const evaluator at the time not making it easy or cheap to check. Now that it is both, this is not accepted anymore.

Since it was undefined behaviour to begin with, this breakage does not violate Rust's stability guarantees. Using undefined behaviour can cause arbitrary behaviour, including failure to build.

invalid-atomic-ordering

The `invalid_atomic_ordering` lint detects passing an `Ordering` to an atomic operation that does not support that ordering.

Example

```

let atom = AtomicU8::new(0);
let value = atom.load(Ordering::Release);

```

This will produce:

```
error: atomic loads cannot have `Release` or `AcqRel` ordering
--> lint_example.rs:4:23
4 | let value = atom.load(Ordering::Release);
   |          ^^^^^^
= help: consider using ordering modes `Acquire`, `SeqCst` or `Relaxed`
= note: `#[deny(invalid_atomic_ordering)]` on by default
```

Explanation

Some atomic operations are only supported for a subset of the `atomic::Ordering` variants. Passing an unsupported variant will cause an unconditional panic at runtime, which is detected by this lint.

This lint will trigger in the following cases: (where `AtomicType` is an atomic type from `core::sync::atomic`, such as `AtomicBool`, `AtomicPtr`, `AtomicUsize`, or any of the other integer atomics).

- Passing `Ordering::Acquire` or `Ordering::AcqRel` to `AtomicType::store`.
- Passing `Ordering::Release` or `Ordering::AcqRel` to `AtomicType::load`.
- Passing `Ordering::Relaxed` to `core::sync::atomic::fence` or `core::sync::atomic::compiler_fence`.
- Passing `Ordering::Release` or `Ordering::AcqRel` as the failure ordering for any of `AtomicType::compare_exchange`, `AtomicType::compare_exchange_weak`, or `AtomicType::fetch_update`.

invalid-type-param-default

The `invalid_type_param_default` lint detects type parameter defaults erroneously allowed in an invalid location.

Example

```
fn foo<T=i32>(t: T) {}
```

This will produce:

```

error: defaults for type parameters are only allowed in `struct`, `enum`, `type`, or `trait` definitions
--> lint_example.rs:2:8
| fn foo<T=i32>(t: T) {}
|     ^^^^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #36887 <https://github.com/rust-lang/rust/issues/36887>
= note: `#[deny(invalid_type_param_default)]` on by default

```

Explanation

Default type parameters were only intended to be allowed in certain situations, but historically the compiler allowed them everywhere. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #36887](#) for more details.

let-underscore-lock

The `let_underscore_lock` lint checks for statements which don't bind a mutex to anything, causing the lock to be released immediately instead of at end of scope, which is typically incorrect.

Example

```

use std::sync::{Arc, Mutex};
use std::thread;
let data = Arc::new(Mutex::new(0));

thread::spawn(move || {
    // The lock is immediately released instead of at the end of the
    // scope, which is probably not intended.
    let _ = data.lock().unwrap();
    println!("doing some work");
    let mut lock = data.lock().unwrap();
    *lock += 1;
});

```

This will produce:

```
error: non-binding let on a synchronization lock
--> lint_example.rs:9:9
9 |     let _ = data.lock().unwrap();
   |     ^^^^^^^^^^^^^^^^^^^^^^ this binding will immediately drop the value assigned to
it
   |
   |         this lock is not assigned to a binding and is immediately dropped
   |
   |= note: `#[deny(let_underscore_lock)]` on by default
help: consider binding to an unused variable to avoid immediately dropping the value
9 |     let _unused = data.lock().unwrap();
   |     ~~~~~
help: consider immediately dropping the value
9 |     drop(data.lock().unwrap());
   |     ~~~~ +
```

Explanation

Statements which assign an expression to an underscore causes the expression to immediately drop instead of extending the expression's lifetime to the end of the scope. This is usually unintended, especially for types like `MutexGuard`, which are typically used to lock a mutex for the duration of an entire scope.

If you want to extend the expression's lifetime to the end of the scope, assign an underscore-prefixed name (such as `_foo`) to the expression. If you do actually want to drop the expression immediately, then calling `std::mem::drop` on the expression is clearer and helps convey intent.

macro-expanded-macro-exports-accessed-by-absolute-paths

The `macro_expanded_macro_exports_accessed_by_absolute_paths` lint detects macro-expanded `macro_export` macros from the current crate that cannot be referred to by absolute paths.

Example

```
macro_rules! define_exported {
    () => {
        #[macro_export]
        macro_rules! exported {
            () => {};
        }
    };
}

define_exported!();

fn main() {
    crate::exported!();
}
```

This will produce:

```
error: macro-expanded `macro_export` macros from the current crate cannot be referred to by
absolute paths
--> lint_example.rs:13:5
13 |     crate::exported!();
|     ^^^^^^^^^^^^^^^^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #52234 <https://github.com/rust-lang/rust/issues/52234>
note: the macro is defined here
--> lint_example.rs:4:9
4 | /         macro_rules! exported {
5 | |             () => {};
6 | |         }
| |_____^
...
10 |     define_exported!();
|     ----- in this macro invocation
= note: `#[deny(macro_expanded_macro_exports_accessed_by_absolute_paths)]` on by default
= note: this error originates in the macro `define_exported` (in Nightly builds, run with -Z
macro-backtrace for more info)
```

Explanation

The intent is that all macros marked with the `#[macro_export]` attribute are made available in the root of the crate. However, when a `macro_rules!` definition is generated by another macro, the macro expansion is unable to uphold this rule. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #53495](#) for more details.

missing-fragment-specifier

The `missing_fragment_specifier` lint is issued when an unused pattern in a `macro_rules!` macro definition has a meta-variable (e.g. `$e`) that is not followed by a fragment specifier (e.g. `:expr`).

This warning can always be fixed by removing the unused pattern in the `macro_rules!` macro definition.

Example

```
macro_rules! foo {
    () => {};
    ($name) => { };
}

fn main() {
    foo!();
}
```

This will produce:

```
error: missing fragment specifier
--> lint_example.rs:3:5
3 |     ($name) => { };
   |     ^^^^^^
   |
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #40107 <https://github.com/rust-lang/rust-lang/issues/40107>
= note: `#[deny(missing_fragment_specifier)]` on by default
```

Explanation

To fix this, remove the unused pattern from the `macro_rules!` macro definition:

```
macro_rules! foo {
    () => {};
}
fn main() {
    foo!();
}
```

mutable-transmutes

The `mutable_transmutes` lint catches transmuting from `&T` to `&mut T` because it is [undefined behavior](#).

Example

```
unsafe {
    let y = std::mem::transmute::<&i32, &mut i32>(&5);
}
```

This will produce:

```
error: transmuting &T to &mut T is undefined behavior, even if the reference is unused,
consider instead using an UnsafeCell
--> lint_example.rs:3:13
  |
3 |     let y = std::mem::transmute::<&i32, &mut i32>(&5);
  |     ^^^^^^^^^^
  |
  = note: #[deny(mutable_transmutes)]` on by default
```

Explanation

Certain assumptions are made about aliasing of data, and this transmute violates those assumptions. Consider using `UnsafeCell` instead.

named-asm-labels

The `named_asm_labels` lint detects the use of named labels in the inline `asm!` macro.

Example

```
use std::arch::asm;

fn main() {
    unsafe {
        asm!("foo: bar");
    }
}
```

This will produce:

```
error: avoid using named labels in inline assembly
--> lint_example.rs:6:15
6 |         asm!("foo: bar");
   |         ^^^
   |
= help: only local labels of the form `<number>:` should be used in inline asm
= note: see the asm section of Rust By Example <https://doc.rust-lang.org/nightly/rust-by-example/unsafe/asm.html#labels> for more information
= note: `#[deny(named_asm_labels)]` on by default
```

Explanation

LLVM is allowed to duplicate inline assembly blocks for any reason, for example when it is in a function that gets inlined. Because of this, GNU assembler [local labels](#) must be used instead of labels with a name. Using named labels might cause assembler or linker errors.

See the explanation in [Rust By Example](#) for more details.

no-mangle-const-items

The `no_mangle_const_items` lint detects any `const` items with the [no_mangle](#) attribute.

Example

```
#[no_mangle]
const FOO: i32 = 5;
```

This will produce:

```
error: const items should never be `#[no_mangle]`
--> lint_example.rs:3:1
3 | const FOO: i32 = 5;
   | -----^^^^^^^^^^^^^
   |
   | help: try a static value: `pub static`
   |
= note: `#[deny(no_mangle_const_items)]` on by default
```

Explanation

Constants do not have their symbols exported, and therefore, this probably means you meant to use a `static`, not a `const`.

order-dependent-trait-objects

The `order_dependent_trait_objects` lint detects a trait coherency violation that would allow creating two trait impls for the same dynamic trait object involving marker traits.

Example

```
pub trait Trait {}

impl Trait for dyn Send + Sync {}
impl Trait for dyn Sync + Send {}
```

This will produce:

```
error: conflicting implementations of trait `Trait` for type `(dyn Send + Sync + 'static)`:
(E0119)
--> lint_example.rs:5:1
4 | impl Trait for dyn Send + Sync {}
| ----- first implementation here
5 | impl Trait for dyn Sync + Send {}
| ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^ conflicting implementation for `(dyn Send + Sync +
'static)`
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #56484 <https://github.com/rust-lang/rust-lang/rust/issues/56484>
= note: `#[deny(order_dependent_trait_objects)]` on by default
```

Explanation

A previous bug caused the compiler to interpret traits with different orders (such as `Send + Sync` and `Sync + Send`) as distinct types when they were intended to be treated the same. This allowed code to define separate trait implementations when there should be a coherence error. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #56484](#) for more details.

overflowing-literals

The `overflowing_literals` lint detects literal out of range for its type.

Example

```
let x: u8 = 1000;
```

This will produce:

```
error: literal out of range for `u8`
--> lint_example.rs:2:13
2 | let x: u8 = 1000;
   |      ^^^^
   |
= note: the literal `1000` does not fit into the type `u8` whose range is `0..=255`
= note: `#[deny(overflowing_literals)]` on by default
```

Explanation

It is usually a mistake to use a literal that overflows the type where it is used. Either use a literal that is within range, or change the type to be within the range of the literal.

patterns-in-fns-without-body

The `patterns_in_fns_without_body` lint detects `mut` identifier patterns as a parameter in functions without a body.

Example

```
trait Trait {
    fn foo(mut arg: u8);
}
```

This will produce:

```
error: patterns aren't allowed in functions without bodies
--> lint_example.rs:3:12
3 |     fn foo(mut arg: u8);
   |         ^^^^^^^ help: remove `mut` from the parameter: `arg`
   |
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #35203 <https://github.com/rust-lang/rust-lang/issues/35203>
= note: `#[deny(patterns_in_fns_without_body)]` on by default
```

Explanation

To fix this, remove `mut` from the parameter in the trait definition; it can be used in the implementation. That is, the following is OK:

```
trait Trait {  
    fn foo(arg: u8); // Removed `mut` here  
}  
  
impl Trait for i32 {  
    fn foo(mut arg: u8) { // `mut` here is OK  
    }  
}
```

Trait definitions can define functions without a body to specify a function that implementors must define. The parameter names in the body-less functions are only allowed to be `_` or an [identifier](#) for documentation purposes (only the type is relevant). Previous versions of the compiler erroneously allowed [identifier patterns](#) with the `mut` keyword, but this was not intended to be allowed. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #35203](#) for more details.

proc-macro-back-compat

The `proc_macro_back_compat` lint detects uses of old versions of certain proc-macro crates, which have hardcoded workarounds in the compiler.

Example

```
use time_macros_impl::impl_macros;  
struct Foo;  
impl_macros!(Foo);
```

This will produce:

```
warning: using an old version of `time-macros-impl`  
::: $DIR/group-compat-hack.rs:27:5  
|  
LL |     impl_macros!(Foo);  
|----- in this macro invocation  
|  
= note: `#[warn(proc_macro_back_compat)]` on by default  
= warning: this was previously accepted by the compiler but is being phased out; it will  
become a hard error in a future release!  
= note: for more information, see issue #83125 <https://github.com/rust-lang/rust/issu...>  
= note: the `time-macros-impl` crate will stop compiling in futures version of Rust. Please  
update to the latest version of the `time` crate to avoid breakage  
= note: this warning originates in a macro (in Nightly builds, run with -Z macro-backtrace  
for more info)
```

Explanation

Eventually, the backwards-compatibility hacks present in the compiler will be removed, causing older versions of certain crates to stop compiling. This is a [future-incompatible](#) lint to ease the transition to an error. See [issue #83125](#) for more details.

proc-macro-derive-resolution-fallback

The `proc_macro_derive_resolution_fallback` lint detects proc macro derives using inaccessible names from parent modules.

Example

```
// foo.rs  
#![crate_type = "proc-macro"]  
  
extern crate proc_macro;  
  
use proc_macro::*;

#[proc_macro_derive(Foo)]
pub fn foo1(a: TokenStream) -> TokenStream {
    drop(a);
    "mod __bar { static mut BAR: Option<Something> = None; }".parse().unwrap()
}
```

```
// bar.rs
#[macro_use]
extern crate foo;

struct Something;

#[derive(Foo)]
struct Another;

fn main() {}
```

This will produce:

```
warning: cannot find type `Something` in this scope
--> src/main.rs:8:10
  |
8 | #[derive(Foo)]
  |          ^^^ names from parent modules are not accessible without an explicit import
  |
  = note: `#[warn(proc_macro_derive_resolution_fallback)]` on by default
  = warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
  = note: for more information, see issue #50504 <https://github.com/rust-lang/rust/issues/50504>
```

Explanation

If a proc-macro generates a module, the compiler unintentionally allowed items in that module to refer to items in the crate root without importing them. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #50504](#) for more details.

pub-use-of-private-extern-crate

The `pub_use_of_private_extern_crate` lint detects a specific situation of re-exporting a private `extern crate`.

Example

```
extern crate core;
pub use core as reexported_core;
```

This will produce:

```
error: extern crate `core` is private, and cannot be re-exported (error E0365), consider
declaring with `pub`
--> lint_example.rs:3:9
| pub use core as reexported_core;
|          ^^^^^^^^^^^^^^^^^^^^^^^^
|
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #34537 <https://github.com/rust-lang/rust/issues/34537>
= note: `#[deny(pub_use_of_private_extern_crate)]` on by default
```

Explanation

A public `use` declaration should not be used to publicly re-export a private `extern crate`. `pub extern crate` should be used instead.

This was historically allowed, but is not the intended behavior according to the visibility rules. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #34537](#) for more details.

soft-unstable

The `soft_unstable` lint detects unstable features that were unintentionally allowed on stable.

Example

```
#[cfg(test)]
extern crate test;

#[bench]
fn name(b: &mut test::Bencher) {
    b.iter(|| 123)
}
```

This will produce:

```
error: use of unstable library feature 'test': `bench` is a part of custom test frameworks
which are unstable
--> lint_example.rs:5:3
  |
5 | #[bench]
  | ^^^^^^
  |
= warning: this was previously accepted by the compiler but is being phased out; it will
become a hard error in a future release!
= note: for more information, see issue #64266 <https://github.com/rust-lang/rust/issues/64266>
= note: `#[deny(soft_unstable)]` on by default
```

Explanation

The `bench` attribute was accidentally allowed to be specified on the [stable release channel](#). Turning this to a hard error would have broken some projects. This lint allows those projects to continue to build correctly when `--cap-lints` is used, but otherwise signal an error that `#[bench]` should not be used on the stable channel. This is a [future-incompatible](#) lint to transition this to a hard error in the future. See [issue #64266](#) for more details.

test-unstable-lint

The `test_unstable_lint` lint tests unstable lints and is perma-unstable.

Example

```
#![allow(test_unstable_lint)]
```

This will produce:

```
warning: unknown lint: `test_unstable_lint`
--> lint_example.rs:1:1
  |
1 | #![allow(test_unstable_lint)]
  | ^^^^^^^^^^^^^^^^^^^^^^^^^^
  |
= note: the `test_unstable_lint` lint is unstable
= help: add `#[feature(test_unstable_lint)]` to the crate attributes to enable
= note: `#[warn(unknown_lints)]` on by default
```

Explanation

In order to test the behavior of unstable lints, a permanently-unstable lint is required. This lint can be used to trigger warnings and errors from the compiler related to unstable lints.

text-direction-codepoint-in-comment

The `text_direction_codepoint_in_comment` lint detects Unicode codepoints in comments that change the visual representation of text on screen in a way that does not correspond to their on memory representation.

Example

```
#![deny(text_direction_codepoint_in_comment)]
fn main() {
    println!("{:?}"); // ';'(
}
```

This will produce:

```
error: unicode codepoint changing visible direction of text present in comment
--> lint_example.rs:3:23
3 |     println!("{:?}"); // '';
   |     ^^^^^-^^
   |     |
   |     '\u{202e}'
   |         this comment contains an invisible unicode text flow control
codepoint
|
= note: these kind of unicode codepoints change the way text flows on applications that
support them, but can cause confusion because they change the order of characters on the
screen
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #![deny(text_direction_codepoint_in_comment)]
   | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
= help: if their presence wasn't intentional, you can remove them
```

Explanation

Unicode allows changing the visual flow of text on screen in order to support scripts that are written right-to-left, but a specially crafted comment can make code that will be compiled appear to be part of a comment, depending on the software used to read the code. To avoid potential problems or confusion, such as in CVE-2021-42574, by default we deny their use.

text-direction-codepoint-in-literal

The `text_direction_codepoint_in_literal` lint detects Unicode codepoints that change the visual representation of text on screen in a way that does not correspond to their on memory representation.

Explanation

The unicode characters `\u{202A}`, `\u{202B}`, `\u{202D}`, `\u{202E}`, `\u{2066}`, `\u{2067}`, `\u{2068}`, `\u{202C}` and `\u{2069}` make the flow of text on screen change its direction on software that supports these codepoints. This makes the text "abc" display as "cba" on screen. By leveraging software that supports these, people can write specially crafted literals that make the surrounding code seem like it's performing one action, when in reality it is performing another. Because of this, we proactively lint against their presence to avoid surprises.

Example

```
#![deny(text_direction_codepoint_in_literal)]
fn main() {
    println!("{}{:?}", 'a', 'b');
}
```

This will produce:

```
error: unicode codepoint changing visible direction of text present in literal
--> lint_example.rs:3:22
3 |     println!("{}{:?}", 'a', 'b');
   |          ^-
   |          ||
   |          | '\u{202e}'
   |          this literal contains an invisible unicode text flow control
codepoint
|
= note: these kind of unicode codepoints change the way text flows on applications that
support them, but can cause confusion because they change the order of characters on the
screen
note: the lint level is defined here
--> lint_example.rs:1:9
1 | #![deny(text_direction_codepoint_in_literal)]
   |          ^^^^^^^^^^^^^^^^^^^^^^
= help: if their presence wasn't intentional, you can remove them
help: if you want to keep them but make them visible in your source code, you can escape them
3 |     println!("{}{:?}", 'a', '\u{202e}' );
   |          ~~~~~~
```

unconditional-panic

The `unconditional_panic` lint detects an operation that will cause a panic at runtime.

Example

```
let x = 1 / 0;
```

This will produce:

```
error: this operation will panic at runtime
--> lint_example.rs:3:9
   |
3 | let x = 1 / 0;
   |      ^^^^^ attempt to divide `1_i32` by zero
   |
= note: `#[deny(unconditional_panic)]` on by default
```

Explanation

This lint detects code that is very likely incorrect because it will always panic, such as division by zero and out-of-bounds array accesses. Consider adjusting your code if this is a bug, or using the `panic!` or `unreachable!` macro instead in case the panic is intended.

unknown-crate-types

The `unknown_crate_types` lint detects an unknown crate type found in a `crate_type` attribute.

Example

```
#![crate_type="lol"]
fn main() {}
```

This will produce:

```
error: invalid `crate_type` value
--> lint_example.rs:1:15
1 | #[crate_type="lol"]
  ^^^^^^
|
= note: `#[deny(unknown_crate_types)]` on by default
```

Explanation

An unknown value give to the `crate_type` attribute is almost certainly a mistake.

useless-deprecated

The `useless_deprecated` lint detects deprecation attributes with no effect.

Example

```
struct X;

#[deprecated = "message"]
impl Default for X {
    fn default() -> Self {
        X
    }
}
```

This will produce:

```
error: this `#[deprecated]` annotation has no effect
--> lint_example.rs:4:1
4 | #[deprecated = "message"]
  ^^^^^^^^^^^^^^^^^^^^^^^ help: remove the unnecessary deprecation attribute
|
= note: `#[deny(useless_deprecated)]` on by default
```

Explanation

Deprecation attributes have no effect on trait implementations.

JSON Output

This chapter documents the JSON structures emitted by `rustc`. JSON may be enabled with the `--error-format=json` flag. Additional options may be specified with the `--json` flag which can change which messages are generated, and the format of the messages.

JSON messages are emitted one per line to stderr.

If parsing the output with Rust, the `cargo_metadata` crate provides some support for parsing the messages.

When parsing, care should be taken to be forwards-compatible with future changes to the format. Optional values may be `null`. New fields may be added. Enumerated fields like "level" or "suggestion_applicability" may add new values.

Diagnostics

Diagnostic messages provide errors or possible concerns generated during compilation. `rustc` provides detailed information about where the diagnostic originates, along with hints and suggestions.

Diagnostics are arranged in a parent/child relationship where the parent diagnostic value is the core of the diagnostic, and the attached children provide additional context, help, and information.

Diagnostics have the following format:

```
{
    /* The primary message. */
    "message": "unused variable: `x`",
    /* The diagnostic code.
       Some messages may set this value to null.
    */
    "code": {
        /* A unique string identifying which diagnostic triggered. */
        "code": "unused_variables",
        /* An optional string explaining more detail about the diagnostic code. */
        "explanation": null
    },
    /* The severity of the diagnostic.
       Values may be:
       - "error": A fatal error that prevents compilation.
       - "warning": A possible error or concern.
       - "note": Additional information or context about the diagnostic.
       - "help": A suggestion on how to resolve the diagnostic.
       - "failure-note": A note attached to the message for further information.
       - "error: internal compiler error": Indicates a bug within the compiler.
    */
    "level": "warning",
    /* An array of source code locations to point out specific details about
       where the diagnostic originates from. This may be empty, for example
       for some global messages, or child messages attached to a parent.

       Character offsets are offsets of Unicode Scalar Values.
    */
    "spans": [
        {
            /* The file where the span is located.
               Note that this path may not exist. For example, if the path
               points to the standard library, and the rust src is not
               available in the sysroot, then it may point to a non-existent
               file. Beware that this may also point to the source of an
               external crate.
            */
            "file_name": "lib.rs",
            /* The byte offset where the span starts (0-based, inclusive). */
            "byte_start": 21,
            /* The byte offset where the span ends (0-based, exclusive). */
            "byte_end": 22,
            /* The first line number of the span (1-based, inclusive). */
            "line_start": 2,
            /* The last line number of the span (1-based, inclusive). */
            "line_end": 2,
            /* The first character offset of the line_start (1-based, inclusive). */
            "column_start": 9,
            /* The last character offset of the line_end (1-based, exclusive). */
            "column_end": 10,
            /* Whether or not this is the "primary" span.
            */
        }
    ]
}
```

This indicates that this span is the focal point of the diagnostic.

There are rare cases where multiple spans may be marked as primary. For example, "immutable borrow occurs here" and "mutable borrow ends here" can be two separate primary spans.

The top (parent) message should always have at least one

```

primary span, unless it has zero spans. Child messages may have
zero or more primary spans.
*/
"is_primary": true,
/* An array of objects showing the original source code for this
span. This shows the entire lines of text where the span is
located. A span across multiple lines will have a separate
value for each line.
*/
"text": [
{
    /* The entire line of the original source code. */
    "text": "    let x = 123;",
    /* The first character offset of the line of
       where the span covers this line (1-based, inclusive). */
    "highlight_start": 9,
    /* The last character offset of the line of
       where the span covers this line (1-based, exclusive). */
    "highlight_end": 10
}
],
/* An optional message to display at this span location.
   This is typically null for primary spans.
*/
"label": null,
/* An optional string of a suggested replacement for this span to
   solve the issue. Tools may try to replace the contents of the
   span with this text.
*/
"suggested_replacement": null,
/* An optional string that indicates the confidence of the
   "suggested_replacement". Tools may use this value to determine
   whether or not suggestions should be automatically applied.

Possible values may be:
- "MachineApplicable": The suggestion is definitely what the
  user intended. This suggestion should be automatically
  applied.
- "MaybeIncorrect": The suggestion may be what the user
  intended, but it is uncertain. The suggestion should result
  in valid Rust code if it is applied.
- "HasPlaceholders": The suggestion contains placeholders like
  `(...)` . The suggestion cannot be applied automatically
  because it will not result in valid Rust code. The user will
  need to fill in the placeholders.
- "Unspecified": The applicability of the suggestion is unknown.
*/
"suggestion_applicability": null,
/* An optional object indicating the expansion of a macro within
   this span.

If a message occurs within a macro invocation, this object will
provide details of where within the macro expansion the message
is located.
*/
"expansion": {
    /* The span of the macro invocation.
       Uses the same span definition as the "spans" array.
    */
    "span": {/*...*/}
    /* Name of the macro, such as "foo!" or "#[derive(Eq)]". */
}

```

```
"macro_decl_name": "some_macro!",
/* Optional span where the relevant part of the macro is
defined. */
"def_site_span": {/*...*/},
}
],
/* Array of attached diagnostic messages.
This is an array of objects using the same format as the parent
message. Children are not nested (children do not themselves
contain "children" definitions).
*/
"children": [
{
  "message": "`#[warn(unused_variables)]` on by default",
  "code": null,
  "level": "note",
  "spans": [],
  "children": [],
  "rendered": null
},
{
  "message": "if this is intentional, prefix it with an underscore",
  "code": null,
  "level": "help",
  "spans": [
    {
      "file_name": "lib.rs",
      "byte_start": 21,
      "byte_end": 22,
      "line_start": 2,
      "line_end": 2,
      "column_start": 9,
      "column_end": 10,
      "is_primary": true,
      "text": [
        {
          "text": "let x = 123;",
          "highlight_start": 9,
          "highlight_end": 10
        }
      ],
      "label": null,
      "suggested_replacement": "_x",
      "suggestion_applicability": "MachineApplicable",
      "expansion": null
    }
  ],
  "children": [],
  "rendered": null
}
],
/* Optional string of the rendered version of the diagnostic as displayed
by rustc. Note that this may be influenced by the `--json` flag.
*/
"rendered": "warning: unused variable: `x`\n→ lib.rs:2:9\n |\n2 | let x = 123;\n| ^ help: if this is intentional, prefix it with an underscore: `_x`\n|\n = note:\n`#[warn(unused_variables)]` on by default\n"
}
```

Artifact notifications

Artifact notifications are emitted when the `--json=artifacts` flag is used. They indicate that a file artifact has been saved to disk. More information about emit kinds may be found in the `--emit` flag documentation.

```
{  
    /* The filename that was generated. */  
    "artifact": "libfoo.rlib",  
    /* The kind of artifact that was generated. Possible values:  
     - "link": The generated crate as specified by the crate-type.  
     - "dep-info": The `.` file with dependency information in a Makefile-like syntax.  
     - "metadata": The Rust `.` file containing metadata about the crate.  
    */  
    "emit": "link"  
}
```

Future-incompatible reports

If the `--json=future-incompat` flag is used, then a separate JSON structure will be emitted if the crate may stop compiling in the future. This contains diagnostic information about the particular warnings that may be turned into a hard error in the future. This will include the diagnostic information, even if the diagnostics have been suppressed (such as with an `#[allow]` attribute or the `--cap-lints` option).

```
{  
    /* An array of objects describing a warning that will become a hard error  
     in the future.  
    */  
    "future_incompat_report":  
    [  
        {  
            /* A diagnostic structure as defined in  
             https://doc.rust-lang.org/rustc/json.html#diagnostics  
            */  
            "diagnostic": {...},  
        }  
    ]  
}
```

Tests

`rustc` has a built-in facility for building and running tests for a crate. More information about writing and running tests may be found in the [Testing Chapter](#) of the Rust Programming Language book.

Tests are written as free functions with the `#[test]` attribute. For example:

```
#[test]
fn it_works() {
    assert_eq!(2 + 2, 4);
}
```

Tests "pass" if they return without an error. They "fail" if they `panic`, or return a type such as `Result` that implements the `Termination` trait with a non-zero value.

By passing the `--test` option to `rustc`, the compiler will build the crate in a special mode to construct an executable that will run the tests in the crate. The `--test` flag will make the following changes:

- The crate will be built as a `bin` crate type, forcing it to be an executable.
- Links the executable with `libtest`, the test harness that is part of the standard library, which handles running the tests.
- Synthesizes a `main` function which will process command-line arguments and run the tests. This new `main` function will replace any existing `main` function as the entry point of the executable, though the existing `main` will still be compiled.
- Enables the `test cfg` option, which allows your code to use conditional compilation to detect if it is being built as a test.
- Enables building of functions annotated with the `test` and `bench` attributes, which will be run by the test harness.

After the executable is created, you can run it to execute the tests and receive a report on what passes and fails. If you are using `Cargo` to manage your project, it has a built-in `cargo test` command which handles all of this automatically. An example of the output looks like this:

```
running 4 tests
test it_works ... ok
test check_valid_args ... ok
test invalid_characters ... ok
test walks_the_dog ... ok

test result: ok. 4 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out; finished in 0.00s
```

Note: Tests must be built with the `unwind` panic strategy. This is because all tests run in the same process, and they are intended to catch panics, which is not possible with the `abort` strategy. See the unstable `-Z panic-abort-tests`

option for experimental support of the `abort` strategy by spawning tests in separate processes.

Test attributes

Tests are indicated using attributes on free functions. The following attributes are used for testing, see the linked documentation for more details:

- `#[test]` – Indicates a function is a test to be run.
- `#[bench]` – Indicates a function is a benchmark to be run. Benchmarks are currently unstable and only available in the nightly channel, see the [unstable docs](#) for more details.
- `#[should_panic]` – Indicates that the test function will only pass if the function [panics](#).
- `#[ignore]` – Indicates that the test function will be compiled, but not run by default. See the `--ignored` and `--include-ignored` options to run these tests.

CLI arguments

The libtest harness has several command-line arguments to control its behavior.

Note: When running with `cargo test`, the libtest CLI arguments must be passed after the `--` argument to differentiate between flags for Cargo and those for the harness. For example: `cargo test -- --nocapture`

Filters

Positional arguments (those without a `-` prefix) are treated as filters which will only run tests whose name matches one of those strings. The filter will match any substring found in the full path of the test function. For example, if the test function `it_works` is located in the module `utils::paths::tests`, then any of the filters `works`, `path`, `utils::`, or `utils::paths::tests::it_works` will match that test.

See [Selection options](#) for more options to control which tests are run.

Action options

The following options perform different actions other than running tests.

--list

Prints a list of all tests and benchmarks. Does not run any of the tests. [Filters](#) can be used to list only matching tests.

-h, --help

Displays usage information and command-line options.

Selection options

The following options change how tests are selected.

--test

This is the default mode where all tests will be run as well as running all benchmarks with only a single iteration (to ensure the benchmark works, without taking the time to actually perform benchmarking). This can be combined with the **--bench** flag to run both tests and perform full benchmarking.

--bench

This runs in a mode where tests are ignored, and only runs benchmarks. This can be combined with **--test** to run both benchmarks and tests.

--exact

This forces [filters](#) to match the full path of the test exactly. For example, if the test `it_works` is in the module `utils::paths::tests`, then only the string `utils::paths::tests::it_works` will match that test.

--skip FILTER

Skips any tests whose name contains the given *FILTER* string. This flag may be passed multiple times.

--ignored

Runs only tests that are marked with the [ignore](#) attribute.

--include-ignored

Runs both [ignored](#) and non-ignored tests.

--exclude-should-panic

Excludes tests marked with the `should_panic` attribute.

⚠️ This option is `unstable`, and requires the `-Z unstable-options` flag. See [tracking issue #82348](#) for more information.

Execution options

The following options affect how tests are executed.

--test-threads NUM_THREADS

Sets the number of threads to use for running tests in parallel. By default, uses the amount of concurrency available on the hardware as indicated by `available_parallelism`.

This can also be specified with the `RUST_TEST_THREADS` environment variable.

--force-run-in-process

Forces the tests to run in a single process when using the `abort` panic strategy.

⚠️ This only works with the unstable `-Z panic-abort-tests` option, and requires the `-Z unstable-options` flag. See [tracking issue #67650](#) for more information.

--ensure-time

⚠️ This option is `unstable`, and requires the `-Z unstable-options` flag. See [tracking issue #64888](#) and the [unstable docs](#) for more information.

--shuffle

Runs the tests in random order, as opposed to the default alphabetical order.

This may also be specified by setting the `RUST_TEST_SHUFFLE` environment variable to anything but `0`.

The random number generator seed that is output can be passed to `--shuffle-seed` to run the tests in the same order again.

Note that `--shuffle` does not affect whether the tests are run in parallel. To run the tests in random order sequentially, use `--shuffle --test-threads 1`.

⚠️ This option is `unstable`, and requires the `-Z unstable-options` flag. See [tracking issue #89583](#) for more information.

--shuffle-seed SEED

Like `--shuffle`, but seeds the random number generator with *SEED*. Thus, calling the test harness with `--shuffle-seed` *SEED* twice runs the tests in the same order both times.

SEED is any 64-bit unsigned integer, for example, one produced by `--shuffle`.

This can also be specified with the `RUST_TEST_SHUFFLE_SEED` environment variable.

⚠  This option is [unstable](#), and requires the `-Z unstable-options` flag. See [tracking issue #89583](#) for more information.

Output options

The following options affect the output behavior.

`-q`, `--quiet`

Displays one character per test instead of one line per test. This is an alias for `-format=terse`.

`--nocapture`

Does not capture the `stdout` and `stderr` of the test, and allows tests to print to the console. Usually the output is captured, and only displayed if the test fails.

This may also be specified by setting the `RUST_TEST_NOCAPTURE` environment variable to anything but `0`.

`--show-output`

Displays the `stdout` and `stderr` of successful tests after all tests have run.

Contrast this with `--nocapture` which allows tests to print *while they are running*, which can cause interleaved output if there are multiple tests running in parallel, `--show-output` ensures the output is contiguous, but requires waiting for all tests to finish.

`--color COLOR`

Control when colored terminal output is used. Valid options:

- `auto`: Colorize if `stdout` is a `tty` and `--nocapture` is not used. This is the default.
- `always`: Always colorize the output.
- `never`: Never colorize the output.

--format FORMAT

Controls the format of the output. Valid options:

- `pretty`: This is the default format, with one line per test.
- `terse`: Displays only a single character per test. `--quiet` is an alias for this option.
- `json`: Emits JSON objects, one per line.  This option is `unstable`, and requires the `-Z unstable-options` flag. See [tracking issue #49359](#) for more information.

--logfile PATH

Writes the results of the tests to the given file.

--report-time

 This option is `unstable`, and requires the `-Z unstable-options` flag. See [tracking issue #64888](#) and the `unstable` docs for more information.

Unstable options

Some CLI options are added in an "unstable" state, where they are intended for experimentation and testing to determine if the option works correctly, has the right design, and is useful. The option may not work correctly, break, or change at any time. To signal that you acknowledge that you are using an unstable option, they require passing the `-Z unstable-options` command-line flag.

Benchmarks

The libtest harness supports running benchmarks for functions annotated with the `# [bench]` attribute. Benchmarks are currently unstable, and only available on the `nightly` channel. More information may be found in the [unstable book](#).

Custom test frameworks

Experimental support for using custom test harnesses is available on the `nightly` channel. See [tracking issue #50297](#) and the `custom_test_frameworks` documentation for more information.

Platform Support

Support for different platforms ("targets") are organized into three tiers, each with a different set of guarantees. For more information on the policies for targets at each tier, see the [Target Tier Policy](#).

Targets are identified by their "target triple" which is the string to inform the compiler what kind of output should be produced.

Component availability is tracked [here](#).

Tier 1 with Host Tools

Tier 1 targets can be thought of as "guaranteed to work". The Rust project builds official binary releases for each tier 1 target, and automated testing ensures that each tier 1 target builds and passes tests after each change.

Tier 1 targets with host tools additionally support running tools like `rustc` and `cargo` natively on the target, and automated testing ensures that tests pass for the host tools as well. This allows the target to be used as a development platform, not just a compilation target. For the full requirements, see [Tier 1 with Host Tools](#) in the Target Tier Policy.

All tier 1 targets with host tools support the full standard library.

target	notes
<code>aarch64-unknown-linux-gnu</code>	ARM64 Linux (kernel 4.1, glibc 2.17+) ¹
<code>i686-pc-windows-gnu</code>	32-bit MinGW (Windows 7+) ²
<code>i686-pc-windows-msvc</code>	32-bit MSVC (Windows 7+) ²
<code>i686-unknown-linux-gnu</code>	32-bit Linux (kernel 3.2+, glibc 2.17+)
<code>x86_64-apple-darwin</code>	64-bit macOS (10.7+, Lion+)
<code>x86_64-pc-windows-gnu</code>	64-bit MinGW (Windows 7+) ²
<code>x86_64-pc-windows-msvc</code>	64-bit MSVC (Windows 7+) ²
<code>x86_64-unknown-linux-gnu</code>	64-bit Linux (kernel 3.2+, glibc 2.17+)

¹ Stack probes support is missing on `aarch64-unknown-linux-gnu`, but it's planned to be implemented in the near future. The implementation is tracked on [issue #77071](#).

² Only Windows 10 currently undergoes automated testing. Earlier versions of Windows rely on testing and support from the community.

Tier 1

Tier 1 targets can be thought of as "guaranteed to work". The Rust project builds official binary releases for each tier 1 target, and automated testing ensures that each tier 1 target builds and passes tests after each change. For the full requirements, see [Tier 1 target policy](#) in the Target Tier Policy.

At this time, all Tier 1 targets are [Tier 1 with Host Tools](#).

Tier 2 with Host Tools

Tier 2 targets can be thought of as "guaranteed to build". The Rust project builds official binary releases for each tier 2 target, and automated builds ensure that each tier 2 target builds after each change. Automated tests are not always run so it's not guaranteed to produce a working build, but tier 2 targets often work to quite a good degree and patches are always welcome!

Tier 2 targets with host tools additionally support running tools like `rustc` and `cargo` natively on the target, and automated builds ensure that the host tools build as well. This allows the target to be used as a development platform, not just a compilation target. For the full requirements, see [Tier 2 with Host Tools](#) in the Target Tier Policy.

All tier 2 targets with host tools support the full standard library.

NOTE: The `rust-docs` component is not usually built for tier 2 targets, so Rustup may install the documentation for a similar tier 1 target instead.

target	notes
<code>aarch64-apple-darwin</code>	ARM64 macOS (11.0+, Big Sur+)
<code>aarch64-pc-windows-msvc</code>	ARM64 Windows MSVC
<code>aarch64-unknown-linux-musl</code>	ARM64 Linux with MUSL
<code>arm-unknown-linux-gnueabi</code>	ARMv6 Linux (kernel 3.2, glibc 2.17)
<code>arm-unknown-linux-gnueabihf</code>	ARMv6 Linux, hardfloat (kernel 3.2, glibc 2.17)
<code>armv7-unknown-linux-gnueabihf</code>	ARMv7 Linux, hardfloat (kernel 3.2, glibc 2.17)
<code>mips-unknown-linux-gnu</code>	MIPS Linux (kernel 4.4, glibc 2.23)
<code>mips64-unknown-linux-gnuabi64</code>	MIPS64 Linux, n64 ABI (kernel 4.4, glibc 2.23)
<code>mips64el-unknown-linux-gnuabi64</code>	MIPS64 (LE) Linux, n64 ABI (kernel 4.4, glibc 2.23)
<code>mipsel-unknown-linux-gnu</code>	MIPS (LE) Linux (kernel 4.4, glibc 2.23)
<code>powerpc-unknown-linux-gnu</code>	PowerPC Linux (kernel 3.2, glibc 2.17)
<code>powerpc64-unknown-linux-gnu</code>	PPC64 Linux (kernel 3.2, glibc 2.17)

target	notes
powerpc64le-unknown-linux-gnu	PPC64LE Linux (kernel 3.10, glibc 2.17)
riscv64gc-unknown-linux-gnu	RISC-V Linux (kernel 4.20, glibc 2.29)
s390x-unknown-linux-gnu	S390x Linux (kernel 3.2, glibc 2.17)
x86_64-unknown-freebsd	64-bit FreeBSD
x86_64-unknown-illumos	illumos
x86_64-unknown-linux-musl	64-bit Linux with MUSL
x86_64-unknown-netbsd	NetBSD/amd64

Tier 2

Tier 2 targets can be thought of as "guaranteed to build". The Rust project builds official binary releases for each tier 2 target, and automated builds ensure that each tier 2 target builds after each change. Automated tests are not always run so it's not guaranteed to produce a working build, but tier 2 targets often work to quite a good degree and patches are always welcome! For the full requirements, see [Tier 2 target policy](#) in the Target Tier Policy.

The `std` column in the table below has the following meanings:

- ✓ indicates the full standard library is available.
- * indicates the target only supports `no_std` development.

NOTE: The `rust-docs` component is not usually built for tier 2 targets, so Rustup may install the documentation for a similar tier 1 target instead.

target	std	notes
aarch64-apple-ios	✓	ARM64 iOS
aarch64-apple-ios-sim	✓	Apple iOS Simulator on ARM64
aarch64-fuchsia	✓	Alias for <code>aarch64-unknown-fuchsia</code>
aarch64-unknown-fuchsia	✓	ARM64 Fuchsia
aarch64-linux-android	✓	ARM64 Android
aarch64-unknown-none-softfloat	*	Bare ARM64, softfloat
aarch64-unknown-none	*	Bare ARM64, hardfloat
aarch64-unknown-uefi	*	ARM64 UEFI
arm-linux-androideabi	✓	ARMv7 Android
arm-unknown-linux-musleabi	✓	ARMv6 Linux with MUSL
arm-unknown-linux-musleabihf	✓	ARMv6 Linux with MUSL, hardfloat
armeabi-v7r-none-eabi	*	Bare ARMv7-R, Big Endian
armeabi-v7r-none-eabihf	*	Bare ARMv7-R, Big Endian, hardfloat

target	std	notes
armv5te-unknown-linux-gnueabi	✓	ARMv5TE Linux (kernel 4.4, glibc 2.23)
armv5te-unknown-linux-musleabi	✓	ARMv5TE Linux with MUSL
armv7-linux-androideabi	✓	ARMv7a Android
armv7-unknown-linux-gnueabi	✓	ARMv7 Linux (kernel 4.15, glibc 2.27)
armv7-unknown-linux-musleabi	✓	ARMv7 Linux with MUSL
armv7-unknown-linux-musleabihf	✓	ARMv7 Linux with MUSL, hardfloat
armv7a-none-eabi	*	Bare ARMv7-A
armv7r-none-eabi	*	Bare ARMv7-R
armv7r-none-eabihf	*	Bare ARMv7-R, hardfloat
asmjs-unknown-emscripten	✓	asm.js via Emscripten
i586-pc-windows-msvc	*	32-bit Windows w/o SSE
i586-unknown-linux-gnu	✓	32-bit Linux w/o SSE (kernel 3.2, glibc 2.17)
i586-unknown-linux-musl	✓	32-bit Linux w/o SSE, MUSL
i686-linux-android	✓	32-bit x86 Android
i686-unknown-freebsd	✓	32-bit FreeBSD
i686-unknown-linux-musl	✓	32-bit Linux with MUSL
i686-unknown-uefi	*	32-bit UEFI
mips-unknown-linux-musl	✓	MIPS Linux with MUSL
mips64-unknown-linux-muslabi64	✓	MIPS64 Linux, n64 ABI, MUSL
mips64el-unknown-linux-muslabi64	✓	MIPS64 (LE) Linux, n64 ABI, MUSL
mipsel-unknown-linux-musl	✓	MIPS (LE) Linux with MUSL
nvptx64-nvidia-cuda	*	--emit=asm generates PTX code that runs on NVIDIA GPUs
riscv32i-unknown-none-elf	*	Bare RISC-V (RV32I ISA)
riscv32imac-unknown-none-elf	*	Bare RISC-V (RV32IMAC ISA)
riscv32imc-unknown-none-elf	*	Bare RISC-V (RV32IMC ISA)
riscv64gc-unknown-none-elf	*	Bare RISC-V (RV64IMAFDC ISA)
riscv64imac-unknown-none-elf	*	Bare RISC-V (RV64IMAC ISA)
sparc64-unknown-linux-gnu	✓	SPARC Linux (kernel 4.4, glibc 2.23)
sparcv9-sun-solaris	✓	SPARC Solaris 10/11, illumos
thumbv6m-none-eabi	*	Bare Cortex-M0, M0+, M1
thumbv7em-none-eabi	*	Bare Cortex-M4, M7
thumbv7em-none-eabihf	*	Bare Cortex-M4F, M7F, FPU, hardfloat
thumbv7m-none-eabi	*	Bare Cortex-M3

target	std	notes
thumbv7neon-linux-androideabi	✓	Thumb2-mode ARMv7a Android with NEON
thumbv7neon-unknown-linux-gnueabihf	✓	Thumb2-mode ARMv7a Linux with NEON (kernel 4.4, glibc 2.23)
thumbv8m.base-none-eabi	*	ARMv8-M Baseline
thumbv8m.main-none-eabi	*	ARMv8-M Mainline
thumbv8m.main-none-eabihf	*	ARMv8-M Mainline, hardfloat
wasm32-unknown-emscripten	✓	WebAssembly via Emscripten
wasm32-unknown-unknown	✓	WebAssembly
wasm32-wasi	✓	WebAssembly with WASI
x86_64-apple-ios	✓	64-bit x86 iOS
x86_64-fortanix-unknown-sgx	✓	Fortanix ABI for 64-bit Intel SGX
x86_64-fuchsia	✓	Alias for x86_64-unknown-fuchsia
x86_64-unknown-fuchsia	✓	64-bit Fuchsia
x86_64-linux-android	✓	64-bit x86 Android
x86_64-pc-solaris	✓	64-bit Solaris 10/11, illumos
x86_64-unknown-linux-gnu32	✓	64-bit Linux (x32 ABI) (kernel 4.15, glibc 2.27)
x86_64-unknown-none	*	Freestanding/bare-metal x86_64, softfloat
x86_64-unknown-redox	✓	Redox OS
x86_64-unknown-uefi	*	64-bit UEFI

Tier 3

Tier 3 targets are those which the Rust codebase has support for, but which the Rust project does not build or test automatically, so they may or may not work. Official builds are not available. For the full requirements, see [Tier 3 target policy](#) in the Target Tier Policy.

The `std` column in the table below has the following meanings:

- ✓ indicates the full standard library is available.
- * indicates the target only supports `no_std` development.
- ? indicates the standard library support is unknown or a work-in-progress.

The `host` column indicates whether the codebase includes support for building host tools.

target	std	host	notes
aarch64-apple-ios-macabi	?		Apple Catalyst on ARM64

target	std	host	notes
aarch64-apple-tvos	*		ARM64 tvOS
aarch64-apple-watchos-sim	✓		ARM64 Apple WatchOS Simulator
aarch64-kmc-solid_asp3	✓		ARM64 SOLID with TOPPERS/ASP3
aarch64-nintendo-switch-freestanding	*		ARM64 Nintendo Switch, Horizon
aarch64-pc-windows-gnullvm	✓	✓	
aarch64-unknown-linux-ohos	✓		ARM64 OpenHarmony
aarch64-unknown-nto-qnx710	✓		ARM64 QNX Neutrino 7.1 RTOS
aarch64-unknown-freebsd	✓	✓	ARM64 FreeBSD
aarch64-unknown-hermit	✓		ARM64 HermitCore
aarch64-unknown-linux-gnu_ilp32	✓	✓	ARM64 Linux (ILP32 ABI)
aarch64-unknown-netbsd	✓	✓	
aarch64-unknown-openbsd	✓	✓	ARM64 OpenBSD
aarch64-unknown-redox	?		ARM64 Redox OS
aarch64-uwp-windows-msvc	?		
aarch64-wrs-vxworks	?		
aarch64_be-unknown-linux-gnu_ilp32	✓	✓	ARM64 Linux (big-endian, ILP32 ABI)
aarch64_be-unknown-linux-gnu	✓	✓	ARM64 Linux (big-endian)
arm64_32-apple-watchos	✓		ARM Apple WatchOS 64-bit with 32-bit pointers
armeabi-unknown-linux-gnueabi	✓	?	ARM BE8 the default ARM big-endian architecture since ARMv6.
armv4t-none-eabi	*		ARMv4T A32
armv4t-unknown-linux-gnueabi	?		
armv5te-none-eabi	*		ARMv5TE A32
armv5te-unknown-linux-uclibceabi	?		ARMv5TE Linux with uClibc
armv6-unknown-freebsd	✓	✓	ARMv6 FreeBSD
armv6-unknown-netbsd-eabihf	?		
armv6k-nintendo-3ds	?		ARMv6K Nintendo 3DS, Horizon (Requires devkitARM toolchain)
armv7-apple-ios	✓		ARMv7 iOS, Cortex-a8
armv7-sony-vita-newlibeabihf	?		ARM Cortex-A9 Sony PlayStation Vita (requires VITASDK toolchain)

target	std	host	notes
armv7-unknown-linux-ohos	✓		ARMv7 OpenHarmony
armv7-unknown-linux-uclibceabi	✓	✓	ARMv7 Linux with uClibc, softfloat
armv7-unknown-linux-uclibceabihf	✓	?	ARMv7 Linux with uClibc, hardfloat
armv7-unknown-freebsd	✓	✓	ARMv7 FreeBSD
armv7-unknown-netbsd-eabihf	✓	✓	
armv7-wrs-vxworks-eabihf	?		
armv7a-kmc-solid_asp3-eabi	✓		ARM SOLID with TOPPERS/ASP3
armv7a-kmc-solid_asp3-eabihf	✓		ARM SOLID with TOPPERS/ASP3, hardfloat
armv7a-none-eabihf	*		ARM Cortex-A, hardfloat
armv7k-apple-watchos	✓		ARM Apple WatchOS
armv7s-apple-ios	✓		
avr-unknown-gnu-atmega328	*		AVR. Requires -Z build-std=core
bpfbe-unknown-none	*		BPF (big endian)
bpfel-unknown-none	*		BPF (little endian)
hexagon-unknown-linux-musl	?		
i386-apple-ios	✓		32-bit x86 iOS
i586-pc-nto-qnx700	*		32-bit x86 QNX Neutrino 7.0 RTOS
i686-apple-darwin	✓	✓	32-bit macOS (10.7+, Lion+)
i686-pc-windows-msvc	*		32-bit Windows XP support
i686-unknown-haiku	✓	✓	32-bit Haiku
i686-unknown-netbsd	✓	✓	NetBSD/i386 with SSE2
i686-unknown-openbsd	✓	✓	32-bit OpenBSD
i686-uwp-windows-gnu	?		
i686-uwp-windows-msvc	?		
i686-wrs-vxworks	?		
loongarch64-unknown-linux-gnu	?		LoongArch64 Linux (LP64D ABI)
m68k-unknown-linux-gnu	?		Motorola 680x0 Linux
mips-unknown-linux-uclibc	✓		MIPS Linux with uClibc
mips64-openwrt-linux-musl	?		MIPS64 for OpenWrt Linux MUSL
mipsel-sony-psp	*		MIPS (LE) Sony PlayStation Portable (PSP)

target	std	host	notes
<code>mipsel-sony-psx</code>	*		MIPS (LE) Sony PlayStation 1 (PSX)
<code>mipsel-unknown-linux-uclibc</code>	✓		MIPS (LE) Linux with uClibc
<code>mipsel-unknown-none</code>	*		Bare MIPS (LE) softfloat
<code>mipsisa32r6-unknown-linux-gnu</code>	?		
<code>mipsisa32r6el-unknown-linux-gnu</code>	?		
<code>mipsisa64r6-unknown-linux-gnuabi64</code>	?		
<code>mipsisa64r6el-unknown-linux-gnuabi64</code>	?		
<code>msp430-none-elf</code>	*		16-bit MSP430 microcontrollers
<code>powerpc-unknown-linux-gnuspe</code>	✓		PowerPC SPE Linux
<code>powerpc-unknown-linux-musl</code>	?		
<code>powerpc-unknown-netbsd</code>	✓	✓	
<code>powerpc-unknown-openbsd</code>	?		
<code>powerpc-wrs-vxworks-spe</code>	?		
<code>powerpc-wrs-vxworks</code>	?		
<code>powerpc64-unknown-freebsd</code>	✓	✓	PPC64 FreeBSD (ELFv1 and ELFv2)
<code>powerpc64le-unknown-freebsd</code>			PPC64LE FreeBSD
<code>powerpc-unknown-freebsd</code>			PowerPC FreeBSD
<code>powerpc64-unknown-linux-musl</code>	?		
<code>powerpc64-wrs-vxworks</code>	?		
<code>powerpc64le-unknown-linux-musl</code>	?		
<code>powerpc64-unknown-openbsd</code>	✓	✓	OpenBSD/powerpc64
<code>powerpc64-ibm-aix</code>	?		64-bit AIX (7.2 and newer)
<code>riscv32gc-unknown-linux-gnu</code>			RISC-V Linux (kernel 5.4, glibc 2.33)
<code>riscv32gc-unknown-linux-musl</code>			RISC-V Linux (kernel 5.4, musl + RISCV32 support patches)
<code>riscv32im-unknown-none-elf</code>	*		Bare RISC-V (RV32IM ISA)
<code>riscv32imac-unknown-xous-elf</code>	?		RISC-V Xous (RV32IMAC ISA)
<code>riscv32imc-esp-espifdf</code>	✓		RISC-V ESP-IDF
<code>riscv64gc-unknown-freebsd</code>			RISC-V FreeBSD
<code>riscv64gc-unknown-fuchsia</code>			RISC-V Fuchsia
<code>riscv64gc-unknown-linux-musl</code>			RISC-V Linux (kernel 4.20, musl 1.2.0)

target	std	host	notes
riscv64gc-unknown-openbsd	✓	✓	OpenBSD/riscv64
s390x-unknown-linux-musl			S390x Linux (kernel 3.2, MUSL)
sparc-unknown-linux-gnu	✓		32-bit SPARC Linux
sparc64-unknown-netbsd	✓	✓	NetBSD/sparc64
sparc64-unknown-openbsd	✓	✓	OpenBSD/sparc64
thumbv4t-none-eabi	*		ARMv4T T32
thumbv5te-none-eabi	*		ARMv5TE T32
thumbv7a-pc-windows-msvc	?		
thumbv7a-uwp-windows-msvc	✓		
thumbv7neon-unknown-linux-musleabihf	?		Thumb2-mode ARMv7a Linux with NEON, MUSL
wasm64-unknown-unknown	?		WebAssembly
x86_64-apple-ios-macabi	✓		Apple Catalyst on x86_64
x86_64-apple-tvos	*		x86 64-bit tvOS
x86_64-apple-watchos-sim	✓		x86 64-bit Apple WatchOS simulator
x86_64-pc-nto-qnx710	✓		x86 64-bit QNX Neutrino 7.1 RTOS
x86_64-pc-windows-gnullvm	✓	✓	
x86_64-pc-windows-msvc	*		64-bit Windows XP support
x86_64-sun-solaris	?		Deprecated target for 64-bit Solaris 10/11, illumos
x86_64-unknown-dragonfly	✓	✓	64-bit DragonFlyBSD
x86_64-unknown-haiku	✓	✓	64-bit Haiku
x86_64-unknown-hermit	✓		HermitCore
x86_64-unknown-l4re-uclibc	?		
x86_64-unknown-openbsd	✓	✓	64-bit OpenBSD
x86_64-uwp-windows-gnu	✓		
x86_64-uwp-windows-msvc	✓		
x86_64-wrs-vxworks	?		

Target Tier Policy

Table of Contents

- General
- Adding a new target
- Tier 3 target policy
- Tier 2 target policy
 - Tier 2 with host tools
- Tier 1 target policy
 - Tier 1 with host tools

General

Rust provides three tiers of target support:

- Rust provides no guarantees about tier 3 targets; they exist in the codebase, but may or may not build.
- Rust's continuous integration checks that tier 2 targets will always build, but they may or may not pass tests.
- Rust's continuous integration checks that tier 1 targets will always build and pass tests.

Adding a new tier 3 target imposes minimal requirements; we focus primarily on avoiding disruption to other ongoing Rust development.

Tier 2 and tier 1 targets place work on Rust project developers as a whole, to avoid breaking the target. The broader Rust community may also feel more inclined to support higher-tier targets in their crates (though they are not obligated to do so). Thus, these tiers require commensurate and ongoing efforts from the maintainers of the target, to demonstrate value and to minimize any disruptions to ongoing Rust development.

This policy defines the requirements for accepting a proposed target at a given level of support.

Each tier builds on all the requirements from the previous tier, unless overridden by a stronger requirement. Targets at tier 2 and tier 1 may also provide *host tools* (such as `rustc` and `cargo`); each of those tiers includes a set of supplementary requirements that must be met if supplying host tools for the target. A target at tier 2 or tier 1 is not required to supply host tools, but if it does, it must meet the corresponding additional requirements for host tools.

The policy for each tier also documents the Rust governance teams that must approve the addition of any target at that tier. Those teams are responsible for reviewing and evaluating the target, based on these requirements and their own judgment. Those teams may apply additional requirements, including subjective requirements, such as to deal with issues not foreseen by this policy. (Such requirements may subsequently motivate additions to this policy.)

While these criteria attempt to document the policy, that policy still involves human judgment. Targets must fulfill the spirit of the requirements as well, as determined by the judgment of the approving teams. Reviewers and team members evaluating targets and target-specific patches should always use their own best judgment regarding the quality of work, and the suitability of a target for the Rust project. Neither this policy nor any decisions made regarding targets shall create any binding agreement or estoppel by any party.

Before filing an issue or pull request (PR) to introduce or promote a target, the target should already meet the corresponding tier requirements. This does not preclude an existing target's maintainers using issues (on the Rust repository or otherwise) to track requirements that have not yet been met, as appropriate; however, before officially proposing the introduction or promotion of a target, it should meet all of the necessary requirements. A target proposal must quote the corresponding requirements verbatim and respond to them as part of explaining how the target meets those requirements. (For the requirements that simply state that the target or the target developers must not do something, it suffices to acknowledge the requirement.)

For a list of all supported targets and their corresponding tiers ("tier 3", "tier 2", "tier 2 with host tools", "tier 1", or "tier 1 with host tools"), see [platform support](#).

Several parts of this policy require providing target-specific documentation. Such documentation should typically appear in a subdirectory of the platform-support section of this rustc manual, with a link from the target's entry in [platform support](#). Use [TEMPLATE.md](#) as a base, and see other documentation in that directory for examples.

Note that a target must have already received approval for the next lower tier, and spent a reasonable amount of time at that tier, before making a proposal for promotion to the next higher tier; this is true even if a target meets the requirements for several tiers at once. This policy leaves the precise interpretation of "reasonable amount of time" up to the approving teams; those teams may scale the amount of time required based on their confidence in the target and its demonstrated track record at its current tier. At a minimum, multiple stable releases of Rust should typically occur between promotions of a target.

The availability or tier of a target in stable Rust is not a hard stability guarantee about the future availability or tier of that target. Higher-level target tiers are an increasing commitment to the support of a target, and we will take that commitment and potential disruptions into account when evaluating the potential demotion or removal of a target that has been part of a stable release. The

promotion or demotion of a target will not generally affect existing stable releases, only current development and future releases.

In this policy, the words "must" and "must not" specify absolute requirements that a target must meet to qualify for a tier. The words "should" and "should not" specify requirements that apply in almost all cases, but for which the approving teams may grant an exception for good reason. The word "may" indicates something entirely optional, and does not indicate guidance or recommendations. This language is based on [IETF RFC 2119](#).

Adding a new target

New targets typically start as Tier 3 and then can be promoted later. To propose addition of a new target, open a pull request on [rust-lang/rust](#):

- Copy the [Tier 3 target policy](#) to the description and fill it out, see [example](#).
- Add a new description for the target in `src/doc/rustc/src/platform-support` using the [template](#).
- Add the target to the `SUMMARY.md` (allows wildcards) and `platform-support.md` (must name all targets verbatim). Link to the created description page.
- Ensure the pull request is assigned to a member of the [Rust compiler team](#) by commenting:

```
r? compiler-team
```

Tier 3 target policy

At this tier, the Rust project provides no official support for a target, so we place minimal requirements on the introduction of targets.

A proposed new tier 3 target must be reviewed and approved by a member of the compiler team based on these requirements. The reviewer may choose to gauge broader compiler team consensus via a [Major Change Proposal \(MCP\)](#).

A proposed target or target-specific patch that substantially changes code shared with other targets (not just target-specific code) must be reviewed and approved by the appropriate team for that shared code before acceptance.

- A tier 3 target must have a designated developer or developers (the "target maintainers") on record to be CCed when issues arise regarding the target. (The mechanism to track and CC such developers may evolve over time.)
- Targets must use naming consistent with any existing targets; for instance, a target for the same CPU or OS as an existing Rust target should use the same name for that CPU or OS. Targets should normally use the same names and naming conventions as used elsewhere in the broader ecosystem beyond Rust (such as in

other toolchains), unless they have a very good reason to diverge. Changing the name of a target can be highly disruptive, especially once the target reaches a higher tier, so getting the name right is important even for a tier 3 target.

- Target names should not introduce undue confusion or ambiguity unless absolutely necessary to maintain ecosystem compatibility. For example, if the name of the target makes people extremely likely to form incorrect beliefs about what it targets, the name should be changed or augmented to disambiguate it.
- If possible, use only letters, numbers, dashes and underscores for the name. Periods (.) are known to cause issues in Cargo.
- Tier 3 targets may have unusual requirements to build or use, but must not create legal issues or impose onerous legal terms for the Rust project or for Rust developers or users.
 - The target must not introduce license incompatibilities.
 - Anything added to the Rust repository must be under the standard Rust license (MIT OR Apache-2.0).
 - The target must not cause the Rust tools or libraries built for any other host (even when supporting cross-compilation to the target) to depend on any new dependency less permissive than the Rust licensing policy. This applies whether the dependency is a Rust crate that would require adding new license exceptions (as specified by the `tidy` tool in the `rust-lang/rust` repository), or whether the dependency is a native library or binary. In other words, the introduction of the target must not cause a user installing or running a version of Rust or the Rust tools to be subject to any new license requirements.
 - Compiling, linking, and emitting functional binaries, libraries, or other code for the target (whether hosted on the target itself or cross-compiling from another target) must not depend on proprietary (non-FOSS) libraries. Host tools built for the target itself may depend on the ordinary runtime libraries supplied by the platform and commonly used by other applications built for the target, but those libraries must not be required for code generation for the target; cross-compilation to the target must not require such libraries at all. For instance, `rustc` built for the target may depend on a common proprietary C runtime library or console output library, but must not depend on a proprietary code generation library or code optimization library. Rust's license permits such combinations, but the Rust project has no interest in maintaining such combinations within the scope of Rust itself, even at tier 3.
 - "onerous" here is an intentionally subjective term. At a minimum, "onerous" legal/licensing terms include but are *not* limited to: non-disclosure requirements, non-compete requirements, contributor license agreements (CLAs) or equivalent, "non-commercial"/"research-only"/etc terms, requirements conditional on the employer or employment of any particular Rust developers, revocable terms, any requirements that create liability for the Rust project or its developers or users, or any requirements that adversely affect the livelihood or prospects of the Rust project or its developers or users.

- Neither this policy nor any decisions made regarding targets shall create any binding agreement or estoppel by any party. If any member of an approving Rust team serves as one of the maintainers of a target, or has any legal or employment requirement (explicit or implicit) that might affect their decisions regarding a target, they must recuse themselves from any approval decisions regarding the target's tier status, though they may otherwise participate in discussions.
 - This requirement does not prevent part or all of this policy from being cited in an explicit contract or work agreement (e.g. to implement or maintain support for a target). This requirement exists to ensure that a developer or team responsible for reviewing and approving a target does not face any legal threats or obligations that would prevent them from freely exercising their judgment in such approval, even if such judgment involves subjective matters or goes beyond the letter of these requirements.
- Tier 3 targets should attempt to implement as much of the standard libraries as possible and appropriate (`core` for most targets, `alloc` for targets that can support dynamic memory allocation, `std` for targets with an operating system or equivalent layer of system-provided functionality), but may leave some code unimplemented (either unavailable or stubbed out as appropriate), whether because the target makes it impossible to implement or challenging to implement. The authors of pull requests are not obligated to avoid calling any portions of the standard library on the basis of a tier 3 target not implementing those portions.
- The target must provide documentation for the Rust community explaining how to build for the target, using cross-compilation if possible. If the target supports running binaries, or running tests (even if they do not pass), the documentation must explain how to run such binaries or tests for the target, using emulation if possible or dedicated hardware if necessary.
- Tier 3 targets must not impose burden on the authors of pull requests, or other developers in the community, to maintain the target. In particular, do not post comments (automated or manual) on a PR that derail or suggest a block on the PR based on a tier 3 target. Do not send automated messages or notifications (via any medium, including via `@`) to a PR author or others involved with a PR regarding a tier 3 target, unless they have opted into such messages.
 - Backlinks such as those generated by the issue/PR tracker when linking to an issue or PR are not considered a violation of this policy, within reason. However, such messages (even on a separate repository) must not generate notifications to anyone involved with a PR who has not requested such notifications.
- Patches adding or updating tier 3 targets must not break any existing tier 2 or tier 1 target, and must not knowingly break another tier 3 target without approval of either the compiler team or the maintainers of the other tier 3 target.
 - In particular, this may come up when working on closely related targets, such as variations of the same architecture with different features. Avoid introducing unconditional uses of features that another variation of the target may not have; use conditional compilation or runtime detection, as appropriate, to let each target run code supported by that target.

If a tier 3 target stops meeting these requirements, or the target maintainers no longer have interest or time, or the target shows no signs of activity and has not built for some time, or removing the target would improve the quality of the Rust codebase, we may post a PR to remove it; any such PR will be CCed to the target maintainers (and potentially other people who have previously worked on the target), to check potential interest in improving the situation.

Tier 2 target policy

At this tier, the Rust project guarantees that a target builds, and will reject patches that fail to build on a target. Thus, we place requirements that ensure the target will not block forward progress of the Rust project.

A proposed new tier 2 target must be reviewed and approved by the compiler team based on these requirements. Such review and approval may occur via a [Major Change Proposal \(MCP\)](#).

In addition, the infrastructure team must approve the integration of the target into Continuous Integration (CI), and the tier 2 CI-related requirements. This review and approval may take place in a PR adding the target to CI, or simply by an infrastructure team member reporting the outcome of a team discussion.

- A tier 2 target must have value to people other than its maintainers. (It may still be a niche target, but it must not be exclusively useful for an inherently closed group.)
- A tier 2 target must have a designated team of developers (the "target maintainers") available to consult on target-specific build-breaking issues, or if necessary to develop target-specific language or library implementation details. This team must have at least 2 developers.
 - The target maintainers should not only fix target-specific issues, but should use any such issue as an opportunity to educate the Rust community about portability to their target, and enhance documentation of the target.
- The target must not place undue burden on Rust developers not specifically concerned with that target. Rust developers are expected to not gratuitously break a tier 2 target, but are not expected to become experts in every tier 2 target, and are not expected to provide target-specific implementations for every tier 2 target.
- The target must provide documentation for the Rust community explaining how to build for the target using cross-compilation, and explaining how to run tests for the target. If at all possible, this documentation should show how to run Rust programs and tests for the target using emulation, to allow anyone to do so. If the target cannot be feasibly emulated, the documentation should explain how to obtain and work with physical hardware, cloud systems, or equivalent.
- The target must document its baseline expectations for the features or versions of CPUs, operating systems, libraries, runtime environments, and similar.
- If introducing a new tier 2 or higher target that is identical to an existing Rust target except for the baseline expectations for the features or versions

of CPUs, operating systems, libraries, runtime environments, and similar, then the proposed target must document to the satisfaction of the approving teams why the specific difference in baseline expectations provides sufficient value to justify a separate target.

- Note that in some cases, based on the usage of existing targets within the Rust community, Rust developers or a target's maintainers may wish to modify the baseline expectations of a target, or split an existing target into multiple targets with different baseline expectations. A proposal to do so will be treated similarly to the analogous promotion, demotion, or removal of a target, according to this policy, with the same team approvals required.

- For instance, if an OS version has become obsolete and unsupported, a target for that OS may raise its baseline expectations for OS version (treated as though removing a target corresponding to the older versions), or a target for that OS may split out support for older OS versions into a lower-tier target (treated as though demoting a target corresponding to the older versions, and requiring justification for a new target at a lower tier for the older OS versions).

- Tier 2 targets must not leave any significant portions of `core` or the standard library unimplemented or stubbed out, unless they cannot possibly be supported on the target.
 - The right approach to handling a missing feature from a target may depend on whether the target seems likely to develop the feature in the future. In some cases, a target may be co-developed along with Rust support, and Rust may gain new features on the target as that target gains the capabilities to support those features.
 - As an exception, a target identical to an existing tier 1 target except for lower baseline expectations for the OS, CPU, or similar, may propose to qualify as tier 2 (but not higher) without support for `std` if the target will primarily be used in `no_std` applications, to reduce the support burden for the standard library. In this case, evaluation of the proposed target's value will take this limitation into account.
- The code generation backend for the target should not have deficiencies that invalidate Rust safety properties, as evaluated by the Rust compiler team. (This requirement does not apply to arbitrary security enhancements or mitigations provided by code generation backends, only to those properties needed to ensure safe Rust code cannot cause undefined behavior or other unsoundness.) If this requirement does not hold, the target must clearly and prominently document any such limitations as part of the target's entry in the target tier list, and ideally also via a failing test in the testsuite. The Rust compiler team must be satisfied with the balance between these limitations and the difficulty of implementing the necessary features.
 - For example, if Rust relies on a specific code generation feature to ensure that safe code cannot overflow the stack, the code generation for the target should support that feature.
 - If the Rust compiler introduces new safety properties (such as via new capabilities of a compiler backend), the Rust compiler team will determine if they consider those new safety properties a best-effort improvement for

specific targets, or a required property for all Rust targets. In the latter case, the compiler team may require the maintainers of existing targets to either implement and confirm support for the property or update the target tier list with documentation of the missing property.

- If the target supports C code, and the target has an interoperable calling convention for C code, the Rust target must support that C calling convention for the platform via `extern "C"`. The C calling convention does not need to be the default Rust calling convention for the target, however.
- The target must build reliably in CI, for all components that Rust's CI considers mandatory.
- The approving teams may additionally require that a subset of tests pass in CI, such as enough to build a functional "hello world" program, `./x.py test --no-run`, or equivalent "smoke tests". In particular, this requirement may apply if the target builds host tools, or if the tests in question provide substantial value via early detection of critical problems.
- Building the target in CI must not take substantially longer than the current slowest target in CI, and should not substantially raise the maintenance burden of the CI infrastructure. This requirement is subjective, to be evaluated by the infrastructure team, and will take the community importance of the target into account.
- Tier 2 targets should, if at all possible, support cross-compiling. Tier 2 targets should not require using the target as the host for builds, even if the target supports host tools.
- In addition to the legal requirements for all targets (specified in the tier 3 requirements), because a tier 2 target typically involves the Rust project building and supplying various compiled binaries, incorporating the target and redistributing any resulting compiled binaries (e.g. built libraries, host tools if any) must not impose any onerous license requirements on any members of the Rust project, including infrastructure team members and those operating CI systems. This is a subjective requirement, to be evaluated by the approving teams.
 - As an exception to this, if the target's primary purpose is to build components for a Free and Open Source Software (FOSS) project licensed under "copyleft" terms (terms which require licensing other code under compatible FOSS terms), such as kernel modules or plugins, then the standard libraries for the target may potentially be subject to copyleft terms, as long as such terms are satisfied by Rust's existing practices of providing full corresponding source code. Note that anything added to the Rust repository itself must still use Rust's standard license terms.
- Tier 2 targets must not impose burden on the authors of pull requests, or other developers in the community, to ensure that tests pass for the target. In particular, do not post comments (automated or manual) on a PR that derail or suggest a block on the PR based on tests failing for the target. Do not send automated messages or notifications (via any medium, including via `@`) to a PR author or others involved with a PR regarding the PR breaking tests on a tier 2 target, unless they have opted into such messages.
 - Backlinks such as those generated by the issue/PR tracker when linking to an issue or PR are not considered a violation of this policy, within

reason. However, such messages (even on a separate repository) must not generate notifications to anyone involved with a PR who has not requested such notifications.

- The target maintainers should regularly run the testsuite for the target, and should fix any test failures in a reasonably timely fashion.
- All requirements for tier 3 apply.

A tier 2 target may be demoted or removed if it no longer meets these requirements. Any proposal for demotion or removal will be CCed to the target maintainers, and will be communicated widely to the Rust community before being dropped from a stable release. (The amount of time between such communication and the next stable release may depend on the nature and severity of the failed requirement, the timing of its discovery, whether the target has been part of a stable release yet, and whether the demotion or removal can be a planned and scheduled action.)

In some circumstances, especially if the target maintainers do not respond in a timely fashion, Rust teams may land pull requests that temporarily disable some targets in the nightly compiler, in order to implement a feature not yet supported by those targets. (As an example, this happened when introducing the 128-bit types `u128` and `i128`.) Such a pull request will include notification and coordination with the maintainers of such targets, and will ideally happen towards the beginning of a new development cycle to give maintainers time to update their targets. The maintainers of such targets will then be expected to implement the corresponding target-specific support in order to re-enable the target. If the maintainers of such targets cannot provide such support in time for the next stable release, this may result in demoting or removing the targets.

Tier 2 with host tools

Some tier 2 targets may additionally have binaries built to run on them as a host (such as `rustc` and `cargo`). This allows the target to be used as a development platform, not just a compilation target.

A proposed new tier 2 target with host tools must be reviewed and approved by the compiler team based on these requirements. Such review and approval may occur via a [Major Change Proposal \(MCP\)](#).

In addition, the infrastructure team must approve the integration of the target's host tools into Continuous Integration (CI), and the CI-related requirements for host tools. This review and approval may take place in a PR adding the target's host tools to CI, or simply by an infrastructure team member reporting the outcome of a team discussion.

- Depending on the target, its capabilities, its performance, and the likelihood of use for any given tool, the host tools provided for a tier 2 target may include only `rustc` and `cargo`, or may include additional tools such as `clippy` and `rustfmt`.

- Approval of host tools will take into account the additional time required to build the host tools, and the substantial additional storage required for the host tools.
- The host tools must have direct value to people other than the target's maintainers. (It may still be a niche target, but the host tools must not be exclusively useful for an inherently closed group.) This requirement will be evaluated independently from the corresponding tier 2 requirement.
 - The requirement to provide "direct value" means that it does not suffice to argue that having host tools will help the target's maintainers more easily provide the target to others. The tools themselves must provide value to others.
- There must be a reasonable expectation that the host tools will be used, for purposes other than to prove that they can be used.
- The host tools must build and run reliably in CI (for all components that Rust's CI considers mandatory), though they may or may not pass tests.
- Building host tools for the target must not take substantially longer than building host tools for other targets, and should not substantially raise the maintenance burden of the CI infrastructure.
- The host tools must provide a substantively similar experience as on other targets, subject to reasonable target limitations.
 - Adding a substantively different interface to an existing tool, or a target-specific interface to the functionality of an existing tool, requires design and implementation approval (e.g. RFC/MCP) from the appropriate approving teams for that tool.
 - Such an interface should have a design that could potentially work for other targets with similar properties.
 - This should happen separately from the review and approval of the target, to simplify the target review and approval processes, and to simplify the review and approval processes for the proposed new interface.
 - By way of example, a target that runs within a sandbox may need to modify the handling of files, tool invocation, and similar to meet the expectations and conventions of the sandbox, but must not introduce a separate "sandboxed compilation" interface separate from the CLI interface without going through the normal approval process for such an interface. Such an interface should take into account potential other targets with similar sandboxes.
- If the host tools for the platform would normally be expected to be signed or equivalent (e.g. if running unsigned binaries or similar involves a "developer mode" or an additional prompt), it must be possible for the Rust project's automated builds to apply the appropriate signature process, without any manual intervention by either Rust developers, target maintainers, or a third party. This process must meet the approval of the infrastructure team.
 - This process may require one-time or semi-regular manual steps by the infrastructure team, such as registration or renewal of a signing key. Any such manual process must meet the approval of the infrastructure team.
 - This process may require the execution of a legal agreement with the signature provider. Such a legal agreement may be revocable, and may

potentially require a nominal fee, but must not be otherwise onerous. Any such legal agreement must meet the approval of the infrastructure team.

(The infrastructure team is not expected or required to sign binding legal agreements on behalf of the Rust project; this review and approval exists to ensure no terms are onerous or cause problems for infrastructure, especially if such terms may impose requirements or obligations on people who have access to target-specific infrastructure.)

- Changes to this process, or to any legal agreements involved, may cause a target to stop meeting this requirement.
- This process involved must be available under substantially similar non-onerous terms to the general public. Making it available exclusively to the Rust project does not suffice.
- This requirement exists to ensure that Rust builds, including nightly builds, can meet the necessary requirements to allow users to smoothly run the host tools.
- Providing host tools does not exempt a target from requirements to support cross-compilation if at all possible.
- All requirements for tier 2 apply.

A target may be promoted directly from tier 3 to tier 2 with host tools if it meets all the necessary requirements, but doing so may introduce substantial additional complexity. If in doubt, the target should qualify for tier 2 without host tools first.

Tier 1 target policy

At this tier, the Rust project guarantees that a target builds and passes all tests, and will reject patches that fail to build or pass the testsuite on a target. We hold tier 1 targets to our highest standard of requirements.

A proposed new tier 1 target must be reviewed and approved by the compiler team based on these requirements. In addition, the release team must approve the viability and value of supporting the target. For a tier 1 target, this will typically take place via a full RFC proposing the target, to be jointly reviewed and approved by the compiler team and release team.

In addition, the infrastructure team must approve the integration of the target into Continuous Integration (CI), and the tier 1 CI-related requirements. This review and approval may take place in a PR adding the target to CI, by an infrastructure team member reporting the outcome of a team discussion, or by including the infrastructure team in the RFC proposing the target.

- Tier 1 targets must have substantial, widespread interest within the developer community, and must serve the ongoing needs of multiple production users of Rust across multiple organizations or projects. These requirements are subjective, and determined by consensus of the approving teams. A tier 1 target

may be demoted or removed if it becomes obsolete or no longer meets this requirement.

- The target maintainer team must include at least 3 developers.
- The target must build and pass tests reliably in CI, for all components that Rust's CI considers mandatory.
 - The target must not disable an excessive number of tests or pieces of tests in the testsuite in order to do so. This is a subjective requirement.
 - If the target does not have host tools support, or if the target has low performance, the infrastructure team may choose to have CI cross-compile the testsuite from another platform, and then run the compiled tests either natively or via accurate emulation. However, the approving teams may take such performance considerations into account when determining the viability of the target or of its host tools.
- The target must provide as much of the Rust standard library as is feasible and appropriate to provide. For instance, if the target can support dynamic memory allocation, it must provide an implementation of `alloc` and the associated data structures.
- Building the target and running the testsuite for the target must not take substantially longer than other targets, and should not substantially raise the maintenance burden of the CI infrastructure.
 - In particular, if building the target takes a reasonable amount of time, but the target cannot run the testsuite in a timely fashion due to low performance of either native code or accurate emulation, that alone may prevent the target from qualifying as tier 1.
- If running the testsuite requires additional infrastructure (such as physical systems running the target), the target maintainers must arrange to provide such resources to the Rust project, to the satisfaction and approval of the Rust infrastructure team.
 - Such resources may be provided via cloud systems, via emulation, or via physical hardware.
 - If the target requires the use of emulation to meet any of the tier requirements, the approving teams for those requirements must have high confidence in the accuracy of the emulation, such that discrepancies between emulation and native operation that affect test results will constitute a high-priority bug in either the emulation or the implementation of the target.
 - If it is not possible to run the target via emulation, these resources must additionally be sufficient for the Rust infrastructure team to make them available for access by Rust team members, for the purposes of development and testing. (Note that the responsibility for doing target-specific development to keep the target well maintained remains with the target maintainers. This requirement ensures that it is possible for other Rust developers to test the target, but does not obligate other Rust developers to make target-specific fixes.)
 - Resources provided for CI and similar infrastructure must be available for continuous exclusive use by the Rust project. Resources provided for access by Rust team members for development and testing must be available on an

exclusive basis when in use, but need not be available on a continuous basis when not in use.

- Tier 1 targets must not have a hard requirement for signed, verified, or otherwise "approved" binaries. Developers must be able to build, run, and test binaries for the target on systems they control, or provide such binaries for others to run. (Doing so may require enabling some appropriate "developer mode" on such systems, but must not require the payment of any additional fee or other consideration, or agreement to any onerous legal agreements.)
 - The Rust project may decide to supply appropriately signed binaries if doing so provides a smoother experience for developers using the target, and a tier 2 target with host tools already requires providing appropriate mechanisms that enable our infrastructure to provide such signed binaries. However, this additional tier 1 requirement ensures that Rust developers can develop and test Rust software for the target (including Rust itself), and that development or testing for the target is not limited.
- All requirements for tier 2 apply.

A tier 1 target may be demoted if it no longer meets these requirements but still meets the requirements for a lower tier. Any proposal for demotion of a tier 1 target requires a full RFC process, with approval by the compiler and release teams. Any such proposal will be communicated widely to the Rust community, both when initially proposed and before being dropped from a stable release. A tier 1 target is highly unlikely to be directly removed without first being demoted to tier 2 or tier 3. (The amount of time between such communication and the next stable release may depend on the nature and severity of the failed requirement, the timing of its discovery, whether the target has been part of a stable release yet, and whether the demotion or removal can be a planned and scheduled action.)

Raising the baseline expectations of a tier 1 target (such as the minimum CPU features or OS version required) requires the approval of the compiler and release teams, and should be widely communicated as well, but does not necessarily require a full RFC.

Tier 1 with host tools

Some tier 1 targets may additionally have binaries built to run on them as a host (such as `rustc` and `cargo`). This allows the target to be used as a development platform, not just a compilation target.

A proposed new tier 1 target with host tools must be reviewed and approved by the compiler team based on these requirements. In addition, the release team must approve the viability and value of supporting host tools for the target. For a tier 1 target, this will typically take place via a full RFC proposing the target, to be jointly reviewed and approved by the compiler team and release team.

In addition, the infrastructure team must approve the integration of the target's host tools into Continuous Integration (CI), and the CI-related requirements for host tools. This review and approval may take place in a PR adding the target's host

tools to CI, by an infrastructure team member reporting the outcome of a team discussion, or by including the infrastructure team in the RFC proposing the target.

- Tier 1 targets with host tools should typically include all of the additional tools such as `clippy` and `rustfmt`, unless there is a target-specific reason why a tool cannot possibly make sense for the target.
 - Unlike with tier 2, for tier 1 we will not exclude specific tools on the sole basis of them being less likely to be used; rather, we'll take that into account when considering whether the target should be at tier 1 with host tools. In general, on any tier 1 target with host tools, people should be able to expect to find and install all the same components that they would for any other tier 1 target with host tools.
- Approval of host tools will take into account the additional time required to build the host tools, and the substantial additional storage required for the host tools.
- Host tools for the target must have substantial, widespread interest within the developer community, and must serve the ongoing needs of multiple production users of Rust across multiple organizations or projects. These requirements are subjective, and determined by consensus of the approving teams. This requirement will be evaluated independently from the corresponding tier 1 requirement; it is possible for a target to have sufficient interest for cross-compilation, but not have sufficient interest for native compilation. The host tools may be dropped if they no longer meet this requirement, even if the target otherwise qualifies as tier 1.
- The host tools must build, run, and pass tests reliably in CI, for all components that Rust's CI considers mandatory.
 - The target must not disable an excessive number of tests or pieces of tests in the testsuite in order to do so. This is a subjective requirement.
- Building the host tools and running the testsuite for the host tools must not take substantially longer than other targets, and should not substantially raise the maintenance burden of the CI infrastructure.
 - In particular, if building the target's host tools takes a reasonable amount of time, but the target cannot run the testsuite in a timely fashion due to low performance of either native code or accurate emulation, that alone may prevent the target from qualifying as tier 1 with host tools.
- Providing host tools does not exempt a target from requirements to support cross-compilation if at all possible.
- All requirements for tier 2 targets with host tools apply.
- All requirements for tier 1 apply.

A target seeking promotion to tier 1 with host tools should typically either be tier 2 with host tools or tier 1 without host tools, to reduce the number of requirements to simultaneously review and approve.

In addition to the general process for demoting a tier 1 target, a tier 1 target with host tools may be demoted (including having its host tools dropped, or being demoted to tier 2 with host tools) if it no longer meets these requirements but still meets the requirements for a lower tier. Any proposal for demotion of a tier 1 target (with or without host tools) requires a full RFC process, with approval by

the compiler and release teams. Any such proposal will be communicated widely to the Rust community, both when initially proposed and before being dropped from a stable release.

target-name-here

Tier: 3

One-sentence description of the target (e.g. CPU, OS)

Target maintainers

- Some Person, email@example.org, <https://github.com/> ...

Requirements

Does the target support host tools, or only cross-compilation? Does the target support std, or alloc (either with a default allocator, or if the user supplies an allocator)?

Document the expectations of binaries built for the target. Do they assume specific minimum features beyond the baseline of the CPU/environment/etc? What version of the OS or environment do they expect?

Are there notable `#[target_feature(...)]` or `-C target-feature=` values that programs may wish to use?

What calling convention does `extern "C"` use on the target?

What format do binaries use by default? ELF, PE, something else?

Building the target

If Rust doesn't build the target by default, how can users build it? Can users just add it to the `target` list in `config.toml`?

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target. To compile for this target, you will either need to build Rust with the target enabled (see "Building the target" above), or build your own copy of `core` by using `build-std` or similar.

Testing

Does the target support running binaries, or do binaries have varying expectations that prevent having a standard way to run them? If users can run binaries, can they do so in some common emulator, or do they need native hardware? Does the target support running the Rust testsuite?

Cross-compilation toolchains and C code

Does the target support C code? If so, what toolchain target should users use to build compatible C code? (This may match the target triple, or it may be a toolchain for a different target triple, potentially with specific options or caveats.)

aarch64-apple-ios-sim

Tier: 2

Apple iOS Simulator on ARM64.

Designated Developers

- [@badboy](#)
- [@deg4uss3r](#)

Requirements

This target is cross-compiled. To build this target Xcode 12 or higher on macOS is required.

Building

The target can be built by enabling it for a `rustc` build:

```
[build]
build-stage = 1
target = ["aarch64-apple-ios-sim"]
```

Cross-compilation

This target can be cross-compiled from `x86_64` or `aarch64` macOS hosts.

Other hosts are not supported for cross-compilation, but might work when also providing the required Xcode SDK.

Testing

Currently there is no support to run the `rustc` test suite for this target.

Building Rust programs

Note: Building for this target requires the corresponding iOS SDK, as provided by Xcode 12+.

From Rust Nightly 1.56.0 (2021-08-03) on the artifacts are shipped pre-compiled:

```
rustup target add aarch64-apple-ios-sim --toolchain nightly
```

Rust programs can be built for that target:

```
rustc --target aarch64-apple-ios-sim your-code.rs
```

There is no easy way to run simple programs in the iOS simulator. Static library builds can be embedded into iOS applications.

*-apple-watchos

- arm64_32-apple-watchos
- armv7k-apple-watchos
- aarch64-apple-watchos-sim
- x86_64-apple-watchos-sim

Tier: 3

Apple WatchOS targets:

- Apple WatchOS on Arm 64_32
- Apple WatchOS on Arm v7k
- Apple WatchOS Simulator on arm64
- Apple WatchOS Simulator on x86_64

Target maintainers

- [@deg4uss3r](#)
- [@vladimir-ea](#)

Requirements

These targets are cross-compiled. To build these targets Xcode 12 or higher on macOS is required.

Building the target

The targets can be built by enabling them for a `rustc` build, for example:

```
[build]
build-stage = 1
target = ["aarch64-apple-watchos-sim"]
```

Building Rust programs

Note: Building for this target requires the corresponding WatchOS SDK, as provided by Xcode 12+.

Rust programs can be built for these targets, if `rustc` has been built with support for them, for example:

```
rustc --target aarch64-apple-watchos-sim your-code.rs
```

Testing

There is no support for running the Rust testsuite on WatchOS or the simulators.

There is no easy way to run simple programs on WatchOS or the WatchOS simulators. Static library builds can be embedded into WatchOS applications.

Cross-compilation toolchains and C code

This target can be cross-compiled from `x86_64` or `aarch64` macOS hosts.

Other hosts are not supported for cross-compilation, but might work when also providing the required Xcode SDK.

aarch64-nintendo-switch-freestanding

Tier: 3

Nintendo Switch with pure-Rust toolchain.

Designated Developers

- [@leo60228](#)
- [@jam1garner](#)

Requirements

This target is cross-compiled. It has no special requirements for the host.

Building

The target can be built by enabling it for a `rustc` build:

```
[build]
build-stage = 1
target = ["aarch64-nintendo-switch-freestanding"]
```

Cross-compilation

This target can be cross-compiled from any host.

Testing

Currently there is no support to run the `rustc` test suite for this target.

Building Rust programs

If `rustc` has support for that target and the library artifacts are available, then Rust programs can be built for that target:

```
rustc --target aarch64-nintendo-switch-freestanding your-code.rs
```

To generate binaries in the NRO format that can be easily run on-device, you can use [cargo-nx](#):

```
cargo nx --triple=aarch64-nintendo-switch-freestanding
```

armeb-unknown-linux-gnueabi

Tier: 3

Target for cross-compiling Linux user-mode applications targeting the ARM BE8 architecture.

Overview

BE8 architecture retains the same little-endian ordered code-stream used by conventional little endian ARM systems, however the data accesses are in big-endian. BE8 is used primarily in high-performance networking applications where the ability to read packets in their native "Network Byte Order" is important (many network protocols transmit data in big-endian byte order for their wire formats).

History

BE8 architecture is the default big-endian architecture for ARM since [ARMv6](#). Its predecessor, used for ARMv4 and ARMv5 devices was [BE32](#). On ARMv6 architecture, endianness can be configured via [system registers](#). However, BE32 was withdrawn for [ARMv7](#) onwards.

Target Maintainers

- [@WorksButNotTested](#)

Requirements

The target is cross-compiled. This target supports `std` in the normal way (indeed only nominal changes are required from the standard ARM configuration).

Target definition

The target definition can be seen [here](#). In particular, it should be noted that the `features` specify that this target is built for the ARMv8 core. Though this can likely be modified as required.

Building the target

Because it is Tier 3, rust does not yet ship pre-compiled artifacts for this target.

Therefore, you can build Rust with support for the target by adding it to the target list in config.toml, a sample configuration is shown below. It is expected that the user already have a working GNU compiler toolchain and update the paths accordingly.

```
[llvm]
download-ci-llvm = false
optimize = true
ninja = true
targets = "ARM;X86"
clang = false

[build]
target = ["x86_64-unknown-linux-gnu", "armeb-unknown-linux-gnueabi"]
docs = false
docs-minification = false
compiler-docs = false
[install]
prefix = "/home/user/x-tools/rust/"

[rust]
debug-logging=true
backtrace = true
incremental = true

[target.x86_64-unknown-linux-gnu]

[dist]

[target.armeb-unknown-linux-gnueabi]
cc = "/home/user/x-tools/armeb-unknown-linux-gnueabi/bin/armeb-unknown-linux-gnueabi-gcc"
cxx = "/home/user/x-tools/armeb-unknown-linux-gnueabi/bin/armeb-unknown-linux-gnueabi-g++"
ar = "/home/user/x-tools/armeb-unknown-linux-gnueabi/bin/armeb-unknown-linux-gnueabi-ar"
ranlib = "/home/user/x-tools/armeb-unknown-linux-gnueabi/bin/armeb-unknown-linux-gnueabi-ranlib"
linker = "/home/user/x-tools/armeb-unknown-linux-gnueabi/bin/armeb-unknown-linux-gnueabi-gcc"
llvm-config = "/home/user/x-tools/clang/bin/llvm-config"
llvm-filecheck = "/home/user/x-tools/clang/bin/FileCheck"
```

Building Rust programs

The following `.cargo/config` is needed inside any project directory to build for the BE8 target:

```
[build]
target = "armeb-unknown-linux-gnueabi"

[target.armeb-unknown-linux-gnueabi]
linker = "armeb-unknown-linux-gnueabi-gcc"
```

Note that it is expected that the user has a suitable linker from the GNU toolchain.

armv4t-none-eabi

Tier 3

Bare-metal target for any cpu in the ARMv4T architecture family, supporting ARM/Thumb code interworking (aka `a32/t32`), with ARM code as the default code generation.

In particular this supports the Gameboy Advance (GBA), but there's nothing GBA specific with this target, so any ARMv4T device should work fine.

Target Maintainers

- [@Lokathor](#)

Requirements

The target is cross-compiled, and uses static linking.

This target doesn't provide a linker script, you'll need to bring your own according to the specific device you want to target. Pass `-Clink-arg=-Tyour_script.ld` as a rustc argument to make the linker use `your_script.ld` during linking.

Building Rust Programs

Because it is Tier 3, rust does not yet ship pre-compiled artifacts for this target.

Just use the `build-std` nightly cargo feature to build the `core` library. You can pass this as a command line argument to cargo, or your `.cargo/config.toml` file might include the following lines:

```
[unstable]
build-std = ["core"]
```

Most of `core` should work as expected, with the following notes:

- the target is "soft float", so `f32` and `f64` operations are emulated in software.
- integer division is also emulated in software.
- the target is old enough that it doesn't have atomic instructions.

Rust programs are output as ELF files.

For running on hardware, you'll generally need to extract the "raw" program code out of the ELF and into a file of its own. The `objcopy` program provided as part of the GNU Binutils can do this:

```
arm-none-eabi-objcopy --output-target binary [in_file] [out_file]
```

Testing

This is a cross-compiled target that you will need to emulate during testing.

Because this is a device-agnostic target, and the exact emulator that you'll need depends on the specific device you want to run your code on.

For example, when programming for the Gameboy Advance, the `mgba-test-runner` program could be used to make a normal set of rust tests be run within the `mgba` emulator.

armv5te-none-eabi

Tier: 3

Bare-metal target for any cpu in the ARMv5TE architecture family, supporting ARM/Thumb code interworking (aka `a32/t32`), with `a32` code as the default code generation.

The `thumbv5te-none-eabi` target is the same as this one, but the instruction set defaults to `t32`.

Target Maintainers

- [@QuinnPainter](#)

Requirements

The target is cross-compiled, and uses static linking.

By default, the `lld` linker included with Rust will be used.

However, you may want to use the `arm-none-eabi-ld` linker instead. This can be obtained for Windows/Mac/Linux from the [ARM Developer Website](#), or possibly from your OS's package manager. To use it, add the following to your `.cargo/config.toml`:

```
[target.armv5te-none-eabi]
linker = "arm-none-eabi-ld"
```

This target doesn't provide a linker script, you'll need to bring your own according to the specific device you want to target. Pass `-Clink-arg=-Tyour_script.ld` as a rustc argument to make the linker use `your_script.ld` during linking.

Building Rust Programs

Because it is Tier 3, rust does not yet ship pre-compiled artifacts for this target.

Just use the `build-std` nightly cargo feature to build the `core` library. You can pass this as a command line argument to cargo, or your `.cargo/config.toml` file might include the following lines:

```
[unstable]
build-std = ["core"]
```

Most of `core` should work as expected, with the following notes:

- the target is "soft float", so `f32` and `f64` operations are emulated in software.
- integer division is also emulated in software.
- the target is old enough that it doesn't have atomic instructions.

`alloc` is also supported, as long as you provide your own global allocator.

Rust programs are output as ELF files.

Testing

This is a cross-compiled target that you will need to emulate during testing.

Because this is a device-agnostic target, and the exact emulator that you'll need depends on the specific device you want to run your code on.

For example, when programming for the DS, you can use one of the several available DS emulators, such as [melonDS](#).

armv6k-nintendo-3ds

Tier: 3

The Nintendo 3DS platform, which has an ARMv6K processor, and its associated operating system ([horizon](#)).

Rust support for this target is not affiliated with Nintendo, and is not derived from nor used with any official Nintendo SDK.

Target maintainers

- [@Meziu](#)
- [@AzureMarker](#)
- [@ian-h-chamberlain](#)

Requirements

This target is cross-compiled. Dynamic linking is not supported.

`#![no_std]` crates can be built using `build-std` to build `core` and optionally `alloc`, and either `panic_abort` or `panic_unwind`.

`std` is partially supported, but mostly works. Some APIs are unimplemented and will simply return an error, such as `std::process`. An allocator is provided by default.

In order to support some APIs, binaries must be linked against `libc` written for the target, using a linker for the target. These are provided by the devkitARM toolchain. See [Cross-compilation toolchains and C code](#) for more details.

Additionally, some helper crates provide implementations of some `libc` functions used by `std` that may otherwise be missing. These, or an alternate implementation of the relevant functions, are required to use `std`:

- `pthread-3ds` provides pthread APIs for `std::thread`.
- `linker-fix-3ds` fulfills some other missing `libc` APIs.

Binaries built for this target should be compatible with all variants of the 3DS (and 2DS) hardware and firmware, but testing is limited and some versions may not work correctly.

This target generates binaries in the ELF format.

Building the target

You can build Rust with support for the target by adding it to the `target` list in `config.toml` and providing paths to the devkitARM toolchain.

```
[build]
build-stage = 1
target = ["armv6k-nintendo-3ds"]

[target.armv6k-nintendo-3ds]
cc = "/opt/devkitpro/devkitARM/bin/arm-none-eabi-gcc"
cxx = "/opt/devkitpro/devkitARM/bin/arm-none-eabi-g++"
ar = "/opt/devkitpro/devkitARM/bin/arm-none-eabi-ar"
ranlib = "/opt/devkitpro/devkitARM/bin/arm-none-eabi-ranlib"
linker = "/opt/devkitpro/devkitARM/bin/arm-none-eabi-gcc"
```

Also, to build `compiler_builtins` for the target, export these flags before building the Rust toolchain:

```
export CFLAGS_armv6k_nintendo_3ds="-mfloat-abi=hard -mtune=mpcore -mtp=soft -march=armv6k"
```

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target.

The recommended way to build binaries is by using the `cargo-3ds` tool, which uses `build-std` and provides commands that work like the usual `cargo run`, `cargo build`, etc.

You can also build Rust with the target enabled (see [Building the target](#) above).

As mentioned in [Requirements](#), programs that use `std` must link against both the devkitARM toolchain and libraries providing the `libc` APIs used in `std`. There is a general-purpose utility crate for working with nonstandard APIs provided by the OS: `ctrus-rs`. Add it to `Cargo.toml` to use it in your program:

```
[dependencies]
ctrus-rs = { git = "https://github.com/Meizu/ctrus-rs.git" }
```

Using this library's `init()` function ensures the symbols needed to link against `std` are present (as mentioned in [Requirements](#) above), as well as providing a runtime suitable for `std`:

```
fn main() {
    ctrus::init();
}
```

Testing

Binaries built for this target can be run in an emulator (most commonly [Citra](#)), or sent to a device through the use of a tool like devkitARM's `3dslink`. They may also simply be copied to an SD card to be inserted in the device.

The `cargo-3ds` tool mentioned in [Building Rust programs](#) supports the use of `3dslink` with `cargo 3ds run`. The default Rust test runner is not supported, but [custom test frameworks](#) can be used with `cargo 3ds test` to run unit tests on a device.

The Rust test suite for `library/std` is not yet supported.

Cross-compilation toolchains and C code

C code can be built for this target using the [devkitARM toolchain](#). This toolchain provides `arm-none-eabi-gcc` as the linker used to link Rust programs as well.

The toolchain also provides a `libc` implementation, which is required by `std` for many of its APIs, and a helper library `libctru` which is used by several of the helper crates listed in [Requirements](#). This toolchain does not, however, include all of the APIs expected by `std`, and the remaining APIs are implemented by `pthread-3ds` and `linker-fix-3ds`.

armv7-sony-vita-newlibeabihf

armv7-unknown-linux-uclibceabi

Tier: 3

This target supports ARMv7 softfloat CPUs and uses the uclibc-ng standard library. This is a common configuration on many consumer routers (e.g., Netgear R7000, Asus RT-AC68U).

Target maintainers

- [@lancethepants](#)

Requirements

This target is cross compiled, and requires a cross toolchain.

This target supports host tools and std.

Building the target

You will need to download or build a `'c'` cross toolchain that targets ARMv7 softfloat and that uses the uclibc-ng standard library. If your target hardware is something like a router or an embedded device, keep in mind that manufacturer supplied SDKs for this class of CPU could be outdated and potentially unsuitable for bootstrapping rust.

Here is a sample toolchain that is built using `buildroot`. It uses modern toolchain components, older thus universal kernel headers (2.6.36.4), and is used for a project called `Tomatoware`. This toolchain is patched so that its sysroot is located at /mmc (e.g., /mmc/bin, /mmc/lib, /mmc/include). This is useful in scenarios where the root filesystem is read-only but you are able attach external storage loaded with user applications. Tomatoware is an example of this that even allows you to run various compilers and developer tools natively on the target device.

Utilizing the Tomatoware toolchain this target can be built for cross compilation and native compilation (host tools) with project

`rust-bootstrap-armv7-unknown-linux-uclibceabi`.

Here is a sample config if using your own toolchain.

```
[build]
build-stage = 2
target = ["armv7-unknown-linux-uclibcabi"]

[target.armv7-unknown-linux-uclibcabi]
cc = "/path/to/arm-unknown-linux-uclibcgnueabi-gcc"
cxx = "/path/to/arm-unknown-linux-uclibcgnueabi-g++"
ar = "path/to/arm-unknown-linux-uclibcgnueabi-ar"
ranlib = "path/to/arm-unknown-linux-uclibcgnueabi-ranlib"
linker = "/path/to/arm-unknown-linux-uclibcgnueabi-gcc"
```

Building Rust programs

The following assumes you are using the Tomatoware toolchain and environment. Adapt if you are using your own toolchain.

Native compilation

Since this target supports host tools, you can natively build rust applications directly on your target device. This can be convenient because it removes the complexities of cross compiling and you can immediately test and deploy your binaries. One downside is that compiling on your ARMv7 CPU will probably be much slower than cross compilation on your x86 machine.

To setup native compilation:

- Download Tomatoware to your device using the latest nightly release found [here](#).
- Extract `tar zxvf arm-soft mmc.tgz -C /mmc`
- Add `/mmc/bin:/mmc:sbin/` to your PATH, or `source /mmc/etc/profile`
- `apt update && apt install rust`

If you bootstrap rust on your own using the project above, it will create a .deb file that you then can install with

```
dpkg -i rust_1.xx.x-x_arm.deb
```

After completing these steps you can use rust normally in a native environment.

Cross Compilation

To cross compile, you'll need to:

- Build the rust cross toolchain using [rust-bootstrap-armv7-unknown-linux-uclibcabi](#) or your own built toolchain.
- Link your built toolchain with

```
rustup toolchain link stage2 \
${HOME}/rust-bootstrap-armv7-unknown-linux-uclibcabi/src/rust/rust/build/x86_64-
unknown-linux-gnu/stage2
```

- Build with:

```
CC_armv7_unknown_linux_uclibcabi=/opt/tomatoware/arm-soft-mmcs/bin/arm-linux-gcc \
CXX_armv7_unknown_linux_uclibcabi=/opt/tomatoware/arm-soft-mmcs/bin/arm-linux-g++ \
AR_armv7_unknown_linux_uclibcabi=/opt/tomatoware/arm-soft-mmcs/bin/arm-linux-ar \
CFLAGS_armv7_unknown_linux_uclibcabi="-march=armv7-a -mtune=cortex-a9" \
CXXFLAGS_armv7_unknown_linux_uclibcabi="-march=armv7-a -mtune=cortex-a9" \
CARGO_TARGET_ARMV7_UNKNOWN_LINUX_UCLIBCEABI_LINKER=/opt/tomatoware/arm-soft-mmcs/bin/arm-
linux-gcc \
CARGO_TARGET_ARMV7_UNKNOWN_LINUX_UCLIBCEABI_RUSTFLAGS='-Clink-arg=-s -Clink-arg=-Wl,--
dynamic-linker=/mmc/lib/ld-uClibc.so.1 -Clink-arg=-Wl,-rpath,/mmc/lib' \
cargo +stage2 \
build \
--target armv7-unknown-linux-uclibcabi \
--release
```

- Copy the binary to your target device and run.

We specify `CC`, `CXX`, `AR`, `CFLAGS`, and `CXXFLAGS` environment variables because sometimes a project or a subproject requires the use of your '`'c'`' cross toolchain. Since Tomatoware has a modified sysroot we also pass via `RUSTFLAGS` the location of the dynamic-linker and `rpath`.

Test with QEMU

To test a cross-compiled binary on your build system follow the instructions for [Cross Compilation](#), install `qemu-arm-static`, and run with the following.

```
CC_armv7_unknown_linux_uclibcabi=/opt/tomatoware/arm-soft-mmcs/bin/arm-linux-gcc \
CXX_armv7_unknown_linux_uclibcabi=/opt/tomatoware/arm-soft-mmcs/bin/arm-linux-g++ \
AR_armv7_unknown_linux_uclibcabi=/opt/tomatoware/arm-soft-mmcs/bin/arm-linux-ar \
CFLAGS_armv7_unknown_linux_uclibcabi="-march=armv7-a -mtune=cortex-a9" \
CXXFLAGS_armv7_unknown_linux_uclibcabi="-march=armv7-a -mtune=cortex-a9" \
CARGO_TARGET_ARMV7_UNKNOWN_LINUX_UCLIBCEABI_LINKER=/opt/tomatoware/arm-soft-mmcs/bin/arm-linux-
gcc \
CARGO_TARGET_ARMV7_UNKNOWN_LINUX_UCLIBCEABI_RUNNER="qemu-arm-static -L /opt/tomatoware/arm-
soft-mmcs/arm-tomatoware-linux-uclibcgnueabi/sysroot/" \
cargo +stage2 \
run \
--target armv7-unknown-linux-uclibcabi \
--release
```

Run in a chroot

It's also possible to build in a chroot environment. This is a convenient way to work without needing to access the target hardware.

To build the chroot:

- `sudo debootstrap --arch armel bullseye $HOME/debian`
- `sudo chroot $HOME/debian/ /bin/bash`
- `mount proc /proc -t proc`
- `mount -t sysfs /sys sys/`
- `export PATH=/mmc/bin:/mmc/sbin:$PATH`

From here you can setup your environment (e.g., add user, install wget).

- Download Tomatoware to the chroot environment using the latest nightly release found [here](#).
- Extract `tar zxvf arm-soft mmc.tgz -C /mmc`
- Add `/mmc/bin:/mmc/sbin/` to your PATH, or `source /mmc/etc/profile`
- `sudo /mmc/bin/apt update && sudo /mmc/bin/apt install rust`

After completing these steps you can use rust normally in a chroot environment.

Remember when using `sudo` the root user's PATH could differ from your user's PATH.

armv7-unknown-linux-uclibceabihf

Tier: 3

This tier supports the ARMv7 processor running a Linux kernel and uClibc-ng standard library. It provides full support for rust and the rust standard library.

Designated Developers

- [@skrap](#)

Requirements

This target is cross compiled, and requires a cross toolchain. You can find suitable pre-built toolchains at [bootlin](#) or build one yourself via [buildroot](#).

Building

Get a C toolchain

Compiling rust for this target has been tested on `x86_64` linux hosts. Other host types have not been tested, but may work, if you can find a suitable cross compilation toolchain for them.

If you don't already have a suitable toolchain, download one [here](#), and unpack it into a directory.

Configure rust

The target can be built by enabling it for a `rustc` build, by placing the following in `config.toml`:

```
[build]
target = ["armv7-unknown-linux-uclibceabihf"]
stage = 2

[target.armv7-unknown-linux-uclibceabihf]
# ADJUST THIS PATH TO POINT AT YOUR TOOLCHAIN
cc = "/TOOLCHAIN_PATH/bin/arm-buildroot-linux-uclibcgnueabihf-gcc"
```

Build

```
# in rust dir  
./x.py build --stage 2
```

Building and Running Rust Programs

To test cross-compiled binaries on a `x86_64` system, you can use the `qemu-arm` [userspace emulation](#) program. This avoids having a full emulated ARM system by doing dynamic binary translation and dynamic system call translation. It lets you run ARM programs directly on your `x86_64` kernel. It's very convenient!

To use:

- Install `qemu-arm` according to your distro.
- Link your built toolchain via:
 - `rustup toolchain link stage2 ${RUST}/build/x86_64-unknown-linux-gnu/stage2`
- Create a test program

```
cargo new hello_world  
cd hello_world
```

- Build and run

```
CARGO_TARGET_ARMV7_UNKNOWN_LINUX_UCLIBCEABIHF_RUNNER="qemu-arm -L ${TOOLCHAIN}/arm-buildroot-  
linux-uclibcgnueabihf/sysroot/" \  
CARGO_TARGET_ARMV7_UNKNOWN_LINUX_UCLIBCEABIHF_LINKER=${TOOLCHAIN}/bin/arm-buildroot-linux-  
uclibcgnueabihf-gcc \  
cargo +stage2 run --target armv7-unknown-linux-uclibcabihi
```

*-linux-android and *-linux-androideabi

Tier: 2

Android is a mobile operating system built on top of the Linux kernel.

Target maintainers

- Chris Wailes ([@chriswailes](#))
- Matthew Maurer ([@maurer](#))
- Martin Geisler ([@mgeisler](#))

Requirements

This target is cross-compiled from a host environment. Development may be done from the [source tree](#) or using the Android NDK.

Android targets support std. Generated binaries use the ELF file format.

NDK/API Update Policy

Rust will support the most recent Long Term Support (LTS) Android Native Development Kit (NDK). By default Rust will support all API levels supported by the NDK, but a higher minimum API level may be required if deemed necessary.

Building the target

To build Rust binaries for Android you'll need a copy of the most recent LTS edition of the [Android NDK](#). Supported Android targets are:

- aarch64-linux-android
- arm-linux-androideabi
- armv7-linux-androideabi
- i686-linux-android
- thumbv7neon-linux-androideabi
- x86_64-linux-android

A list of all supported targets can be found [here](#)

*-unknown-linux-ohos

Tier: 3

Targets for the [OpenHarmony](#) operating system.

Target maintainers

- Amanieu d'Antras ([@Amanieu](#))

Setup

The OpenHarmony SDK doesn't currently support Rust compilation directly, so some setup is required.

First, you must obtain the OpenHarmony SDK from [this page](#). Select the version of OpenHarmony you are developing for and download the "Public SDK package for the standard system".

Create the following shell scripts that wrap Clang from the OpenHarmony SDK:

aarch64-unknown-linux-ohos-clang.sh

```
#!/bin/sh
exec /path/to/ohos-sdk/linux/native/llvm/bin/clang \
-target aarch64-linux-ohos \
--sysroot=/path/to/ohos-sdk/linux/native/sysroot \
-D__MUSL__ \
"$@"
```

aarch64-unknown-linux-ohos-clang++.sh

```
#!/bin/sh
exec /path/to/ohos-sdk/linux/native/llvm/bin/clang++ \
-target aarch64-linux-ohos \
--sysroot=/path/to/ohos-sdk/linux/native/sysroot \
-D__MUSL__ \
"$@"
```

armv7-unknown-linux-ohos-clang.sh

```
#!/bin/sh
exec /path/to/ohos-sdk/linux/native/llvm/bin/clang \
-ttarget arm-linux-ohos \
--sysroot=/path/to/ohos-sdk/linux/native/sysroot \
-D__MUSL__ \
-march=armv7-a \
-mfloating-abi=softfp \
-mtune=generic-armv7-a \
-mthumb \
"$@"
```

`armv7-unknown-linux-ohos-clang++.sh`

```
#!/bin/sh
exec /path/to/ohos-sdk/linux/native/llvm/bin/clang++ \
-ttarget arm-linux-ohos \
--sysroot=/path/to/ohos-sdk/linux/native/sysroot \
-D__MUSL__ \
-march=armv7-a \
-mfloating-abi=softfp \
-mtune=generic-armv7-a \
-mthumb \
"$@"
```

Future versions of the OpenHarmony SDK will avoid the need for this process.

Building the target

To build a rust toolchain, create a `config.toml` with the following contents:

```
profile = "compiler"
changelog-seen = 2

[build]
sanitizers = true
profiler = true

[target.aarch64-unknown-linux-ohos]
cc = "/path/to/aarch64-unknown-linux-ohos-clang.sh"
cxx = "/path/to/aarch64-unknown-linux-ohos-clang++.sh"
ar = "/path/to/ohos-sdk/linux/native/llvm/bin/llvm-ar"
ranlib = "/path/to/ohos-sdk/linux/native/llvm/bin/llvm-ranlib"
linker = "/path/to/aarch64-unknown-linux-ohos-clang.sh"

[target.armv7-unknown-linux-ohos]
cc = "/path/to/armv7-unknown-linux-ohos-clang.sh"
cxx = "/path/to/armv7-unknown-linux-ohos-clang++.sh"
ar = "/path/to/ohos-sdk/linux/native/llvm/bin/llvm-ar"
ranlib = "/path/to/ohos-sdk/linux/native/llvm/bin/llvm-ranlib"
linker = "/path/to/armv7-unknown-linux-ohos-clang.sh"
```

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target. To compile for this target, you will either need to build Rust with the target enabled (see "Building the target" above), or build your own copy of `core` by using `build-std` or similar.

You will need to configure the linker to use in `~/.cargo/config`:

```
[target.aarch64-unknown-linux-ohos]
ar = "/path/to/ohos-sdk/linux/native/llvm/bin/llvm-ar"
linker = "/path/to/aarch64-unknown-linux-ohos-clang.sh"

[target.armv7-unknown-linux-ohos]
ar = "/path/to/ohos-sdk/linux/native/llvm/bin/llvm-ar"
linker = "/path/to/armv7-unknown-linux-ohos-clang.sh"
```

Testing

Running the Rust testsuite is possible, but currently difficult due to the way the OpenHarmony emulator is set up (no networking).

Cross-compilation toolchains and C code

You can use the shell scripts above to compile C code for the target.

aarch64-unknown-fuchsia and x86_64-unknown-fuchsia

Tier: 2

Fuchsia is a modern open source operating system that's simple, secure, updatable, and performant.

Target maintainers

The Fuchsia team:

- Tyler Mandry ([@tmandry](#))
- Dan Johnson ([@computerdruid](#))
- David Koloski ([@djkoloski](#))
- Joseph Ryan ([@P1n3appl3](#))

As the team evolves over time, the specific members listed here may differ from the members reported by the API. The API should be considered to be authoritative if this occurs. Instead of pinging individual members, use `@rustbot ping fuchsia` to contact the team on GitHub.

Table of contents

1. Requirements
2. Walkthrough structure
3. Compiling a Rust binary targeting Fuchsia
 1. Targeting Fuchsia with rustup and cargo
 2. Targeting Fuchsia with a compiler built from source
4. Creating a Fuchsia package
 1. Creating a Fuchsia component
 2. Building a Fuchsia package
5. Publishing a Fuchsia package
 1. Creating a Fuchsia package repository
 2. Publishing Fuchsia package to repository
6. Running a Fuchsia component on an emulator
 1. Starting the Fuchsia emulator
 2. Watching emulator logs
 3. Serving a Fuchsia package
 4. Running a Fuchsia component
7. `.gitignore` extensions
8. Testing
 1. Running unit tests

- 2. Running the compiler test suite
- 9. Debugging
 - 1. `zxdb`
 - 2. Attaching `zxdb`
 - 3. Using `zxdb`
 - 4. Displaying source code in `zxdb`

Requirements

This target is cross-compiled from a host environment. You will need a recent copy of the [Fuchsia SDK](#), which provides the tools, libraries, and binaries required to build and link programs for Fuchsia.

Development may also be done from the [source tree](#).

Fuchsia targets support `std` and follow the `sysv64` calling convention on `x86_64`. Fuchsia binaries use the ELF file format.

Walkthrough structure

This walkthrough will cover:

1. Compiling a Rust binary targeting Fuchsia.
2. Building a Fuchsia package.
3. Publishing and running a Fuchsia package to a Fuchsia emulator.

For the purposes of this walkthrough, we will only target `x86_64-unknown-fuchsia`.

Compiling a Rust binary targeting Fuchsia

Today, there are two main ways to build a Rust binary targeting Fuchsia using the Fuchsia SDK:

1. Allow `rustup` to handle the installation of Fuchsia targets for you.
2. Build a toolchain locally that can target Fuchsia.

Targeting Fuchsia with `rustup` and `cargo`

The easiest way to build a Rust binary targeting Fuchsia is by allowing `rustup` to handle the installation of Fuchsia targets for you. This can be done by issuing the following commands:

```
rustup target add x86_64-unknown-fuchsia
rustup target add aarch64-unknown-fuchsia
```

After installing our Fuchsia targets, we can now compile a Rust binary that targets Fuchsia.

To create our Rust project, we can use `cargo` as follows:

From base working directory

```
cargo new hello_fuchsia
```

The rest of this walkthrough will take place from `hello_fuchsia`, so we can change into that directory now:

```
cd hello_fuchsia
```

Note: From this point onwards, all commands will be issued from the `hello_fuchsia/` directory, and all `hello_fuchsia/` prefixes will be removed from references for sake of brevity.

We can edit our `src/main.rs` to include a test as follows:

src/main.rs

```
fn main() {
    println!("Hello Fuchsia!");
}

#[test]
fn it_works() {
    assert_eq!(2 + 2, 4);
}
```

In addition to the standard workspace created, we will want to create a `.cargo/config.toml` file to link necessary libraries during compilation:

.cargo/config.toml

```
[target.x86_64-unknown-fuchsia]

rustflags = [
    "-Lnative=<SDK_PATH>/arch/x64/lib",
    "-Lnative=<SDK_PATH>/arch/x64/sysroot/lib"
]
```

Note: Make sure to fill out `<SDK_PATH>` with the path to the downloaded [Fuchsia SDK](#).

These options configure the following:

- `-Lnative=${SDK_PATH}/arch/${ARCH}/lib`: Link against Fuchsia libraries from the SDK

- `-Lnative=${SDK_PATH}/arch/${ARCH}/sysroot/lib`: Link against Fuchsia sysroot libraries from the SDK

In total, our new project will look like:

Current directory structure

```
hello_fuchsia/
└── src/
    └── main.rs
└── Cargo.toml
└── .cargo/
    └── config.toml
```

Finally, we can build our rust binary as:

```
cargo build --target x86_64-unknown-fuchsia
```

Now we have a Rust binary at `target/x86_64-unknown-fuchsia/debug/hello_fuchsia`, targeting our desired Fuchsia target.

Current directory structure

```
hello_fuchsia/
└── src/
    └── main.rs
└── target/
    └── x86_64-unknown-fuchsia/
        └── debug/
            └── hello_fuchsia
└── Cargo.toml
└── .cargo/
    └── config.toml
```

Targeting Fuchsia with a compiler built from source

An alternative to the first workflow is to target Fuchsia by using `rustc` built from source.

Before building Rust for Fuchsia, you'll need a clang toolchain that supports Fuchsia as well. A recent version (14+) of clang should be sufficient to compile Rust for Fuchsia.

x86-64 and AArch64 Fuchsia targets can be enabled using the following configuration in `config.toml`:

```
[build]
target = ["<host_platform>", "aarch64-unknown-fuchsia", "x86_64-unknown-fuchsia"]

[rust]
lld = true

[llvm]
download-ci-llvm = false

[target.x86_64-unknown-fuchsia]
cc = "clang"
cxx = "clang++"

[target.aarch64-unknown-fuchsia]
cc = "clang"
cxx = "clang++"
```

Though not strictly required, you may also want to use `clang` for your host target as well:

```
[target.<host_platform>]
cc = "clang"
cxx = "clang++"
```

By default, the Rust compiler installs itself to `/usr/local` on most UNIX systems. You may want to install it to another location (e.g. a local `install` directory) by setting a custom prefix in `config.toml`:

```
[install]
# Make sure to use the absolute path to your install directory
prefix = "<RUST_SRC_PATH>/install"
```

Next, the following environment variables must be configured. For example, using a script we name `config-env.sh`:

```
# Configure this environment variable to be the path to the downloaded SDK
export SDK_PATH=<SDK path goes here>

export CFLAGS_aarch64_unknown_fuchsia="--target=aarch64-unknown-fuchsia --
sysroot=${SDK_PATH}/arch/arm64/sysroot -I${SDK_PATH}/pkg/fdio/include"
export CXXFLAGS_aarch64_unknown_fuchsia="--target=aarch64-unknown-fuchsia --
sysroot=${SDK_PATH}/arch/arm64/sysroot -I${SDK_PATH}/pkg/fdio/include"
export LDFLAGS_aarch64_unknown_fuchsia="--target=aarch64-unknown-fuchsia --
sysroot=${SDK_PATH}/arch/arm64/sysroot -L${SDK_PATH}/arch/arm64/lib"
export CARGO_TARGET_AARCH64_UNKNOWN_FUCHSIA_RUSTFLAGS="-C link-arg=--
sysroot=${SDK_PATH}/arch/arm64/sysroot -Lnative=${SDK_PATH}/arch/arm64/sysroot/lib -
Lnative=${SDK_PATH}/arch/arm64/lib"
export CFLAGS_x86_64_unknown_fuchsia="--target=x86_64-unknown-fuchsia --
sysroot=${SDK_PATH}/arch/x64/sysroot -I${SDK_PATH}/pkg/fdio/include"
export CXXFLAGS_x86_64_unknown_fuchsia="--target=x86_64-unknown-fuchsia --
sysroot=${SDK_PATH}/arch/x64/sysroot -I${SDK_PATH}/pkg/fdio/include"
export LDFLAGS_x86_64_unknown_fuchsia="--target=x86_64-unknown-fuchsia --
sysroot=${SDK_PATH}/arch/x64/sysroot -L${SDK_PATH}/arch/x64/lib"
export CARGO_TARGET_X86_64_UNKNOWN_FUCHSIA_RUSTFLAGS="-C link-arg=--
sysroot=${SDK_PATH}/arch/x64/sysroot -Lnative=${SDK_PATH}/arch/x64/sysroot/lib -
Lnative=${SDK_PATH}/arch/x64/lib"
```

Finally, the Rust compiler can be built and installed:

```
(source config-env.sh && ./x.py install)
```

Once `rustc` is installed, we can create a new working directory to work from, `hello_fuchsia` along with `hello_fuchsia/src`:

```
mkdir hello_fuchsia
cd hello_fuchsia
mkdir src
```

Note: From this point onwards, all commands will be issued from the `hello_fuchsia/` directory, and all `hello_fuchsia/` prefixes will be removed from references for sake of brevity.

There, we can create a new file named `src/hello_fuchsia.rs`:

src/hello_fuchsia.rs

```
fn main() {
    println!("Hello Fuchsia!");
}

#[test]
fn it_works() {
    assert_eq!(2 + 2, 4);
}
```

Current directory structure

```
hello_fuchsia/
└ src/
    └ hello_fuchsia.rs
```

Using your freshly installed `rustc`, you can compile a binary for Fuchsia using the following options:

- `--target x86_64-unknown-fuchsia`/`--target aarch64-unknown-fuchsia`: Targets the Fuchsia platform of your choice
- `-Lnative ${SDK_PATH}/arch/${ARCH}/lib`: Link against Fuchsia libraries from the SDK
- `-Lnative ${SDK_PATH}/arch/${ARCH}/sysroot/lib`: Link against Fuchsia sysroot libraries from the SDK

Putting it all together:

```
# Configure these for the Fuchsia target of your choice
TARGET_ARCH="<x86_64-unknown-fuchsia|aarch64-unknown-fuchsia>"
ARCH="<x64|aarch64>"

rustc \
  --target ${TARGET_ARCH} \
  -Lnative=${SDK_PATH}/arch/${ARCH}/lib \
  -Lnative=${SDK_PATH}/arch/${ARCH}/sysroot/lib \
  --out-dir bin src/hello_fuchsia.rs
```

Current directory structure

```
hello_fuchsia/
├ src/
│ └ hello_fuchsia.rs
└ bin/
    └ hello_fuchsia
```

Creating a Fuchsia package

Before moving on, double check your directory structure:

Current directory structure

```
hello_fuchsia/
├ src/                                (if using rustc)
│ └ hello_fuchsia.rs
├ bin/                                 ...
│ └ hello_fuchsia
└ src/                                (if using cargo)
    └ main.rs
    target/
        └ x86_64-unknown-fuchsia/
            └ debug/
                └ hello_fuchsia
                    ...
```

With our Rust binary built, we can move to creating a Fuchsia package. On Fuchsia, a package is the unit of distribution for software. We'll need to create a new package directory where we will place files like our finished binary and any data it may need.

To start, make the `pkg`, and `pkg/meta` directories:

```
mkdir pkg  
mkdir pkg/meta
```

Current directory structure

```
hello_fuchsia/  
└ pkg/  
    └ meta/
```

Now, create the following files inside:

`pkg/meta/package`

```
{  
    "name": "hello_fuchsia",  
    "version": "0"  
}
```

The `package` file describes our package's name and version number. Every package must contain one.

`pkg/hello_fuchsia.manifest` if using cargo

```
bin/hello_fuchsia=target/x86_64-unknown-fuchsia/debug/hello_fuchsia  
lib/ld.so.1=<SDK_PATH>/arch/x64/sysroot/dist/lib/ld.so.1  
lib/libfdio.so=<SDK_PATH>/arch/x64/dist/libfdio.so  
meta/package(pkg/meta/package  
meta/hello_fuchsia.cm(pkg/meta/hello_fuchsia.cm
```

`pkg/hello_fuchsia.manifest` if using rustc

```
bin/hello_fuchsia=bin/hello_fuchsia  
lib/ld.so.1=<SDK_PATH>/arch/x64/sysroot/dist/lib/ld.so.1  
lib/libfdio.so=<SDK_PATH>/arch/x64/dist/libfdio.so  
meta/package(pkg/meta/package  
meta/hello_fuchsia.cm(pkg/meta/hello_fuchsia.cm
```

Note: Relative manifest paths are resolved starting from the working directory of `pm`. Make sure to fill out `<SDK_PATH>` with the path to the downloaded SDK.

The `.manifest` file will be used to describe the contents of the package by relating their location when installed to their location on the file system. The `bin/hello_fuchsia=` entry will be different depending on how your Rust binary was built, so choose accordingly.

Current directory structure

```
hello_fuchsia/
└ pkg/
  └ meta/
    └ package
      hello_fuchsia.manifest
```

Creating a Fuchsia component

On Fuchsia, components require a component manifest written in Fuchsia's markup language called CML. The Fuchsia devsite contains an [overview of CML](#) and a [reference for the file format](#). Here's a basic one that can run our single binary:

`pkg/hello_fuchsia.cml`

```
{
  include: [ "syslog/client.shard.cml" ],
  program: {
    runner: "elf",
    binary: "bin/hello_fuchsia",
  },
}
```

Current directory structure

```
hello_fuchsia/
└ pkg/
  └ meta/
    └ package
      hello_fuchsia.manifest
      hello_fuchsia.cml
```

Now we can compile that CML into a component manifest:

```
 ${SDK_PATH}/tools/${ARCH}/cmc compile \
  pkg/hello_fuchsia.cml \
  --includepath ${SDK_PATH}/pkg \
  -o pkg/meta/hello_fuchsia.cm
```

Note: `--includepath` tells the compiler where to look for `includes` from our CML. In our case, we're only using `syslog/client.shard.cml`.

Current directory structure

```
hello_fuchsia/
└ pkg/
  └ meta/
    └ package
      hello_fuchsia.cm
      hello_fuchsia.manifest
      hello_fuchsia.cml
```

Building a Fuchsia package

Next, we'll build a package manifest as defined by our manifest:

```
`${SDK_PATH}/tools/${ARCH}/pm \
    -api-level $($SDK_PATH/tools/${ARCH}/ffx version -v | grep "api-level" | head -1 | awk
-F ' ' '{print $2}') \
    -o pkg/hello_fuchsia_manifest \
    -m pkg/hello_fuchsia.manifest \
    build \
    -output-package-manifest pkg/hello_fuchsia_package_manifest
```

This will produce `pkg/hello_fuchsia_manifest/` which is a package manifest we can publish directly to a repository.

Current directory structure

```
hello_fuchsia/
└── pkg/
    ├── meta/
    │   └── package
    │       └── hello_fuchsia.cm
    ├── hello_fuchsia_manifest/
    │   └── ...
    ├── hello_fuchsia.manifest
    ├── hello_fuchsia.cml
    └── hello_fuchsia_package_manifest
```

We are now ready to publish the package.

Publishing a Fuchsia package

With our package and component manifests setup, we can now publish our package. The first step will be to create a Fuchsia package repository to publish to.

Creating a Fuchsia package repository

We can set up our repository with:

```
`${SDK_PATH}/tools/${ARCH}/pm newrepo \
    -repo pkg/repo
```

Current directory structure

```
hello_fuchsia/
└ pkg/
   └ meta/
      └ package
         └ hello_fuchsia.cm
   └ hello_fuchsia_manifest/
      └ ...
   └ repo/
      └ ...
   └ hello_fuchsia.manifest
   └ hello_fuchsia.cml
   └ hello_fuchsia_package_manifest
```

Publishing Fuchsia package to repository

We can publish our new package to that repository with:

```
 ${SDK_PATH}/tools/${ARCH}/pm publish \
 -repo pkg/repo \
 -lp -f <(echo "pkg/hello_fuchsia_package_manifest")
```

Then we can add the repository to `ffx`'s package server as `hello-fuchsia` using:

```
 ${SDK_PATH}/tools/${ARCH}/ffx repository add-from-pm \
 pkg/repo \
 -r hello-fuchsia
```

Running a Fuchsia component on an emulator

At this point, we are ready to run our Fuchsia component. For reference, our final directory structure will look like:

Final directory structure

```
hello_fuchsia/
└── src/           (if using rustc)
    └── hello_fuchsia.rs
    ...
    bin/
    └── hello_fuchsia
    ...
    src/           (if using cargo)
    └── main.rs
    target/
    └── x86_64-unknown-fuchsia/
        └── debug/
            └── hello_fuchsia
            ...
    pkg/
    └── meta/
        └── package
            └── hello_fuchsia.cm
    hello_fuchsia_manifest/
    └── ...
    repo/
    └── ...
    hello_fuchsia.manifest
    hello_fuchsia.cml
    hello_fuchsia_package_manifest
```

Starting the Fuchsia emulator

Start a Fuchsia emulator in a new terminal using:

```
 ${SDK_PATH}/tools/${ARCH}/ffx product-bundle get workstation_eng.qemu-${ARCH}
 ${SDK_PATH}/tools/${ARCH}/ffx emu start workstation_eng.qemu-${ARCH} --headless
```

Watching emulator logs

Once the emulator is running, open a separate terminal to watch the emulator logs:

In separate terminal

```
 ${SDK_PATH}/tools/${ARCH}/ffx log \
    --since now
```

Serving a Fuchsia package

Now, start a package repository server to serve our package to the emulator:

```
 ${SDK_PATH}/tools/${ARCH}/ffx repository server start
```

Once the repository server is up and running, register it with the target Fuchsia system running in the emulator:

```
`${SDK_PATH}/tools/${ARCH}/ffx target repository register \
--repository hello-fuchsia
```

Running a Fuchsia component

Finally, run the component:

```
`${SDK_PATH}/tools/${ARCH}/ffx component run \
/core/ffx-laboratory:hello_fuchsia \
fuchsia-pkg://hello-fuchsia/hello_fuchsia_manifest#meta/hello_fuchsia.cm
```

On reruns of the component, the `--recreate` argument may also need to be passed.

```
`${SDK_PATH}/tools/${ARCH}/ffx component run \
--recreate \
/core/ffx-laboratory:hello_fuchsia \
fuchsia-pkg://hello-fuchsia/hello_fuchsia_manifest#meta/hello_fuchsia.cm
```

.gitignore extensions

Optionally, we can create/extend our `.gitignore` file to ignore files and directories that are not helpful to track:

```
pkg/repo
pkg/meta/hello_fuchsia.cm
pkg/hello_fuchsia_manifest
pkg/hello_fuchsia_package_manifest
```

Testing

Running unit tests

Tests can be run in the same way as a regular binary.

- If using `cargo`, you can simply pass `test --no-run` to the `cargo` invocation and then repackage and rerun the Fuchsia package. From our previous example, this would look like `cargo test --target x86_64-unknown-fuchsia --no-run`, and moving the executable binary path found from the line `Executable unittests src/main.rs (target/x86_64-unknown-fuchsia/debug/deps/hello_fuchsia-<HASH>)` into `pkg/hello_fuchsia.manifest`.
- If using the compiled `rustc`, you can simply pass `--test` to the `rustc` invocation and then repackage and rerun the Fuchsia package.

The test harness will run the applicable unit tests.

Often when testing, you may want to pass additional command line arguments to your binary. Additional arguments can be set in the component manifest:

pkg/hello_fuchsia.cml

```
{  
    include: [ "syslog/client.shard.cml" ],  
    program: {  
        runner: "elf",  
        binary: "bin/hello_fuchsia",  
        args: ["it_works"],  
    },  
}
```

This will pass the argument `it_works` to the binary, filtering the tests to only those tests that match the pattern. There are many more configuration options available in CML including environment variables. More documentation is available on the [Fuchsia devsite](#).

Running the compiler test suite

The commands in this section assume that they are being run from inside your local Rust source checkout:

```
cd ${RUST_SRC_PATH}
```

To run the Rust test suite on an emulated Fuchsia device, you must install the Rust compiler locally. See "[Targeting Fuchsia with a compiler built from source](#)" for the steps to build locally.

You'll also need to download a copy of the Fuchsia SDK. The current minimum supported SDK version is [10.20221207.2.89](#).

Fuchsia's test runner interacts with the Fuchsia emulator and is located at `src/ci/docker/scripts/fuchsia-test-runner.py`. We can use it to start our test environment with:

```
src/ci/docker/scripts/fuchsia-test-runner.py start  
--rust ${RUST_SRC_PATH}/install  
--sdk ${SDK_PATH}  
--target {x86_64-unknown-fuchsia|aarch64-unknown-fuchsia}
```

Where `${RUST_SRC_PATH}/install` is the `prefix` set in `config.toml` and `${SDK_PATH}` is the path to the downloaded and unzipped SDK.

Once our environment is started, we can run our tests using `x.py` as usual. The test runner script will run the compiled tests on an emulated Fuchsia device. To run the full `tests/ui` test suite:

```
( \
    source config-env.sh &&
    ./x.py
    --config config.toml
    --stage=2
    test tests/ui
    --target x86_64-unknown-fuchsia
    --run=always
    --test-args --target-rustcflags
    --test-args -Lnative=${SDK_PATH}/arch/{x64|arm64}/sysroot/lib
    --test-args --target-rustcflags
    --test-args -Lnative=${SDK_PATH}/arch/{x64|arm64}/lib
    --test-args --target-rustcflags
    --test-args -Clink-arg=--undefined-version
    --test-args --remote-test-client
    --test-args src/ci/docker/scripts/fuchsia-test-runner.py
)
```

By default, `x.py` compiles test binaries with `panic=unwind`. If you built your Rust toolchain with `-Cpanic=abort`, you need to tell `x.py` to compile test binaries with `panic=abort` as well:

```
--test-args --target-rustcflags
--test-args -Cpanic=abort
--test-args --target-rustcflags
--test-args -Zpanic_abort_tests
```

When finished testing, the test runner can be used to stop the test environment:

```
src/ci/docker/scripts/fuchsia-test-runner.py stop
```

Debugging

`zxdb`

Debugging components running on a Fuchsia emulator can be done using the console-mode debugger: `zxdb`. We will demonstrate attaching necessary symbol paths to debug our `hello-fuchsia` component.

Attaching `zxdb`

In a separate terminal, issue the following command from our `hello_fuchsia` directory to launch `zxdb`:

`In separate terminal`

```
`${SDK_PATH}/tools/${ARCH}/ffx debug connect -- \
--symbol-path target/x86_64-unknown-fuchsia/debug
```

- `--symbol-path` gets required symbol paths, which are necessary for stepping through your program.

The "displaying source code in `zxdb`" section describes how you can display Rust and/or Fuchsia source code in your debugging session.

Using `zxdb`

Once launched, you will be presented with the window:

```
Connecting (use "disconnect" to cancel)...
Connected successfully.
👉 To get started, try "status" or "help".
[zxdb]
```

To attach to our program, we can run:

```
[zxdb] attach hello_fuchsia
```

Expected output

```
Waiting for process matching "hello_fuchsia".
Type "filter" to see the current filters.
```

Next, we can create a breakpoint at `main` using "`b main`":

```
[zxdb] b main
```

Expected output

```
Created Breakpoint 1 @ main
```

Finally, we can re-run the "`hello_fuchsia`" component from our original terminal:

```
`${SDK_PATH}/tools/${ARCH}/ffx component run \
--recreate \
fuchsia-pkg://hello-fuchsia/hello_fuchsia_manifest#meta/hello_fuchsia.cm
```

Once our component is running, our `zxdb` window will stop execution in our `main` as desired:

Expected output

```
Breakpoint 1 now matching 1 addrs for main
① on bp 1 hello_fuchsia::main() • main.rs:2
  1 fn main() {
▶ 2     println!("Hello Fuchsia!");
  3 }
  4
[zxdb]
```

[`zxdb`] has similar commands to other debuggers like `gdb`. To list the available commands, run "help" in the `zxdb` window or visit [the zxdb documentation](#).

```
[zxdb] help
```

Expected output

```
Help!

Type "help <command>" for command-specific help.

Other help topics (see "help <topic>")
...
```

Displaying source code in `zxdb`

By default, the debugger will not be able to display source code while debugging. For our user code, we displayed source code by pointing our debugger to our debug binary via the `--symbol-path` arg. To display library source code in the debugger, you must provide paths to the source using `--build-dir`. For example, to display the Rust and Fuchsia source code:

```
 ${SDK_PATH}/tools/${ARCH}/ffx debug connect -- \
  --symbol-path target/x86_64-unknown-fuchsia/debug \
  --build-dir ${RUST_SRC_PATH}/rust \
  --build-dir ${FUCHSIA_SRC_PATH}/fuchsia/out/default
```

- `--build-dir` links against source code paths, which are not strictly necessary for debugging, but is a nice-to-have for displaying source code in `zxdb`.

Linking to a Fuchsia checkout can help with debugging Fuchsia libraries, such as `fdio`.

Debugging the compiler test suite

Debugging the compiler test suite requires some special configuration:

First, we have to properly configure `zxdb` so it will be able to find debug symbols and source information for our test. The test runner can do this for us with:

```
src/ci/docker/scripts/fuchsia-test-runner.py debug
  --rust-src ${RUST_SRC_PATH}
  --fuchsia-src ${FUCHSIA_SRC_PATH}
  --test ${TEST}
```

where `${TEST}` is relative to Rust's `tests` directory (e.g. `ui/abi/...`).

This will start a `zxdb` session that is properly configured for the specific test being run. All three arguments are optional, so you can omit `--fuchsia-src` if you don't have it downloaded. Now is a good time to set any desired breakpoints, like `b main` .

Next, we have to tell `x.py` not to optimize or strip debug symbols from our test suite binaries. We can do this by passing some new arguments to `rustc` through our `x.py` invocation. The full invocation is:

```
( \
  source config-env.sh &&
  ./x.py
  --config config.toml
  --stage=2
  test tests/${TEST}
  --target x86_64-unknown-fuchsia
  --run=always
  --test-args --target-rustcflags
  --test-args -Lnative=${SDK_PATH}/arch/{x64|arm64}/sysroot/lib
  --test-args --target-rustcflags
  --test-args -Lnative=${SDK_PATH}/arch/{x64|arm64}/lib
  --test-args --target-rustcflags
  --test-args -Clink-arg=-undefined-version
  --test-args --target-rustcflags
  --test-args -Cdebuginfo=2
  --test-args --target-rustcflags
  --test-args -Copt-level=0
  --test-args --target-rustcflags
  --test-args -Cstrip=none
  --test-args --remote-test-client
  --test-args src/ci/docker/scripts/fuchsia-test-runner.py
)
```

If you built your Rust toolchain with `panic=abort` , make sure to include the previous flags so your test binaries are also compiled with `panic=abort` .

Upon running this command, the test suite binary will be run and `zxdb` will attach and load any relevant debug symbols.

-kmc-solid_

Tier: 3

SOLID embedded development platform by Kyoto Microcomputer Co., Ltd.

The target names follow this format: `$ARCH-kmc-solid_$KERNEL-$ABI`, where `$ARCH` specifies the target processor architecture, `$KERNEL` the base kernel, and `$ABI` the target ABI (optional). The following targets are currently defined:

Target name	target_arch	target_vendor	target_os
aarch64-kmc-solid_asp3	aarch64	kmc	solid_asp3
armv7a-kmc-solid_asp3-eabi	arm	kmc	solid_asp3
armv7a-kmc-solid_asp3-eabihf	arm	kmc	solid_asp3

Designated Developers

- @kawadakk

Requirements

This target is cross-compiled. A platform-provided C compiler toolchain is required, though it can be substituted by [GNU Arm Embedded Toolchain](#) for the purpose of building Rust and functional binaries.

Building

The target can be built by enabling it for a `rustc` build.

```
[build]
target = ["aarch64-kmc-solid_asp3"]
```

Make sure `aarch64-kmc-elf-gcc` is included in `$PATH`. Alternatively, you can use GNU Arm Embedded Toolchain by adding the following to `config.toml`:

```
[target.aarch64-kmc-solid_asp3]
cc = "arm-none-eabi-gcc"
```

Cross-compilation

This target can be cross-compiled from any hosts.

Testing

Currently there is no support to run the `rustc` test suite for this target.

Building Rust programs

Building executables is not supported yet.

If `rustc` has support for that target and the library artifacts are available, then Rust static libraries can be built for that target:

```
$ rustc --target aarch64-kmc-solid_asp3 your-code.rs --crate-type staticlib  
$ ls libyour_code.a
```

On Rust Nightly it's possible to build without the target artifacts available:

```
cargo build -Z build-std --target aarch64-kmc-solid_asp3
```

loongarch*-unknown-linux-*

Tier: 3

[LoongArch](#) is a new RISC ISA developed by Loongson Technology Corporation Limited.

The target name follow this format: `<machine>-<vendor>-<os><abi-suffix>`, where `<machine>` specifies the CPU family/model, `<vendor>` specifies the vendor and `<os>` the operating system name. While the integer base ABI is implied by the machine field, the floating point base ABI type is encoded into the os field of the specifier using the string suffix `<abi-suffix>`.

<abi-suffix>	Description		
ABI type(Base ABI/ABI extension)	c library	kernel	target tuple
lp64d/base	glibc	linux	loongarch64-unknown-linux-gnu
lp64f/base	glibc	linux	loongarch64-unknown-linux-gnuf32
lp64s/base	glibc	linux	loongarch64-unknown-linux-gnusf
lp64d/base	musl libc	linux	loongarch64-unknown-linux-musl
lp64f/base	musl libc	linux	loongarch64-unknown-linux-muslf32
lp64s/base	musl libc	linux	loongarch64-unknown-linux-muslsf

Target maintainers

- [Zhai Xiaojuan](#) `zhaixiaojuan@loongson.cn`
- [Wang Rui](#) `wangrui@loongson.cn`
- [Zhai Xiang](#) `zhaixiang@loongson.cn`
- [Wang Xuerui](#) `git@xen0n.name`

Requirements

This target is cross-compiled. A GNU toolchain for LoongArch target is required. It can be downloaded from <https://github.com/loongson/build-tools/releases>, or built from the source code of GCC (12.1.0 or later) and Binutils (2.40 or later).

Building the target

The target can be built by enabling it for a `rustc` build.

```
[build]
target = ["loongarch64-unknown-linux-gnu"]
```

Make sure `loongarch64-unknown-linux-gnu-gcc` can be searched from the directories specified in `$PATH`. Alternatively, you can use GNU LoongArch Toolchain by adding the following to `config.toml`:

```
[target.loongarch64-unknown-linux-gnu]
# ADJUST THIS PATH TO POINT AT YOUR TOOLCHAIN
cc = "/TOOLCHAIN_PATH/bin/loongarch64-unknown-linux-gnu-gcc"
cxx = "/TOOLCHAIN_PATH/bin/loongarch64-unknown-linux-gnu-g++"
ar = "/TOOLCHAIN_PATH/bin/loongarch64-unknown-linux-gnu-ar"
ranlib = "/TOOLCHAIN_PATH/bin/loongarch64-unknown-linux-gnu-ranlib"
linker = "/TOOLCHAIN_PATH/bin/loongarch64-unknown-linux-gnu-gcc"
```

Cross-compilation

This target can be cross-compiled on a `x86_64-unknown-linux-gnu` host. Cross-compilation on other hosts may work but is not tested.

Testing

To test a cross-compiled binary on your build system, install the `qemu` binary that supports the LoongArch architecture and execute the following commands.

```
CC_loongarch64_unknown_linux_gnu=/TOOLCHAIN_PATH/bin/loongarch64-unknown-linux-gnu-gcc \
CXX_loongarch64_unknown_linux_gnu=/TOOLCHAIN_PATH/bin/loongarch64-unknown-linux-gnu-g++ \
AR_loongarch64_unknown_linux_gnu=/TOOLCHAIN_PATH/bin/loongarch64-unknown-linux-gnu-gcc-ar \
CARGO_TARGET_LOONGARCH64_UNKNOWN_LINUX_GNUN_LINKER=/TOOLCHAIN_PATH/bin/loongarch64-unknown-
linux-gnu-gcc \
# SET TARGET SYSTEM LIBRARY PATH
CARGO_TARGET_LOONGARCH64_UNKNOWN_LINUX_GNUN_RUNNER="qemu-loongarch64 -L
/TOOLCHAIN_PATH/TARGET_LIBRAY_PATH" \
cargo run --target loongarch64-unknown-linux-gnu --release
```

Tested on x86 architecture, other architectures not tested.

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target. To compile for this target, you will either need to build Rust with the target enabled (see "Building the target" above), or build your own copy of `std` by using `build-std` or similar.

If `rustc` has support for that target and the library artifacts are available, then Rust static libraries can be built for that target:

```
$ rustc --target loongarch64-unknown-linux-gnu your-code.rs --crate-type staticlib  
$ ls libyour_code.a
```

On Rust Nightly it's possible to build without the target artifacts available:

```
cargo build -Z build-std --target loongarch64-unknown-linux-gnu
```

m68k-unknown-linux-gnu

Tier: 3

Motorola 680x0 Linux

Designated Developers

- [@glaubitz](#)
- [@ricky26](#)

Requirements

This target requires a Linux/m68k build environment for cross-compilation which is available on Debian and Debian-based systems, openSUSE and other distributions.

On Debian, it should be sufficient to install a g++ cross-compiler for the m68k architecture which will automatically pull in additional dependencies such as the glibc cross development package:

```
# apt install g++-m68k-linux-gnu
```

Binaries can be run using QEMU user emulation. On Debian-based systems, it should be sufficient to install the package `qemu-user-static` to be able to run simple static binaries:

```
# apt install qemu-user-static
```

To run more complex programs, it will be necessary to set up a Debian/m68k chroot with the help of the command `debootstrap`:

```
# apt install debootstrap debian-ports-archive-keyring
# debootstrap --keyring=/usr/share/keyrings/debian-ports-archive-keyring.gpg --arch=m68k
unstable debian-68k http://ftp.ports.debian.org/debian-ports
```

This chroot can then seamlessly entered using the normal `chroot` command thanks to QEMU user emulation:

```
# chroot /path/to/debian-68k
```

To get started with native builds, which are currently untested, a native Debian/m68k system can be installed either on real hardware such as 68k-based Commodore Amiga or Atari systems or emulated environments such as QEMU version 4.2 or newer or ARAnyM.

ISO images for installation are provided by the Debian Ports team and can be obtained from the Debian CD image server available at:

<https://cdimage.debian.org/cdimage/ports/current>

Documentation for Debian/m68k is available on the Debian Wiki at:

<https://wiki.debian.org/M68k>

Support is available either through the `debian-68k` mailing list:

<https://lists.debian.org/debian-68k/>

or the `#debian-68k` IRC channel on OFTC network.

Building

The codegen for this target should be built by default. However, core and std are currently missing but are being worked on and should become available in the near future.

Cross-compilation

This target can be cross-compiled from a standard Debian or Debian-based, openSUSE or any other distribution which has a basic m68k cross-toolchain available.

Testing

Currently there is no support to run the `rustc` test suite for this target.

Building Rust programs

Rust programs can be built for that target:

```
rustc --target m68k-unknown-linux-gnu your-code.rs
```

Very simple programs can be run using the `qemu-m68k-static` program:

```
$ qemu-m68k-static your-code
```

For more complex applications, a chroot or native (emulated) Debian/m68k system are required for testing.

mips64-openwrt-linux-musl

Tier: 3

Target maintainers

- Donald Hoskins grommish@gmail.com, <https://github.com/Itus-Shield>

Requirements

This target is cross-compiled. There is no support for `std`. There is no default allocator, but it's possible to use `alloc` by supplying an allocator.

By default, Rust code generated for this target uses `-msoft-float` and is dynamically linked.

This target generates binaries in the ELF format.

Building the target

This target is built exclusively within the `OpenWrt` build system via the `rust-lang` HOST package

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target. To compile for this target, you will either need to build Rust with the target enabled (see "Building the target" above).

Testing

As `mips64-openwrt-linux-musl` supports a variety of different environments and does not support `std`, this target does not support running the Rust testsuite at this time.

mipsel-sony-psx

Tier: 3

Sony PlayStation 1 (psx)

Designated Developer

- [@ayrtonm](#)

Requirements

This target is cross-compiled. It has no special requirements for the host.

Building

The target can be built by enabling it for a `rustc` build:

```
[build]
build-stage = 1
target = ["mipsel-sony-psx"]
```

Cross-compilation

This target can be cross-compiled from any host.

Testing

Currently there is no support to run the `rustc` test suite for this target.

Building Rust programs

Since it is Tier 3, `rust` doesn't ship pre-compiled artifacts for this target.

Just use the `build-std` nightly cargo feature to build the `core` and `alloc` libraries:

```
cargo build -Zbuild-std=core,alloc --target mipsel-sony-psx
```

The command above generates an ELF. To generate binaries in the PSEXEX format that emulators run, you can use [cargo-psx](#):

```
cargo psx build
```

or use `-Clink-arg=-oformat=binary` to produce a flat binary.

nvptx64-nvidia-cuda

Tier: 2

This is the target meant for deploying code for Nvidia® accelerators based on their CUDA platform.

Target maintainers

- Riccardo D'Ambrosio, <https://github.com/RDambrosio016>
- Kjetil Kjeka, <https://github.com/kjetilkjeka>

riscv32imac-unknown-xous-elf

Tier: 3

Xous microkernel, message-based operating system that powers devices such as Precursor and Betrusted. The operating system is written entirely in Rust, so no additional software is required to compile programs for Xous.

Target maintainers

- [@xobs](#)

Requirements

Building the target itself requires a RISC-V compiler that is supported by [cc-rs](#). For example, you can use the prebuilt [xPack](#) toolchain.

Cross-compiling programs does not require any additional software beyond the toolchain. Prebuilt versions of the toolchain are available [from Betrusted](#).

Building the target

The target can be built by enabling it for a `rustc` build.

```
[build]
target = ["riscv32imac-unknown-xous-elf"]
```

Make sure your C compiler is included in `$PATH`, then add it to the `config.toml`:

```
[target.riscv32imac-unknown-xous-elf]
cc = "riscv-none-elf-gcc"
ar = "riscv-none-elf-ar"
```

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target. To compile for this target, you will need to do one of the following:

- Build Rust with the target enabled (see "Building the target" above)
- Build your own copy of `core` by using `build-std` or similar

- Download a prebuilt toolchain from [Betrusted](#)

Cross-compilation

This target can be cross-compiled from any host.

Testing

Currently there is no support to run the `rustc` test suite for this target.

*-pc-windows-gnu LLVM

Tier: 3

Windows targets similar to `*-pc-windows-gnu` but using UCRT as the runtime and various LLVM tools/libraries instead of GCC/Binutils.

Target triples available so far:

- `aarch64-pc-windows-gnu LLVM`
- `x86_64-pc-windows-gnu LLVM`

Target maintainers

- `@mati865`

Requirements

The easiest way to obtain these targets is cross-compilation but native build from `x86_64-pc-windows-gnu` is possible with few hacks which I don't recommend. Std support is expected to be on par with `*-pc-windows-gnu`.

Binaries for this target should be at least on par with `*-pc-windows-gnu` in terms of requirements and functionality.

Those targets follow Windows calling convention for `extern "C"`.

Like with any other Windows target created binaries are in PE format.

Building the target

For cross-compilation I recommend using `llvm-mingw` toolchain, one change that seems necessary beside configuring cross compilers is disabling experimental `m86k` target. Otherwise LLVM build fails with `multiple definition ...` errors. Native bootstrapping builds require rather fragile hacks until host artifacts are available so I won't describe them here.

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target. To compile for this target, you will either need to build Rust with the target enabled (see "Building the target" above), or build your own copy of `core` by using `build-std` or similar.

Testing

Created binaries work fine on Windows or Wine using native hardware. Testing AArch64 on `x86_64` is problematic though and requires spending some time with QEMU. Once these targets bootstrap themselves on native hardware they should pass Rust testsuite.

Cross-compilation toolchains and C code

Compatible C code can be built with Clang's `aarch64-pc-windows-gnu` and `x86_64-pc-windows-gnu` targets as long as LLVM based C toolchains are used. Those include:

- `llvm-mingw`
- MSYS2 with CLANG* environment

nto-qnx

Tier: 3

QNX® Neutrino (nto) Real-time operating system. The support has been implemented jointly by [Elektrobit Automotive GmbH](#) and [Blackberry QNX](#).

Target maintainers

- Florian Bartels, Florian.Bartels@elektrobit.com, <https://github.com/flba-eb>
- Tristan Roach, TRoach@blackberry.com, <https://github.com/gh-tr>

Requirements

Currently, the following QNX Neutrino versions and compilation targets are supported:

QNX Neutrino Version	Target Architecture	Full support	<code>no_std</code> support
7.1	AArch64	✓	✓
7.1	x86_64	✓	✓
7.0	x86		✓

Adding other architectures that are supported by QNX Neutrino is possible.

In the table above, 'full support' indicates support for building Rust applications with the full standard library. '`no_std` support' indicates that only `core` and `alloc` are available.

For building or using the Rust toolchain for QNX Neutrino, the [QNX Software Development Platform \(SDP\)](#) must be installed and initialized. Initialization is usually done by sourcing `qnxsdः-env.sh` (this will be installed as part of the SDP, see also installation instruction provided with the SDP). Afterwards `qcc` (QNX C/C++ compiler) should be available (in the `$PATH` variable). `qcc` will be called e.g. for linking executables.

When linking `no_std` applications, they must link against `libc.so` (see example). This is required because applications always link against the `crt` library and `crt` depends on `libc.so`. This is done automatically when using the standard library.

Small example application

Small `no_std` example is shown below. Applications using the standard library work as well.

```
#![no_std]
#![no_main]
#![feature(lang_items)]

// We must always link against libc, even if no external functions are used
// "extern C" - Block can be empty but must be present
#[link(name = "c")]
extern "C" {
    pub fn printf(format: *const core::ffi::c_char, ...) -> core::ffi::c_int;
}

#[no_mangle]
pub extern "C" fn main(_argc: isize, _argv: *const *const u8) -> isize {
    const HELLO: &'static str = "Hello World, the answer is %d\n\0";
    unsafe {
        printf(HELLO.as_ptr() as *const _, 42);
    }
    0
}

use core::panic::PanicInfo;

#[panic_handler]
fn panic(_panic: &PanicInfo<'_>) -> ! {
    loop {}
}

#[lang = "eh_personality"]
#[no_mangle]
pub extern "C" fn rust_eh_personality() {}
```

The QNX Neutrino support of Rust has been tested with QNX Neutrino 7.0 and 7.1.

There are no further known requirements.

Conditional compilation

For conditional compilation, following QNX Neutrino specific attributes are defined:

- `target_os` = `"nto"`
- `target_env` = `"nto71"` (for QNX Neutrino 7.1)
- `target_env` = `"nto70"` (for QNX Neutrino 7.0)

Building the target

1. Create a `config.toml`

Example content:

```
profile = "compiler"
changelog-seen = 2
```

2. Compile the Rust toolchain for an `x86_64-unknown-linux-gnu` host (for both `aarch64` and `x86_64` targets)

Compiling the Rust toolchain requires the same environment variables used for compiling C binaries. Refer to the [QNX developer manual](#).

To compile for QNX Neutrino (aarch64 and x86_64) and Linux (x86_64):

```
export build_env='
    CC_aarch64-unknown-nto-qnx710=qcc
    CFLAGS_aarch64-unknown-nto-qnx710=-Vgcc_ntoaarch64le_cxx
    CXX_aarch64-unknown-nto-qnx710=qcc
    AR_aarch64_unknown_nto_qnx710=ntoaarch64-ar
    CC_x86_64-pc-nto-qnx710=qcc
    CFLAGS_x86_64-pc-nto-qnx710=-Vgcc_ntox86_64_cxx
    CXX_x86_64-pc-nto-qnx710=qcc
    AR_x86_64_pc_nto_qnx710=ntox86_64-ar'

env $build_env \
    ./x.py build \
        --target aarch64-unknown-nto-qnx710 \
        --target x86_64-pc-nto-qnx710 \
        --target x86_64-unknown-linux-gnu \
        rustc library/core library/alloc
```

Running the Rust test suite

The test suites of the Rust compiler and standard library can be executed much like other Rust targets. The environment for testing should match the one used during compiler compilation (refer to `build_env` and `qcc`/`PATH` above) with the addition of the `TEST_DEVICE_ADDR` environment variable. The `TEST_DEVICE_ADDR` variable controls the remote runner and should point to the target, despite `localhost` being shown in the following example. Note that some tests are failing which is why they are currently excluded by the target maintainers which can be seen in the following example.

To run all tests on a `x86_64` QNX Neutrino target:

```

export TEST_DEVICE_ADDR="localhost:12345" # must address the test target, can be a SSH tunnel
export build_env='
    CC_aarch64-unknown-nto-qnx710=qcc
    CFLAGS_aarch64-unknown-nto-qnx710=-Vgcc_ntoaarch64le_cxx
    CXX_aarch64-unknown-nto-qnx710=qcc
    AR_aarch64_unknown_nto_qnx710=ntoaarch64-ar
    CC_x86_64-pc-nto-qnx710=qcc
    CFLAGS_x86_64-pc-nto-qnx710=-Vgcc_ntox86_64_cxx
    CXX_x86_64-pc-nto-qnx710=qcc
    AR_x86_64_pc_nto_qnx710=ntox86_64-ar'

# Disable tests that only work on the host or don't make sense for this target.
# See also:
# - src/ci/docker/host-x86_64/i686-gnu/Dockerfile
# - https://rust-lang.zulipchat.com/#narrow/stream/182449-t-compiler.2Fhelp/topic/Running.20tests.20on.20remote.20target
# - .github/workflows/ci.yml
export exclude_tests='
    --exclude src/bootstrap
    --exclude src/tools/error_index_generator
    --exclude src/tools/linkchecker
    --exclude tests/ui-fulldeps
    --exclude rustc
    --exclude rustdoc
    --exclude tests/run-make-fulldeps'

env $build_env \
    ./x.py test -j 1 \
        $exclude_tests \
        --stage 1 \
        --target x86_64-pc-nto-qnx710

```

Currently, only one thread can be used when testing due to limitations in `libc::fork` and `libc::posix_spawnnp`. See [fork documentation](#) (error section) for more information. This can be achieved by using the `-j 1` parameter in the `x.py` call. This issue is being researched and we will try to allow parallelism in the future.

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target. To compile for this target, you must either build Rust with the target enabled (see "Building the target" above), or build your own copy of `core` by using `build-std` or similar.

Testing

Compiled executables can run directly on QNX Neutrino.

Rust std library test suite

The target needs sufficient resources to execute all tests. The commands below assume that a QEMU image is used.

- Ensure that the temporary directory used by `remote-test-server` has enough free space and inodes. 5GB of free space and 40000 inodes are known to be sufficient (the test will create more than 32k files). To create a QEMU image in an empty directory, run this command inside the directory:

```
mkqnximage --type=qemu --ssh-ident=$HOME/.ssh/id_ed25519.pub --data-size=5000 --data-inodes=40000
```

`/data` should have enough free resources. Set the `TMPDIR` environment variable accordingly when running `remote-test-server`, e.g.:

```
TMPDIR=/data/tmp/rust remote-test-server --bind 0.0.0.0:12345
```

- Ensure the TCP stack can handle enough parallel connections (default is 200, should be 300 or higher). After creating an image (see above), edit the file `output/build/startup.sh`:

1. Search for `io-pkt-v6-hc`
2. Add the parameter `-ptcpip threads_max=300`, e.g.:

```
io-pkt-v6-hc -U 33:33 -d e1000 -ptcpip threads_max=300
```

3. Update the image by running `mkqnximage` again with the same parameters as above for creating it.

- Running and stopping the virtual machine

To start the virtual machine, run inside the directory of the VM:

```
mkqnximage --run=h
```

To stop the virtual machine, run inside the directory of the VM:

```
mkqnximage --stop
```

- Ensure local networking

Ensure that 'localhost' is getting resolved to 127.0.0.1. If you can't ping the localhost, some tests may fail. Ensure it's appended to /etc/hosts (if first `ping` command fails). Commands have to be executed inside the virtual machine!

```
$ ping localhost
ping: Cannot resolve "localhost" (Host name lookup failure)

$ echo "127.0.0.1 localhost" >> /etc/hosts

$ ping localhost
PING localhost (127.0.0.1): 56 data bytes
64 bytes from 127.0.0.1: icmp_seq=0 ttl=255 time=1 ms
```

Cross-compilation toolchains and C code

Compiling C code requires the same environment variables to be set as compiling the Rust toolchain (see above), to ensure `qcc` is used with proper arguments. To ensure compatibility, do not specify any further arguments that for example change calling conventions or memory layout.

*-unknown-openbsd

Tier: 3

OpenBSD multi-platform 4.4BSD-based UNIX-like operating system.

The target names follow this format: `$ARCH-unknown-openbsd`, where `$ARCH` specifies the target processor architecture. The following targets are currently defined:

Target name	C++ library	OpenBSD Platform
<code>aarch64-unknown-openbsd</code>	<code>libc++</code>	64-bit ARM systems
<code>i686-unknown-openbsd</code>	<code>libc++</code>	Standard PC and clones based on the Intel i386 architecture and compatible processors
<code>powerpc64-unknown-openbsd</code>	<code>libc++</code>	IBM POWER-based PowerNV systems
<code>riscv64gc-unknown-openbsd</code>	<code>libc++</code>	64-bit RISC-V systems
<code>sparc64-unknown-openbsd</code>	<code>estdc++</code>	Sun UltraSPARC and Fujitsu SPARC64 systems
<code>x86_64-unknown-openbsd</code>	<code>libc++</code>	AMD64-based systems

Note that all OS versions are *major* even if using X.Y notation (`6.8` and `6.9` are different major versions) and could be binary incompatibles (with breaking changes).

Designated Developers

- `@semarie`, `semarie@openbsd.org`
- `lang/rust` maintainer (see `MAINTAINER` variable)

Fallback to `ports@openbsd.org`, OpenBSD third parties public mailing-list (with openbsd developers readers)

Requirements

These targets are natively compiled and could be cross-compiled. C compiler toolchain is required for the purpose of building Rust and functional binaries.

Building

The target can be built by enabling it for a `rustc` build.

```
[build]
target = ["$ARCH-unknown-openbsd"]

[target.$ARCH-unknown-openbsd]
cc = "$ARCH-openbsd-cc"
```

Cross-compilation

These targets can be cross-compiled, but LLVM might not build out-of-box.

Testing

The Rust testsuite could be run natively.

Building Rust programs

Rust does not yet ship pre-compiled artifacts for these targets.

*-unknown-uefi

Tier: 2

Unified Extensible Firmware Interface (UEFI) targets for application, driver, and core UEFI binaries.

Available targets:

- `aarch64-unknown-uefi`
- `i686-unknown-uefi`
- `x86_64-unknown-uefi`

Target maintainers

- David Rheinsberg ([@dvdhram](#))
- Nicholas Bishop ([@nicholasbishop](#))

Requirements

All UEFI targets can be used as `no-std` environments via cross-compilation. Support for `std` is missing, but actively worked on. `alloc` is supported if an allocator is provided by the user. No host tools are supported.

The UEFI environment resembles the environment for Microsoft Windows, with some minor differences. Therefore, cross-compiling for UEFI works with the same tools as cross-compiling for Windows. The target binaries are PE32+ encoded, the calling convention is different for each architecture, but matches what Windows uses (if the architecture is supported by Windows). The special `efiapi` Rust calling-convention chooses the right ABI for the target platform (`extern "C"` is incorrect on Intel targets at least). The specification has an elaborate section on the different supported calling-conventions, if more details are desired.

MMX, SSE, and other FP-units are disabled by default, to allow for compilation of core UEFI code that runs before they are set up. This can be overridden for individual compilations via `rustc` command-line flags. Not all firmwares correctly configure those units, though, so careful inspection is required.

As native to PE32+, binaries are position-dependent, but can be relocated at runtime if their desired location is unavailable. The code must be statically linked. Dynamic linking is not supported. Code is shared via UEFI interfaces, rather than dynamic linking. Additionally, UEFI forbids running code on anything but the boot CPU/thread, nor is interrupt-usage allowed (apart from the timer interrupt). Device drivers are required to use polling methods.

UEFI uses a single address-space to run all code in. Multiple applications can be loaded simultaneously and are dispatched via cooperative multitasking on a single stack.

By default, the UEFI targets use the `link`-flavor of the LLVM linker `lld` to link binaries into the final PE32+ file suffixed with `*.efi`. The PE subsystem is set to `EFI_APPLICATION`, but can be modified by passing `/subsystem:<...>` to the linker. Similarly, the entry-point is to `efi_main` but can be changed via `/entry:<...>`. The panic-strategy is set to `abort`,

The UEFI specification is available online for free: [UEFI Specification Directory](#)

Building rust for UEFI targets

Rust can be built for the UEFI targets by enabling them in the `rustc` build configuration. Note that you can only build the standard libraries. The compiler and host tools currently cannot be compiled for UEFI targets. A sample configuration would be:

```
[build]
build-stage = 1
target = ["x86_64-unknown-uefi"]
```

Building Rust programs

Starting with Rust 1.67, precompiled artifacts are provided via `rustup`. For example, to use `x86_64-unknown-uefi`:

```
# install cross-compile toolchain
rustup target add x86_64-unknown-uefi
# target flag may be used with any cargo or rustc command
cargo build --target x86_64-unknown-uefi
```

Testing

UEFI applications can be copied into the ESP on any UEFI system and executed via the firmware boot menu. The `qemu` suite allows emulating UEFI systems and executing UEFI applications as well. See its documentation for details.

The `uefi-run` rust tool is a simple wrapper around `qemu` that can spawn UEFI applications in `qemu`. You can install it via `cargo install uefi-run` and execute `qemu` applications as `uefi-run ./application.efi`.

Cross-compilation toolchains and C code

There are 3 common ways to compile native C code for UEFI targets:

- Use the official SDK by Intel: [Tianocore/EDK2](#). This supports a multitude of platforms, comes with the full specification transposed into C, lots of examples and build-system integrations. This is also the only officially supported platform by Intel, and is used by many major firmware implementations. Any code compiled via the SDK is compatible to rust binaries compiled for the UEFI targets. You can link them directly into your rust binaries, or call into each other via UEFI protocols.
- Use the [GNU-EFI](#) suite. This approach is used by many UEFI applications in the Linux/OSS ecosystem. The GCC compiler is used to compile ELF binaries, and linked with a pre-loader that converts the ELF binary to PE32+ at runtime. You can combine such binaries with the rust UEFI targets only via UEFI protocols. Linking both into the same executable will fail, since one is an ELF executable, and one a PE32+. If linking to [GNU-EFI](#) executables is desired, you must compile your rust code natively for the same GNU target as [GNU-EFI](#) and use their pre-loader. This requires careful consideration about which calling-convention to use when calling into native UEFI protocols, or calling into linked [GNU-EFI](#) code (similar to how these differences need to be accounted for when writing [GNU-EFI](#) C code).
- Use native Windows targets. This means compiling your C code for the Windows platform as if it was the UEFI platform. This works for static libraries, but needs adjustments when linking into an UEFI executable. You can, however, link such static libraries seamlessly into rust code compiled for UEFI targets. Be wary of any includes that are not specifically suitable for UEFI targets (especially the C standard library includes are not always compatible). Freestanding compilations are recommended to avoid incompatibilities.

Ecosystem

The rust language has a long history of supporting UEFI targets. Many crates have been developed to provide access to UEFI protocols and make UEFI programming more ergonomic in rust. The following list is a short overview (in alphabetical ordering):

- [efi](#): *Ergonomic Rust bindings for writing UEFI applications.* Provides *rustified* access to UEFI protocols, implements allocators and a safe environment to write UEFI applications.
- [r-efi](#): *UEFI Reference Specification Protocol Constants and Definitions.* A pure transpose of the UEFI specification into rust. This provides the raw definitions from the specification, without any extended helpers or *rustification*. It serves as baseline to implement any more elaborate rust UEFI layers.

- **uefi-rs**: Safe and easy-to-use wrapper for building UEFI apps. An elaborate library providing safe abstractions for UEFI protocols and features. It implements allocators and provides an execution environment to UEFI applications written in rust.
- **uefi-run**: Run UEFI applications. A small wrapper around `qemu` to spawn UEFI applications in an emulated `x86_64` machine.

Example: Freestanding

The following code is a valid UEFI application returning immediately upon execution with an exit code of 0. A panic handler is provided. This is executed by rust on panic. For simplicity, we simply end up in an infinite loop.

This example can be compiled as binary crate via `cargo`:

```
cargo build --target x86_64-unknown-uefi
```

```
#![no_main]
#![no_std]

#[panic_handler]
fn panic_handler(_info: &core::panic::PanicInfo) -> ! {
    loop {}
}

#[export_name = "efi_main"]
pub extern "C" fn main(_h: *mut core::ffi::c_void, _st: *mut core::ffi::c_void) -> usize {
    0
}
```

Example: Hello World

This is an example UEFI application that prints "Hello World!", then waits for key input before it exits. It serves as base example how to write UEFI applications without any helper modules other than the standalone UEFI protocol definitions provided by the `r-efi` crate.

This extends the "Freestanding" example and builds upon its setup. See there for instruction how to compile this as binary crate.

Note that UEFI uses UTF-16 strings. Since rust literals are UTF-8, we have to use an open-coded, zero-terminated, UTF-16 array as argument to `output_string()`. Similarly to the panic handler, real applications should rather use UTF-16 modules.

```
#![no_main]
#![no_std]

use r_efi::efi;

#[panic_handler]
fn panic_handler(_info: &core::panic::PanicInfo) -> ! {
    loop {}
}

#[export_name = "efi_main"]
pub extern "C" fn main(_h: efi::Handle, st: *mut efi::SystemTable) -> efi::Status {
    let s = [
        0x0048u16, 0x0065u16, 0x006cu16, 0x006cu16, 0x006fu16, // "Hello"
        0x0020u16, // "
        0x0057u16, 0x006fu16, 0x0072u16, 0x006cu16, 0x0064u16, // "World"
        0x0021u16, // !
        0x000au16, // "\n"
        0x0000u16, // NUL
    ];
    // Print "Hello World!".
    let r =
        unsafe { ((*(*st).con_out).output_string)((*st).con_out, s.as_ptr() as *mut efi::Char16) };
    if r.is_error() {
        return r;
    }

    // Wait for key input, by waiting on the `wait_for_key` event hook.
    let r = unsafe {
        let mut x: usize = 0;
        ((*(*st).boot_services).wait_for_event)(1, &mut (*(*st).con_in).wait_for_key, &mut x);
    };
    if r.is_error() {
        return r;
    }

    efi::Status::SUCCESS
}
```

wasm64-unknown-unknown

Tier: 3

WebAssembly target which uses 64-bit memories, relying on the [memory64](#) WebAssembly proposal.

Target maintainers

- Alex Crichton, <https://github.com/alexcrichton>

Requirements

This target is cross-compiled. The target supports `std` in the same manner as the `wasm32-unknown-unknown` target which is to say that it comes with the standard library but many I/O functions such as `std::fs` and `std::net` will simply return error. Additionally I/O operations like `println!` don't actually do anything and the prints aren't routed anywhere. This is the same as the `wasm32-unknown-unknown` target. This target comes by default with an allocator, currently `dmalloc` which is [ported to rust](#).

The difference of this target with `wasm32-unknown-unknown` is that it's compiled for 64-bit memories instead of 32-bit memories. This means that `usize` is 8-bytes large as well as pointers. The tradeoff, though, is that the maximum memory size is now the full 64-bit address space instead of the 4GB as limited by the 32-bit address space for `wasm32-unknown-unknown`.

This target is not a stable target. The [memory64](#) WebAssembly proposal is still in-progress and not standardized. This means that there are not many engines which implement the `memory64` feature and if they do they're likely behind a flag, for example:

- Nodejs - `--experimental-wasm-memory64`
- Wasmtime - `--wasm-features memory64`

Also note that at this time the `wasm64-unknown-unknown` target assumes the presence of other merged wasm proposals such as (with their LLVM feature flags):

- [Bulk memory](#) - `+bulk-memory`
- [Mutable imported globals](#) - `+mutable-globals`
- [Sign-extending operations](#) - `+sign-ext`
- [Non-trapping fp-to-int operations](#) - `+nontrapping-fptoint`

The `wasm64-unknown-unknown` target intends to match the default Clang targets for its "C" ABI, which is likely to be the same as Clang's `wasm32-unknown-unknown` largely.

Note: due to the relatively early-days nature of this target when working with this target you may encounter LLVM bugs. If an assertion hit or a bug is found it's recommended to open an issue either with `rust-lang/rust` or ideally with LLVM itself.

This target does not support `panic=unwind` at this time.

Building the target

You can build Rust with support for the target by adding it to the `target` list in `config.toml`, and the target also requires `lld` to be built to work.

```
[build]
target = ["wasm64-unknown-unknown"]

[rust]
lld = true
```

Building Rust programs

Rust does not yet ship pre-compiled artifacts for this target. To compile for this target, you will either need to build Rust with the target enabled (see "Building the target" above), or build your own copy of `std` by using `build-std` or similar.

Note that the following `cfg` directives are set for `wasm64-unknown-unknown`:

- `cfg(target_arch = "wasm64")`
- `cfg(target_family = "wasm")`

Testing

Currently testing is not well supported for `wasm64-unknown-unknown` and the Rust project doesn't run any tests for this target. Testing support sort of works but without `println!` it's not the most exciting tests to run.

Cross-compilation toolchains and C code

Compiling Rust code with C code for `wasm64-unknown-unknown` is theoretically possible, but there are no known toolchains to do this at this time. At the time of this writing there is no known "libc" for wasm that works with `wasm64-unknown-unknown`, which means that mixing C & Rust with this target effectively cannot be done.

x86_64-fortanix-unknown-sgx

Tier: 2

Secure enclaves using Intel Software Guard Extensions (SGX) based on the ABI defined by Fortanix for the Enclave Development Platform (EDP).

Target maintainers

The EDP team at Fortanix.

- Jethro Beekman [@jethrogb](#)
- Raoul Strackx [@raoulstrackx](#)
- Mohsen Zohrevandi [@mzohreva](#)

Requirements

The target supports `std` with a default allocator. Only cross compilation is supported.

Binaries support all CPUs that include Intel SGX. Only 64-bit mode is supported.

Not all `std` features are supported, see [Using Rust's std](#) for details.

The `extern "C"` calling convention is the System V AMD64 ABI.

The supported ABI is the [fortanix-sgx-abi](#).

The compiler output is ELF, but the native format for the platform is the SGX stream (SGXS) format. A converter like `ftxsgx-elf2sgxs` is needed.

Programs in SGXS format adhering to the Fortanix SGX ABI can be run with any compatible runner, such as [ftxsgx-runner](#).

See the [EDP installation guide](#) for recommendations on how to setup a development and runtime environment.

Building the target

As a tier 2 target, the target is built by the Rust project.

You can configure rustbuild like so:

```
[build]
build-stage = 1
target = ["x86_64-fortanix-unknown-sgx"]
```

Building Rust programs

Standard build flows using `cargo` or `rustc` should work.

Testing

The Rust test suite as well as custom unit and integration tests will run on hardware that has Intel SGX enabled if a cargo runner is configured correctly, see the requirements section.

Cross-compilation toolchains and C code

C code is not generally supported, as there is no libc. C code compiled for x86-64 in freestanding mode using the System V AMD64 ABI may work. The `rs-libc` crate contains a subset of libc that's known to work with this target.

x86_64-unknown-none

Tier: 2

Freestanding/bare-metal x86-64 binaries in ELF format: firmware, kernels, etc.

Target maintainers

- Harald Hoyer harald@profian.com, <https://github.com/haraldh>
- Mike Leany, <https://github.com/mikeleany>

Requirements

This target is cross-compiled. There is no support for `std`. There is no default allocator, but it's possible to use `alloc` by supplying an allocator.

By default, Rust code generated for this target does not use any vector or floating-point registers (e.g. SSE, AVX). This allows the generated code to run in environments, such as kernels, which may need to avoid the use of such registers or which may have special considerations about the use of such registers (e.g. saving and restoring them to avoid breaking userspace code using the same registers). You can change code generation to use additional CPU features via the `-C target-feature=` codegen options to `rustc`, or via the `#[target_feature]` mechanism within Rust code.

By default, code generated with this target should run on any `x86_64` hardware; enabling additional target features may raise this baseline.

Code generated with this target will use the `kernel` code model by default. You can change this using the `-C code-model=` option to `rustc`.

On `x86_64-unknown-none`, `extern "C"` uses the standard System V calling convention, without red zones.

This target generates binaries in the ELF format. Any alternate formats or special considerations for binary layout will require linker options or linker scripts.

Building the target

You can build Rust with support for the target by adding it to the `target` list in `config.toml`:

```
[build]
build-stage = 1
target = ["x86_64-unknown-none"]
```

Building Rust programs

Starting with Rust 1.62, precompiled artifacts are provided via `rustup`:

```
# install cross-compile toolchain
rustup target add x86_64-unknown-none
# target flag may be used with any cargo or rustc command
cargo build --target x86_64-unknown-none
```

Testing

As `x86_64-unknown-none` supports a variety of different environments and does not support `std`, this target does not support running the Rust test suite.

Cross-compilation toolchains and C code

If you want to compile C code along with Rust (such as for Rust crates with C dependencies), you will need an appropriate `x86_64` toolchain.

Rust *may* be able to use an `x86_64-linux-gnu-` toolchain with appropriate standalone flags to build for this toolchain (depending on the assumptions of that toolchain, see below), or you may wish to use a separate `x86_64-unknown-none` (or `x86_64-elf-`) toolchain.

On some `x86_64` hosts that use ELF binaries, you *may* be able to use the host C toolchain, if it does not introduce assumptions about the host environment that don't match the expectations of a standalone environment. Otherwise, you may need a separate toolchain for standalone/freestanding development, just as when cross-compiling from a non-`x86_64` platform.

Targets

`rustc` is a cross-compiler by default. This means that you can use any compiler to build for any architecture. The list of *targets* are the possible architectures that you can build for. See the [Platform Support](#) page for a detailed list of targets, or [Built-in Targets](#) for instructions on how to view what is available for your version of `rustc`.

To see all the options that you can set with a target, see the docs [here](#).

To compile to a particular target, use the `--target` flag:

```
$ rustc src/main.rs --target=wasm32-unknown-unknown
```

Target Features

`x86`, and `ARMv8` are two popular CPU architectures. Their instruction sets form a common baseline across most CPUs. However, some CPUs extend these with custom instruction sets, e.g. vector (`AVX`), bitwise manipulation (`BMI`) or cryptographic (`AES`).

Developers, who know on which CPUs their compiled code is going to run can choose to add (or remove) CPU specific instruction sets via the `-C target-feature=val` flag.

Please note, that this flag is generally considered as unsafe. More details can be found in [this section](#).

Built-in Targets

`rustc` ships with the ability to compile to many targets automatically, we call these "built-in" targets, and they generally correspond to targets that the team is supporting directly. To see the list of built-in targets, you can run `rustc --print target-list`.

Typically, a target needs a compiled copy of the Rust standard library to work. If using `rustup`, then check out the documentation on [Cross-compilation](#) on how to download a pre-built standard library built by the official Rust distributions. Most targets will need a system linker, and possibly other things.

Custom Targets

If you'd like to build for a target that is not yet supported by `rustc`, you can use a "custom target specification" to define a target. These target specification files are JSON. To see the JSON for the host target, you can run:

```
rustc +nightly -Z unstable-options --print target-spec-json
```

To see it for a different target, add the `--target` flag:

```
rustc +nightly -Z unstable-options --target=wasm32-unknown-unknown --print target-spec-json
```

To use a custom target, see the (unstable) [build-std](#) feature of [cargo](#).

Known Issues

This section informs you about known "gotchas". Keep in mind, that this section is (and always will be) incomplete. For suggestions and amendments, feel free to [contribute](#) to this guide.

Target Features

Most target-feature problems arise, when mixing code that have the target-feature *enabled* with code that have it *disabled*. If you want to avoid undefined behavior, it is recommended to build *all code* (including the standard library and imported crates) with a common set of target-features.

By default, compiling your code with the `-C target-feature` flag will not recompile the entire standard library and/or imported crates with matching target features. Therefore, target features are generally considered as unsafe. Using `#[target_feature]` on individual functions makes the function unsafe.

Examples:

Target-Feature	Issue	Seen on	Description	Details
<code>+soft-float</code> and <code>-sse</code>	Segfaults and ABI mismatches	<code>x86</code> and <code>x86-64</code>	<p>The <code>x86</code> and <code>x86_64</code> architecture uses SSE registers (aka <code>xmm</code>) for floating point operations. Using software emulated floats ("soft-floats") disables usage of <code>xmm</code> registers, but parts of Rust's core libraries (e.g. <code>std::f32</code> or <code>std::f64</code>) are compiled without soft-floats and expect parameters to be passed in <code>xmm</code> registers. This leads to ABI mismatches.</p> <p>Attempting to compile with disabled SSE causes the same error, too.</p>	#63466

Profile-guided Optimization

`rustc` supports doing profile-guided optimization (PGO). This chapter describes what PGO is, what it is good for, and how it can be used.

What Is Profiled-Guided Optimization?

The basic concept of PGO is to collect data about the typical execution of a program (e.g. which branches it is likely to take) and then use this data to inform optimizations such as inlining, machine-code layout, register allocation, etc.

There are different ways of collecting data about a program's execution. One is to run the program inside a profiler (such as `perf`) and another is to create an instrumented binary, that is, a binary that has data collection built into it, and run that. The latter usually provides more accurate data and it is also what is supported by `rustc`.

Usage

Generating a PGO-optimized program involves following a workflow with four steps:

1. Compile the program with instrumentation enabled (e.g. `rustc -Cprofile-generate=/tmp/pgo-data main.rs`)
2. Run the instrumented program (e.g. `./main`) which generates a `default_<id>.profraw` file
3. Convert the `.profraw` file into a `.profdata` file using LLVM's `llvm-profdata` tool
4. Compile the program again, this time making use of the profiling data (for example `rustc -Cprofile-use=merged.profdata main.rs`)

An instrumented program will create one or more `.profraw` files, one for each instrumented binary. E.g. an instrumented executable that loads two instrumented dynamic libraries at runtime will generate three `.profraw` files. Running an instrumented binary multiple times, on the other hand, will re-use the respective `.profraw` files, updating them in place.

These `.profraw` files have to be post-processed before they can be fed back into the compiler. This is done by the `llvm-profdata` tool. This tool is most easily installed via

```
rustup component add llvm-tools-preview
```

Note that installing the `llvm-tools-preview` component won't add `llvm-profdata` to the `PATH`. Rather, the tool can be found in:

```
~/.rustup/toolchains/<toolchain>/lib/rustlib/<target-triple>/bin/
```

Alternatively, an `llvm-propdata` coming with a recent LLVM or Clang version usually works too.

The `llvm-propdata` tool merges multiple `.profraw` files into a single `.profdata` file that can then be fed back into the compiler via `-Cprofile-use`:

```
# STEP 1: Compile the binary with instrumentation
rustc -Cprofile-generate=/tmp/pgo-data -O ./main.rs

# STEP 2: Run the binary a few times, maybe with common sets of args.
#           Each run will create or update `profraw` files in /tmp/pgo-data
./main mydata1.csv
./main mydata2.csv
./main mydata3.csv

# STEP 3: Merge and post-process all the `profraw` files in /tmp/pgo-data
llvm-propdata merge -o ./merged.profdata /tmp/pgo-data

# STEP 4: Use the merged `profdata` file during optimization. All `rustc`
#           flags have to be the same.
rustc -Cprofile-use=./merged.profdata -O ./main.rs
```

A Complete Cargo Workflow

Using this feature with Cargo works very similar to using it with `rustc` directly. Again, we generate an instrumented binary, run it to produce data, merge the data, and feed it back into the compiler. Some things of note:

- We use the `RUSTFLAGS` environment variable in order to pass the PGO compiler flags to the compilation of all crates in the program.
- We pass the `--target` flag to Cargo, which prevents the `RUSTFLAGS` arguments to be passed to Cargo build scripts. We don't want the build scripts to generate a bunch of `.profraw` files.
- We pass `--release` to Cargo because that's where PGO makes the most sense. In theory, PGO can also be done on debug builds but there is little reason to do so.
- It is recommended to use *absolute paths* for the argument of `-Cprofile-generate` and `-Cprofile-use`. Cargo can invoke `rustc` with varying working directories, meaning that `rustc` will not be able to find the supplied `.profdata` file. With absolute paths this is not an issue.
- It is good practice to make sure that there is no left-over profiling data from previous compilation sessions. Just deleting the directory is a simple way of doing so (see `STEP 0` below).

This is what the entire workflow looks like:

```
# STEP 0: Make sure there is no left-over profiling data from previous runs
rm -rf /tmp/pgo-data

# STEP 1: Build the instrumented binaries
RUSTFLAGS="-Cprofile-generate=/tmp/pgo-data" \
cargo build --release --target=x86_64-unknown-linux-gnu

# STEP 2: Run the instrumented binaries with some typical data
./target/x86_64-unknown-linux-gnu/release/myprogram mydata1.csv
./target/x86_64-unknown-linux-gnu/release/myprogram mydata2.csv
./target/x86_64-unknown-linux-gnu/release/myprogram mydata3.csv

# STEP 3: Merge the ``.profraw` files into a ``.profdata` file
llvm-propdata merge -o /tmp/pgo-data/merged.profdata /tmp/pgo-data

# STEP 4: Use the ``.profdata` file for guiding optimizations
RUSTFLAGS="-Cprofile-use=/tmp/pgo-data/merged.profdata" \
cargo build --release --target=x86_64-unknown-linux-gnu
```

Troubleshooting

- It is recommended to pass `-cllvm-args=-pgo-warn-missing-function` during the `-Cprofile-use` phase. LLVM by default does not warn if it cannot find profiling data for a given function. Enabling this warning will make it easier to spot errors in your setup.
- There is a [known issue](#) in Cargo prior to version 1.39 that will prevent PGO from working correctly. Be sure to use Cargo 1.39 or newer when doing PGO.

Further Reading

`rustc`'s PGO support relies entirely on LLVM's implementation of the feature and is equivalent to what Clang offers via the `-fprofile-generate` / `-fprofile-use` flags. The [Profile Guided Optimization](#) section in Clang's documentation is therefore an interesting read for anyone who wants to use PGO with Rust.

Instrumentation-based Code Coverage

Introduction

The Rust compiler includes two code coverage implementations:

- A GCC-compatible, gcov-based coverage implementation, enabled with `-Z profile`, which derives coverage data based on `DebugInfo`.
- A source-based code coverage implementation, enabled with `-C instrument-coverage`, which uses LLVM's native, efficient coverage instrumentation to generate very precise coverage data.

This document describes how to enable and use the LLVM instrumentation-based coverage, via the `-C instrument-coverage` compiler flag.

How it works

When `-C instrument-coverage` is enabled, the Rust compiler enhances rust-based libraries and binaries by:

- Automatically injecting calls to an LLVM intrinsic (`llvm.instrprof.increment`), at functions and branches in compiled code, to increment counters when conditional sections of code are executed.
- Embedding additional information in the data section of each library and binary (using the [LLVM Code Coverage Mapping Format Version 5](#), if compiling with LLVM 12, or [Version 6](#), if compiling with LLVM 13 or higher), to define the code regions (start and end positions in the source code) being counted.

When running a coverage-instrumented program, the counter values are written to a `profraw` file at program termination. LLVM bundles tools that read the counter results, combine those results with the coverage map (embedded in the program binary), and generate coverage reports in multiple formats.

Note: `-C instrument-coverage` also automatically enables `-C symbol-mangling-version=v0` (tracking issue [#60705](#)). The `v0` symbol mangler is strongly recommended. The `v0` demangler can be overridden by explicitly adding `-Z unstable-options -C symbol-mangling-version=legacy`.

Enable coverage profiling in the Rust compiler

Rust's source-based code coverage requires the Rust "profiler runtime". Without it, compiling with `-C instrument-coverage` generates an error that the profiler runtime is missing.

The Rust `nightly` distribution channel includes the profiler runtime, by default.

Important: If you are building the Rust compiler from the source distribution, the profiler runtime is *not* enabled in the default `config.example.toml`. Edit your `config.toml` file and ensure the `profiler` feature is set it to `true` (either under the `[build]` section, or under the settings for an individual `[target.<triple>]`):

```
# Build the profiler runtime (required when compiling with options that depend
# on this runtime, such as `'-C profile-generate` or `'-C instrument-coverage`).
profiler = true
```

Building the demangler

LLVM coverage reporting tools generate results that can include function names and other symbol references, and the raw coverage results report symbols using the compiler's "mangled" version of the symbol names, which can be difficult to interpret. To work around this issue, LLVM coverage tools also support a user-specified symbol name demangler.

One option for a Rust demangler is `rustfilt`, which can be installed with:

```
cargo install rustfilt
```

Another option, if you are building from the Rust compiler source distribution, is to use the `rust-demangler` tool included in the Rust source distribution, which can be built with:

```
$ ./x.py build rust-demangler
```

Compiling with coverage enabled

Set the `-C instrument-coverage` compiler flag in order to enable LLVM source-based code coverage profiling.

The default option generates coverage for all functions, including unused (never called) functions and generics. The compiler flag supports an optional value to tailor this behavior. (See `-C instrument-coverage=<options>`, below.)

With `cargo`, you can instrument your program binary and dependencies at the same time.

For example (if your project's `Cargo.toml` builds a binary by default):

```
$ cd your-project  
$ cargo clean  
$ RUSTFLAGS="-C instrument-coverage" cargo build
```

If `cargo` is not configured to use your `profiler`-enabled version of `rustc`, set the path explicitly via the `RUSTC` environment variable. Here is another example, using a `stage1` build of `rustc` to compile an `example` binary (from the `json5format` crate):

```
$ RUSTC=$HOME/rust/build/x86_64-unknown-linux-gnu/stage1/bin/rustc \  
RUSTFLAGS="-C instrument-coverage" \  
cargo build --example formatjson5
```

Note: that some compiler options, combined with `-C instrument-coverage`, can produce LLVM IR and/or linked binaries that are incompatible with LLVM coverage maps. For example, coverage requires references to actual functions in LLVM IR. If any covered function is optimized out, the coverage tools may not be able to process the coverage results. If you need to pass additional options, with coverage enabled, test them early, to confirm you will get the coverage results you expect.

Running the instrumented binary to generate raw coverage profiling data

In the previous example, `cargo` generated the coverage-instrumented binary `formatjson5`:

```
$ echo "{some: 'thing'}" | target/debug/examples/formatjson5 -
```

```
{  
    some: "thing",  
}
```

After running this program, a new file named like `default_11699812450447639123_0_20944` should be in the current working directory. A new, unique file name will be generated each time the program is run to avoid overwriting previous data.

```
$ echo "{some: 'thing'}" | target/debug/examples/formatjson5 -  
***  
$ ls default_*.profraw  
default_11699812450447639123_0_20944.profraw
```

You can also set a specific file name or path for the generated `.profraw` files by using the environment variable `LLVM_PROFILE_FILE`:

```
$ echo "{some: 'thing'}" \
| LLVM_PROFILE_FILE="formatjson5.profraw" target/debug/examples/formatjson5 -
...
$ ls formatjson5.profraw
formatjson5.profraw
```

If `LLVM_PROFILE_FILE` contains a path to a non-existent directory, the missing directory structure will be created. Additionally, the following special pattern strings are rewritten:

- `%p` - The process ID.
- `%h` - The hostname of the machine running the program.
- `%t` - The value of the `TMPDIR` environment variable.
- `%Nm` - the instrumented binary's signature: The runtime creates a pool of `N` raw profiles, used for on-line profile merging. The runtime takes care of selecting a raw profile from the pool, locking it, and updating it before the program exits. `N` must be between `1` and `9`, and defaults to `1` if omitted (with simply `%m`).
- `%c` - Does not add anything to the filename, but enables a mode (on some platforms, including Darwin) in which profile counter updates are continuously synced to a file. This means that if the instrumented program crashes, or is killed by a signal, perfect coverage information can still be recovered.

In the first example above, the value `11699812450447639123_0` in the generated filename is the instrumented binary's signature, which replaced the `%m` pattern and the value `20944` is the process ID of the binary being executed.

Installing LLVM coverage tools

LLVM's supplies two tools—`llvm-propdata` and `llvm-cov`—that process coverage data and generate reports. There are several ways to find and/or install these tools, but note that the coverage mapping data generated by the Rust compiler requires LLVM version 12 or higher, and processing the *raw* data may require exactly the LLVM version used by the compiler. (`llvm-cov --version` typically shows the tool's LLVM version number, and `rustc --verbose --version` shows the version of LLVM used by the Rust compiler.)

- You can install compatible versions of these tools via the `rustup` component `llvm-tools-preview`. This component is the recommended path, though the specific tools available and their interface is not currently subject to Rust's usual stability guarantees. In this case, you may also find `cargo-binutils` useful as a wrapper around these tools.
- You can install a compatible version of LLVM tools from your operating system distribution, or from your distribution of LLVM.

- If you are building the Rust compiler from source, you can optionally use the bundled LLVM tools, built from source. Those tool binaries can typically be found in your build platform directory at something like: `rust/build/x86_64-unknown-linux-gnu/llvm/bin/llvm-*`.

The examples in this document show how to use the `llvm` tools directly.

Creating coverage reports

Raw profiles have to be indexed before they can be used to generate coverage reports. This is done using `llvm-propdata merge`, which can combine multiple raw profiles and index them at the same time:

```
$ llvm-propdata merge -sparse formatjson5.profraw -o formatjson5.propdata
```

Finally, the `.propdata` file is used, in combination with the coverage map (from the program binary) to generate coverage reports using `llvm-cov report`, for a coverage summaries; and `llvm-cov show`, to see detailed coverage of lines and regions (character ranges) overlaid on the original source code.

These commands have several display and filtering options. For example:

```
$ llvm-cov show -Xdemangler=rustfilt target/debug/examples/formatjson5 \
  -instr-profile=formatjson5.propdata \
  -show-line-counts-or-regions \
  -show-instantiations \
  -name=add_quoted_string
```

```

<json5format::parser::Parser>::add_quoted_string:
439| 110|     match captured {
440| 110|         Some(unquoted) => {
441| 110|             if self.is_in_object()
442| 110|                 && !self.with_object(|object| object.has_pending_property()))
443| 110|                     ^34^0
444| 0|             {
445| 0|                 let captured = self.colon_capturer.capture(self.remaining);
446| 0|                 if self.consume_if_matched(captured) {
447| 0|                     if matches_unquoted_property_name(&unquoted) {
448| 0|                         self.set_pending_property(unquoted)
449| 0|                     } else {
450| 0|                         self.set_pending_property(&format!("{}{}{}", quote, &unquoted, quote))
451| 0|                     }
452| 0|                 } else {
453| 0|                     return Err(self.error("Property name separator (:) missing"));
454| 0|                 }
455| 110|             } else {
456| 110|                 let comments = self.take_pending_comments(); ^0
457| 110|                 self.add_value(Primitive::new(
458| 110|                     format!("{}{}{}", quote, &unquoted, quote),
459| 110|                     comments,
460| 0|                 ))
461| 0|             }
462| 0|         None => return Err(self.error("Unclosed string")),
463| 110|     }
464| 110| }
<json5format::parser::Parser>::add_quoted_string::{closure#0}:
442| 34|     && !self.with_object(|object| object.has_pending_property())?
<json5format::parser::Parser>::with_object::<<json5format::parser::Parser>::add_quoted_string::{closure#0}, bool>:
290| 34|     match &mut *self.current_scope().borrow_mut() {
291| 34|         Value::Object { val, .. } => f(val),
292| 0|         unexpected => Err(self.error(format!(
293| 0|             "Invalid Object token found while parsing an {}:{} (mismatched braces?)",
294| 0|             unexpected
295| 0|         )));
296| 0|     }
297| 34|

```

Some of the more notable options in this example include:

- `--Xdemangler=rustfilt` - the command name or path used to demangle Rust symbols (`rustfilt` in the example, but this could also be a path to the `rust-demangler` tool)
- `target/debug/examples/formatjson5` - the instrumented binary (from which to extract the coverage map)
- `--instr-profile=<path-to-file>.profdata` - the location of the `.profdata` file created by `llvm-profdata merge` (from the `.profraw` file generated by the instrumented binary)
- `--name=<exact-function-name>` - to show coverage for a specific function (or, consider using another filter option, such as `--name-regex=<pattern>`)

Note: Coverage can also be disabled on an individual function by annotating the function with the `[no_coverage]` attribute (which requires the feature flag `#![feature(no_coverage)]`).

Interpreting reports

There are four statistics tracked in a coverage summary:

- Function coverage is the percentage of functions that have been executed at least once. A function is considered to be executed if any of its instantiations are executed.
- Instantiation coverage is the percentage of function instantiations that have been executed at least once. Generic functions and functions generated from macros are two kinds of functions that may have multiple instantiations.
- Line coverage is the percentage of code lines that have been executed at least once. Only executable lines within function bodies are considered to be code lines.
- Region coverage is the percentage of code regions that have been executed at least once. A code region may span multiple lines: for example, in a large function body with no control flow. In other cases, a single line can contain multiple code regions: `return x || (y && z)` has countable code regions for `x` (which may resolve the expression, if `x` is `true`), `|| (y && z)` (executed only if `x` was `false`), and `return` (executed in either situation).

Of these four statistics, function coverage is usually the least granular while region coverage is the most granular. The project-wide totals for each statistic are listed in the summary.

Test coverage

A typical use case for coverage analysis is test coverage. Rust's source-based coverage tools can both measure your tests' code coverage as percentage, and pinpoint functions and branches not tested.

The following example (using the `json5format` crate, for demonstration purposes) show how to generate and analyze coverage results for all tests in a crate.

Since `cargo test` both builds and runs the tests, we set the additional `RUSTFLAGS`, to add the `-C instrument-coverage` flag.

```
$ RUSTFLAGS="-C instrument-coverage" \
cargo test --tests
```

Note: The default for `LLVM_PROFILE_FILE` is `default_%m_%p.profraw`. Versions prior to 1.65 had a default of `default.profraw`, so if using those earlier versions, it is recommended to explicitly set `LLVM_PROFILE_FILE="default_%m_%p.profraw"` to avoid having multiple tests overwrite the `.profraw` files.

Make note of the test binary file paths, displayed after the word "Running" in the test output:

```
...
Compiling json5format v0.1.3 ($HOME/json5format)
Finished test [unoptimized + debuginfo] target(s) in 14.60s

    Running target/debug/deps/json5format-fececd4653271682
running 25 tests
...
test result: ok. 25 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out

    Running target/debug/deps/lib-30768f9c53506dc5
running 31 tests
...
test result: ok. 31 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out
```

You should have one or more `.profraw` files now, one for each test binary. Run the `profdata` tool to merge them:

```
$ llvm-profdata merge -sparse default_*.profraw -o json5format.profdata
```

Then run the `cov` tool, with the `profdata` file and all test binaries:

```
$ llvm-cov report \
--use-color --ignore-filename-regex='/.cargo/registry' \
--instr-profile=json5format.profdata \
--object target/debug/deps/lib-30768f9c53506dc5 \
--object target/debug/deps/json5format-fececd4653271682
$ llvm-cov show \
--use-color --ignore-filename-regex='/.cargo/registry' \
--instr-profile=json5format.profdata \
--object target/debug/deps/lib-30768f9c53506dc5 \
--object target/debug/deps/json5format-fececd4653271682 \
--show-instantiations --show-line-counts-or-regions \
--Xdemangler=rustfilt | less -R
```

Note: If overriding the default `profraw` file name via the `LLVM_PROFILE_FILE` environment variable, it's highly recommended to use the `%m` and `%p` special pattern strings to generate unique file names in the case of more than a single test binary being executed.

Note: The command line option `--ignore-filename-regex=/.cargo/registry`, which excludes the sources for dependencies from the coverage results._

Tips for listing the binaries automatically

For `bash` users, one suggested way to automatically complete the `cov` command with the list of binaries is with a command like:

```
$ llvm-cov report \
$( \
    for file in \
        $( \
            RUSTFLAGS="-C instrument-coverage" \
            cargo test --tests --no-run --message-format=json \
                | jq -r "select(.profile.test == true) | .filenames[]" \
                | grep -v dSYM - \
        ); \
    do \
        printf "%s %s " -object $file; \
    done \
) \
--instr-profile=json5format.profdata --summary-only # and/or other options
```

Adding `--no-run --message-format=json` to the same `cargo test` command used to run the tests (including the same environment variables and flags) generates output in a JSON format that `jq` can easily query.

The `printf` command takes this list and generates the `--object <binary>` arguments for each listed test binary.

Including doc tests

The previous examples run `cargo test` with `--tests`, which excludes doc tests.¹

To include doc tests in the coverage results, drop the `--tests` flag, and apply the `-C instrument-coverage` flag, and some doc-test-specific options in the `RUSTDOCFLAGS` environment variable. (The `llvm-profdata` command does not change.)

```
$ RUSTFLAGS="-C instrument-coverage" \
RUSTDOCFLAGS="-C instrument-coverage -Z unstable-options --persist-doctests
target/debug/doctestbins" \
cargo test
$ llvm-profdata merge -sparse default_*.*.profraw -o json5format.profdata
```

The `-Z unstable-options --persist-doctests` flag is required, to save the test binaries (with their coverage maps) for `llvm-cov`.

```
$ llvm-cov report \
$( \
    for file in \
    $( \
        RUSTFLAGS="-C instrument-coverage" \
        RUSTDOCFLAGS="-C instrument-coverage -Z unstable-options --persist-doctests
target/debug/doctestbins" \
        cargo test --no-run --message-format=json \
            | jq -r "select(.profile.test == true) | .filenames[]" \
            | grep -v dSYM - \
    ) \
        target/debug/doctestbins/*/rust_out; \
do \
    [[ -x $file ]] && printf "%s %s" -object $file; \
done \
) \
--instr-profile=json5format.propdata --summary-only # and/or other options
```

Note: The differences in this `llvm-cov` invocation, compared with the version without doc tests, include:

- The `cargo test ... --no-run` command is updated with the same environment variables and flags used to build the tests, *including* the doc tests.
- The file glob pattern `target/debug/doctestbins/*/rust_out` adds the `rust_out` binaries generated for doc tests (note, however, that some `rust_out` files may not be executable binaries).
- `[[-x $file]] &&` filters the files passed on to the `printf`, to include only executable binaries.

1

There is ongoing work to resolve a known issue
[(#79417)](<https://github.com/rust-lang/rust/issues/79417>) that doc test coverage generates incorrect source line numbers in `llvm-cov show` results.

-C instrument-coverage=<options>

- `-C instrument-coverage=all`: Instrument all functions, including unused functions and unused generics. (This is the same as `-C instrument-coverage`, with no value.)
- `-C instrument-coverage=off`: Do not instrument any functions. (This is the same as simply not including the `-C instrument-coverage` option.)
- `-Zunstable-options -C instrument-coverage=except-unused-generics`: Instrument all functions except unused generics.
- `-Zunstable-options -C instrument-coverage=except-unused-functions`: Instrument only used (called) functions and instantiated generic functions.

Other references

Rust's implementation and workflow for source-based code coverage is based on the same library and tools used to implement [source-based code coverage in Clang](#). (This document is partially based on the Clang guide.)

Linker-plugin-based LTO

The `-C linker-plugin-lto` flag allows for deferring the LTO optimization to the actual linking step, which in turn allows for performing interprocedural optimizations across programming language boundaries if all the object files being linked were created by LLVM based toolchains. The prime example here would be linking Rust code together with Clang-compiled C/C++ code.

Usage

There are two main cases how linker plugin based LTO can be used:

- compiling a Rust `staticlib` that is used as a C ABI dependency
- compiling a Rust binary where `rustc` invokes the linker

In both cases the Rust code has to be compiled with `-C linker-plugin-lto` and the C/C++ code with `-flto` or `-flto=thin` so that object files are emitted as LLVM bitcode.

Rust `staticlib` as dependency in C/C++ program

In this case the Rust compiler just has to make sure that the object files in the `staticlib` are in the right format. For linking, a linker with the LLVM plugin must be used (e.g. LLD).

Using `rustc` directly:

```
# Compile the Rust staticlib
rustc --crate-type=staticlib -Clinker-plugin-lto -Copt-level=2 ./lib.rs
# Compile the C code with `'-flto=thin'
clang -c -O2 -flto=thin -o cmain.o ./cmain.c
# Link everything, making sure that we use an appropriate linker
clang -flto=thin -fuse-ld=lld -L . -l"name-of-your-rust-lib" -o main -O2 ./cmain.o
```

Using `cargo`:

```
# Compile the Rust staticlib
RUSTFLAGS="-Clinker-plugin-lto" cargo build --release
# Compile the C code with `'-flto=thin'
clang -c -O2 -flto=thin -o cmain.o ./cmain.c
# Link everything, making sure that we use an appropriate linker
clang -flto=thin -fuse-ld=lld -L . -l"name-of-your-rust-lib" -o main -O2 ./cmain.o
```

C/C++ code as a dependency in Rust

In this case the linker will be invoked by `rustc`. We again have to make sure that an appropriate linker is used.

Using `rustc` directly:

```
# Compile C code with `'-flto'
clang ./clib.c -flto=thin -c -o ./clib.o -O2
# Create a static library from the C code
ar crus ./libxyz.a ./clib.o

# Invoke `rustc` with the additional arguments
rustc -Clinker-plugin-lto -L. -Copt-level=2 -Clinker=clang -Clink-arg=-fuse-ld=lld ./main.rs
```

Using `cargo` directly:

```
# Compile C code with `'-flto'
clang ./clib.c -flto=thin -c -o ./clib.o -O2
# Create a static library from the C code
ar crus ./libxyz.a ./clib.o

# Set the linking arguments via RUSTFLAGS
RUSTFLAGS="-Clinker-plugin-lto -Clinker=clang -Clink-arg=-fuse-ld=lld" cargo build --release
```

Explicitly specifying the linker plugin to be used by `rustc`

If one wants to use a linker other than LLD, the LLVM linker plugin has to be specified explicitly. Otherwise the linker cannot read the object files. The path to the plugin is passed as an argument to the `-Clinker-plugin-lto` option:

```
rustc -Clinker-plugin-lto="/path/to/LLVMgold.so" -L. -Copt-level=2 ./main.rs
```

Usage with `clang-cl` and `x86_64-pc-windows-msvc`

Cross language LTO can be used with the `x86_64-pc-windows-msvc` target, but this requires using the `clang-cl` compiler instead of the MSVC `cl.exe` included with Visual Studio Build Tools, and linking with `lld-link`. Both `clang-cl` and `lld-link` can be downloaded from [LLVM's download page](#). Note that most crates in the ecosystem are likely to assume you are using `cl.exe` if using this target and that some things, like for example `vcpkg`, [don't work very well with clang-cl](#).

You will want to make sure your rust major LLVM version matches your installed LLVM tooling version, otherwise it is likely you will get linker errors:

```
rustc -V --verbose
clang-cl --version
```

If you are compiling any proc-macros, you will get this error:

```
error: Linker plugin based LTO is not supported together with `'-C prefer-dynamic` when targeting Windows-like targets
```

This is fixed if you explicitly set the target, for example `cargo build --target x86_64-pc-windows-msvc`. Without an explicit `--target` the flags will be passed to all compiler invocations (including build scripts and proc macros), see [cargo docs on rustflags](#)

If you have dependencies using the `cc` crate, you will need to set these environment variables:

```
set CC=clang-cl
set CXX=clang-cl
set CFLAGS=/clang:-flto=thin /clang:-fuse-ld=lld-link
set CXXFLAGS=/clang:-flto=thin /clang:-fuse-ld=lld-link
REM Needed because msvc's lib.exe crashes on LLVM LTO .obj files
set AR=llvm-lib
```

If you are specifying lld-link as your linker by setting `linker = "lld-link.exe"` in your cargo config, you may run into issues with some crates that compile code with separate cargo invocations. You should be able to get around this problem by setting `-Clinker=lld-link` in RUSTFLAGS

Toolchain Compatibility

In order for this kind of LTO to work, the LLVM linker plugin must be able to handle the LLVM bitcode produced by both `rustc` and `clang`.

Best results are achieved by using a `rustc` and `clang` that are based on the exact same version of LLVM. One can use `rustc -vV` in order to view the LLVM used by a given `rustc` version. Note that the version number given here is only an approximation as Rust sometimes uses unstable revisions of LLVM. However, the approximation is usually reliable.

The following table shows known good combinations of toolchain versions.

Rust Version	Clang Version
1.34 - 1.37	8
1.38 - 1.44	9
1.45 - 1.46	10
1.47 - 1.51	11
1.52 - 1.55	12
1.56 - 1.59	13
1.60 - 1.64	14
1.65	15

Note that the compatibility policy for this feature might change in the future.

Exploit Mitigations

This chapter documents the exploit mitigations supported by the Rust compiler, and is by no means an extensive survey of the Rust programming language's security features.

This chapter is for software engineers working with the Rust programming language, and assumes prior knowledge of the Rust programming language and its toolchain.

Introduction

The Rust programming language provides memory[1] and thread[2] safety guarantees via its ownership[3], references and borrowing[4], and slice types[5] features. However, Unsafe Rust[6] introduces unsafe blocks, unsafe functions and methods, unsafe traits, and new types that are not subject to the borrowing rules.

Parts of the Rust standard library are implemented as safe abstractions over unsafe code (and historically have been vulnerable to memory corruption[7]). Furthermore, the Rust code and documentation encourage creating safe abstractions over unsafe code. This can cause a false sense of security if unsafe code is not properly reviewed and tested.

Unsafe Rust introduces features that do not provide the same memory and thread safety guarantees. This causes programs or libraries to be susceptible to memory corruption (CWE-119)[8] and concurrency issues (CWE-557)[9]. Modern C and C++ compilers provide exploit mitigations to increase the difficulty to exploit vulnerabilities resulting from these issues. Therefore, the Rust compiler must also support these exploit mitigations in order to mitigate vulnerabilities resulting from the use of Unsafe Rust. This chapter documents these exploit mitigations and how they apply to Rust.

This chapter does not discuss the effectiveness of these exploit mitigations as they vary greatly depending on several factors besides their design and implementation, but rather describe what they do, so their effectiveness can be understood within a given context.

Exploit mitigations

This section documents the exploit mitigations applicable to the Rust compiler when building programs for the Linux operating system on the AMD64 architecture and equivalent.¹

The Rust Programming Language currently has no specification. The Rust compiler (i.e., `rustc`) is the language reference implementation. All references to "the Rust

compiler” in this chapter refer to the language reference implementation.

Table I

Summary of exploit mitigations supported by the Rust compiler when building programs for the Linux operating system on the AMD64 architecture and equivalent.

Exploit mitigation	Supported and enabled by default	Since
Position-independent executable	Yes	0.12.0 (2014-10-09)
Integer overflow checks	Yes (enabled when debug assertions are enabled, and disabled when debug assertions are disabled)	1.1.0 (2015-06-25)
Non-executable memory regions	Yes	1.8.0 (2016-04-14)
Stack clashing protection	Yes	1.20.0 (2017-08-31)
Read-only relocations and immediate binding	Yes	1.21.0 (2017-10-12)
Heap corruption protection	Yes	1.32.0 (2019-01-17) (via operating system default or specified allocator)
Stack smashing protection	No	
Forward-edge control flow protection	Yes	Nightly
Backward-edge control flow protection (e.g., shadow and safe stack)	No	

- See https://github.com/rust-lang/rust/tree/master/compiler/rustc_target/src/spec for a list of targets and their default options. ↴

Position-independent executable

Position-independent executable increases the difficulty of the use of code reuse exploitation techniques, such as return-oriented programming (ROP) and variants, by generating position-independent code for the executable, and instructing the dynamic linker to load it similarly to a shared object at a random load address, thus also benefiting from address-space layout randomization (ASLR). This is also referred to as “full ASLR”.

The Rust compiler supports position-independent executable, and enables it by default since version 0.12.0 (2014-10-09)[10]-[13].

```
$ readelf -h target/release/hello-rust | grep Type:
  Type: DYN (Shared object file)
```

Fig. 1. Checking if an executable is a position-independent executable.

An executable with an object type of `ET_DYN` (i.e., shared object) and not `ET_EXEC` (i.e., executable) is a position-independent executable (see Fig. 1).

Integer overflow checks

Integer overflow checks protects programs from undefined and unintended behavior (which may cause vulnerabilities) by checking for results of signed and unsigned integer computations that cannot be represented in their type, resulting in an overflow or wraparound.

The Rust compiler supports integer overflow checks, and enables it when debug assertions are enabled since version 1.1.0 (2015-06-25)[14]-[20].

```
fn main() {
    let u: u8 = 255;
    println!("u: {}", u + 1);
}
```

Fig. 2. hello-rust-integer program.

```
$ cargo run
Compiling hello-rust-integer v0.1.0 (/home/rcvalle/hello-rust-integer)
Finished dev [unoptimized + debuginfo] target(s) in 0.23s
    Running `target/debug/hello-rust-integer`
thread 'main' panicked at 'attempt to add with overflow', src/main.rs:3:23
note: run with `RUST_BACKTRACE=1` environment variable to display a backtrace.
```

Fig. 3. Build and execution of hello-rust-integer with debug assertions enabled.

```
$ cargo run --release
Compiling hello-rust-integer v0.1.0 (/home/rcvalle/hello-rust-integer)
Finished release [optimized] target(s) in 0.23s
    Running `target/release/hello-rust-integer`
u: 0
```

Fig. 4. Build and execution of hello-rust-integer with debug assertions disabled.

Integer overflow checks are enabled when debug assertions are enabled (see Fig. 3), and disabled when debug assertions are disabled (see Fig. 4). To enable integer overflow checks independently, use the option to control integer overflow checks, scoped attributes, or explicit checking methods such as `checked_add`².

It is recommended that explicit wrapping methods such as `wrapping_add` be used when wrapping semantics are intended, and that explicit checking and wrapping methods always be used when using Unsafe Rust.

2. See the [u32 docs](#) for more information on the checked, overflowing, saturating, and wrapping methods (using u32 as an example). ↩

Non-executable memory regions

Non-executable memory regions increase the difficulty of exploitation by limiting the memory regions that can be used to execute arbitrary code. Most modern processors provide support for the operating system to mark memory regions as non executable, but it was previously emulated by software, such as in grsecurity/PaX's [PAGEEXEC](#) and [SEGMEXEC](#), on processors that did not provide support for it. This is also known as "No Execute (NX) Bit", "Execute Disable (XD) Bit", "Execute Never (XN) Bit", and others.

The Rust compiler supports non-executable memory regions, and enables it by default since its initial release, version 0.1 (2012-01-20)[21], [22], but has regressed since then[23]-[25], and enforced by default since version 1.8.0 (2016-04-14)[25].

```
$ readelf -l target/release/hello-rust | grep -A 1 GNU_STACK
GNU_STACK      0x0000000000000000 0x0000000000000000 0x0000000000000000
                0x0000000000000000 0x0000000000000000 RW      0x10
```

Fig. 5. Checking if non-executable memory regions are enabled for a given binary.

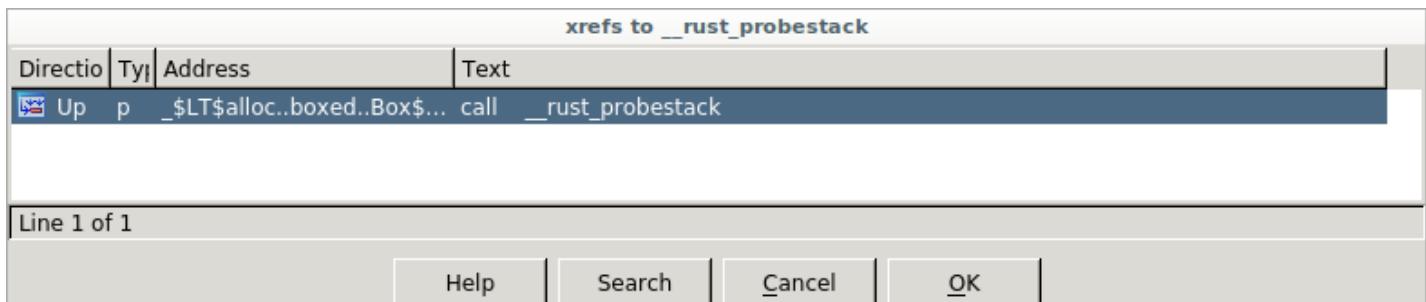
The presence of an element of type `PT_GNU_STACK` in the program header table with the `PF_X` (i.e., executable) flag unset indicates non-executable memory regions³ are enabled for a given binary (see Fig. 5). Conversely, the presence of an element of type `PT_GNU_STACK` in the program header table with the `PF_X` flag set or the absence of an element of type `PT_GNU_STACK` in the program header table indicates non-executable memory regions are not enabled for a given binary.

3. See the Appendix section for more information on why it affects other memory regions besides the stack. ↩

Stack clashing protection

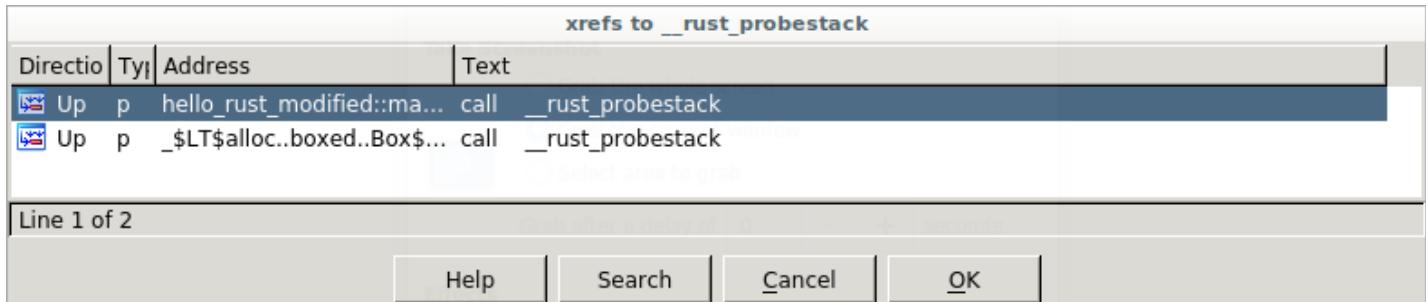
Stack clashing protection protects the stack from overlapping with another memory region—allowing arbitrary data in both to be overwritten using each other—by reading from the stack pages as the stack grows to cause a page fault when attempting to read from the guard page/region. This is also referred to as "stack probes" or "stack probing".

The Rust compiler supports stack clashing protection via stack probing, and enables it by default since version 1.20.0 (2017-08-31)[26]-[29].

Fig. 6. IDA Pro listing cross references to `__rust_probestack` in `hello-rust`.

```
fn hello() {
    println!("Hello, world!");
}

fn main() {
    let _: [u64; 1024] = [0; 1024];
    hello();
}
```

Fig 7. Modified `hello-rust`.Fig. 8. IDA Pro listing cross references to `__rust_probestack` in modified `hello-rust`.

To check if stack clashing protection is enabled for a given binary, search for cross references to `__rust_probestack`. The `__rust_probestack` is called in the prologue of functions whose stack size is larger than a page size (see Fig. 6), and can be forced for illustration purposes by modifying the `hello-rust` example as seen in Fig. 7 and Fig. 8.

Read-only relocations and immediate binding

Read-only relocations protect segments containing relocations and relocation information (i.e., `.init_array`, `.fini_array`, `.dynamic`, and `.got`) from being overwritten by marking these segments read only. This is also referred to as “partial RELRO”.

The Rust compiler supports read-only relocations, and enables it by default since version 1.21.0 (2017-10-12)[30], [31].

```
$ readelf -l target/release/hello-rust | grep GNU_RELRO
GNU_RELRO      0x000000000002ee00 0x0000000000002fe00 0x0000000000002fe00
```

Fig. 9. Checking if read-only relocations is enabled for a given binary.

The presence of an element of type `PT_GNU_RELRO` in the program header table indicates read-only relocations are enabled for a given binary (see Fig. 9). Conversely, the absence of an element of type `PT_GNU_RELRO` in the program header table indicates read-only relocations are not enabled for a given binary.

Immediate binding protects additional segments containing relocations (i.e., `.got.plt`) from being overwritten by instructing the dynamic linker to perform all relocations before transferring control to the program during startup, so all segments containing relocations can be marked read only (when combined with read-only relocations). This is also referred to as “full RELRO”.

The Rust compiler supports immediate binding, and enables it by default since version 1.21.0 (2017-10-12)[30], [31].

```
$ readelf -d target/release/hello-rust | grep BIND_NOW
0x000000000000001e (FLAGS)
          BIND_NOW
```

Fig. 10. Checking if immediate binding is enabled for a given binary.

The presence of an element with the `DT_BIND_NOW` tag and the `DF_BIND_NOW` flag⁴ in the dynamic section indicates immediate binding is enabled for a given binary (see Fig. 10). Conversely, the absence of an element with the `DT_BIND_NOW` tag and the `DF_BIND_NOW` flag in the dynamic section indicates immediate binding is not enabled for a given binary.

The presence of both an element of type `PT_GNU_RELRO` in the program header table and of an element with the `DT_BIND_NOW` tag and the `DF_BIND_NOW` flag in the dynamic section indicates full RELRO is enabled for a given binary (see Fig. 9 and Fig. 10).

4. And the `DF_1_NOW` flag for some link editors. ↩

Heap corruption protection

Heap corruption protection protects memory allocated dynamically by performing several checks, such as checks for corrupted links between list elements, invalid pointers, invalid sizes, double/multiple “frees” of the same memory allocated, and many corner cases of these. These checks are implementation specific, and vary per allocator.

[ARM Memory Tagging Extension \(MTE\)](#), when available, will provide hardware assistance for a probabilistic mitigation to detect memory safety violations by tagging memory allocations, and automatically checking that the correct tag is used on every memory access.

Rust’s default allocator has historically been `jemalloc`, and it has long been the cause of issues and the subject of much discussion[32]-[38]. Consequently, it has

been removed as the default allocator in favor of the operating system's standard C library default allocator⁵ since version 1.32.0 (2019-01-17)[39].

```
fn main() {
    let mut x = Box::new([0; 1024]);

    for i in 0..1026 {
        unsafe {
            let elem = x.get_unchecked_mut(i);
            *elem = 0x4141414141414141u64;
        }
    }
}
```

Fig. 11. hello-rust-heap program.

```
$ cargo run
Compiling hello-rust-heap v0.1.0 (/home/rcvalle/hello-rust-heap)
Finished dev [unoptimized + debuginfo] target(s) in 0.25s
    Running `target/debug/hello-rust-heap`
free(): invalid next size (normal)
Aborted
```

Fig. 12. Build and execution of hello-rust-heap with debug assertions enabled.

```
$ cargo run --release
Compiling hello-rust-heap v0.1.0 (/home/rcvalle/hello-rust-heap)
Finished release [optimized] target(s) in 0.25s
    Running `target/release/hello-rust-heap`
free(): invalid next size (normal)
Aborted
```

Fig. 13. Build and execution of hello-rust-heap with debug assertions disabled.

Heap corruption checks are being performed when using the default allocator (i.e., the GNU Allocator) as seen in Fig. 12 and Fig. 13.

⁵ Linux's standard C library default allocator is the GNU Allocator, which is derived from ptmalloc (pthread malloc) by Wolfram Gloger, which in turn is derived from dlmalloc (Doug Lea malloc) by Doug Lea. ↩

Stack smashing protection

Stack smashing protection protects programs from stack-based buffer overflows by inserting a random guard value between local variables and the saved return instruction pointer, and checking if this value has changed when returning from a function. This is also known as "Stack Protector" or "Stack Smashing Protector (SSP)".

The Rust compiler does not support stack smashing protection. However, more comprehensive alternatives to stack smashing protection exist, such as shadow and safe stack (see backward-edge control flow protection).

xrefs to <code>__stack_chk_fail</code>				
Directio	Ty	Address	Text	
[D]	p	fileline_initialize:loc_101...	call	<code>__stack_chk_fail</code>
[D]	p	_rdos_backtrace_creat...	call	<code>__stack_chk_fail</code>
[D]	p	elf_zlib_inflate_table:loc...	call	<code>__stack_chk_fail</code>
[D]	p	elf_try_debugfile:loc_10...	call	<code>__stack_chk_fail</code>
[D]	p	elf_zlib_inflate_and_verif...	call	<code>__stack_chk_fail</code>
[D]	p	elf_add:loc_12A1E	call	<code>__stack_chk_fail</code>
[D]	p	phdr_callback:loc_1318E	call	<code>__stack_chk_fail</code>
[D]	p	_rdos_backtrace_initiali...	call	<code>__stack_chk_fail</code>
[D]	p	dwarf_buf_error:loc_135...	call	<code>__stack_chk_fail</code>
[D]	p	find_address_ranges:loc...	call	<code>__stack_chk_fail</code>
[D]	p	read_referenced_name:l...	call	<code>__stack_chk_fail</code>
[D]	p	read_function_entry:loc...	call	<code>__stack_chk_fail</code>
[D]	p	dwarf_lookup_pc:loc_16...	call	<code>__stack_chk_fail</code>
[D]	p	dwarf_fileline:loc_16A69	call	<code>__stack_chk_fail</code>
[D]	p	_rdos_backtrace_dwarf...	call	<code>__stack_chk_fail</code>

Line 1 of 15

[Help](#)[Search](#)[Cancel](#)[OK](#)Fig. 14. IDA Pro listing cross references to `__stack_chk_fail` in hello-rust.

To check if stack smashing protection is enabled for a given binary, search for cross references to `__stack_chk_fail`. The only cross references to `__stack_chk_fail` in hello-rust are from the statically-linked libbacktrace library (see Fig. 14).

Forward-edge control flow protection

Forward-edge control flow protection protects programs from having its control flow changed/hijacked by performing checks to ensure that destinations of indirect branches are one of their valid destinations in the control flow graph. The comprehensiveness of these checks vary per implementation. This is also known as “forward-edge control flow integrity (CFI)”.

Newer processors provide hardware assistance for forward-edge control flow protection, such as ARM Branch Target Identification (BTI), ARM Pointer Authentication, and Intel Indirect Branch Tracking (IBT) as part of Intel Control-flow Enforcement Technology (CET). However, ARM BTI and Intel IBT -based implementations are less comprehensive than software-based implementations such as [LLVM ControlFlowIntegrity \(CFI\)](#), and the commercially available [grsecurity/PaX Reuse Attack Protector \(RAP\)](#).

The Rust compiler supports forward-edge control flow protection on nightly builds[40]-[41] ⁶.

```
$ readelf -s -W target/debug/rust-cfi | grep "\.cfi"
 12: 0000000000005170    46 FUNC    LOCAL  DEFAULT   14
_RNvCsjaOHoaNjor6_8rust_cfi7add_one.cfi
 15: 00000000000051a0    16 FUNC    LOCAL  DEFAULT   14
_RNvCsjaOHoaNjor6_8rust_cfi7add_two.cfi
 17: 0000000000005270   396 FUNC    LOCAL  DEFAULT   14
_RNvCsjaOHoaNjor6_8rust_cfi4main.cfi
...
...
```

Fig. 15. Checking if LLVM CFI is enabled for a given binary[41].

The presence of symbols suffixed with ".cfi" or the `__cfi_init` symbol (and references to `__cfi_check`) indicates that LLVM CFI (i.e., forward-edge control flow protection) is enabled for a given binary. Conversely, the absence of symbols suffixed with ".cfi" or the `__cfi_init` symbol (and references to `__cfi_check`) indicates that LLVM CFI is not enabled for a given binary (see Fig. 15).

6. It also supports Control Flow Guard (CFG) on Windows (see <https://github.com/rust-lang/rust/issues/68793>). ↩

Backward-edge control flow protection

Shadow stack protects saved return instruction pointers from being overwritten by storing a copy of them on a separate (shadow) stack, and using these copies as authoritative values when returning from functions. This is also known as “ShadowCallStack” and “Return Flow Guard”, and is considered an implementation of backward-edge control flow protection (or “backward-edge CFI”).

Safe stack protects not only the saved return instruction pointers, but also register spills and some local variables from being overwritten by storing unsafe variables, such as large arrays, on a separate (unsafe) stack, and using these unsafe variables on the separate stack instead. This is also known as “SafeStack”, and is also considered an implementation of backward-edge control flow protection.

Both shadow and safe stack are intended to be a more comprehensive alternatives to stack smashing protection as they protect the saved return instruction pointers (and other data in the case of safe stack) from arbitrary writes and non-linear out-of-bounds writes.

Newer processors provide hardware assistance for backward-edge control flow protection, such as ARM Pointer Authentication, and Intel Shadow Stack as part of Intel CET.

The Rust compiler does not support shadow or safe stack. There is work currently ongoing to add support for the sanitizers[40], which may or may not include support for safe stack⁷.

```
$ readelf -s target/release/hello-rust | grep __safestack_init
```

Fig. 16. Checking if LLVM SafeStack is enabled for a given binary.

The presence of the `__safestack_init` symbol indicates that LLVM SafeStack is enabled for a given binary. Conversely, the absence of the `__safestack_init` symbol indicates that LLVM SafeStack is not enabled for a given binary (see Fig. 16).

7. The shadow stack implementation for the AMD64 architecture and equivalent in LLVM was removed due to performance and security issues. ↩

Appendix

As of the latest version of the [Linux Standard Base \(LSB\) Core Specification](#), the `PT_GNU_STACK` program header indicates whether the stack should be executable, and the absence of this header indicates that the stack should be executable. However, the Linux kernel currently sets the `READ_IMPLIES_EXEC` personality upon loading any executable with the `PT_GNU_STACK` program header and the `PF_X` flag set or with the absence of this header, resulting in not only the stack, but also all readable virtual memory mappings being executable.

An attempt to fix this [was made in 2012](#), and another [was made in 2020](#). The former never landed, and the latter partially fixed it, but introduced other issues—the absence of the `PT_GNU_STACK` program header still causes not only the stack, but also all readable virtual memory mappings to be executable in some architectures, such as IA-32 and equivalent (or causes the stack to be non-executable in some architectures, such as AMD64 and equivalent, contradicting the LSB).

The `READ_IMPLIES_EXEC` personality needs to be completely separated from the `PT_GNU_STACK` program header by having a separate option for it (or `setarch -X` could just be used whenever `READ_IMPLIES_EXEC` is needed), and the absence of the `PT_GNU_STACK` program header needs to have more secure defaults (unrelated to `READ_IMPLIES_EXEC`).

References

1. D. Hosfelt. “Fearless security: memory safety.” Mozilla Hacks.
<https://hacks.mozilla.org/2019/01/fearless-security-memory-safety/>.
2. D. Hosfelt. “Fearless security: thread safety.” Mozilla Hacks.
<https://hacks.mozilla.org/2019/02/fearless-security-thread-safety/>.
3. S. Klabnik and C. Nichols. “What Is Ownership?.” The Rust Programming Language.
<https://doc.rust-lang.org/book/ch04-01-what-is-ownership.html>.
4. S. Klabnik and C. Nichols. “References and Borrowing.” The Rust Programming Language. <https://doc.rust-lang.org/book/ch04-02-references-and-borrowing.html>.
5. S. Klabnik and C. Nichols. “The Slice Type.” The Rust Programming Language. <https://doc.rust-lang.org/book/ch04-03-slices.html>.

6. S. Klabnik and C. Nichols. "Unsafe Rust." The Rust Programming Language. <https://doc.rust-lang.org/book/ch19-01-unsafe-rust.html>.
7. S. Davidoff. "How Rust's standard library was vulnerable for years and nobody noticed." Medium. <https://medium.com/@shnatsel/how-rusts-standard-library-was-vulnerable-for-years-and-nobody-noticed-aebf0503c3d6>.
8. "Improper restriction of operations within the bounds of a memory buffer (CWE-119)." MITRE CWE List. <https://cwe.mitre.org/data/definitions/119.html>.
9. "Concurrency issues (CWE-557)." MITRE CWE List. <https://cwe.mitre.org/data/definitions/557.html>.
10. K. McAllister. "Memory exploit mitigations #15179." GitHub. <https://github.com/rust-lang/rust/issues/15179>.
11. K. McAllister. "RFC: Memory exploit mitigation #145." GitHub. <https://github.com/rust-lang/rfcs/pull/145>.
12. K. McAllister. "RFC: Memory exploit mitigation." GitHub. <https://github.com/kmcallister/rfcs/blob/hardening/active/0000-memory-exploit-mitigation.md>.
13. D. Micay. "Enable PIE by default on Linux for full ASLR #16340." GitHub. <https://github.com/rust-lang/rust/pull/16340>.
14. N. Matsakis. "Integer overflow #560." GitHub. <https://github.com/rust-lang/rfcs/pull/560>.
15. G. Lehel and N. Matsakis. "Integer overflow." GitHub. <https://rust-lang.github.io/rfcs/0560-integer-overflow.html>.
16. A. Turon. "Tracking issue for integer overflow (RFC 560) #22020." GitHub. <https://github.com/rust-lang/rust/issues/22020>.
17. H. Wilson. "Myths and legends about integer overflow in Rust." Huon on the Internet. <http://huonw.github.io/blog/2016/04/myths-and-legends-about-integer-overflow-in-rust/>.
18. B. Anderson. "Stabilize -C overflow-checks #1535." GitHub. <https://github.com/rust-lang/rfcs/pull/1535>.
19. B. Anderson. "Stable overflow checks." GitHub. <https://github.com/brson/rfcs/blob/overflow/text/0000-stable-overflow-checks.md>.
20. N. Froyd. "Add -C overflow-checks option #40037." GitHub. <https://github.com/rust-lang/rust/pull/40037>.
21. R. Á. de Espíndola. "rustc requires executable stack #798." GitHub. <https://github.com/rust-lang/rust/issues/798>.

22. A. Seipp. "Make sure librustrt.so is linked with a non-executable stack. #1066." GitHub. <https://github.com/rust-lang/rust/pull/1066>.
23. D. Micay. "Rust binaries should not have an executable stack #5643." GitHub. <https://github.com/rust-lang/rust/issues/5643>.
24. D. Micay. "Mark the assembly object stacks as non-executable #5647." GitHub. <https://github.com/rust-lang/rust/pull/5647>.
25. A. Clark. "Explicitly disable stack execution on linux and bsd #30859." GitHub. <https://github.com/rust-lang/rust/pull/30859>.
26. "Replace stack overflow checking with stack probes #16012." GitHub. <https://github.com/rust-lang/rust/issues/16012>.
27. B. Striegel. "Extend stack probe support to non-tier-1 platforms, and clarify policy for mitigating LLVM-dependent unsafety #43241." GitHub. <https://github.com/rust-lang/rust/issues/43241>.
28. A. Crichton. "rustc: Implement stack probes for x86 #42816." GitHub. <https://github.com/rust-lang/rust/pull/42816>.
29. A. Crichton. "Add __rust_probestack intrinsic #175." GitHub. <https://github.com/rust-lang/compiler-builtins/pull/175>.
30. B. Anderson. "Consider applying -Wl,-z,relro or -Wl,-z,relro,-z,now by default #29877." GitHub. <https://github.com/rust-lang/rust/issues/29877>.
31. J. Löthberg. "Add support for full RELRO #43170." GitHub. <https://github.com/rust-lang/rust/pull/43170>.
32. N. Matsakis. "Allocators in Rust." Baby Steps. <http://smallcultfollowing.com/babysteps/blog/2014/11/14/allocators-in-rust/>.
33. A. Crichton. "RFC: Allow changing the default allocator #1183." GitHub. <https://github.com/rust-lang/rfcs/pull/1183>.
34. A. Crichton. "RFC: Swap out jemalloc." GitHub. <https://rust-lang.github.io/rfcs/1183-swap-out-jemalloc.html>.
35. A. Crichton. "Tracking issue for changing the global, default allocator (RFC 1974) #27389." GitHub. <https://github.com/rust-lang/rust/issues/27389>.
36. S. Fackler. "Prepare global allocators for stabilization #1974." GitHub. <https://github.com/rust-lang/rfcs/pull/1974>.
37. A. Crichton. "RFC: Global allocators." GitHub. <https://rust-lang.github.io/rfcs/1974-global-allocators.html>.
38. B. Anderson. "Switch the default global allocator to System, remove alloc_jemalloc, use jemallocator in rustc #36963." GitHub. <https://github.com/rust-lang/rust/issues/36963>.

39. A. Crichton. “Remove the alloc_jemalloc crate #55238.” GitHub. <https://github.com/rust-lang/rust/pull/55238>.
40. R. de C Valle. “Tracking Issue for LLVM Control Flow Integrity (CFI) Support for Rust #89653.” GitHub. <https://github.com/rust-lang/rust/issues/89653>.
41. “ControlFlowIntegrity.” The Rust Unstable Book. <https://doc.rust-lang.org/unstable-book/compiler-flags/sanitizer.html#controlflowintegrity>.

Contributing to rustc

We'd love to have your help improving `rustc`! To that end, we've written [a whole book](#) on its internals, how it works, and how to get started working on it. To learn more, you'll want to check that out.

If you would like to contribute to *this* book, you can find its source in the `rustc` source at [src/doc/rustc](#).